



For workspace administrators

Administrator Guide

Aperture Chat

The complete User Guide, plus the Admin console: accounts, groups, model access, connections, single sign-on, analytics, token budgets, the tenant audit trail, and alerts. No prior knowledge assumed.

Contents

PART 1 · GETTING STARTED

1	Welcome	4
2	Signing in	4
3	Finding your way around	5
4	Organizing chats: folders, pins, and recent	6
5	Searching everything	7

PART 2 · CHAT

6	Your first chat	8
7	Reading and acting on replies	9
8	Composer symbol shortcuts	10
9	Attaching files, web pages, and connected sources	10
10	Send options: Knowledge, Web, Agent, and Reasoning	11
11	Session details and sources	12

PART 3 · DRAFTS, DECKS, AND AGENTS

12	Drafting documents	13
13	Building a slide deck	14
14	Agent profiles	16

PART 4 · KNOWLEDGE/TOOLS AND AUTOMATIONS

15	Knowledge bases	17
16	The Tools library	17
17	Scheduled automations	18

PART 5 · APPEARANCE, YOUR ACCOUNT, AND HELP

18	Light and dark mode	20
19	Your account	20
20	Personalization memory	21
21	Install the app on your phone	22
22	Getting help	22

PART 6 · THE ADMIN CONSOLE

23	Opening the Admin console	23
----	---------------------------------	----

24	Users	23
25	Groups	24
26	Model Access	25
27	Connections: connectors and response actions	26
28	Automations: preparing scheduled runs	27
29	SSO: single sign-on for your tenant	28
30	Analytics	28
31	Policies: service availability and memory governance	29
32	Audit	30
33	Alerts	31

Getting started

1. Welcome

What Aperture Chat is and how to read this guide.

Aperture Chat is your organization's secure workspace assistant. You can chat with approved AI models, draft documents and slide decks, search your organization's knowledge, and schedule recurring work — all in one place, with every answer traceable back to its sources.

This guide assumes nothing. Every instruction spells out where to look and what to click. Words in bold — like **Save version** — are the exact labels you will see on screen. If a step says “hover”, rest your mouse pointer on the item without clicking; a short explanation appears for every clickable control in the platform.

GUIDE	WHO IT IS FOR	WHAT IT COVERS
User Guide	Everyone	Chat , cross-session personalization memory, Drafts , slide decks, Agents/Automations , Knowledge/Tools , search, folders, appearance, and installing the app on your phone.
Administrator Guide	Workspace admins	Everything in the User Guide, plus the Admin console : users, groups, model access, connections, SSO , analytics, policies and memory governance, audit, and alerts.

TIP

The platform is honest by design. If a feature is not configured — a connector without credentials, a model that is not enabled — the screen says so plainly instead of pretending it works. When you see “**Needs credentials**” or “not configured”, that is real, and your administrator can fix it.

2. Signing in

Reach the sign-in screen and get into your workspace.

- 1 Open the web address your organization gave you for Aperture Chat in a modern browser (Chrome, Edge, Safari, or Firefox).
- 2 On the sign-in screen, type your work email address in the field marked `you@company.com`.
- 3 If your organization uses single sign-on (**SSO**), the screen recognizes your email domain and offers **Continue with SSO**. Click it — you are sent to your organization's identity provider (for example Microsoft or Google), sign in there as usual, and return to Aperture Chat already signed in.

- 4 If your account uses a password instead, type it in the **Password** field and click the sign-in button.

GOOD TO KNOW

If both **SSO** and a password are available, either works. If your organization enforces **SSO** for your email domain, password sign-in is blocked on purpose — use the **SSO** button.

GOOD TO KNOW

See a message you do not expect? The sign-in screen reports real problems plainly (for example, a domain that is not allowed for **SSO**). **Copy** the message and send it to your administrator.

3. Finding your way around

The sidebar, the main area, and what each navigation button opens.

The screen has two areas: a sidebar on the left and the main work area on the right. The sidebar is how you move between the platform's views and your saved chats.

SIDEBAR BUTTON	WHAT IT OPENS
New chat	A fresh conversation with your workspace assistant. This is the default view.
Drafts	A document and slide-deck editor with an AI assistant beside it.
Agents/Automations	Reusable agent profiles, plus an Automations tab for scheduled runs.
Knowledge/Tools	Two tabs: Knowledge (searchable document collections) and Tools (connections, prompts, and skills).

Below the navigation buttons, a **Chats** row expands into your organized chat history — **Folders**, **Pinned**, and **Recent** sections, covered in the next section.

Resizing and collapsing the sidebar

- Click the small chevron (◀) near the top of the sidebar to collapse it and give your work more room. Click the chevron again (▶) to bring it back.
- Drag the sidebar's right edge to make it wider or narrower. Double-click that edge to collapse it.
- On a narrow window or a phone, the sidebar hides behind a menu button (three stacked lines, ≡) in the top-left corner. Tap it to open the menu; tap the dimmed area to close it.

The bottom of the sidebar

- **Search** — opens a search box over your past work: chats (including archived ones), agents, drafts, and indexed documents. The keyboard shortcut is Ctrl+K (Windows) or ⌘K (Mac).
- **Help** — opens the guided walkthrough videos and this downloadable guide.

- **Dark mode / Light mode** — switches the appearance instantly (see “Light and dark mode” later in this guide).
- **Install app** — appears on phones and tablets, and adds the workspace to your home screen (see “Install the app on your phone”).
- Your account card — your initials or photo, name, and role. Click it to open the account drawer, including Personalization memory when your organization has enabled it.

Where the consoles live

Administrators open the **Admin console** from the account drawer: click your account card, expand the **Management** section, and choose **Admin console**. Regular users simply do not have a **Management** section. The account drawer also holds **View as** role previews, **Usage this month**, and your **Archived chats** — all covered in “Your account”.

4. Organizing chats: folders, pins, and recent

Keep the chats you care about easy to find.

Click the **Chats** row in the sidebar to expand your organized history. Inside it are three sections — **Folders**, **Pinned**, and **Recent** — and each one expands or collapses with a click.

Folders

- 1 Expand the **Folders** section and click the folder icon with a plus sign next to its heading.
- 2 Type a name for the folder — a client, a matter, a project — and click **Create**.
- 3 To file a chat into the folder, hover the chat's row in the sidebar and click the folder-plus icon that appears, then pick the folder (or create a new one right there).
- 4 Click a folder to expand or collapse the chats inside it.

GOOD TO KNOW

Deleting a folder does not delete its chats — they move back to the **Recent** list.

Pinned chats

- 1 Hover any chat row and click the pin icon that appears.
- 2 The chat moves to the **Pinned** section, where it stays until you unpin it (click the pin icon again).

Recent chats

- The **Recent** section lists your latest conversations. Click one to reopen it exactly where you left off.
- Hover a row to reveal three quick actions: add to folder (folder-plus icon), pin (pin icon), and archive (box icon).

- Archiving moves a chat out of the sidebar without deleting it — use it to tidy up. **Archived chats** are listed in your account drawer, where you can restore or permanently delete them.
- Click **View all chats** at the bottom of the list to browse your full history, including everything that no longer fits in the sidebar.

5. Searching everything

One search box over chats, agents, drafts, and indexed documents.

- 1 Click **Search** near the bottom of the sidebar, or press Ctrl+K (Windows) or ⌘K (Mac) from anywhere.
- 2 The **Search past work** box opens. Type a few words from what you remember — a chat title, a phrase from a message, an agent or draft name, or text from an indexed document.
- 3 Results are grouped by kind. **Chat** results lead, and archived conversations are labeled explicitly so you know where a result will reopen.
- 4 Click a result to jump straight to it.

GOOD TO KNOW

Search covers every chat title and message you can access — archived chats included — plus agents, drafts, and documents indexed in your knowledge bases. Sections with no matches simply do not appear; nothing is padded with placeholders.

6. Your first chat

Send a message, pick a model, and control web search.

- 1 Click **New chat** in the sidebar. A fresh chat greets you by the time of day — “Good morning”, “Good afternoon”, or “Good evening” with your first name (late at night it asks if you are burning the midnight oil) — above the tagline “Your approved models, your sources, your guardrails — ask Aperture Chat anything.” (The name reflects your organization's branding.)
- 2 Click into the message box (the composer) — it reads **Ask anything...** — and type your question.
- 3 Press Enter to send. Press **Shift+Enter** when you want a new line without sending.

Picking the model

The model selector sits in the top bar of every chat. Click it to choose which AI model answers this conversation. Each chat remembers its own choice, and the list only ever shows models your workspace has approved for you.

Web search and the active-tools chip

- When the selected model supports it, public web search starts turned on. A chip at the bottom of the composer shows what is active: with one tool on it names it — for example **Web search** — and with several on it reads **Tools** with a count.
- Click the small x on that chip to turn off every active tool — **Knowledge**, **Web**, and **Agent** — for your next message. Every new chat starts with web search on again when the model supports it.
- Replies that used the web come back with citations you can click, and every source is listed in the session details panel (covered below).

Sending

- The paper-plane button sends your message with the current settings.
- The chevron (v) beside it opens send options — per-reply switches for **Knowledge**, **Web**, and **Agent**, plus the **Reasoning** level (covered in their own section).

Dictate instead of typing

- Click the microphone button in the send controls at the bottom-right of the composer to start dictating. A live waveform shows that the platform is hearing you.
- Click the button again — it becomes a stop square — and your words are transcribed and inserted into the composer as editable text. Nothing sends until you press Enter.

- If the microphone is blocked or no speech was heard, the composer says so plainly so you can fix the problem and try again. The **Drafts** workspace has the same dictation button for drafting instructions.

7. Reading and acting on replies

The work trace, citations, and the actions under every message.

The work trace

- While a reply is being produced, a single status line shows the current step, a live timer, and how many steps remain (for example “4s · step 2 of 5”).
- Click that line at any time — while running or after — to expand the full trace: which model was routed, what context was prepared, whether web search ran, and how the answer was finalized.
- While a reply is still streaming, a stop button appears — its label is **Stop this response**. Clicking it halts the reply and keeps the conversation.

Actions under each response

- **Copy** — copies the response as clean, readable text.
- Thumbs up / thumbs down — sends feedback your administrators can review.
- **Branch** — starts a new chat that continues from that response, leaving the original untouched.
- **Regenerate response** — asks the model to answer the same prompt again. Version arrows appear on the response with a counter (for example “2 / 3”) so you can flip between the versions it has produced.
- **Transfer to Drafts** — turns the response into an editable document in the **Drafts** workspace, where you can format, version, and export it.

Actions under your own messages

- **Edit message** — reopens your prompt for editing. Sending the edit removes the later replies and asks again from that point.
- **Load prompt in new chat** — copies the prompt and its attachments into a fresh chat's composer, ready to adjust and resend.

Diagrams in replies

When a reply contains a diagram (a flowchart, sequence, or similar), it renders as a real figure, not a code block. The figure's toolbar offers **Copy** for the diagram source, **PNG** and **SVG** to download it as an image file, and **Code** to flip between the drawing and its source text.

Generated images

- When the selected model can generate images, just ask for one — the finished images appear directly in the reply.
- Click **Download** under an image to save it as a file.

Citations and token usage

- When a reply used the web or your knowledge bases, numbered citations appear with it. Click a citation to open the original source.
- A small token count appears under a reply only when the model provider actually reported one — the platform never shows an estimate dressed up as a real number.

8. Composer symbol shortcuts

Five characters that insert prompts, agents, knowledge, skills, and automations.

Start a word in the composer with one of five symbols and a menu opens with matching items. Hover the composer to see this cheat sheet at any time.

TYPE	WHAT IT OPENS
/	Your saved prompts and enabled MCP connections.
@	Agent profiles — the reply routes through the one you pick.
#	Knowledge bases and the files inside them, to ground your question.
\$	Saved skill files, inserted into the message.
>	Automations — the one you pick queues and runs when you send.

- 1 Type the symbol, then keep typing to filter the list.
- 2 Use the ↑ and ↓ arrow keys to move through the menu.
- 3 Press Enter to insert the highlighted item, Tab to complete it, or Escape to dismiss the menu.

9. Attaching files, web pages, and connected sources

Upload from your computer, fetch a page by link, or pull from your organization's drives.

- 1 Click the paperclip button at the bottom-left of the composer.
- 2 Click **Upload from computer** to attach files from this device. The model reads them as context for your next message.
- 3 Or click **Web page by link** and paste one or more web addresses — the platform fetches up to 3 public pages and attaches them as cited sources for this message.
- 4 Attached items appear as chips above the composer. Click the x on a chip to remove one before sending.

Attach from source

Below those options, under the **Attach from source** heading, the menu lists your organization's workspace sources: Google Drive, OneDrive, SharePoint, Box, and iManage. Which of these actually work depends on what your administrators have connected.

Connecting your own account

- Cloud sources read from your own account. The first time you pick one, the picker may ask you to connect it — click the Connect button for that source and a sign-in window opens.
- Approve read access in that window. The file list refreshes on its own once access is granted, and only you can see files from your account.
- The picker shows files at the top level of the drive or the folder your administrators configured.

GOOD TO KNOW

If a source is not configured, the platform says so honestly when you try it — nothing silently fails. Ask your administrator to configure the connector (the Administrator Guide covers exactly how).

10. Send options: Knowledge, Web, Agent, and Reasoning

Control what the model may use, one reply at a time.

- 1 Click the chevron (v) beside the paper-plane send button.
- 2 The menu opens with **Send now** at the top — that simply sends with the current settings.
- 3 Below it are three switches you can turn on or off for the next reply, and a **Reasoning** slider at the bottom.

SWITCH	WHAT IT DOES WHEN ON
Knowledge	The reply searches your enabled knowledge bases and cites what it finds, with links back to the sources. A picker lets you choose which bases to use.
Web	The reply uses public web search. Results come back as clickable citations, and the sources are listed in session details. Available when the selected model supports it.
Agent	The reply may use the tools your workspace has enabled. Pair it with an agent profile (type @) to route through a purpose-built configuration.

Reasoning level

The **Reasoning** group at the bottom of the menu is a three-position slider from **Fast** to **Smart**. **Fast** favors quicker answers; **Smart** makes the model think longer for more detailed output; the middle is a balanced default. The slider is only active when the selected model supports reasoning levels — otherwise it is disabled and says so.

Whenever tools are on, the active-tools chip appears next to the paperclip — the single tool's name, or **Tools** with a count. Click the x on that chip to turn **Knowledge**, **Web**, and **Agent** all off for the next message.

11. Session details and sources

Real token usage, active tools, and every source this chat gathered.

- 1 Click the round information button (i) at the right end of the chat's top bar.
 - 2 The session details panel opens for the current chat.
- **Messages and model** — how many messages the chat holds and which model is active.
 - **Token usage** — shown only when the model provider reported real numbers. Otherwise the panel says “Not reported by the provider”.
 - **Tools** — exactly what is switched on for your next message: web search, knowledge, agent mode, plus any automations or MCP connections you added in the composer. “Off” means nothing is active.
 - **Sources gathered** — every web source (as clickable links) and workspace citation collected during this conversation.

Drafts, decks, and agents

12. Drafting documents

Generate, edit, version, and export documents with the drafting assistant.

Click **Drafts** in the sidebar. The view pairs a full document editor with an AI assistant rail on the left that writes into the editor for you. A **Draft format** switch at the top of the editor has two buttons — **Document** and **Deck**. This section covers document mode; “Building a slide deck” covers deck mode.

Generating a draft

- 1 Look at the assistant rail first: it shows which context sources are on — workspace files, web, and templates — before you generate anything.
- 2 In the assistant's message box, describe the document you need: a memo, an engagement letter, a summary. Be specific about audience and tone if it matters.
- 3 The assistant writes the document directly into the editor. You can also send any chat response here with **Transfer to Drafts**.
- 4 Use the model selector in the toolbar area to choose which approved model writes and edits this draft.

Editing like a document

The editor toolbar covers real document formatting: undo and redo arrows, block styles (paragraph, headings, quote), a **Text size** selector, bold and italics, **Superscript** and **Subscript**, colors, a **Highlight** color, bulleted and numbered lists, links, citations, and **Clear formatting** to strip styling from the selected text.

- The plus (+) insert button opens an insert menu: **Web image** finds a real image from the web, **Chart** drops in an editable chart, **Table** inserts a fillable table, **Divider** adds a horizontal rule, and **Page break** starts a new page at your cursor.
- **Inline AI edit** rewrites just the text you have selected: select a passage, click the AI pen, type an instruction, and review the change.

Versions and comparing

- Click **Save version** whenever the draft reaches a good state. You can restore an earlier version at any time from the draft's history.
- **Compare versions** opens a read-only visual redline of two saved versions, so you can see exactly what changed between them. It needs two genuinely different saved versions before it activates.

Exporting

- Click **Export** to open the export panel. In document mode it offers a **Word document** (an editable Word file with preview page breaks and embedded images), **Markdown** (best for plain text or web publishing), and **Print / Save as PDF**, which opens your browser's print dialog with the saved version — choose “Save as PDF” there to keep a PDF copy.
- Exports always run from a saved version. If you have unsaved edits, the panel says **Save your edits first** and offers a one-click **Save version and export** (or **Save version and print**) button.

TIP

On a narrow window, the assistant rail becomes a slide-out drawer so the editor gets the full screen. Open it with the floating pen button on the left edge; press Escape or click the dimmed area to close it.

13. Building a slide deck

Turn a draft into slides, edit them, brand them, and export real PowerPoint.

Deck mode turns the **Drafts** workspace into a slide editor with a real PowerPoint export. Click the **Deck** button on the **Draft format** switch. If the draft already has document content, a dialog asks “Turn this draft into slides?” — choose **Start a blank deck** to keep the document untouched, or **Convert into slides** to split its headings and text into slides. Your document versions are kept either way.

The filmstrip

- Slide thumbnails run along the side. Click one to select it, or drag it to reorder the deck.
- Each thumbnail has its own actions: move-up and move-down arrows, a duplicate button, and a delete button (undo restores a deleted slide).
- **Add slide** at the end of the filmstrip opens the layout menu for a new slide.

Slide layouts

Seven layouts are available today: **Title**, **Title + bullets**, **Two columns**, **Image + caption**, **Quote**, **Section**, and **Closing**. (An eighth, **Chart**, is defined in the deck format but not yet offered in the editor.) The layout button in the toolbar switches the selected slide's layout at any time.

Editing slide text

- Click any text region on a slide to edit it in place. The deck toolbar carries the same rich formatting controls as documents — undo and redo, bold, italics, and more.
- Inside bullet lists, press Tab to indent a bullet one level and Shift+Tab to outdent it.
- **Edit selection with AI** rewrites highlighted slide text: select the text, click the AI pen, type what should change, and press **Apply AI edit**.

Starter templates and your brand template

- The templates drawer (the slides icon in the assistant rail) offers five starter structures: **Pitch deck**, **Quarterly review**, **Project kickoff**, **Client proposal**, and **Training session**. Pick one and start it to get scaffold slides you replace.
- A **Use templates in chat** toggle (off by default) lets the selected template guide the deck assistant when it drafts for you.
- **Upload brand template** accepts your firm's own file — “.pptx or .potx — brand colors, fonts, logo, and every slide's text are extracted.” The extracted theme is stored only on this device.
- **Once** a brand theme is loaded, **Apply to deck** restyles the current deck with its colors, fonts, and logo, and a “Load all N slides” button can replace the deck with every slide from the uploaded template.

Backgrounds and AI images

- The background button in the toolbar opens a per-slide menu: **Upload background...** (a PNG, JPEG, or WebP from this device), **Use on every slide**, **Remove from this slide**, and **Clear from all slides**.
- **Generate AI slide image** writes an image prompt for you, prefilled from the slide's content; press **Generate image** and the result becomes that slide's background. This needs an image-generation model enabled for your workspace — the button says so honestly when none is.
- In the assistant rail, **Toggle AI slide images** makes the assistant also generate an image for every slide whenever it drafts a deck.

Speaker notes

Every slide has a **Speaker notes** field below the stage. **Notes** export with the deck: they land in each slide's notes pane in the PowerPoint file, they appear in the **Markdown outline**, and they follow you into presentation mode.

Presenting from Aperture

- 1 Click the **Present deck** button at the end of the deck toolbar (the monitor icon). The deck opens full screen.
- 2 Advance by clicking anywhere on the slide, or with the arrow keys; move back with the left arrow.
- 3 Keep your speaker notes open in the panel beneath the slide and read from them as you go — the **Notes** button (or the N key) shows and hides the panel.
- 4 Leave the presentation with Escape or the **Exit** button in the top corner.

Saving and exporting

- Decks are stored on this device, not on the server — the save state chip says so, and the **Save version** tooltip spells it out: “Save a version on this device — decks stay local until you export them.”

- Click **Export** to open the **Export deck** panel: **PowerPoint deck** produces an editable .pptx that mirrors the slides on screen — speaker notes included in each slide’s notes pane — and **Markdown outline** exports slide titles, bullets, and speaker notes as text.
- Limits: a deck holds up to 100 slides and 8 MB of content. The editor tells you plainly if a deck exceeds them.

14. Agent profiles

Reusable bundles of model, prompts, knowledge, and tools.

An agent profile bundles a model route, meta prompts, knowledge bases, and MCP tools into one reusable configuration. Instead of switching all of those on by hand, you route a reply through the agent and get its whole setup at once.

Creating an agent

- 1 Click **Agents**/Automations in the sidebar, then **New Agent**.
- 2 Give it a name and a short description of what it is for.
- 3 Work through the editor tabs: **Profile** (name, model, and description), **Knowledge** (assign knowledge bases), **Tools** (select MCP tools), **Prompts & Skills** (attach system prompts and skill files), **Access** (who can use it), and **Hermes** (the optional learning companion).
- 4 Click the save button. The agent is now available to every chat.

Using an agent in chat

- Type @ in the composer and pick the profile, or
- Turn on **Agent** in the send options menu.
- Either way, the reply routes through the agent's full configuration — its model, prompts, knowledge, and tools.

From the **Agents** view you can also open a profile to edit it, jump straight into a chat with it, or delete it. The Automations tab at the top of this view holds your scheduled runs — covered in “Scheduled automations”.

Knowledge/Tools and automations

15. Knowledge bases

Give chats grounded, citable access to your documents, web pages, and APIs.

Knowledge lives in the **Library**: click **Knowledge/Tools** in the sidebar, then the **Knowledge** tab. The **Knowledge Bases** panel lists every collection with its status, security posture, and whether it is enabled.

- 1 Click **Add Knowledge Base** and name the collection.
- 2 Choose what feeds it. A knowledge base can index documents you upload, a public web page (**Add a web link** — enter the address and an optional note), or an API (**Connect an API** — enter the endpoint and its authentication).
- 3 **Create** the base, then check its row in the table: it shows the status, the security posture, and whether the base is enabled.

GOOD TO KNOW

Only enabled knowledge bases can be searched from chat, and access control is enforced per source — people only ever retrieve what they are allowed to see.

Using knowledge in a conversation

- Type # in the composer to reference a knowledge base or a specific file inside one.
- Or turn on the **Knowledge** switch in send options to let the reply search your enabled bases.
- Replies grounded in knowledge return citations that link straight back to the source documents.

16. The Tools library

MCP connections, saved prompts, and skill files inside the Library.

Click **Knowledge/Tools** in the sidebar, then the **Tools** tab. The **Tools Library and Connectors** panel has three sections:

TAB	WHAT LIVES THERE	WHERE IT APPEARS IN CHAT
Connections	MCP tools and connections available to models and agents.	Type / in the composer.
Prompts	Saved prompt templates, including <code>{{variables}}</code> you fill in when using them.	Type / in the composer.
Skills	Skill files — reusable instructions for recurring workflows.	Type \$ in the composer.

- Click **New Prompt** or **New Skill** to create one: give it a name, a category, a description, and the content itself, then save.
- Click any existing card to review or edit it; the trash icon deletes it permanently.
- Every tool row shows its real status — draft, approval required, or enabled — so you always know whether a connection is actually live before relying on it.

17. Scheduled automations

Recurring chat or drafting runs, as one model call or a multi-step chain.

Automations are saved workflows that run a chat or drafting task on a schedule. They are useful for recurring digests, weekly status checks, report refreshes, policy reviews, and any workflow where the same model chain should run the same way each time. Enabled schedules run automatically in the background (times are UTC), and each scheduled run delivers its output as a new chat thread in your sidebar.

Creating an automation

- 1 Click **Agents/Automations** in the sidebar, then the Automations tab, then **New automation**.
- 2 Name it — for example “Monday client digest”.
- 3 Choose what it runs against: chat or draft.
- 4 Pick a trigger: **Weekly** (a day and a time), **Once** (a specific date and time), or **Cron expression** (such as `0 9 * * 1` for 9:00 every Monday, in UTC).
- 5 Write the initial input — what the first step should work on.
- 6 Build the model chain: each step has a model and an instruction, and each step's output feeds into the next step. Use **Add step** for multi-step chains and the x to remove a step.
- 7 Click **Save automation**.

What to check before saving

- **Models** — every step must use a model you can actually access. If a model is missing, your admin may need to approve it for your group.

- Context — write the initial input as if the automation will run without you watching. Include the matter, audience, date range, source expectations, and desired output format.
- Cadence — use **Once** for a single future run, **Weekly** for normal recurring work, and a cron expression only when your administrator gave you an exact one. All times are UTC.
- Ownership — name the automation so another person can tell what it does later.

Running and managing

- Each automation card shows its schedule, its steps, and its last run with an honest result — succeeded or failed, with the reason.
- **Run now** executes the whole chain immediately and shows a transcript of every step's output.
- The toggle pauses or resumes the schedule; the pencil edits it; the trash deletes it.
- You can also queue an automation into any chat by typing > in the composer — it runs when you send.
- The list shows the automations you created. Organization administrators can review the automations in their workspace.

TIP

Always use **Run now** once after creating or editing an automation. If the run fails, the card shows the reason so you can fix the model choice, input, or instructions before depending on the saved schedule.

Appearance, your account, and help

18. Light and dark mode

Switch the whole platform's appearance with one click.

- 1 Find the appearance row near the bottom of the sidebar — it reads **Dark mode** with a moon icon in light mode, and **Light mode** with a sun icon in dark mode.
- 2 Click it. The entire platform switches immediately — no reload, nothing to save.
- 3 Click it again to switch back.

19. Your account

Your profile, personalization memory, password, role previews, usage, archives, and signing out.

Click your account card (your initials or photo) at the very bottom of the sidebar. The account drawer opens with everything about you in one place.

- **Profile** — click the card with your name and the pencil icon to edit your display name, firm, website, bio, phone number, and photo (upload an image up to 5 MB, or paste a URL). Click **Save profile** when done, or **Cancel** to discard.
- **Management** — organization administrators see this section; expanding it opens the **Admin console**. Regular users do not have it.
- **Personalization memory** — opens your private memory manager when the feature is enabled. Use it to control, review, add, correct, pin, forget, or clear what the assistant remembers about you.
- **Password** — accounts that sign in with a password can change it here (click the pencil, enter the current password, then the new one twice — at least 12 characters). Accounts that sign in through **SSO** manage their password with the **SSO** provider instead, and the panel says so.
- **Role and organization** — your role, your organization, any personal token caps that apply to you, and how you sign in.
- **View as** — if your role allows it, preview the workspace exactly as a lower role would see it. A note reminds you which role you are previewing; switch back the same way.
- **Usage this month** — your prompts, responses, and estimated tokens.
- **Archived chats** — every chat you have archived, with buttons to restore each one to your sidebar or delete it permanently.
- **Sign out** — ends your session and returns to the sign-in screen.

20. Personalization memory

Save and recall private preferences in plain English across chats and sessions.

Personalization memory belongs to your account, not to one conversation. When service policy and your organization administrator allow it, administrators and regular users can each build their own private memory. A saved item can influence a new chat immediately and still be available after you sign out and return in a later session.

Save a memory in ordinary English

- 1 In any chat, state what should persist in natural language — for example: “Remember that I prefer a short summary before the detail,” “Please remember that I work in commercial litigation,” or “Remember my project codename is Silver Horizon.”
- 2 Send the message. The assistant confirms the memory after it is saved; you do not need a slash command or special syntax.
- 3 Start a new chat, or return in a later sign-in session, and ask a natural question such as “What do you remember about me?”, “What are my writing preferences?”, or “What is my project codename?” The same account memory is available across those sessions.

GOOD TO KNOW

Memory is different from the context window shown in Session details. The context window describes how much of the current chat the selected model can hold at once; personalization memory is durable account context that follows you into other chats and sessions.

Open and control the memory manager

- 1 Click your account card — your name and role at the bottom-left of the sidebar.
- 2 Click Personalization memory. The dialog titled “What the assistant remembers about you” opens.
- 3 Use memory in my chats controls whether saved items are applied to answers. Learn from my conversations controls whether the assistant may notice durable preferences automatically. Turning a switch off does not expose or silently erase existing items.
- 4 To add something directly, type it under Add something you want remembered, choose Standing instructions, Preferences, About you, Your work, or Other details, and click Add.
- 5 Review the list below. Click a memory's wording to correct it; pin an important item; use Forget to remove one item; or choose Forget everything to clear the list.

GOOD TO KNOW

Only you can read your memory content. Administrators can see counts and purge memories for compliance, but the platform never shows them what an individual memory says.

TIP

If Personalization memory is missing from your account drawer, the feature is not currently available at the platform, organization, or group level. Ask your administrator; refreshing the page cannot override that policy.

CAREFUL

Do not use memory as a password or secret vault. Credentials, keys, and sensitive identifiers are rejected on purpose.

21. Install the app on your phone

Put your workspace on the home screen with its own icon and name.

On a phone or tablet, an **Install app** row appears near the bottom of the sidebar. Tap it and a dialog titled with “Add ... to your home screen” opens, previewing exactly the icon and name that will be installed — your organization's own branding, not a generic one.

- On Android, when the browser supports one-tap install, an **Install app** button opens the phone's native install sheet directly. Otherwise the dialog shows the honest manual route: open your browser's menu and choose “Add to Home screen” or “**Install app**”.
- On iPhone and iPad, Apple only allows installs through the browser's share menu, so the dialog walks you through it: tap the **Share** button, scroll the share sheet, choose “Add to Home Screen”, then tap Add.
- **Once** installed, the app opens full-screen from its own home-screen icon, signed in to the same workspace.

22. Getting help

Walkthrough videos and this guide, always one click away.

- Click **Help** at the bottom of the sidebar to open the guided walkthroughs — short narrated videos of real platform screens, with captions and transcripts, covering everything in this guide.
- Choose Personalization memory in the **Help** playlist for the narrated account-card path, the two memory switches, plain-English save and recall examples, and cross-session behavior.
- The same panel offers this guide as a downloadable PDF, so you can keep it, print it, or share it with a new teammate.
- Administrators have a **Documentation** button inside the **Admin console** with role-specific videos and a printable guide.

The Admin console

23. Opening the Admin console

Where the console lives and what its nine tenant-governance tabs control.

- 1 Click your account card at the bottom of the sidebar to open the account drawer.
- 2 Expand the **Management** section and click **Admin console**. (This section appears only for workspace administrators.)
- 3 The console opens with nine tabs across the top: **Users**, **Groups**, **Model Access**, **Connections**, **SSO**, **Analytics**, Policies, **Audit**, and **Alerts**. Policies is always present between **Analytics** and **Audit**; service-wide availability determines which organization controls are active inside it.

Everything you change here writes through the admin API immediately — changes persist across refreshes and restarts, and every action lands in the tenant audit trail. Status messages under the header tell you honestly whether an action synced or failed.

TIP

The **Documentation** button at the top of the console opens narrated video walkthroughs of every tab, plus this guide as a PDF.

24. Users

Create tenant accounts, manage passwords, and remove leavers safely.

The **Users** tab lists the organization accounts you administer, with each person's role, groups, sign-in method, and status.

Adding a user

- 1 Click **Add User**.
- 2 Enter the person's name and work email, then pick from the roles allowed by current organization policy: User, Power User, Auditor, and **Agent** Approver, plus Admin when administrator creation is available.
- 3 Optionally pick a starting group, then create the account.

Reading and filtering the list

- Use the group filter to work one team at a time.

- Watch the auth and status columns: users who arrived through **SSO** stay **Pending** until you assign them to a group — that is the signal they are waiting on you.

Per-row actions

Every row has an Actions column with three controls:

- **Password** — opens the password dialog for that person: type a password or click **Generate** for a strong random one, optionally mark it a **Temporary password** (they must choose their own at first sign-in), and click **Set password**. The password is shown only here — share it over a safe channel.
- **Deactivate / Activate** — ends or restores sign-in access immediately, keeping the account's audit history intact. You can also select several accounts with their checkboxes and use the **Deactivate** button at the top.
- **Delete** (trash icon) — permanently deletes an eligible account and its chat history. The tooltip spells it out per person: “Permanently delete ... and their chat history”. Administrator-account actions that are unavailable under current service policy remain disabled with an explanation.

GOOD TO KNOW

Administrative continuity rules are enforced by the service. When an account action would violate them, the console blocks the action and explains that it is restricted by administrative continuity policy.

25. Groups

Groups carry permissions and model access — everything flows through them.

Groups are how access flows in the tenant: model grants, knowledge access controls, and permissions all attach to groups, not to individual people. New **SSO** users stay **Pending** until you put them in one.

- 1 Click **Add Group** and name it for the team or matter it represents.
- 2 Select the group to open its editor. It has three tabs: **Users**, **Permissions**, and **Import**.
- 3 On the **Users** tab, tick the people who belong in it.
- 4 On the **Import** tab — headed “**Import Users to**” the group's name — paste addresses into the **User emails** box and click **Add users to group** to bulk-add them.
- 5 On the **Permissions** tab, switch each capability on or off.

The seven permission toggles

TOGGLE	WHAT IT GRANTS
Can use chat	Start conversations and use assigned models.
Can use knowledge	Query approved knowledge bases.
Can use agents	Run approved agent workspaces.
Can use tools	Invoke enabled tools and MCP actions.
Can use API	Create a personal API key for approved models. Service policy must make downstream API access available; otherwise the toggle says so and stays off.
Can use Hermes companion	Build and run agent profiles with the Hermes learning companion. Off until approved.
Can build agents	Create and edit their own private agent profiles from models available to the organization. Off by default, and it only takes effect when service policy permits user-built agents. Publishing agents to the organization stays admin-only.

CAREFUL

Deleting a group removes its members' access that flowed through it. Check what the group grants before deleting.

26. Model Access

Decide which synced models users can see, and which groups can use them.

Model Access starts from the catalog available to your organization. You decide what portion of it your users actually see.

- 1 Click **Sync models** to refresh the catalog whenever service availability changes.
- 2 Read the counters: the All, Enabled, and Disabled status counters carry live totals alongside how many groups are available for scoping.
- 3 Use the search box and the funnel filters in the column headers to cut through the catalog: check off providers, check off model labs, or type text to match runtime routes, and Clear filter resets one. The table has seven columns: **Model**, **Provider**, **User Access**, **Groups**, **Filters**, **Knowledge**, and **Tools**. Everything starts hidden until you decide otherwise.
- 4 Flip the **User Access** toggle on a row to make that model visible, and use **Choose groups** to narrow it to specific teams.

- 5 Use the **Filters** column to attach per-model content filters, and the **Knowledge** and **Tools** columns to scope what each model may search and call.

GOOD TO KNOW

Because model access flows through groups, a tenant with no groups yet sees the guard up front — “**Create** a group before enabling models.” — and the access switches stay off until one exists.

GOOD TO KNOW

Newly synced models arrive disabled until service availability and organization policy permit them. If a model you expect is missing, it is not currently available to this organization.

27. Connections: connectors and response actions

Configure real credentials for each source, and add admin-approved chat output actions.

The **Connections** tab holds the workspace-wide switches for the sources and tools your tenant offers. Every toggle applies to every user in the workspace: turning a connector off removes the capability everywhere users would meet it — chat attach pickers, the / command palette, and the **Knowledge/Tools** library — and the API refuses it too, so off really is off. Credential connectors also carry an honest pill with three states: **Credentials saved** (configured and enabled), **Saved · disabled** (configured but switched off), and **Needs credentials** (not configured yet). Switch-style rows — MCP Servers, Prompt **Library**, **Knowledge** Ingestion, **Document** Templates, and **Audit** and **Analytics Export** — need no credentials; their toggle is the whole story.

Configuring a connector

- 1 Click **Configure** on the connector's card. The form shows that vendor's real fields:

CONNECTOR	WHAT IT NEEDS
OneDrive / SharePoint	Microsoft Graph client credentials: tenant ID, client ID, and client secret from your Azure app registration.
Box	A client-credentials grant with your Box enterprise ID.
Google Drive	A Google OAuth client — save it, then click Connect Google to authorize.
iManage	The server base URL plus a service account username and password.
Web Search	Nothing — it works keyless out of the box. You can choose the engine (DuckDuckGo or a self-hosted SearXNG URL) and the number of results.

- 1 Save the configuration, then click **Test connection**. The test calls the real vendor API and reports the genuine result — including the vendor's own error message if the credentials are wrong.

- 2 Only enable the connector once its test passes.

GOOD TO KNOW

The credentials you save here power knowledge sync and workspace-level access. For chat attachments, users connect their own Google, Microsoft, or Box account from the composer's attach menu and only ever see their own files — the connector card's setup notes list the extra redirect URI and delegated permissions that flow requires.

Chat output actions

The second panel on this tab, **Chat output actions**, adds admin-approved buttons to assistant responses — export, format, or handoff actions. It does not add MCP servers or model-callable tools; those are managed from the **Tools** tab under **Knowledge/Tools** in the sidebar (**Connections**), where each MCP connection and its transport is configured.

28. Automations: preparing scheduled runs

What admins need to configure before users can rely on recurring model chains.

Users create and run automations from the Automations tab of the **Agents** view, and enabled schedules fire automatically in the background (UTC), delivering each run's output as a new chat thread. Admins control whether those runs have the models, groups, and connectors they need. Treat every automation failure as a setup signal first: the run is usually missing model access, source access, or a connector credential.

Admin readiness checklist

- 1 Open **Model Access** and confirm the automation's model is visible to the user's group.
- 2 Open **Groups** and confirm the user belongs to the group that receives the model grant.
- 3 Open **Connections** and configure any connector the automation expects to use, then run **Test connection** before enabling it.
- 4 If the automation uses knowledge, confirm the knowledge base is enabled and its source ACL allows the user's group.
- 5 Ask the user to press **Run now** after saving the automation; the transcript and last-run status will show whether the setup is complete.

GOOD TO KNOW

Admins do not need to expose every model to make automations work. Grant the smallest approved model set that fits the scheduled workflow, then expand only when a run proves it needs more capability.

29. SSO: single sign-on for your tenant

Connect an identity provider, provision on first sign-in, and map IdP groups.

GOOD TO KNOW

If organization policy does not permit **SSO** management in this console, the tab is read-only and says so plainly.

- 1 Open the **SSO** tab and click **Add SSO configuration**.
- 2 Pick a preset — Microsoft Entra ID, Google Workspace, Okta — or choose custom OIDC and enter any issuer. Presets prefill the issuer URL; for Entra, replace the `{tenant-id}` placeholder with your directory ID.
- 3 Paste the **Client ID** and **Client secret** from your identity provider's app registration. The secret is vaulted server-side and never shown again.
- 4 List the email domains allowed to sign in through this provider.
- 5 Decide on just-in-time provisioning: when on, the first sign-in from an allowed domain creates the account automatically with the User role. Remember — they stay **Pending** until you assign a group.
- 6 **Copy** the redirect URI shown at the bottom of the form and register it with your identity provider. Sign-ins cannot complete until the provider trusts that exact callback address.
- 7 Click create. Enforcement always starts off, so nothing can lock the tenant out.
- 8 On the new configuration's card, click **Test connection** — it performs a real discovery and key check against the provider.
- 9 Map identity-provider group values to tenant groups on the card, so JIT users land with the right access.
- 10 Only enforce tenant sign-in after the test passes.

CAREFUL

Enforcement blocks password sign-in for the allowed domains. Treat it as the very last step: if the provider is misconfigured, enforcement is what locks people out.

30. Analytics

Runtime metadata, feedback, per-user usage, and token budgets for this tenant.

The **Analytics** tab measures how the organization actually uses the platform. Sections start collapsed behind descriptive headers — click a header to expand one — and each section carries its own filter row pairing a person picker (defaulting to **All admins and users**) with a date range and its presets (All, Today, Week, 30 days). Scoping one section never narrows another.

- **Runtime Clock Metadata** — organization-scoped execution timestamps captured from admin and user **Chat** and **Draft** completion events. The cards total runtime events, chat completions, and draft calls. Its **CSV** includes actor_id and actor_name columns, and when one user is selected the filename gains that user's suffix so exports stay unambiguous.
- **Chat Feedback Analytics** — response-level thumbs up and thumbs down signals submitted by tenant admins and users, with totals for overall, positive, and negative feedback.
- **Model Activity** — saved prompt volume by model, date, and person for the selected range, drawn as “**Prompts** by model”, “Prompt trend”, and “**Users** by prompt activity”.
- **User Usage** — durable per-user usage for this organization's admins and users across chat, drafts, agents, automations, and the API gateway. Token counts are provider-reported only and stay blank when a provider reported none. The panel has its own **Filter usage by user** selector for drilling into one person.
- **Workspace Usage Budget** — a read-only view of the organization's service-managed ceiling and its current UTC accounting period. This panel shows where usage stands without exposing service-level configuration.
- **Token Allocations** — per-user and per-group token caps beneath the workspace ceiling, with Per day, Per week, or Per month UTC resets. The most restrictive applicable cap wins. Choose a user or group, enter a cap, pick the period, and click **Set cap**. Until you add one, the table says “No allocations yet. Everyone shares the workspace ceiling.”

GOOD TO KNOW

Every number here comes from saved audit and usage events, and each **CSV** export button opens its own small date-range popover so a file contains exactly the rows you chose. If a panel has nothing to show, it says so plainly instead of showing sample data.

31. Policies: service availability and memory governance

Apply downstream tenant defaults, configure memory, and govern by count without reading content.

The Policies tab is always present. **Policy Controls**, Personalization Memory, and Memory by User all start collapsed, so the page stays easy to scan. Expand only the panel you need; organization administrators can narrow available capabilities, while anything unavailable under service policy remains locked.

Policy Controls: read service availability

- Administrator accounts — shows whether this console can create and manage administrators or whether new administrator accounts require service approval.
- Admin sign-in policy — shows whether admins must use **SSO** or may use an explicitly provisioned local account.
- **SSO** configuration — shows whether tenant **SSO** mappings are delegated or read-only.

- New model defaults — shows whether newly available models begin with **Default Users** or require explicit grants.

Below those status rows are three downstream switches for the protected **Default Users** group: personal API keys, private-agent building, and personalization memory. A switch locks when its capability is unavailable under service policy; the saved group grant is preserved rather than silently erased. Use **Groups** for exceptions, **Model Access** for available models, and **Connections** for available connectors.

Personalization Memory

- 1 Expand Personalization Memory and turn on Memory for this organization. This makes the account-level Personalization memory row available to eligible users; existing memories stay saved but are not applied while this switch is off.
- 2 Choose whether to allow Learn from conversations automatically. When off, only explicit requests such as “remember that ...” and direct additions in the memory manager create memories. Every user can still opt out of automatic learning individually.
- 3 Set Retention (days), from 1 through 3650. Older memories retire automatically when they pass this policy.
- 4 Set Maximum memories per user, from 1 through 2000. The default is 200, which is a practical general-purpose limit; raise or lower it only when your retention, compliance, or workload policy calls for a different capacity. When the cap is passed, the least useful unpinned memories retire first.
- 5 Expand Memory by User to review content-free counts. Click **Refresh** to reload the counts, or purge one person's memories when compliance or account cleanup requires it.

CAREFUL

Memory administration never grants reading access. Administrators see policy, counts, and purge controls only; the API and the interface do not return another person's memory content.

32. Audit

Governance signals, prompt monitoring, security alerts, and the tenant trail.

The **Audit** tab is the tenant's governance station. It opens on the **Admin Audit** posture dashboard while the sections below start collapsed — click a section header to expand it. **User Prompt Activity**, **Security Alerts**, and the **Audit Trail** each carry their own filter row pairing a person picker with a date range, so one section's scope never narrows another. The **CSV** export buttons produce files of exactly the filtered rows — including `actor_id` and `actor_name` columns, so every exported row names who acted.

- **Admin Audit** — summary cards for tenant security and governance signals: audit events, the prompt watchlist of active DLP or misuse alerts, prompt volume, active admins, active users, connector issues, and ungrouped models. Cards that need attention are highlighted; hover one to see what it counts.

- **Recent Governance Activity** — the current tenant snapshot for identity, user, model, and connector posture: how many **SSO** configurations are enforced, how many users are active, and how many connectors are enabled.
- **User Prompt Activity** — drill into saved prompts from this organization's admins and users and their model responses by person, thread, model, and timestamp, scoped by the section's own user and date filter.
- **Security Alerts** — DLP and malicious-behavior flags raised from admin and user prompts, shown with redacted snippets for review. Click **Acknowledge** once an alert is handled, or **Reopen** if it needs another look.
- **Audit Trail** — the tenant's append-only transaction log, newest first. Its toolbar has a severity select, an action-category select, and a search box (**Search actions, people, targets...**) that narrow the rows together; the trail's **CSV** exports exactly the rows you are looking at. Click **Refresh** to reload it straight from the admin audit API.

GOOD TO KNOW

Security alerts fire on real prompt content — payment card numbers, shared credentials, prompt-injection attempts — but show only redacted snippets, so reviewing an alert never re-exposes the sensitive value itself.

33. Alerts

Watch rules over audit activity, and honest delivery statuses for every alert.

The **Alerts** tab turns audit activity into notifications. It has three panels: **Email Delivery**, **Alert Rules**, and **Alert Deliveries**. Every alert is always logged in-app regardless of email — email is a delivery channel, not the record.

Email Delivery

For organization administrators this panel is a read-only status: it reports whether email delivery is configured at the service level. If email is not configured, rules still work — their alerts are logged in-app and the delivery log says so honestly.

Alert Rules

- Click **Suspicious-activity template** to start from a sensible security rule, or **New rule** to build one from scratch.
- A rule has: a name; **Action patterns** — exact audit actions or prefixes such as security.* or admin.user_deleted; a **Minimum severity**; an optional **Watched user**; a threshold — how many matches within a time window before it fires; a cooldown between alerts; and email recipients. Leave recipients empty and the alert is in-app only.
- Rules you create here watch this organization's admin and user audit activity.

Alert Deliveries

Every alert trigger is listed with its real delivery status: sent, queued, failed with the actual SMTP error, email not configured, or logged in-app. The log exports to **CSV**, archived deliveries included. Click **Archive** on a delivery to clear it from the default view — its history is kept, and **Show archived** reveals archived deliveries so you can review or **Restore** them. Tenant admins can archive their tenant's deliveries only.

Tenant SSO, and troubleshooting

Administrator Guide

This appendix adds step-by-step identity provider configuration for tenant SSO, the connector permissions each external system requires, and a reference of the exact messages the platform emits.

A.1 How Aperture Chat authenticates

Aperture Chat is an OpenID Connect relying party. Knowing exactly what it sends, and what it requires back, turns every provider configuration in this appendix into a mechanical task.

ELEMENT	WHAT APERTURE CHAT DOES
Flow	Authorization Code with PKCE (<code>code_challenge_method=S256</code>) and a validated <code>nonce</code>
Discovery	Fetches <code>{issuer}/.well-known/openid-configuration</code> and requires <code>issuer</code> , <code>authorization_endpoint</code> , <code>token_endpoint</code> , and <code>jwks_uri</code>
Scopes	<code>openid email profile</code> unless the configuration overrides them
Redirect URI	<code>{API base URL}/api/auth/sso/callback</code> — the console prints the exact value; register it verbatim
Identity claim	First present of <code>email</code> , <code>preferred_username</code> , <code>upn</code>
Group claim	<code>groups</code> by default; configurable per SSO configuration
Token validation	Signature against the published JWKS, plus issuer, audience, expiry, and nonce

The issuer must be publicly reachable

Discovery runs through an egress guard that rejects private, loopback, and link-local addresses. An issuer on an internal-only hostname fails with “OIDC issuer URL is not permitted”. Publish the discovery document, or terminate it on a reachable host.

A.2 Microsoft Entra ID

Before you begin

You need the Application Administrator or Cloud Application Administrator role in Entra ID, and platform owner or tenant administrator access in Aperture Chat. Have the redirect URI from the Aperture Chat SSO form to hand.

In the Microsoft Entra admin center

1. Go to **Entra ID** → **App registrations** → **New registration** and name it for the deployment.
2. Under **Supported account types**, choose single tenant unless guests from other tenants must sign in.
3. Set **Redirect URI** to platform **Web** and paste the exact value Aperture Chat shows, ending in `/api/auth/sso/callback`. A trailing-slash or scheme mismatch is rejected by Microsoft, not by Aperture Chat.
4. Register, then copy the **Application (client) ID** and **Directory (tenant) ID** from Overview.
5. Open **Certificates & secrets** → **New client secret**, choose an expiry your rotation process can meet, and copy the secret *Value* immediately — it is never shown again.
6. Open **API permissions** and confirm the delegated Microsoft Graph permissions `openid`, `email`, and `profile`. Grant admin consent if your tenant requires it.
7. Optional, for group mapping: **Token configuration** → **Add groups claim**. Choose Security groups, and under ID token select **Group ID**.

In Aperture Chat

1. Open the SSO form and choose **Microsoft Entra ID**.
2. Set **Issuer URL** to `https://login.microsoftonline.com/{tenant-id}/v2.0`, replacing the placeholder with your Directory (tenant) ID.
3. Paste **Client ID** and **Client secret**. The secret is vaulted server-side and never returned to a browser.
4. List the **Allowed email domains** this configuration may authenticate.
5. Leave **Provision new users on first sign-in (JIT)** off until you have mapped groups, unless you intend every verified directory user to receive an account.
6. Save, then run the discovery and signing-key test before enforcing anything.

Group IDs, not names

Entra emits group object IDs in the `groups` claim, not display names. Map those GUIDs to tenant groups. Directories that exceed Microsoft's claim limit send a `_claim_names` overage reference instead of the list — reduce the group scope, or assign the application to specific groups.

A.3 Google Workspace, Okta, and other issuers

Google Workspace

1. In the Google Cloud console, under APIs & Services → Credentials, create an **OAuth client ID** of type **Web application**.
2. Add the Aperture Chat callback under **Authorized redirect URIs**, ending in `/api/auth/sso/callback`.
3. Configure the OAuth consent screen as **Internal** so only your Workspace users can complete sign-in.
4. In Aperture Chat set the issuer to `https://accounts.google.com` and list your Workspace domains.

Google does not emit Workspace groups in the ID token by default. Either assign tenant groups in Aperture Chat after provisioning, or drive membership through SCIM.

Okta

1. Create an app integration of type **OIDC – OpenID Connect, Web Application**.
2. Set the sign-in redirect URI to the Aperture Chat callback.
3. Under Assignments, limit the application to the groups that should reach this deployment.
4. Set the issuer to your Okta domain, or to an authorization server such as `https://your-org.okta.com/oauth2/default`. Add a `groups` claim to the ID token to map groups.

Any other OIDC provider

No provider is special-cased. Any issuer works when all four of these hold:

1. It publishes `{issuer}/.well-known/openid-configuration` with the four required fields.
2. It is reachable from the API container over the public internet.
3. It supports the Authorization Code flow with PKCE and a confidential client secret.
4. Its ID token carries `email`, `preferred_username`, or `upn`.

Confirm the issuer string in the discovery document matches what you enter exactly. A mismatch fails ID-token validation even when sign-in appears to succeed at the provider.

Appendix B • Troubleshooting and reference

B.1 Sign-in and SSO messages

These are the exact messages the platform emits. Match on the text rather than on symptom.

MESSAGE	CAUSE AND RESOLUTION
This SSO provider is not fully configured...	Issuer URL, client ID, or client secret is missing. Re-save all three; the secret must be re-entered after a change.
Could not fetch OIDC discovery document from...	The issuer is unreachable, is not publishing discovery, or returned non-JSON. Open the URL from outside your network and confirm it resolves.
OIDC issuer URL is not permitted...	The issuer resolves to a private, loopback, or link-local address and the egress guard rejected it. Use a publicly resolvable issuer.
The identity provider did not return an email claim...	The <code>email</code> scope was not granted, or the directory account has no email. Grant the scope and confirm the account has a mail attribute.
The identity provider reports {email} as unverified...	The provider marked the address unverified. Verify it at the provider; an unverified identity is never accepted.
ID token nonce does not match the login request.	The sign-in was resumed from a stale tab or replayed. Start again from the sign-in screen.
The sign-in link is invalid or has expired...	The authorization round trip took too long, or the state was reused. Start the sign-in again.
The SSO client secret is missing. An administrator must re-save it.	The vaulted secret was cleared or rotated out. Re-enter it in the SSO form.
SSO is enforced for this email domain; local sign-in is disabled.	Working as configured. Use the identity provider, or a break-glass account on an uncovered domain.
Organization policy requires admin accounts to sign in through SSO.	The platform owner enabled the admin-SSO policy. Administrators must authenticate through the provider.
JIT provisioning is disabled for this SSO provider.	The account does not exist and JIT is off. Create the account, or enable JIT for this configuration.
SSO tenant does not match this account.	The email already exists in a different tenant. Resolve the duplicate before signing in.
Too many failed sign-in attempts for this account...	Lockout protection. Wait for the window to pass; it clears without administrator action.

B.2 Connector permissions

Microsoft connectors (OneDrive, SharePoint, Outlook)

- For knowledge sync, grant `Files.Read.All` — and `Sites.Read.All` for SharePoint — as *application* permissions, then choose **Grant admin consent** in Entra.
- For chat attachments, where each user reaches only their own files, add the same scopes as *delegated* permissions and register `{API base URL}/api/connector-oauth/callback` as a Web redirect URI.

Google Drive uses an OAuth client plus per-user authorization; each user sees only what their own account can open. **Box** uses enterprise and client-credential configuration with a delegated path for personal files. **iManage** needs the server URL and a supported service or delegated identity. **Web search** runs keyless against DuckDuckGo, or against an operator-hosted SearXNG instance.

Test before enabling

Every connector has a real connection test. A saved-but-untested connector that fails at query time surfaces as a failed automation run, not as a silent skip.

B.3 Email delivery states

Alert rules always log in-app. Email is an additional channel, and its status is reported literally.

STATE	MEANING
Sent	The SMTP server accepted the message. Delivery to the mailbox is then the mail system's responsibility.
Queued	Accepted for later delivery and not yet completed.
Failed	Shown with the actual SMTP error. Authentication failures, TLS mode mismatches, and rejected From addresses appear verbatim.
Email not configured	No SMTP settings exist. A platform owner configures them in Owner → Alerts.
Logged in-app	The trigger fired and was recorded; no email channel applied.

B.4 External reference index

Authoritative documentation for every external system this platform touches.

Identity and provisioning

- OpenID Connect Core 1.0
https://openid.net/specs/openid-connect-core-1_0.html
- OpenID Connect Discovery 1.0
https://openid.net/specs/openid-connect-discovery-1_0.html
- PKCE (RFC 7636)
<https://datatracker.ietf.org/doc/html/rfc7636>
- SCIM Protocol (RFC 7644)
<https://datatracker.ietf.org/doc/html/rfc7644>
- Entra ID: register an application
<https://learn.microsoft.com/entra/identity-platform/quickstart-register-app>
- Entra ID: configure optional claims
<https://learn.microsoft.com/entra/identity-platform/optional-claims>
- Entra ID: ID token claims reference
<https://learn.microsoft.com/entra/identity-platform/id-token-claims-reference>
- Entra ID: provision with SCIM
<https://learn.microsoft.com/entra/identity/app-provisioning/use-scim-to-provision-users-and-groups>
- Google: OpenID Connect
<https://developers.google.com/identity/openid-connect/openid-connect>
- Okta: OIDC web application
<https://developer.okta.com/docs/guides/sign-into-web-app-redirect/>

Connectors, tools, and operations

- Microsoft Graph permissions reference
<https://learn.microsoft.com/graph/permissions-reference>
- Model Context Protocol
<https://modelcontextprotocol.io/>
- SearXNG documentation
<https://docs.searxng.org/>
- Docker Compose
<https://docs.docker.com/compose/>
- Caddy documentation
<https://caddyserver.com/docs/>

- Aperture Chat source

<https://github.com/Aperture-Chat/Aperture-Chat>