



For platform owners

# Platform Owner Guide

## Aperture Chat

The complete User and Administrator Guides, plus first-run guidance, providers and shared connectors, the API key vault, organization model availability, SSO and MFA policy, branding, search indexing, budgets, Elastic export, releases, analytics, audit, and alerts. No prior knowledge assumed.

# Contents

## PART 1 · GETTING STARTED

|   |                                                    |    |
|---|----------------------------------------------------|----|
| 1 | Welcome                                            | 4  |
| 2 | Signing in                                         | 5  |
| 3 | Request access and finish your first sign-in       | 8  |
| 4 | Two-step verification and recovery codes           | 9  |
| 5 | Finding your way around                            | 14 |
| 6 | Organizing chats: folders, pins, and the chat menu | 16 |
| 7 | Searching everything                               | 19 |

## PART 2 · CHAT

|    |                                                    |    |
|----|----------------------------------------------------|----|
| 8  | Your first chat                                    | 22 |
| 9  | Understand model availability and request access   | 27 |
| 10 | Reading and acting on replies                      | 30 |
| 11 | Symbol shortcuts: / @ # \$ >                       | 32 |
| 12 | Attaching files, web pages, and connected sources  | 35 |
| 13 | Send options: Knowledge, Web, Agent, and Reasoning | 38 |
| 14 | Session details and sources                        | 41 |

## PART 3 · DRAFTS, DECKS, AND AGENTS

|    |                                       |    |
|----|---------------------------------------|----|
| 15 | Drafting documents                    | 44 |
| 16 | Building a slide deck                 | 50 |
| 17 | Recover work kept only on this device | 55 |
| 18 | Agent profiles                        | 57 |

## PART 4 · LIBRARY AND AUTOMATIONS

|    |                       |    |
|----|-----------------------|----|
| 19 | Knowledge bases       | 60 |
| 20 | The Tools library     | 62 |
| 21 | Scheduled automations | 64 |

## PART 5 · APPEARANCE, YOUR ACCOUNT, AND HELP

|    |                               |    |
|----|-------------------------------|----|
| 22 | Light and dark mode           | 69 |
| 23 | Your account                  | 70 |
| 24 | Personalization memory        | 73 |
| 25 | Install the app on your phone | 76 |
| 26 | Getting help                  | 77 |

## PART 6 · THE ADMIN CONSOLE

|    |                                                      |     |
|----|------------------------------------------------------|-----|
| 27 | Opening the Admin console                            | 78  |
| 28 | Users                                                | 79  |
| 29 | Groups                                               | 84  |
| 30 | Model Access                                         | 87  |
| 31 | Resolve model requests and trace access              | 89  |
| 32 | Connections: response actions and connector handoff  | 91  |
| 33 | Automations: preparing scheduled runs                | 93  |
| 34 | SSO: single sign-on for your tenant                  | 94  |
| 35 | Analytics                                            | 98  |
| 36 | Policies: service availability and memory governance | 100 |
| 37 | Data retention and conversation tags                 | 103 |
| 38 | Review feedback and reported issues                  | 106 |
| 39 | Audit                                                | 107 |
| 40 | Alerts                                               | 110 |

## PART 7 · THE PLATFORM OWNER CONSOLE

|    |                                                                |     |
|----|----------------------------------------------------------------|-----|
| 41 | Opening the Platform owner console                             | 114 |
| 42 | First-run setup: from owner account to a working team          | 115 |
| 43 | Providers                                                      | 118 |
| 44 | API Key Vault                                                  | 122 |
| 45 | Models: organization availability                              | 125 |
| 46 | Org Settings: users and role boundaries                        | 127 |
| 47 | Single sign-on: connect any OpenID Connect provider            | 129 |
| 48 | Single sign-on: Microsoft Entra ID, Okta, and Google Workspace | 134 |
| 49 | Single sign-on: groups, MFA, enforcement, and recovery         | 140 |
| 50 | Org Settings: platform branding                                | 142 |
| 51 | Org Settings: policies and budgets                             | 144 |
| 52 | Org Settings: shared connectors and source credentials         | 152 |
| 53 | Org Settings: search index readiness                           | 154 |
| 54 | Review and install a platform release                          | 155 |
| 55 | Organization retention and tagging                             | 156 |
| 56 | Analytics: runtime, feedback, activity, and usage              | 158 |
| 57 | Automation readiness and governance                            | 160 |
| 58 | Audit: owner governance signals                                | 161 |
| 59 | Alerts: email delivery and platform-wide rules                 | 163 |

# Getting started

## 1. Welcome

What Aperture Chat is and how to read this guide.

Aperture Chat is your organization's workspace assistant. You can chat with approved AI models, draft documents and slide decks, search your organization's knowledge, and schedule recurring work. Review the sources and work trace when a response uses connected knowledge or tools.

This guide explains where to look and what to choose. Words in bold — like **Save version** — match the labels on screen. If a step says “hover”, rest your pointer on the item without clicking; controls that offer extra help show it nearby.

### How each task is taught

Each task appears in three matching forms: a narrated walkthrough video in the product, a written procedure in this guide, and the same procedure on the website guide. They come from one source, so they describe the same steps. Every walkthrough is a complete loop, from where you start to a result you can see, and it covers each way the product offers to do the task.

| IN EACH WALKTHROUGH BOX     | WHAT IT GIVES YOU                                                                                                                                                                             |
|-----------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Before you begin</b>     | The role, access, earlier setup, and outside accounts or values you need first.                                                                                                               |
| <b>Step by step</b>         | Every action in order, with the exact labels on screen. Where there are several ways to do the task (an identity provider, a delivery target, a device), each path has its own complete list. |
| <b>Check that it worked</b> | What you should see when the task succeeded.                                                                                                                                                  |
| <b>Troubleshooting</b>      | Messages you may meet, quoted as the product shows them, and what to do about each.                                                                                                           |

Walkthrough screens are real captures of a running workspace with sample accounts. Steps that happen in another product, such as an identity provider's console, are shown as labeled instruction cards rather than imitation screenshots.

| GUIDE                       | WHO IT IS FOR    | WHAT IT COVERS                                                                                                                                                                          |
|-----------------------------|------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>User Guide</b>           | Everyone         | <b>Chat</b> , model access requests, personalization memory, documents and decks, account sync, agents, knowledge and tools, search and commands, appearance, and account security.     |
| <b>Administrator Guide</b>  | Workspace admins | Everything in the User Guide, plus the <b>Admin console</b> : users, groups, model access, response actions, <b>SSO</b> , analytics, policies and memory governance, audit, and alerts. |
| <b>Platform Owner Guide</b> | Platform owners  | Everything in the other guides, plus setup readiness, providers and keys, organization policy, shared connectors, branding, releases, search indexing, and platform governance.         |

#### TIP

If a feature is not configured — a connector without credentials, a model that is not enabled — the screen reports that state. Ask your administrator about account or model access; shared connector settings and credentials are managed by the service team outside tenant administration.

## 2. Signing in

Reach the sign-in screen and get into your workspace.

- Open the web address your organization gave you for Aperture Chat in a modern browser (Chrome, Edge, Safari, or Firefox).
- On the sign-in screen, type your work email address in the field marked `you@company.com`.
- If your organization uses single sign-on (**SSO**), the screen recognizes your email domain and offers **Continue with SSO**. Click it — you are sent to your organization's identity provider (for example Microsoft or Google), sign in there as usual, and return to Aperture Chat already signed in.
- When both methods are offered, choose Organization **SSO** or Email & password. The password field appears in Email & password mode; enter your password and click Sign in. Pressing Enter submits the method you selected.

#### GOOD TO KNOW

If both **SSO** and a password are available, either works. If your organization enforces **SSO** for your email domain, password sign-in is blocked on purpose — use the **SSO** button.

#### GOOD TO KNOW

See a message you do not expect? The sign-in screen reports real problems plainly (for example, a domain that is not allowed for **SSO**). **Copy** the message and send it to your administrator.

## Request access and enter your workspace

Ask for an account, then sign in with a temporary password, your own password, Organization SSO, or an authenticator code.

### Before you begin

- Your workspace address, which your organization gives you (for example <https://your-instance.example>).
- Your work email address.
- For Organization **SSO**: an account at your organization's identity provider, such as Microsoft Entra ID, Okta, or Google.
- For a password account: the temporary password your administrator gives you after approving your request.
- If two-step verification is on for your account: your authenticator app, or one unused recovery code.

### Step by step

- 1 Open your workspace address in a web browser. The sign-in page opens.
- 2 Under New to Aperture Chat?, choose **Request access**.
- 3 Enter First name, Last name, and Work email, then choose Submit access request.
- 4 Read Request received. The form does not send an email or set a password. Contact your administrator, who approves the request and tells you how to sign in.
- 5 When your administrator gives you a temporary password, return to the sign-in page and choose Email & password (it appears beside Organization **SSO** when your workspace offers both).
- 6 Enter your work email and the temporary password, then choose Sign in.
- 7 On Set a new password, type a password of at least 12 characters into New password and Confirm password, then choose **Set password** and continue.
- 8 Check that the welcome card, A good place to begin., says You're signed in.

### Path: Organization SSO

- 1 Open your workspace address. Organization **SSO** is selected first when your workspace offers it.
- 2 Optionally type your work email. It selects your organization's provider, shown as Sign in with followed by the provider's name.
- 3 Choose **Continue with SSO**.
- 4 Sign in on your identity provider's page, including any multifactor check it asks for. Aperture Chat never sees that password.

- 5 You return to the workspace signed in. On your first visit your account is created automatically if your organization allows it; otherwise ask your administrator to add you.

### Path: Email & password

- 1 Choose Email & password.
- 2 Enter your work email and your password, then choose Sign in.
- 3 If two-step verification is on, follow the **Authenticator code** path.
- 4 The workspace opens where you left off.

### Path: Authenticator code

- 1 After Sign in, Two-step verification asks for a code. It shows the attempts remaining and when the challenge expires (about five minutes).
- 2 Open your authenticator app and type the current 6-digit code into **Authenticator code**.
- 3 Choose Verify and continue.
- 4 No phone? Choose **Use a recovery code instead**, type one unused recovery code, and choose Verify and continue. Each code works once.
- 5 If the challenge expires, choose Back to sign in and start again.

### Path: Trouble signing in

- 1 Choose Trouble signing in? under the sign-in button.
- 2 For Organization **SSO**, use your organization's own password recovery.
- 3 For an email-and-password account, ask your workspace administrator for a temporary password.
- 4 If you lost your authenticator and have no recovery codes, ask your administrator to reset two-step verification for your account.

### Check that it worked

- ✓ The welcome card A good place to begin. shows You're signed in.
- ✓ Your name appears on the account card at the bottom of the sidebar.
- ✓ With Organization **SSO**, the identity provider's page handled your password and you returned signed in.

## Troubleshooting

| WHAT YOU SEE                                                                                              | WHAT TO DO                                                                                                                                                                                                    |
|-----------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "Invalid local credentials."                                                                              | The password is wrong. Retype it (Show password helps), or ask your administrator for a new temporary password.                                                                                               |
| "Unknown local account."                                                                                  | There is no active password account for that email yet: the request may still be pending, or your account signs in with Organization <b>SSO</b> . Choose Organization <b>SSO</b> , or ask your administrator. |
| "SSO is enforced for this email domain; local sign-in is disabled. Use the configured identity provider." | Your organization requires single sign-on for your domain. Choose Organization <b>SSO</b> and <b>Continue with SSO</b> .                                                                                      |
| "Use a password with at least 12 characters." or "The passwords do not match."                            | On Set a new password, use 12 or more characters and type the same password in both fields.                                                                                                                   |
| "Choose a new password."                                                                                  | The new password matches the temporary one. Pick a different password.                                                                                                                                        |
| "The MFA code is invalid." or "Too many MFA attempts."                                                    | Wait for the next code in your authenticator app and check that your phone's clock is set automatically. After too many attempts, wait for the timer, or choose Back to sign in.                              |
| "Verification unavailable"                                                                                | The sign-in challenge expired or was used up. Choose Back to sign in and sign in again.                                                                                                                       |
| "Too many failed sign-in attempts for this account. Try again shortly."                                   | Wait a minute before trying again, or use Trouble signing in?.                                                                                                                                                |
| "Too many access requests for this email. Try again shortly."                                             | Your request was already recorded. Contact your administrator instead of submitting again.                                                                                                                    |

### 3. Request access and finish your first sign-in

Submit a request, understand the administrator handoff, and enter the workspace.

- 1 On the sign-in screen, choose **Request access**. Enter your first name, last name, and work email, then submit the form.
- 2 The Request received screen confirms submission for your email. If this email is new to the workspace, an administrator reviews the request. This form does not email you or create a password.



- 3 Contact your administrator for an update. After approval, they provide the workspace address and confirm whether to use organization **SSO** or a temporary password.
- 4 Choose **Back** to sign in. Your submitted email remains filled in. Use the sign-in method your administrator arranged.
- 5 If you received a temporary password, sign in with it and complete **Set a new password**. Enter a new password of at least 12 characters twice. After saving, the app continues with a new authenticated session and revokes your older sessions.
- 6 Complete any required authenticator verification. **Once** in the workspace, select an available model and send a short first message. If no model is available, ask your administrator to check your account, group, and model access.

#### GOOD TO KNOW

Submitting another access request does not approve an account or recover an existing password. Trouble signing in explains the recovery path: contact your administrator for local password access, or use your organization identity provider for **SSO**. If sign-in settings cannot load, choose **Retry connection**.

## 4. Two-step verification and recovery codes

Set up an authenticator, complete required verification, and manage recovery codes.

To add an authenticator after signing in with a local password, open your account drawer, find **Security**, choose **Manage security**, and select **Set up authenticator**. Confirm your current password, add the QR code or setup key to your authenticator, acknowledge that you added it, and choose **Verify authenticator** with a current six-digit code. **Copy** the one-time recovery codes into secure storage, acknowledge that you saved them, and choose **Done**. Copying alone does not store them for you; finish verification and code storage before leaving the panel.

- If sign-in asks you to set up an authenticator, choose **Begin authenticator setup**. Add the displayed QR code or setup secret to your authenticator app.
- Confirm that you added the account, enter the current six-digit **Authenticator code**, and choose **Verify and enable MFA** before the setup expires.
- Use **Copy recovery codes** and save the codes privately when they are shown. They are not saved to a file automatically. Each recovery code can be used once when your authenticator is unavailable; do not include these codes in screenshots or issue reports.
- On later sign-ins, enter your current authenticator code. **Use a recovery code instead** switches to a single-use recovery code when necessary.
- If you cancel or the challenge expires, return to sign-in and start again. A cancelled setup or failed verification does not finish sign-in.

When verification is enabled, the account card shows how many recovery codes remain. **Replace recovery codes** requires a fresh authenticator code or an unused recovery code; after replacement, all previous

recovery codes stop working. Save the new set before leaving the dialog.

**Turn off verification** appears only when your organization's policy allows it. Confirm with an authenticator or recovery code; successful removal signs you out. Organization-required verification cannot be turned off here.

If you lose both the authenticator and all unused recovery codes, contact your administrator for the organization's verified recovery process. A new password does not remove the authenticator requirement. For identity-provider MFA, use that provider's recovery process.

#### TIP

Use the code currently shown for this workspace in your authenticator app. If it expires while you are typing, use the next code. Read the attempts, expiry, and any cooldown shown; wait for a cooldown to finish before trying again. If **Security** cannot load, choose Retry security settings. Keep setup keys and recovery codes out of shared documents, videos, and support attachments.

#### GOOD TO KNOW

**SSO** accounts manage voluntary authenticator setup in their identity-provider settings. Your identity provider may also require its own verification during **SSO**. Follow the provider's screen for that check; these instructions describe Aperture Chat's authenticator screens.

#### WALKTHROUGH · 2:51 · [HELP](#) · [LEARN](#)

### Protect your account and recover access

Turn on two-step verification, sign in with a code or a recovery code, replace recovery codes, and change your password.

#### Before you begin

- An email-and-password account. With Organization **SSO**, your identity provider manages your password and multifactor settings, and the panel says so.
- An authenticator app on your phone, such as **Microsoft Authenticator**, Google Authenticator, or a password manager with one-time codes.
- A safe place for recovery codes, such as a password manager.
- Your current password. If you still have a temporary password, replace it first.

#### Step by step

- 1 Select your account card at the bottom of the sidebar. In the **Security** card, choose **Manage security**.
- 2 Under Two-step verification (Off), choose **Set up authenticator**.
- 3 Enter your Current password and choose Continue setup.
- 4 In your authenticator app, scan the QR code, or type the Setup key shown under it.

- 5 Tick I added this account to my authenticator., type the 6–digit code from the app into **Authenticator code**, and choose **Verify authenticator**.
- 6 Choose **Copy recovery codes** and store them somewhere safe. Tick I stored these recovery codes somewhere safe., then choose **Done**.
- 7 Check that Two–step verification shows On and Your authenticator is connected. 10 recovery codes remain.

#### Path: Sign in with a code

- 1 Sign in with your email and password as usual.
- 2 On Two–step verification, type the current 6–digit code from your authenticator app.
- 3 Choose Verify and continue.

#### Path: Sign in with a recovery code

- 1 On Two–step verification, choose **Use a recovery code instead**.
- 2 Type one unused recovery code into Recovery code and choose Verify and continue.
- 3 Open **Account** › **Manage security**: the count of recovery codes remaining has gone down by one.
- 4 When few codes remain, replace them.

#### Path: Replace recovery codes

- 1 **Account** › **Manage security** › **Replace recovery codes**.
- 2 Enter a current **Authenticator code** (or choose **Use a recovery code instead**).
- 3 Choose **Create** new recovery codes. The old set stops working at once.
- 4 **Copy** and store the new codes, tick I stored these recovery codes somewhere safe., and choose **Done**.

#### Path: Change your password

- 1 **Account** › **Password** card › Edit.
- 2 Enter Current password, then a New password of at least 12 characters, and the same in Confirm new password.
- 3 Choose Update password and wait for **Password** updated.
- 4 You stay signed in here; your older sessions on other devices are signed out.

#### Path: Lost or replaced phone

- 1 Sign in with your password and one recovery code (**Use a recovery code instead**).

- 2 If your organization allows it, choose **Turn off verification**, confirm with another recovery code, sign back in, and set up the new phone.
- 3 If verification is required by your organization, or you have no recovery codes left, ask your administrator to reset two-step verification for your account.
- 4 After the reset, sign in and set up your authenticator again.

### Check that it worked


- ✓ Two-step verification shows On, with Your authenticator is connected. and the number of recovery codes that remain.
- ✓ Your next password sign-in asks for an **Authenticator code**.
- ✓ After signing in with a recovery code, the remaining count is one lower.
- ✓ After a password change, the **Password** card shows **Password** updated.

## Troubleshooting

| WHAT YOU SEE                                                                           | WHAT TO DO                                                                                                                                                         |
|----------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "Current password is incorrect."                                                       | Retype your current password. If you have forgotten it, ask your administrator for a temporary password.                                                           |
| "The MFA code is invalid."                                                             | Wait for the next code and type it promptly. Codes depend on the time, so set your phone's clock to update automatically. A code that was already used is refused. |
| "Change the temporary password before enrolling MFA."                                  | Sign in and choose your own password first, then set up the authenticator.                                                                                         |
| "Finish verification or save your recovery codes before leaving this panel."           | Complete <b>Verify authenticator</b> , or copy and store the recovery codes and choose <b>Done</b> , before closing the panel.                                     |
| "Required by your organization. You cannot turn this off."                             | Your organization requires two-step verification, so <b>Turn off verification</b> is not offered. To move to a new phone, ask your administrator to reset it.      |
| "Your organization manages sign-in. Use your identity provider's security settings..." | You sign in with Organization <b>SSO</b> . Change your password and multifactor settings at your identity provider.                                                |
| "Use a password with at least 12 characters." or "The new passwords do not match."     | Use 12 or more characters and type the same new password in both fields.                                                                                           |
| "Password changes are only available for local password accounts."                     | Your account signs in with Organization <b>SSO</b> . Change the password at your identity provider.                                                                |

## 5. Finding your way around

The sidebar, the main area, and what each navigation button opens.



The chat workspace: navigation on the left, model selection above, and the message composer in the main area. Example workspace with synthetic data.

The screen has two areas: a sidebar on the left and the main work area on the right. The sidebar is how you move between the platform's views and your saved chats.

### Your first visit

The Getting started card explains what is ready for your account and what comes next. Open quick-start guide opens user help. Administrators see Open admin guide and Manage access. Follow the setup action offered for your role when no model is available. A missing model is an access or setup task, not a successful chat connection.

Choose a guide, take the suggested action, or select I'll explore on my own (or Dismiss welcome) when you are ready. Simply loading the page does not dismiss this card. You can return to guides from **Help** or console **Documentation** later.

| SIDEBAR BUTTON  | WHAT IT OPENS                                                                                                                                                      |
|-----------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>New chat</b> | A fresh conversation with your workspace assistant. This is the default view.                                                                                      |
| <b>Search</b>   | <b>Search past work</b> for chats (including archived), documents and decks, agents, and indexed documents, plus commands. Shortcut: Ctrl+K (Windows) or ⌘K (Mac). |
| <b>Drafts</b>   | A document and slide-deck editor with an AI assistant beside it.                                                                                                   |
| <b>Agents</b>   | Reusable agent profiles, plus an Automations tab for scheduled runs.                                                                                               |
| <b>Library</b>  | Two tabs: <b>Knowledge</b> (searchable document collections) and <b>Tools</b> (connections, prompts, and skills).                                                  |

Below the navigation buttons, your chat history is always listed under a small CHATS heading — folders first, then pinned chats, then recent chats. The next section covers organizing them.

## Resizing and collapsing the sidebar

- Click the small chevron (◀) near the top of the sidebar to collapse it and give your work more room. Click the chevron again (▶) to bring it back.
- Drag the sidebar's right edge to make it wider or narrower. Double-click that edge to collapse it.
- On a narrow window or a phone, the sidebar hides behind a menu button (three stacked lines, ≡) in the top-left corner. Tap it to open the menu; tap the dimmed area to close it.

## Using the keyboard

Use Tab and Shift+Tab to move between controls and Enter or Space to activate the focused control. In a dialog, keyboard focus stays with that dialog. Use its Close or **Cancel** button, or Escape when available, to return to the control that opened it. Confirmations explain when an action is permanent before you commit it.

## The bottom of the sidebar

- **Admin console** and Platform console — links shown only to administrators and platform owners (see “Where the consoles live”). Platform owners may also see a Release available notice [here](#).
- **Help** — opens the guided walkthrough videos and this downloadable guide.
- **Dark mode / Light mode** and **Theme schedule** — two small buttons beside **Help**: the moon or sun switches the appearance instantly, and the clock sets a schedule (see “Light and dark mode” later in this guide).
- **Install app** — appears on phones and tablets, and adds the workspace to your home screen (see “Install the app on your phone”).
- Your account card — your initials or photo, name, and role. Click it to open the account drawer, including Personalization memory when your organization has enabled it.

## Where the consoles live

Administrators choose **Admin console** near the bottom of the sidebar; platform owners also see Platform console. The account drawer's **Management** section lists the same consoles. Regular users see neither. The account drawer also holds **View as** role previews, **Usage this month**, and your **Archived chats** — all covered in “Your account”.

## Returning to a view

The browser address follows the workspace view and console tab. Browser Back and Forward return to those views; opening a saved chat or draft link still checks your account's access. A link does not share private content or grant permission to another person.

## 6. Organizing chats: folders, pins, and the chat menu

Keep the chats you care about easy to find.

Your chat history is always visible under the CHATS heading: folders first, then a **Pinned** group when something is pinned, then your recent chats (labeled **Recent** only when a group sits above them). Click the CHATS heading to hide or show the whole list; this is remembered for your account on this device. While hidden, it shows a count and a dot for unread replies, and **Search** still finds every chat.

### Folders

- Click the folder icon with a plus sign next to the CHATS heading.
- Type a name for the folder — a client, a matter, a project — and click **Create**.
- To file a chat into a folder, hover or focus its row, click the ⋮ (**More** actions) button, and choose Move to folder. Pick a folder, choose New folder to create one right there, or choose Remove from folder to take it out.
- Click a folder to expand or collapse the chats inside it.

#### GOOD TO KNOW

Deleting a folder does not delete its chats — they move back to the **Recent** list.

### Pinned chats

- Hover or focus any chat row, click ⋮ (**More** actions), and choose Pin chat.
- The chat moves to the **Pinned** group, just below your folders, where it stays until you choose Unpin chat from the same menu.

### Recent chats

- **Recent** chats — below folders and pinned chats — list your latest conversations. Click one to reopen it exactly where you left off.



- Hover or focus a row to reveal its ⋮ (**More** actions) button. The menu offers Pin chat (or Unpin chat), Move to folder, and **Archive**.
- A dot beside a chat marks a newest assistant reply you have not opened. Read status is saved to your account, so it matches in every browser and device.
- Archiving moves a chat out of the sidebar without deleting it — use it to tidy up. **Archived chats** are listed in your account drawer, where you can restore or permanently delete them.
- Click **View all chats** at the bottom of the list to browse your full history, including everything that no longer fits in the sidebar.

#### WALKTHROUGH · 1:36 · [HELP](#) > [LEARN](#)

### Organize and find your work

Create folders and file chats in them, pin, archive and restore, rename chats, and hide your chat list.

#### Before you begin

- A few chats in your sidebar.
- **Folders** are kept in this browser: they do not follow you to another browser or device.

#### Step by step

- 1 Choose **Create** chat folder (the folder button beside CHATS), type a Folder name, and choose **Create**.
- 2 Point at a chat, choose its **More** actions button (...), then Move to folder, and choose the folder.
- 3 Check that the chat now sits inside the folder, with the folder's count updated.
- 4 To keep a chat at the top, choose ... > Pin chat. It moves to **Pinned**.
- 5 To put a chat away, choose ... > **Archive**. Find it later in your account card > **Archived chats** > View, and choose **Restore**.
- 6 To rename the open chat, choose Rename chat (the pencil beside its title), type the name, and choose Save chat name.
- 7 Choose the CHATS heading to hide your chat list; choose it again to show it.

#### Path: Folders

- 1 **Folders** can hold subfolders, up to four levels: choose **Create** subfolder in ... on a folder.
- 2 ... > Move to folder > Remove from folder takes a chat out again.
- 3 **Delete** ... folder asks first; chats in it stay in **Recent**.

#### Path: Archive and delete

- 1 **Archive** keeps the chat and removes it from the sidebar (and unpins it).

- 2 **Account** card › **Archived chats** › View lists archived chats, with **Restore** and Permanently delete.
- 3 Permanently delete asks first and cannot be undone.

### Path: Rename with AI

- 1 After the first reply, choose Rename chat with AI beside the title for a suggested name.

### Path: See everything

- 1 **View all chats** opens All chats, with Filter chats to search titles.
- 2 View all pinned and View all folders appear when you have more than three.

### Check that it worked

- ✓ The folder shows the filed chat and its count.
- ✓ The pinned chat appears under **Pinned**.
- ✓ An archived chat is listed under **Archived chats** until you restore it.
- ✓ With the list hidden, CHATS shows how many chats are tucked away.

### Troubleshooting

| WHAT YOU SEE                                           | WHAT TO DO                                                                                                                                        |
|--------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>My folders are missing in another browser</b>       | <b>Folders</b> are stored in the browser where you made them. <b>Create</b> them again there; the chats themselves are in your account.           |
| <b>"Chat title cannot be blank."</b>                   | Type a name before choosing Save chat name, or choose <b>Cancel</b> renaming chat.                                                                |
| <b>"The AI could not rename this chat. Try again."</b> | <b>Try again</b> , or type a name yourself with Rename chat.                                                                                      |
| <b>A chat disappeared from the sidebar</b>             | It may be archived (account card › <b>Archived chats</b> ), inside a folder, or hidden with the CHATS heading. <b>Search</b> finds it either way. |

### WALKTHROUGH · 0:34 · [HELP](#) › [LEARN](#)

## Preview chats at a glance

Hover any listed chat to read its prompts and outputs in a compact, scrollable preview.

### Before you begin

- **Chats** in your sidebar, in All chats, or in **Archived chats**.

## Step by step

- 1 Point at a chat name in the sidebar and wait a moment; keyboard users can focus the chat row instead.
- 2 Read the **Chat** preview: the title and time, then Prompt 1 and Output 1, in order.
- 3 Scroll inside the preview to read later prompts and outputs. The footer counts prompts and outputs.
- 4 Move away, or press Escape, to close it without opening the chat.
- 5 Choose the chat if you want to open it.

### Path: In All chats

- 1 Choose **View all chats** under your recent chats.
- 2 Point at any chat in All chats; the preview opens beside the list.
- 3 Filter chats narrows the list by title first.

### Path: In Archived chats

- 1 **Account** card > **Archived chats** > View.
- 2 Point at an archived chat to preview it before you restore or delete it.

## Check that it worked

- ✓ The preview shows **Chat** preview, the chat's title, and Prompt 1.
- ✓ The footer counts the prompts and outputs.

## Troubleshooting

| WHAT YOU SEE                                       | WHAT TO DO                                                                          |
|----------------------------------------------------|-------------------------------------------------------------------------------------|
| "This chat does not have previewable content yet." | The chat has no messages yet.                                                       |
| No preview appears                                 | Rest the pointer on the chat name for a moment, or focus the row with the keyboard. |

## 7. Searching everything

Find saved work, preview a result, or run an available command.

- Click **Search** near the top of the sidebar, just below **New chat**, or press Ctrl+K (Windows) or ⌘K (Mac) from anywhere.

- The **Search past work** box opens. Type a few words from what you remember: a chat title, message text, an agent or draft name, slide text, or text from an indexed document.
- Results are grouped by kind. **Chat** results lead, matching text is highlighted, and archived conversations are labeled explicitly. Saved documents and decks open in their matching editor.
- Hover over or keyboard-focus a supported chat or draft result to preview it. The chat row also offers folder, pin, and archive or restore actions when that conversation is available in the current workspace.
- Click a result to open it, or use Up and Down followed by Enter. Escape closes search. Recently opened results are available when the query is empty.

## Commands and saved work

**Commands** appear in a separate section from search results. Type > at the start of this search box to see all commands available to your role and permissions, then keep typing to narrow them. **Commands** can open workspace views and permitted actions; an unavailable feature does not become accessible through search. The same > symbol inside a chat message has a different purpose: choosing an automation for that message.

### GOOD TO KNOW

**Search** includes only work your account can access. If the panel says Indexing your workspace, results may be incomplete until indexing finishes; try again after it completes. A missing result is not proof the work was deleted. Unsaved device-only changes may need saving to your account before server search can find them.

### WALKTHROUGH · 1:20 · HELP › LEARN

## Search, commands, and workspace links

Search every chat, draft, agent, and document you can access, reopen recent results, and run commands from the keyboard.

### Before you begin

- Some saved work to find. **Archived chats** are included.

### Step by step

- 1 Choose **Search** in the sidebar, or press Ctrl+K (⌘K on a Mac).
- 2 Type words from a title or a message.
- 3 Read the results, grouped by kind: Previous chats, Documents & knowledge, **Agents**, Automations, **Drafts**, and more.
- 4 Use the arrow keys and Enter, or choose a result, to open it.
- 5 Press Ctrl+K again. With an empty box, **Recent** lists what you opened from **Search**.
- 6 Type > to list commands, add a word to narrow them (for example >dark), and press Enter to run one.

### Path: Act on a chat from Search

- 1 Point at a chat result for its buttons: add to a folder, pin, and archive or restore.
- 2 **Archived chats** are marked **Archived** conversation and can be restored from here.

### Path: Commands

- 1 **New chat**; Go to **Chat**, **Drafts**, **Agents**, Automations, **Knowledge**, or **Tools**.
- 2 Switch to dark mode or Switch to light mode; Open **Help**; **Sign out**.
- 3 **Install app** appears on phones and tablets.

### Path: Workspace links

- 1 Each screen has its own address (for example /drafts or /chat/ followed by the chat's id).
- 2 Browser Back and Forward move between the screens you visited.
- 3 A link never shares private content: whoever opens it must sign in and have access.

### Check that it worked

- ✓ Results list the matching chats and other work.
- ✓ Choosing a result opens it.
- ✓ A command such as >dark changes the workspace at once.

### Troubleshooting

| WHAT YOU SEE                                                              | WHAT TO DO                                                                         |
|---------------------------------------------------------------------------|------------------------------------------------------------------------------------|
| "Indexing your workspace... results may be incomplete until it finishes." | <b>Search</b> is still catching up. <b>Try again</b> in a moment.                  |
| "No chats or workspace items found for ..."                               | Try fewer or different words, such as a word from the reply rather than the title. |
| "No commands match ..."                                                   | Type > alone to see every command.                                                 |
| "Search failed."                                                          | Check your connection and search again.                                            |

## 8. Your first chat

Send a message, pick a model, and control web search.

- Click **New chat** in the sidebar. A fresh chat greets you by the time of day — “Good morning”, “Good afternoon”, or “Good evening” with your first name (late at night it asks if you are burning the midnight oil) — above the tagline “Your approved models, your sources, your guardrails — ask Aperture Chat anything.” (The name reflects your organization's branding.)
- Click into the message box (the composer) — it reads **Ask anything...** — and type your question.
- Explore an idea, Compare options, and **Draft** a message offer starting points. Selecting one fills an editable prompt; review it before sending.
- Press Enter to send. Press **Shift+Enter** when you want a new line without sending.

### Picking the model

The model selector sits in the top bar of every chat. Click it to choose which AI model answers this conversation. Each chat remembers its own choice, and the list only ever shows models your workspace has approved for you.

Use the arrow keys, Home, End, or type a model name to move through available models; Enter selects the highlighted model. Each row has its own star: Set [model name] as default model selects that model and remembers it for new chats. The filled star marks the single current default. Selecting a row without its star changes this conversation's model without changing the default. **Drafts** has a separate default drafting model, set with the star beside a row in its own selector.

### Web search and the active-tools chip

- When the selected model supports it, public web search starts turned on. A chip at the bottom of the composer shows what is active: with one tool on it names it — for example **Web search** — and with several on it reads **Tools** with a count.
- Click the small x on that chip to clear the next message's active tools: **Knowledge**, **Web**, and **Agent**, selected MCP connections, and queued automations. Every new chat starts with web search on again when the model supports it.
- Replies that used the web come back with citations you can click, and every source is listed in the session details panel (covered below).

### Sending

- The paper-plane button sends your message with the current settings.

- The chevron (v) beside it opens send options — per-reply switches for **Knowledge**, **Web**, and **Agent**, plus the **Reasoning** level (covered in their own section).

## Dictate instead of typing

- Click the microphone button in the send controls at the bottom-right of the composer to start dictating. A live waveform shows that the platform is hearing you.
- Click the button again — it becomes a stop square — and your words are transcribed and inserted into the composer as editable text. Nothing sends until you press Enter.
- If the microphone is blocked or no speech was heard, the composer says so plainly so you can fix the problem and try again. The **Drafts** workspace has the same dictation button for drafting instructions.

### WALKTHROUGH · 1:38 · [HELP](#) · [LEARN](#)

## Start chatting

Send your first message, watch it work, stop a reply, edit and resend your prompt, and copy the answer.

### Before you begin

- A signed-in account with at least one available model. If the composer says Connect a model provider to start chatting..., see Choose models and request access.

### Step by step

- 1 Choose **New chat** at the top of the sidebar.
- 2 Check the model selector at the top right (**Model:** and a name). Choose another model there if you need one.
- 3 Type your request where it says **Ask anything...**, or choose a starter such as Explore an idea. A starter only fills in the opening words.
- 4 Press Enter, or choose the paper-plane Send message button, to send. **Shift+Enter** adds a new line instead.
- 5 Watch Aperture Chat is working while the reply is prepared; the answer appears under it.
- 6 Read the reply. The work trace above it reads complete, and the chat is listed under **Recent** with a title from your request.
- 7 To keep the answer, choose **Copy** response text under the reply and wait for Copied.

### Path: Stop a reply

- 1 While Aperture Chat is working, choose · stop at the top right of the trace.
- 2 The reply stops and keeps any text already written, with You stopped this response before it finished.
- 3 To ask again, use Edit and resend.

### Path: Edit and resend

- 1 Point at your message and choose **Edit message** (the pencil).
- 2 Change the text. Escape cancels.
- 3 Choose Send edited message. Everything after that message is removed, and the model answers the new text.

### Path: Regenerate

- 1 On a finished reply, choose **Regenerate response** (the circular arrows).
- 2 The model answers the same request again.
- 3 A stopped or failed reply has no actions; use Edit and resend instead.

### Path: Reuse a prompt in a new chat

- 1 Point at your message and choose **Load prompt in new chat** (the paper plane beside the pencil).
- 2 A new chat opens with your prompt and its attachments in the message box, ready to send.

### Check that it worked

- ✓ Your message and the model's reply both appear, and the trace above the reply reads complete.
- ✓ The chat is listed under **Recent** in the sidebar.
- ✓ Copied appears beside the reply's actions after you copy it.



## Troubleshooting

| WHAT YOU SEE                                                                                      | WHAT TO DO                                                                                                                                                                                                     |
|---------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| The message box says "Connect a model provider to start chatting..."                              | No model is available to your account. Open the model selector or <b>Why isn't a model listed?</b> , and request access.                                                                                       |
| "You stopped this response before it finished. Regenerate or resend the message to run it again." | A stopped reply has no Regenerate button. Choose <b>Edit message</b> on your prompt, then Send edited message.                                                                                                 |
| A reply reads "... did not return a completion: ..."                                              | The model service failed. Resend with <b>Edit message</b> . If it repeats, tell your administrator. Some self-hosted models reject Regenerate with "System message must be at the beginning."; resend instead. |
| The pencil says "Wait for the current response to finish before editing"                          | Wait for the reply, or choose · stop first.                                                                                                                                                                    |
| "The edited message could not be sent. Check the text and try again."                             | Make sure the edited text is not empty, then choose Send edited message again.                                                                                                                                 |
| "Copy unavailable"                                                                                | The browser blocked the clipboard. Select the reply text and copy it with the keyboard instead.                                                                                                                |

### WALKTHROUGH · 1:27 · [HELP](#) · [LEARN](#)

## Dictation, images, and diagrams

Dictate a message when your workspace has a transcription model, ask an image model for pictures, and copy, view, and download diagrams.

### Before you begin

- For dictation: a microphone, permission for this site to use it, and a transcription model (a Gemini Flash model) set up by your platform owner.
- For images: an image-generation model available to you in the model selector.
- For diagrams: any available model. Ask for a Mermaid diagram.

### Step by step

- 1 In the message box, choose Dictate a message (the microphone).
- 2 Allow microphone access if the browser asks.
- 3 Speak. A live waveform shows it is recording; recording stops by itself after 2 minutes.

- 4 Choose Stop dictation and transcribe (the square). The transcript is added to the message box as editable text.
- 5 Review the text, then send it yourself; dictation never sends on its own.

### Path: Images

- 1 In the model selector, choose a model that can generate images.
- 2 Describe the picture you want; ask for up to 4.
- 3 The images appear in the reply, labeled Generated image.
- 4 Point at an image and choose **Download** to save it. **Download** links last 7 days.

### Path: Diagrams

- 1 Ask for a diagram, for example: show the stages of a process as a Mermaid flowchart.
- 2 The reply shows a live figure with **Copy**, **PNG**, **SVG**, **Code**, and Edit in its header.
- 3 Choose **Code** to read or check the diagram's source; choose Diagram to switch back.
- 4 Choose **PNG** or **SVG** to download the picture, or **Copy** to copy it.

### Check that it worked

- ✓ Dictation: the transcript appears in the message box, ready to edit.
- ✓ Images: **Download** saves the image as a file.
- ✓ Diagrams: the figure renders in the reply, and **PNG** downloads it as a picture file.

## Troubleshooting

| WHAT YOU SEE                                                                                                                     | WHAT TO DO                                                                                               |
|----------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------|
| "No configured Gemini Flash model is available for dictation. ..."                                                               | Your workspace has no transcription model. Type your message, and ask your platform owner to set one up. |
| "Microphone access was denied. Allow microphone use for this site to dictate."                                                   | Allow the microphone in your browser's site settings, then try again.                                    |
| "Microphone capture is not available in this browser."                                                                           | Use a current browser on a device with a microphone.                                                     |
| "No audio was captured. Check the selected microphone and try again." or "The dictation model heard no speech in the recording." | Check which microphone your device uses and speak closer to it.                                          |
| "... returned no images for this request. Try rephrasing the prompt or selecting a different image model."                       | Rephrase the request, or choose another image model.                                                     |
| "This diagram could not be rendered. Open Code to inspect the source, or Edit to fix it."                                        | The model's diagram source has an error. Ask it to correct the Mermaid code, or fix it yourself.         |

## 9. Understand model availability and request access

Read the server's reason, ask for a model when permitted, and follow its request status.

- Open the chat model selector and choose **Why isn't a model listed?** If no model is usable, the No models available control opens the same explanation.
- **Models in your organization** separates **Usable now** from **Not available to you**. Read the reason on the model row: **Locked** describes an access restriction; **Provider offline** means you already have access but the provider needs attention.
- Choose **Request access** only where it is enabled. The request goes to an administrator for review; sending it does not grant immediate access.
- While a request is pending, the row shows its date and **Withdraw request**. Use **Refresh model access** after your administrator resolves it, then select the usable model in the chat picker.

### GOOD TO KNOW

Your organization may show only models you can already use. Its browsing policy controls whether you can inspect other enabled models and request access. Some restrictions cannot be lifted by a request; ask your administrator about the stated reason. An offline provider needs a platform-owner repair, not another group-access request.

## Choose models and request access

Pick the model for a chat, star your default, find out why a model is missing, and request, withdraw, and use a model your administrator approves.

### Before you begin

- A signed-in account with at least one model available. If the selector reads **No models available**, start at **Why isn't a model listed?**.
- Your organization must allow access requests. If it shows only models you can already use, ask your administrator about others.
- An administrator to review requests. Approval happens in their **Admin console > Model Access**.

### Step by step

- 1 In a chat, open the model selector (**Model:** and the model's name) at the top right.
- 2 Choose a model to use it in this chat. Each chat keeps its own model.
- 3 Choose the star beside a model to make it your default for new chats in this browser.
- 4 Choose **Why isn't a model listed?** at the bottom of the selector. **Models in your organization** lists **Usable now** and **Not available to you**, with the server's reason for each.
- 5 Beside a **Locked** model, choose **Request access**. The row changes to Request pending since and the date.
- 6 Wait for your administrator to approve. Approval adds you to a group that carries the model and signs you out of your current sessions.
- 7 Sign in again. The model is listed under **Usable now** and in the selector. Choose it and send a message.

### Path: Withdraw a request

- 1 Open the model selector > **Why isn't a model listed?**.
- 2 Beside the pending model, choose **Withdraw request**.
- 3 Withdrew the request for the model appears, and **Request access** is offered again.

### Path: Provider offline

- 1 A **Provider offline** badge means the connection behind the model is not working.
- 2 Requesting access does not fix a disconnected provider; the platform owner must reconnect it.
- 3 Choose **Refresh model access** later, or use another model in the meantime.

## Path: After approval

- 1 If the dialog or a chat says Session is invalid or expired. Sign in again., sign in again.
- 2 Open the selector: the new model is listed. Choose it for this chat, or star it to make it your default.
- 3 Send a message; the reply comes from the newly granted model.

## Check that it worked

- ✓ The selector shows **Model:** and your chosen model's name, and the reply comes from that model.
- ✓ The starred model is filled in, and new chats start with it.
- ✓ After approval and a fresh sign-in, the model is listed under **Usable now** with Available.

## Troubleshooting

| WHAT YOU SEE                                                                                       | WHAT TO DO                                                                                                                                     |
|----------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|
| "Not granted to any of your groups. An administrator can add you to a group that has this model. " | Choose <b>Request access</b> , or ask your administrator to add you to a group that has the model.                                             |
| "The provider behind this model is not connected right now." (Provider offline)                    | A request cannot fix this. The platform owner must reconnect the provider. Use another model meanwhile.                                        |
| "Disabled by the platform owner for the whole organization. "                                      | The model is switched off for everyone. Ask your administrator whether it will be enabled.                                                     |
| "Session is invalid or expired. Sign in again." right after a request                              | Your administrator approved the request, which signs you out. Sign in again; the model is then listed.                                         |
| "Model access requests are disabled by organization policy. "                                      | Your organization does not take requests here. Contact your administrator directly.                                                            |
| "Too many access requests. Try again shortly. "                                                    | Wait a minute. A pending request stays pending; you do not need to resend it.                                                                  |
| The selector reads No models available                                                             | No model is usable for your account yet. Choose it to open <b>Models in your organization</b> , request access, or contact your administrator. |

## 10. Reading and acting on replies

The work trace, citations, and the actions under every message.

---

### The work trace

- While a reply is being produced, a single status line shows the current step, a live timer, and how many steps remain (for example “4s · step 2 of 5”).
- Click that line at any time — while running or after — to expand the full trace: which model was routed, what context was prepared, whether web search ran, and how the answer was finalized.
- While a reply is still streaming, a stop button appears — its label is **Stop this response**. Clicking it halts the reply and keeps the conversation.

### Actions under each response

- **Copy** — copies the response as clean, readable text.
- Thumbs up / thumbs down — sends feedback your administrators can review.
- **Branch** — starts a new chat that continues from that response, leaving the original untouched.
- **Regenerate response** — asks the model to answer the same prompt again. Version arrows appear on the response with a counter (for example “2 / 3”) so you can flip between the versions it has produced.
- **Transfer to Drafts** — turns the response into an editable document in the **Drafts** workspace, where you can format, version, and export it.

### Actions under your own messages

- **Edit message** — reopens your prompt for editing. Sending the edit removes the later replies and asks again from that point.
- **Load prompt in new chat** — copies the prompt and its attachments into a fresh chat's composer, ready to adjust and resend.

### Diagrams in replies

When a reply contains a diagram (a flowchart, sequence, or similar), it renders as a real figure, not a code block. The figure's toolbar offers **Copy** for the diagram source, **PNG** and **SVG** to download it as an image file, and **Code** to flip between the drawing and its source text.

### Generated images

- When the selected model can generate images, just ask for one — the finished images appear directly in the reply.
- Click **Download** under an image to save it as a file.

## Citations and token usage

- When a reply used the web or your knowledge bases, numbered citations appear with it. Click a citation to open the original source.
- A small token count appears under a reply only when the model provider actually reported one — the platform never shows an estimate dressed up as a real number.

WALKTHROUGH · 1:43 · [HELP](#) · [LEARN](#)

### Follow the work trace and act on replies

Read the steps behind a reply, rate it with a note, branch it into a new chat, and turn it into an editable document.

#### Before you begin

- A chat with at least one finished reply (see Start chatting).

#### Step by step

- 1 Under a finished reply's heading, find Work trace and its summary, such as complete · 4 steps.
- 2 Choose Work trace to expand it and read every step: Routing request names the model, Preparing context lists the sources, then Generating answer and Finalizing response.
- 3 Choose Work trace again to collapse it.
- 4 Choose Send positive feedback (thumbs up) or Send negative feedback (thumbs down) under the reply.
- 5 Optionally type a note in the box that opens, then choose Send note. Note sent confirms it.
- 6 Choose **Branch** response into new chat to continue from this reply in a separate chat.
- 7 Choose **Transfer to Drafts** to open the reply as an editable document.

#### Path: Rate a reply

- 1 Choose the thumbs up or thumbs down under the reply. Positive feedback sent (or Negative feedback sent) appears.
- 2 In What worked well? or What went wrong?, add an optional note of up to 2,000 characters.
- 3 Choose Send note, or close the box to skip the note.

#### Path: Branch into a new chat

- 1 Choose **Branch** response into new chat under the reply.
- 2 A new chat opens, titled **Branch:** and the original title, holding the prompt and this reply.
- 3 Continue there; the original chat is unchanged.

## Path: Transfer to Drafts

- 1 Choose **Transfer to Drafts** at the right of the reply's actions.
- 2 **Drafts** opens a new document titled after the chat, saved as Version 1, with the reply and any citations.
- 3 Edit it, save versions, and export it (see **Draft** documents).

## Check that it worked

- ✓ The expanded trace lists the model used in Routing request.
- ✓ Note sent appears beside the reply's actions after sending a note.
- ✓ A chat titled **Branch: ...** appears under **Recent**.
- ✓ **Drafts** opens with the reply as a saved document, and the assistant rail says it was transferred.

## Troubleshooting

| WHAT YOU SEE                                                         | WHAT TO DO                                                                                                                                                       |
|----------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| The running trace moves to the next step before that step seems done | While a reply runs, the trace advances through the planned steps on a timer. Treat it as a guide to what the request involves; the finished trace is the record. |
| "Could not branch"                                                   | The branch could not be created, often because the connection dropped. <b>Try again</b> when the chat has saved.                                                 |
| No actions under a reply                                             | Actions appear only on finished replies. A stopped or failed reply has none; resend with <b>Edit message</b> .                                                   |
| Transfer to Drafts is missing                                        | It appears on finished replies with text. Wait for the reply to finish.                                                                                          |

## 11. Symbol shortcuts: / @ # \$ >

Five characters that pull prompts, agents, knowledge, skill files, and automations into a message.

Start a word in the message box with one of five symbols and a menu opens with matching items. You can also choose **Send options** → **Resources** to browse without memorizing a symbol. Forgot one? The bottom of Session details (①) lists all five, and **Help** → Symbol shortcuts: / @ # \$ > shows each menu in action.



| TYPE | WHAT IT OPENS                                                                                            |
|------|----------------------------------------------------------------------------------------------------------|
| /    | Saved prompts (inserted as text) and enabled MCP connections (turned on for this message).               |
| @    | <b>Agent</b> profiles — the reply routes through the one you pick.                                       |
| #    | <b>Knowledge</b> bases and the files inside them. Picking one turns on <b>Knowledge</b> for that source. |
| \$   | Saved skill files. The one you pick is attached to the message.                                          |
| >    | Automations, including paused ones. The one you pick runs on your message when you send.                 |

- Type the symbol, then keep typing to filter the list.
- Use the ↑ and ↓ arrow keys to move through the menu.
- Press Enter to insert the highlighted item, Tab to complete it, or Escape to dismiss the menu.

## Browse resources

On the **Resources** tab, use **Find a resource** to search by name, or select All, **Knowledge** bases, Files in knowledge sources, MCP connections, **Prompts**, **Agents**, Skill files, or Automations. A selected connection or automation shows a check; choose it again to remove it. Read an empty or failed-loading state before retrying. The **Type shortcuts in chat** reference explains each symbol below the results.

### TIP

The expanded composer focuses on editing a longer message. Return to the normal composer to browse resources, adjust send options, and send with the visible controls.

### WALKTHROUGH · 1:49 · HELP > LEARN

## Symbol shortcuts: / @ # \$ >

Type a symbol in the message box to add a saved prompt, an agent, knowledge, a skill file, or an automation, then send.

### Before you begin

- A chat with an available model.
- Items to choose from: saved prompts and skill files are shared by your administrators; agents, knowledge bases, and automations are ones you can use or have built.

### Step by step

- 1 Click in the message box. At the start of a word, type one of the five symbols: / @ # \$ or >.
- 2 A list opens above the message box. Keep typing to narrow it.
- 3 Use the up and down arrow keys to move, and Enter (or a click) to choose. Escape closes the list.

- 4 Check what was added: prompt text in the box, or a chip under it (**Agent**, **Knowledge**, a skill file, or an automation).
- 5 Finish your message and send it.

### Path: / Prompts and MCP connections

- 1 Type / to list saved prompts and enabled MCP connections.
- 2 Choose a prompt to put its text in the message box; edit it before sending.
- 3 Choose a connection to turn it on for this message.

### Path: @ Agents

- 1 Type @ to list the agent profiles you can use.
- 2 Choose one. @ and its name are added, and the **Agent** chip appears.
- 3 The reply runs through that agent's model, instructions, knowledge, and tools.

### Path: # Knowledge

- 1 Type # to list knowledge bases and the files in them.
- 2 Choose a knowledge base, or a single file to search its source.
- 3 The **Knowledge** chip appears; the reply searches it and cites what it uses.

### Path: \$ Skill files

- 1 Type \$ to list saved skill files, such as citation rules or an approval checklist.
- 2 Choose one. It is attached to the message as a file chip (Name.md).
- 3 The model follows it for this message.

### Path: > Automations

- 1 Type > to list your automations, including paused ones.
- 2 Choose one. It is queued in the chip under the message box.
- 3 When you send, the automation runs on your message. x removes it before sending.

### Check that it worked

- ✓ After choosing, the message box holds the prompt text or shows the matching chip.
- ✓ A message sent with # and \$ returns a reply that lists its sources and View N citations →.
- ✓ Session details > Symbol shortcuts lists all five symbols at any time.

## Troubleshooting

| WHAT YOU SEE                                                             | WHAT TO DO                                                                                             |
|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------|
| "No matching commands. Try another name."                                | Nothing matches what you typed after the symbol. <b>Delete</b> some letters, or check the item's name. |
| "No saved prompts or enabled MCP connections yet — add them under Tools" | Nothing is shared with you yet. Your administrator adds prompts and connections.                       |
| "No agent profiles available yet — create one under Agents"              | Build an agent (if your account can) or ask your administrator to share one.                           |
| "No knowledge bases are enabled yet — add one under Knowledge"           | Ask your administrator to share a knowledge base, or create one if your account can.                   |
| "Some files could not be loaded. Retry"                                  | Choose Retry in the list. Sources that already loaded stay usable.                                     |
| The list does not open                                                   | The symbol must start a word: type a space first if it follows other text.                             |

## 12. Attaching files, web pages, and connected sources

Upload from your computer, fetch a page by link, or pull from your organization's drives.

- Click the paperclip button at the bottom-left of the composer.
- Click **Upload from computer** to attach files from this device. The model reads them as context for your next message.
- Or click **Web page by link** and paste one or more web addresses — the platform fetches up to 3 public pages and attaches them as cited sources for this message.
- Attached items appear as chips above the composer. Click the x on a chip to remove one before sending.

### Attach from source

Below those options, under the **Attach from source** heading, the menu lists your organization's workspace sources: Google Drive, OneDrive, SharePoint, Box, and iManage. Availability depends on the service team's connector setup and your own source-account access.

### Connecting your own account

- Cloud sources read from your own account. The first time you pick one, the picker may ask you to connect it — click the Connect button for that source and a sign-in window opens.

- **Approve** read access in that window. The file list refreshes on its own once access is granted, and only you can see files from your account.
- The picker shows files at the top level of the drive or the folder configured for the source, within your account's permissions.

#### GOOD TO KNOW

If a source needs shared configuration, ask your administrator to coordinate with the service team. Shared connector administration is managed outside tenant administration. Your own Connect action in the attach menu remains separate and grants access only through your source account.

#### WALKTHROUGH · 2:08 · HELP › LEARN

### Attach files and sources

Attach a file from your computer, a web page by link, or a file from a connected cloud source, and know what the model actually receives.

#### Before you begin

- A chat with an available model.
- For cloud sources (Google Drive, OneDrive, SharePoint, Box, iManage): your administrator must set up the connector, and you sign in to your own account the first time.
- Files up to 25 MB each. **Web** pages must be public addresses starting with `http://` or `https://`.

#### Step by step

- 1 In the message box, choose Add attachment (the paperclip).
- 2 Choose **Upload from computer** and pick one or more files.
- 3 Wait for the file chip to finish uploading. Remove a file with the x on its chip.
- 4 Type your question about the file and send it.
- 5 Check the reply: it uses the file's content, and the trace includes a step for reading attachments.

#### Path: Web page by link

- 1 Paperclip › **Web page by link**.
- 2 Type or paste the full address into **Web** page address, such as `https://example.com`, and choose Add link (or press Enter).
- 3 The link appears as a chip. Add up to 3 links per message.
- 4 Send your question. The page is fetched when you send and cited in the reply (View 1 citation →).

## Path: Cloud sources

- 1 Paperclip › choose a source under **Attach from source**: Google Drive, OneDrive, SharePoint, Box, or iManage.
- 2 The first time, choose Connect and approve read access in the sign-in window. Only you see files from your account.
- 3 Tick up to 10 files and choose Attach selected.
- 4 If the dialog says the source is not ready, your administrator must finish setting up that connector.

## Path: What the model receives

- 1 Word (.docx) and text files (.txt, .md, .csv, .json, .xml, .html, .rtf, .eml, .log): the opening text, up to about 2,400 characters.
- 2 Images: the picture itself, if the selected model can read images; otherwise the model says it cannot see them.
- 3 Audio and video: a transcript, when your workspace has a transcription model. Media longer than 45 minutes is refused.
- 4 PDF, Excel, and PowerPoint files from your computer: currently only the file name. To ask about their content, add them to a knowledge base, attach them from a cloud source, or open Word files in **Drafts**.

## Check that it worked

- ✓ The file chip finishes uploading without Upload failed.
- ✓ The reply states facts that appear only in your file or page.
- ✓ A web-page reply offers View 1 citation → for the attached page.

## Troubleshooting

| WHAT YOU SEE                                                                  | WHAT TO DO                                                                                                                               |
|-------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------|
| "Attachment exceeds the 25 MB chat upload limit."                             | Split or compress the file, or add it to a knowledge base, which accepts larger files.                                                   |
| "Could not upload the attachment. Check your connection and try again."       | Check your connection, remove the failed chip, and attach the file again.                                                                |
| The reply says it cannot see the content of a PDF, spreadsheet, or slide file | Files of those types uploaded from your computer currently reach the model by name only. Use a knowledge base or a cloud source instead. |
| "Enter a full web address starting with http:// or https://."                 | Paste the complete address, including https://.                                                                                          |
| "You can attach up to 3 web links per message."                               | Send the first three, then attach the rest in a follow-up message.                                                                       |
| "Web source fetch failed for ..."                                             | The page could not be fetched (it may need a sign-in or be offline). <b>Copy</b> the text into your message instead.                     |
| "Google Drive is not ready." (or another source)                              | Your administrator has not finished that connector. Ask them, or upload the file from your computer.                                     |
| "Media longer than 45 minutes cannot be transcribed in this version."         | Split the recording and upload the part you need.                                                                                        |

## 13. Send options: Knowledge, Web, Agent, and Reasoning

Control what the model may use, one reply at a time.

- Click the chevron (v) beside the paper-plane send button.
- The menu opens with **Send now** at the top — that simply sends with the current settings.
- Below it are three switches you can turn on or off for the next reply, and a **Reasoning** slider at the bottom.

| SWITCH           | WHAT IT DOES WHEN ON                                                                                                                                                                                                              |
|------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Knowledge</b> | The reply searches your enabled knowledge bases and cites what it finds, with links back to the sources. A picker lets you choose which bases to use.                                                                             |
| <b>Web</b>       | The reply uses public web search. Results come back as clickable citations, and the sources are listed in session details. Available when the selected model supports it.                                                         |
| <b>Agent</b>     | The reply uses the selected agent profile and its permitted tools. Turning this on can choose the first available profile if none is selected; confirm the <b>Agent profile</b> field before sending, or choose a profile with @. |

## Reasoning level

The **Reasoning** group at the bottom of the menu is a three-position slider from **Fast** to **Smart**. **Fast** favors quicker answers; **Smart** makes the model think longer for more detailed output; the middle is a balanced default. The slider is only active when the selected model supports reasoning levels — otherwise it is disabled and says so.

Whenever tools are on, the active-tools chip appears next to the paperclip: the single tool's name, or **Tools** with a count. Click its x to turn off **Knowledge**, **Web**, and **Agent** and clear selected MCP connections and queued automations for the next message.

On the **Reply settings** tab, **Stream replies** controls whether text appears as the model produces it or only after the reply finishes. The **Resources** tab opens the searchable resource picker; **MCP connections and resources** also opens it from **Reply settings**. These controls change how the next reply is prepared or displayed; they do not supply missing provider or connector access.

### WALKTHROUGH · 2:38 · [HELP](#) · [LEARN](#)

## Knowledge, Web, Agent, and reply settings

Ground a reply in a knowledge base and check its citations, and choose Web, an agent, reasoning, streaming, and resources for a message.

### Before you begin

- A chat with an available model.
- For **Knowledge**: at least one knowledge base shared with you (see **Knowledge** bases).
- For **Web**: your workspace must connect web search to the model you use. Otherwise **Web** is greyed out.
- For **Agent**: an agent profile you can use (see **Agent** profiles).

### Step by step

- 1 In the message box, choose **Send options** (the arrow beside the paper plane). **Send options** opens on **Reply settings**.

- 2 Choose **Knowledge**, then choose the source to search, such as Litigation Playbook. Each source shows how many files it holds.
- 3 Choose Close send options. The **Knowledge** chip under the message box shows it is on.
- 4 Type your question and send it.
- 5 Check the reply: the work trace lists one more step, and View 2 citations → appears under the answer.
- 6 Choose View citations to open Session details › **Sources gathered**, with each source's K number, file name, and the passage used.

### Path: Web

- 1 **Send options** › **Web**. When it is available, the chip under the message box reads **Web search**.
- 2 Send your message. Aperture Chat searches public web sources and gives the results, with their addresses, to the model so the reply can cite them.
- 3 If **Web** is greyed out, your workspace has not connected web search to this model. Ask your administrator.

### Path: Agent

- 1 **Send options** › **Agent**. An **Agent profile** list appears under it.
- 2 Choose the profile you want; the line under it counts its knowledge, files, and tools.
- 3 Send your message. The trace reads **Agent** work trace. See **Agent** profiles to build your own.

### Path: Reasoning and streaming

- 1 Under **Reasoning**, drag the slider toward **Fast** for quicker answers or **Smart** for deeper thinking, on models that support it.
- 2 Clear **Stream replies** to show each reply only when it is complete; leave it ticked to watch the answer arrive.

### Path: Resources

- 1 **Send options** › **Resources** (or **MCP connections and resources**).
- 2 **Search** by name, or filter by **Knowledge** bases, Files in knowledge sources, MCP connections, **Prompts**, **Agents**, Skill files, or Automations.
- 3 Choose an item. A prompt's text fills the message box; other items turn on for this message.

### Path: Turn everything off

- 1 Choose x on the chip under the message box.



- 2 **Knowledge, Web, and Agent** turn off, selected MCP connections clear, and queued automations are removed.

### Check that it worked

- ✓ The chip under the message box names what is on (**Knowledge, Web search, Agent, or Tools** with a count).
- ✓ A knowledge reply shows View N citations →, and **Sources gathered** lists the files and passages.
- ✓ The reply's trace includes a step for each option you turned on.

### Troubleshooting

| WHAT YOU SEE                                                                                        | WHAT TO DO                                                                                                                             |
|-----------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------|
| <b>Web is greyed out: "Web search is turned off for this model by your workspace configuration"</b> | Your workspace has not connected web search to this model. Choose another model, or ask your administrator.                            |
| <b>With Web on, the reply stays on Aperture Chat is working</b>                                     | Some self-hosted models currently reject web-search requests. Choose · stop, turn <b>Web</b> off, and resend; tell your administrator. |
| <b>"Web search (...) is unavailable: ... Turn off Web for this reply to continue without it."</b>   | The search service could not be reached. Turn <b>Web</b> off for this message, or try again later.                                     |
| <b>Knowledge cannot be chosen</b>                                                                   | No knowledge base is shared with you. Ask your administrator, or create one if your account can (see <b>Knowledge bases</b> ).         |
| <b>"Create an agent profile in Agents to enable this mode."</b>                                     | No agent profile is available to you. Build one under <b>Agents</b> , or ask your administrator to share one.                          |
| <b>"... does not support reasoning levels"</b>                                                      | The selected model has no reasoning control. The slider has no effect for it.                                                          |

## 14. Session details and sources

Real token usage, active tools, every source this chat gathered, and the symbol shortcuts.

- Click the round information button (i) at the right end of the chat's top bar.
- The session details panel opens for the current chat.
- **Current chat** and **Model** — the user-message and response counts, and which model is active.
- **Tokens used** — shown only when the model provider reported real numbers. Otherwise the panel says "Not reported by the provider".

- **Tools** — exactly what is switched on for your next message: web search, knowledge, agent mode, plus any automations or MCP connections you added in the composer. “Off” means nothing is active.
- **Sources gathered** — every web source (as clickable links) and workspace citation collected during this conversation.
- Symbol shortcuts — at the bottom of the panel, the five symbols you can type in the message box and what each one pulls in. See “Symbol shortcuts: / @ # \$ >”.

The **Context window** meter describes how much context this conversation uses. For older conversations without provider counts it can show  $\approx$  and estimated from message length. Treat that as an estimate; **Tokens used** remains based on actual provider reports. When the context window fills, start a new chat and include the details it needs to continue reliably.

#### WALKTHROUGH · 0:59 · HELP > LEARN

### Session details and context

Open Session details to check the model, tokens used, how full the context window is, the sources a reply used, and the symbol shortcuts.

#### Before you begin

- A chat with at least one reply. **Sources gathered** appears after a reply that used knowledge or the web.

#### Step by step

- 1 In a chat, choose Session info (the i button at the top right). Session details opens on the right.
- 2 Read Session summary: **Current chat**, **Tokens used**, **Model**, **Knowledge** bases, **Tools**, and **Agent profile**.
- 3 Read **Context window**: the meter and how many of the model's tokens this chat already uses.
- 4 Scroll to **Sources gathered** to see every source a reply used, with its K (knowledge) or W (web) number and the passage.
- 5 Scroll to Symbol shortcuts for the five symbols you can type in the message box.
- 6 Choose Close session details when you are done.

#### Path: From a reply's citations

- 1 Under a reply, choose View N citations →.
- 2 Session details opens at **Sources gathered** for that reply. Hide citations closes it again.

#### Check that it worked

- ✓ **Tokens used** shows the counts the provider reported (in, out, and total), or Not reported by the provider.

- ✓ The context window meter shows a percentage and tokens used of the model's total.
- ✓ **Sources gathered** lists the same number of sources as View N citations.


## Troubleshooting

| WHAT YOU SEE                                            | WHAT TO DO                                                                                                                |
|---------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------|
| <b>Tokens used reads "Not reported by the provider"</b> | The model service did not return usage. Nothing is wrong with your chat; the context meter may use an estimate instead.   |
| <b>The context meter is marked estimated or ≈</b>       | It is based on message length, because the provider did not report usage for some messages.                               |
| <b>The context window is nearly full</b>                | Start a new chat for a fresh window. Older details may otherwise matter less to the answer.                               |
| <b>No Sources gathered section</b>                      | The replies in this chat did not use knowledge or the web. Turn on <b>Knowledge</b> in <b>Send options</b> and ask again. |

## Drafts, decks, and agents

### 15. Drafting documents

Generate, edit, version, and export documents with the drafting assistant.



Drafts combines the document editor with a separate assistant and document controls. Example workspace with synthetic data.

Click **Drafts** in the sidebar. The view pairs a full document editor with an AI assistant rail on the left that writes into the editor for you. A **Draft format** switch at the top of the editor has two buttons — **Document** and **Deck**. This section covers document mode; “Building a slide deck” covers deck mode.

#### GOOD TO KNOW

If the model selector says No models connected, AI drafting, AI rename, inline AI edits, and AI deck prompts are unavailable. You can still edit manually, import, save, use document history, and export. Follow the setup guidance for your role or ask your administrator for access to a working model.

#### Generating a draft

- Look at the Sources, **Web**, and Templates chips at the top of the assistant rail first: they show which context is on before you generate anything.

- In the assistant's message box, describe the document you need: a memo, an engagement letter, a summary. Be specific about audience and tone if it matters. The suggested requests under the heading fill the box for you.
- The assistant writes the document directly into the editor. You can also send any chat response here with **Transfer to Drafts**.
- Use the model selector in the toolbar area to choose which approved model writes and edits this draft.

## Opening a file from this device

Click the paperclip in the assistant's message box. Under From this device are two choices. Attach to chat gives files to the assistant to read with your next request; they become sources, not the document. Open in editor opens one file as the document itself, so you can edit it directly or select text and use **Ask AI**.

- Open in editor accepts Word (.docx, .dotx, .doc), **Markdown** (.md), plain text (.txt), and web page (.html) files. Word and web pages keep their formatting, and **Markdown** keeps its headings, lists, and emphasis. When something could not be carried over, an import note in the assistant rail says what.
- The file opens as a new draft and is saved to your account as one. The draft you had open stays in **Document history**; if it has unsaved changes, the unsaved-changes dialog asks first.
- Open in editor is unavailable while the assistant is still working on a request. In deck mode the same choice opens a PowerPoint file; see “Building a slide deck”.

## Editing like a document

The main toolbar keeps undo and redo, block style, bold, italic, underline, and inline AI editing close at hand. **Text** opens font, size, advanced styles, color, highlighting, and **Clear formatting**. **Paragraph** opens alignment, lists, and quotes. **More** opens **Copy document** and the **AI edit trail**. Select the intended text before opening a panel; Escape returns to the editor.

- The plus (+) **Insert** button opens links and citations alongside **Web image**, **Chart**, **Table**, **Divider**, and **Page break**. A manual page break starts a new page at your cursor; the editor also paginates longer documents automatically.
- Selecting text shows a floating toolbar beside it with **Ask AI**, block style, bold, italic, underline, strikethrough, link, and highlight.
- Type / at the start of a line for the command menu: AI actions such as **Continue writing** and **Summarize the document**, headings, lists, quotes, a table, a divider, a page break, a web image, or today's date. Keep typing to filter and press Enter to insert.
- Typing shortcuts format as you go: start a line with “#”, “##”, or “###” and a space for a title, heading, or subheading; “-” or “1.” for a list; or “>” for a quote. Type “---” and press Enter for a divider, and wrap words in “\*\*”, “\*”, or “`” for bold, italic, or code.
- Click inside a table to insert or delete rows and columns, turn the header row on or off, or delete the table. Click a picture to set its size (Small, Half, Large, or Full width), alignment, and alt text, or to delete it.

- Pasting from Word, Google Docs, or a web page keeps headings, lists, tables, links, pictures, and emphasis but drops the source's fonts, sizes, and colors so the text matches your document. Ctrl+Shift+V (⌘⇧V on a Mac) pastes plain text.

## Editing with AI

- Select a passage and choose **Ask AI** in the floating toolbar, or press Ctrl+J (⌘J on a Mac). With nothing selected, the same command writes new text at the cursor.
- Pick an action — **Improve writing**, **Fix spelling & grammar**, **Make shorter**, **Make longer**, **Simplify language**, **Turn into a bulleted list**, **Turn into a table**, **Summarize**, or **Use active voice** — or a tone or language, or type your own instruction.
- Review the suggestion. **Changes** marks what was added and removed; **Result** shows the finished text. Nothing in the document changes yet.
- Choose **Replace**, **Insert below**, **Try again**, or **Discard**. To adjust the suggestion instead, type a follow-up such as “shorter” or “add a number” in the box underneath.

An accepted AI edit stays highlighted briefly and is recorded with the document; **More → AI edit trail** highlights the recorded edits again so you can review what changed.

## Revise the document you already have

Describe the change in the assistant: shorten a section, change tone, apply a template, or transform the existing paper. Revision requests work from the current document. Ask explicitly to start over, replace, or clear it when that is your intention. The editor keeps instructions separate from the deliverable; review the resulting document and use Undo or a saved version when a change needs correction.

For an existing paper, **Assistant settings → Apply MLA layout** applies double-spaced 12-point Times New Roman, a student heading, centered title, body and reference indents, and a separate Works Cited page without asking a model. It preserves the paper's text and supports Undo. Fill missing student details and verify quotations and sources yourself; formatting is not citation verification.

The page navigator shows the current sheet and total pages outside the document text. Older drafts are paginated when opened, and manual edits reflow after you leave the editor. Page labels do not become part of the exported body.

## Status bar, outline, find, and zoom

- The status bar under the page shows the current page, word and character counts, and reading time. When text is selected, it counts the selected words too.
- **Outline** lists the document's headings; select one to jump to that section.
- Find (Ctrl+F, or ⌘F on a Mac) searches the document, with **Match case** and whole-word options. **Find and replace** (Ctrl+H, or ⌘⇧H on a Mac) replaces one match at a time or all of them.
- The keyboard button lists every editor shortcut (Ctrl+/, or ⌘/ on a Mac). The zoom controls resize the page on screen from 50% to 200% without changing the document or its exports.

## Versions and comparing

- Click **Save version** whenever the draft reaches a good state. You can restore an earlier version at any time from the draft's history.
- **Compare versions** opens a read-only visual redline of two saved versions, so you can see exactly what changed between them. It needs two genuinely different saved versions before it activates.
- Open **Document history** to browse saved documents and decks. Hover or keyboard-focus an entry for a preview. **Archive** moves finished work into **Archived**, where **Unarchive** restores it; **Delete** requires confirmation and permanently removes the saved item and revisions.
- Read the save status: account-backed documents and decks can be reopened after sign-in on another device. A pending or failed save stays a working copy in this browser. If the account version changed elsewhere, resolve the reported conflict before treating the save as complete.

## Leaving with unsaved changes

When you leave an edited draft through workspace navigation, open another chat or draft, or choose **Sign out**, the unsaved-changes dialog lets you Keep editing, Save copy and continue, or **Discard** and continue. The saved copy goes into local document history in this browser. If browser storage fails, the app keeps you in the draft so you can recover it. A forced security sign-out can still interrupt work; save important changes regularly.

## Exporting

- Click **Export** to open the export panel. In document mode it offers a **Word document** (an editable Word file with preview page breaks and embedded images), **Markdown** (best for plain text or web publishing), and **Print / Save as PDF**, which opens your browser's print dialog with the saved version — choose “Save as PDF” there to keep a PDF copy.
- Exports always run from a saved version. If you have unsaved edits, the panel says **Save your edits first** and offers a one-click **Save version and export** (or **Save version and print**) button.
- When your browser supports it, choose **Choose a location** to select a destination, or **Browser downloads** to use its download folder. Word and print retain document formatting; **Markdown** preserves the text without its typography. Exporting a file does not by itself upload pending work to your account.

### TIP

On a narrow window, the assistant rail becomes a slide-out drawer so the editor gets the full screen. Open it with the pen tab on the left edge or Ctrl+. (⌘. on a Mac); resting the mouse on that edge previews it. Press Escape or click the dimmed area to close it.

### WALKTHROUGH · 3:34 · [HELP](#) · [LEARN](#)

## Draft documents

Start a document from a blank page, a chat reply, or a file; have the assistant write it; edit with AI; save a version; and export it.



## Before you begin

- An available model for AI writing and editing. Without one, you can still type, open files, save, and export.
- To open a file in the editor: a Word (.docx or .doc), **Markdown** (.md), text (.txt), or web page (.html) file.

## Step by step

- 1 Choose **Drafts** in the sidebar. A blank document opens, with the **Document** Assistant on the left.
- 2 Type a title in the title field at the top of the page.
- 3 In Ask the document assistant, describe the document you need, or choose a suggestion. Choose Apply instruction (the arrow).
- 4 Wait while the assistant writes on the page. Ready to write returns when it is done.
- 5 Check the title. If the assistant renamed the document, type your own title again.
- 6 Select a passage, choose **Ask AI** in the floating toolbar (or press Ctrl+J, ⌘J on a Mac), and choose an action such as **Make shorter**.
- 7 Review the **Changes** or **Result** view, then choose **Replace**, **Insert below**, **Try again**, and **Discard** are the alternatives.
- 8 Choose **Save version** and wait for Saved.
- 9 Choose **Export**, set Save to to **Browser downloads**, and choose **Word document**. The receipt names the downloaded file.

## Path: From a chat reply

- 1 Under a finished chat reply, choose **Transfer to Drafts**.
- 2 **Drafts** opens a new document with the reply, saved as Version 1.
- 3 Continue from step 6 of the main steps.

## Path: Open a file from this device

- 1 Choose Attach file (the paperclip) in the assistant's message box.
- 2 Under From this device, choose Open in editor, and pick a .docx, .doc, .md, .txt, or .html file.
- 3 The file opens as a new draft with its headings and lists. The draft you had open stays in **Document history**.
- 4 Attach to chat instead gives the file to the assistant to read with your next request.



### Path: Write it yourself

- 1 Click on the page and type. **Markdown** shortcuts work: **#** and a space starts a title, **-** and a space starts a list.
- 2 Type **/** at the start of a line for the command menu: headings, lists, tables, page breaks, and AI actions.
- 3 Use **Text**, **Paragraph**, **More**, and **Insert** above the page for formatting.

### Path: Export formats

- 1 **Word document**: an editable .docx with page breaks and images.
- 2 **Markdown**: a .md file for plain text or web publishing.
- 3 **Print / Save as PDF**: opens the browser's print dialog; choose Save as PDF there.
- 4 Save to: **Browser downloads** saves straight to your downloads; **Choose a location** asks the browser where to save.

### Check that it worked

- ✓ The toolbar shows Saved and the version number after **Save version**.
- ✓ The export receipt names the file, for example synthetic-vendor-onboarding-checklist.docx, and offers **Download** again.
- ✓ An opened file shows a note in the assistant rail: the file opened in the editor.

## Troubleshooting

| WHAT YOU SEE                                                                     | WHAT TO DO                                                                                                                     |
|----------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------|
| "AI drafting is unavailable." or <b>Model: No models connected</b>               | No model is available to you. You can still type, open files, save, and export. Ask your administrator for model access.       |
| <b>The document title changed to a template name, such as Board Update Draft</b> | The assistant can rename a new document after a template it matched. Type your own title back into the title field, then save. |
| "Could not open ... Choose a .docx, .doc, .md, .txt, or .html file."             | Save the file in one of those formats and open it again. For PDFs, use a knowledge base or attach the file to chat.            |
| "Wait for the assistant to finish before opening another document."              | Let the current request finish, then open the file.                                                                            |
| <b>Export says "Save your edits first"</b>                                       | Choose <b>Save version and export</b> (or <b>Save version and print</b> ). Exports use the saved version.                      |
| "Local only — server save failed"                                                | Your changes are kept in this browser. Choose Retry, or export a copy before you leave.                                        |
| "Could not export ... Try again from the export menu."                           | <b>Try again</b> , or choose <b>Browser downloads</b> in Save to.                                                              |

## 16. Building a slide deck

Turn a draft into slides, edit them, brand them, and export real PowerPoint.

**Deck** mode turns the **Drafts** workspace into a slide editor with a real PowerPoint export. Click the **Deck** button on the **Draft format** switch. If the draft already has document content, a dialog asks "Turn this draft into slides?" — choose **Start a blank deck** to keep the document untouched, or **Convert into slides** to split its headings and text into slides. Your document versions are kept either way.

### The filmstrip

- Slide thumbnails run along the side. Click one to select it, or drag it to reorder the deck.
- Hover a thumbnail for its actions: move up, move down, duplicate, and delete (undo restores a deleted slide).
- **Add slide** at the end of the filmstrip opens the layout menu for a new slide.
- **Slide sorter**, at the right of the deck toolbar, lays every slide out in a grid. Drag to reorder, double-click a slide to edit it, and use each slide's duplicate and delete buttons.

## Slide layouts

Above the slide, **Layouts** previews seven arrangements: **Title**, **Title + bullets**, **Two columns**, **Image + caption**, **Section**, **Quote**, and **Closing**. Select one to switch the selected slide; the layout button on the slide and in the toolbar opens the same choices. **Themes** offers five color palettes — Aperture, Midnight, Editorial, Violet, and Forest; a palette change recolors every slide, keeps text, media, and layout, and can be undone.

## Editing slide text

- Click any text region on a slide to edit it in place. The deck toolbar carries undo and redo, bold, italic, underline, and **Text** for font, size, and color; **More** copies the deck outline or lists the keyboard shortcuts.
- Drag a text box by its frame to move it. Guides snap it to the slide's edges and center and to other boxes; arrow keys nudge the selected box (Shift moves ten times as far).
- Inside bullet lists, press Tab to indent a bullet one level and Shift+Tab to outdent it.

## Editing slides with AI

- **Edit slide with AI** (the AI pen in the deck toolbar, or Ctrl+J / ⌘J) reworks the whole slide: **Improve this slide**, **Make it punchier**, **Cut the text in half**, **Write speaker notes**, **Pick a better layout**, **Split into two slides**, **Fix spelling & grammar**, a tone or language, or your own instruction.
- The review shows the slide before and after, plus any new speaker notes. Choose **Apply to slide**, **Try again**, or **Discard**; undo restores the previous slide.
- To rewrite only part of a slide, highlight the text and choose **Ask AI**, then review the suggestion before replacing it.

## Opening a PowerPoint file

- In deck mode, click the paperclip in the assistant's message box and choose Open in editor (Edit slides directly) to open a .pptx file as the deck itself.
- Each slide keeps its title, subtitle, bullets with their indent levels, two-column text, speaker notes, one picture, and its layout artwork as the background. **Text** with no place in the chosen layout moves into that slide's speaker notes instead of disappearing.
- If the presentation is larger than the deck limit, pictures are left out first and then layout artwork, and the assistant rail says so. The opened deck is a new deck; the deck you had open stays in **Document history**.

## Starter templates and your brand template

- **Deck starters & brand themes**, above the slide (or the Templates chip in the assistant rail), opens the templates panel with five starter structures: **Pitch deck**, **Quarterly review**, **Project kickoff**, **Client proposal**, and **Training session**. Pick one and start it to get scaffold slides you replace.
- A **Use templates in chat** toggle (off by default) lets the selected template guide the deck assistant when it drafts for you.
- **Upload brand template** accepts your firm's own file — “.pptx or .potx — brand colors, fonts, logo, and every slide's text are extracted.” The extracted theme is stored only on this device.

- **Once** a brand theme is loaded, **Apply to deck** restyles the current deck with its colors, fonts, and logo, and a “Load all N slides” button can replace the deck with every slide from the uploaded template.

## Backgrounds and AI images

- The background button in the toolbar opens a per-slide menu: **Upload background...** (a PNG, JPEG, or WebP from this device), **Use on every slide**, **Remove from this slide**, and **Clear from all slides**.
- **Generate AI slide image** writes an image prompt for you, prefilled from the slide's content; press **Generate image** and the result becomes that slide's background. This needs an image-generation model enabled for your workspace — the button says so honestly when none is.
- In the assistant rail, **Toggle AI slide images** makes the assistant also generate an image for every slide whenever it drafts a deck.

## Speaker notes

Every slide has a **Speaker notes** field below the stage. **Notes** export with the deck: they land in each slide's notes pane in the PowerPoint file, they appear in the **Markdown outline**, and they follow you into presentation mode.

## Presenting from Aperture

- Click **Present deck** in the deck toolbar (the monitor icon), or press F5 (⌘Enter on a Mac); add Shift to start from the current slide. The deck opens full screen.
- Advance by clicking the slide or with the arrow keys or Space; move back with the left arrow. Type a slide number and press Enter to jump to it.
- The **Notes** button (or N) shows your speaker notes under the slide. **Presenter view** (or P) shows the current slide, the next slide, a timer you can pause or reset, and your notes side by side.
- Press B or W to black or white out the screen; press it again to return. Leave with Escape or the **Exit** button in the top corner.

## Saving and exporting

- **Save version** stores decks to your account with revision history, like documents. Read the status to confirm the save reached the server. A local working copy or pending save remains only in this browser until synchronization succeeds; use the **Only on this device** reminder to find unfinished saves.
- **Document history** includes decks. Preview, reopen, archive, unarchive, and confirmed deletion use the same account controls as documents. A concurrent account edit can prevent a stale save or deletion; read the conflict message before retrying.
- Click **Export** to open the **Export deck** panel: **PowerPoint deck** produces an editable .pptx that mirrors the slides on screen — speaker notes included in each slide's notes pane — and **Markdown outline** exports slide titles, bullets, and speaker notes as text.
- Limits: a deck holds up to 100 slides and 8 MB of content. The editor tells you plainly if a deck exceeds them.

## Build a slide deck

Turn a document into slides, build a deck from an outline or a starter template, refine it, present it, and export a PowerPoint file.

### Before you begin

- An available model for the deck assistant and **Edit slide with AI**. Templates, layouts, notes, presenting, and export work without one.
- Optional: your organization's PowerPoint brand template (.pptx or .potx, up to 30 MB, no macros).
- AI slide images need an image-generation model in your workspace.

### Step by step

- 1 Choose **Drafts** in the sidebar, then choose **Deck** in the **Draft format** switch. A blank deck opens with the **Deck** Assistant.
- 2 Type a title for the deck.
- 3 In Ask the deck assistant, describe the deck: how many slides and what each covers. Choose Apply instruction.
- 4 Wait while Building slides... runs; on smaller models this can take a few minutes. The filmstrip fills and the rail says how many slides were drafted.
- 5 Review each slide. Use **Layouts**, **Themes**, speaker notes, and **Edit slide with AI** to refine it.
- 6 Choose **Save version** and wait for Saved.
- 7 Choose **Present deck**, then press P for presenter view. Escape exits.
- 8 Choose **Export**, set Save to to **Browser downloads**, and choose **PowerPoint deck**. The receipt names the .pptx file.

### Path: Convert a document

- 1 In a document with content, choose **Deck** in the **Draft format** switch.
- 2 Turn this draft into slides? appears. Choose **Convert into slides** (or **Start a blank deck**).
- 3 The headings become slides. Your document versions are kept either way.

### Path: Start from a template

- 1 In a deck, choose **Deck starters & brand themes**.
- 2 Choose a starter: **Pitch deck**, **Quarterly review**, **Project kickoff**, **Client proposal**, or **Training session**.
- 3 Choose Start and the template's name. **Replace** the scaffold text with your content.

### Path: Your brand template

- 1 In **Deck starters & brand themes**, choose Upload under Brand theme and pick your .pptx or .potx file.
- 2 Choose **Apply to deck**. The colors, fonts, and logo carry through the slides.
- 3 Load all N slides brings in every slide design. The brand theme is stored only on this device.

### Path: Present

- 1 Choose **Present deck**. Arrow keys, Space, or a click move forward.
- 2 Press N to show notes, P for presenter view (timer, next slide, notes), and B to black out the screen.
- 3 Press Escape, or choose **Exit**, to leave.

### Path: Export formats

- 1 **PowerPoint deck**: an editable .pptx with your slides and speaker notes.
- 2 **Markdown outline**: slide titles, bullets, and notes as text.

### Check that it worked

- ✓ The filmstrip lists the slides, and the rail reports how many the assistant drafted.
- ✓ **Presenter view** shows the current slide, a timer, the next slide, and your notes.
- ✓ The export receipt names the .pptx file.

## Troubleshooting

| WHAT YOU SEE                                                                                                                          | WHAT TO DO                                                                                                                |
|---------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------|
| Building slides... runs for several minutes                                                                                           | Larger decks take longer, especially on smaller models. Ask for fewer slides, or wait; your draft stays intact meanwhile. |
| Generate AI slide image is greyed out: "No image-generation model is enabled for your workspace"                                      | Your workspace has no image model. Upload your own background from Slide background instead.                              |
| "Could not open ... Choose a PowerPoint .pptx file."                                                                                  | Open in editor accepts .pptx and .potx files. Save other formats as .pptx first.                                          |
| "... exceeds the 30 MB brand-template upload limit." or "Macro-enabled presentations are not accepted. Save as .pptx or .potx first." | Use a smaller template, and save it without macros.                                                                       |
| "Brand template upload failed: ..."                                                                                                   | Check that the file is a valid PowerPoint template and try again.                                                         |

## 17. Recover work kept only on this device

Distinguish a browser working copy from a completed account save.

The sidebar's **Only on this device** reminder appears when work has changes kept in this browser that have not reached your account. It can include chats whose last save failed, documents or decks with unsent changes, and older history from before account sync.

- Open the reminder to review the named items. For chats whose last save failed, choose **Retry all** and check the resulting save state.
- For documents and decks, choose **Open Drafts**, open the item, and save it to your account. Review any load, quota, connection, or conflict error before leaving the browser.
- Older browser history remains separate until you explicitly import each item from **Drafts** → **Document history** with **Import to my account**. Confirm the item belongs in this account before importing.
- **Archive** finished drafts when you want them out of the active list. **Account** archives remain recoverable; archiving a browser-only item does not prove it has been uploaded.

### GOOD TO KNOW

**Clear list** dismisses these reminders without deleting or uploading work. New edits may appear again. **Hide this reminder** turns off this account's sidebar notice in this browser; save status remains in **Drafts** and chats. Downloading an export also leaves pending account saves to be resolved.

## Save, organize, and recover your drafts

Save documents and decks to your account, reopen them from Document history, archive and restore them, return to an earlier version, and recover work kept only on this device.

### Before you begin

- At least one document or deck in **Drafts**.

### Step by step

- 1 In a document or deck, choose **Save version** and wait for Saved in the toolbar.
- 2 Choose **Document history** (the clock at the top of the assistant rail).
- 3 Point at a card to preview it, or choose the card to reopen that draft.
- 4 To put finished work away, point at its card and choose **Archive. Draft** archived. Open **Archived** to restore it. confirms it.
- 5 Choose **Archived** to find it, then choose **Unarchive** to bring it back to Active.
- 6 Under Versions of this draft, choose a version to return to it.
- 7 **Delete** is separate: choose **Delete**, read the confirmation, then choose **Delete** draft or **Cancel**.

### Path: When a save fails

- 1 Local only means the latest changes are kept in this browser but have not reached your account.
- 2 Choose Retry, or export a copy before you leave.
- 3 If **Draft** changed elsewhere appears, keep your local changes safe (export), then choose Reload server copy.

### Path: Only on this device

- 1 The sidebar reminder counts items only on this device. Choose it to open the list.
- 2 **Retry all** retries chat saves; **Open Drafts** takes you to drafts and decks with unsent changes.
- 3 **Clear list** dismisses the current reminders; **Hide this reminder** hides the notice. Neither uploads nor deletes your work.

### Path: Leaving with unsaved edits

- 1 Keep your unsaved edits? appears when you leave or open something else.
- 2 Choose Keep editing, Save copy and continue, or **Discard** and continue.

### Check that it worked

- ✓ Saved appears in the toolbar with the version number.



- ✓ The draft is listed under **Archived** after archiving, and under Active after **Unarchive**.
- ✓ Versions of this draft lists each saved version with its time.

## Troubleshooting

| WHAT YOU SEE                                                              | WHAT TO DO                                                                                                               |
|---------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------|
| "Local only — ... Your changes are kept on this device."                  | The server save failed. Choose Retry, or export a copy before you close the browser.                                     |
| "Not saved — ..., and browser storage could not keep a copy."             | Keep the workspace open, then retry or export your changes.                                                              |
| "this draft exceeds the 2 MB server draft limit" (or the 8 MB deck limit) | Split the document, or remove large images, then save again.                                                             |
| "Draft changed elsewhere"                                                 | Another tab or device saved this draft. <b>Export</b> your local version if you need it, then choose Reload server copy. |
| "Save two different versions of this draft to compare them"               | <b>Compare versions</b> needs two saved versions that differ. Save again after an edit.                                  |

## 18. Agent profiles

Reusable bundles of model, prompts, knowledge, and tools.

An agent profile bundles a model route, meta prompts, knowledge bases, and MCP tools into one reusable configuration. A usable profile needs an approved, enabled model backed by a working provider. A saved profile marked not connected is a configuration example, not proof that it can run.

### Creating an agent

- If your account has agent authoring permission, click **Agents** in the sidebar, then **New Agent**.
- Give it a name and a short description of what it is for.
- Work through the editor tabs: **Profile** (name, model, and description), **Knowledge** (assign knowledge bases), **Tools** (select MCP tools), **Prompts & Skills** (attach system prompts and skill files), **Access** (who can use it), and **Hermes** (the optional learning companion).
- Click the save button. Standard users create private profiles; administrators control group sharing and organization publishing. Check model readiness and access before trying the profile in chat.

### Using an agent in chat

- Type @ in the composer and pick an available profile to use its configured model route, instructions, knowledge, and tools.

- The **Agent** switch in **Send options** can choose the first available profile when none is selected. Check **Agent profile** and select the intended configuration before sending.
- If no profiles appear, ask your administrator to check that a ready profile is available to your account.

From the **Agents** view you can also open a profile to edit it, jump straight into a chat with it, or delete it. The Automations tab at the top of this view holds your scheduled runs — covered in “Scheduled automations”.

WALKTHROUGH · 1:52 · [HELP](#) · [LEARN](#)

## Agent profiles

Build a private agent with its own instructions, model, knowledge, and skill files, then use it in chat and check that the reply follows it.

### Before you begin

- Permission to build agents: your platform owner allows users to build their own agents, and your administrator grants **Can build agents** to one of your groups. Until then, **Agents** says Building agents is not available to you.
- An available model for the agent to use.
- Optional: knowledge bases, tools, prompts, and skill files shared with you, to attach to the agent.

### Step by step

- 1 Choose **Agents** in the sidebar, then **New Agent**.
- 2 On **Profile**, type an **Agent** name and choose the AI model. Optionally start from a template such as Research Analyst.
- 3 In System prompt, write the agent's instructions: its role, what it should produce, and what it must never do. Meta prompt adds standing rules on every turn.
- 4 On **Knowledge**, tick the knowledge bases it should search on every message.
- 5 On **Tools**, tick any tools it may use in agent mode; tools that need approval ask before they run.
- 6 On **Prompts & Skills**, tick saved prompts and skill files to add to its instructions.
- 7 On **Access**, read Private — only you can use this agent.
- 8 Choose **Create** agent and wait for the agent's name followed by created. It is ready to use in chat.
- 9 In a chat, type @ and choose the agent, then type your question and send it.
- 10 Check the reply: the trace reads **Agent** work trace, and the answer follows the agent's instructions and cites its knowledge.

### Path: Use it from Send options

- 1 **Send options** > **Agent**, then choose the profile in **Agent profile**.
- 2 Check the selected profile before sending; turning **Agent** on can select the first one.

### Path: Use it from the Agents page

- 1 On the agent's card or in its editor, choose **Chat** (or Use in **Chat**).
- 2 A chat opens with **Agent** mode on for that agent.

### Path: Edit, duplicate, or delete

- 1 On the agent's card, choose Edit, change the settings, and choose Save **Profile**.
- 2 **More** actions on the card offers Duplicate and **Delete**.

### Path: Share it

- 1 **Agents** you build stay private to your account.
- 2 Ask an administrator to share the agent with a group.

### Check that it worked

- ✓ The agent appears under **Agents** with your chosen model.
- ✓ In chat, the **Agent** chip shows the agent, and the reply's trace reads **Agent** work trace with Preparing agent tools and the knowledge search.
- ✓ The answer follows the instructions you wrote (here, three bullet points from the playbook) and offers its citations.

### Troubleshooting

| WHAT YOU SEE                                                              | WHAT TO DO                                                                                                                 |
|---------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------|
| "Building agents is not available to you."                                | Your groups lack <b>Can build agents</b> , or the service policy does not allow user-built agents. Ask your administrator. |
| "Select a base AI model and enter an agent name before saving."           | Type a name and choose an available model in AI model.                                                                     |
| "... is not connected, so the agent cannot answer."                       | The agent's model has no working provider. Choose another model in <b>Profile</b> and save.                                |
| "Knowledge base ... is turned off, so the agent answers without it."      | Ask the knowledge base's owner to turn it back on, or untick it.                                                           |
| "No agent profiles available yet — create one under Agents" when typing @ | You have no agents yet, or none is shared with you. Build one, or ask your administrator.                                  |
| The reply ignores part of the instructions                                | Smaller models follow short, specific instructions best. Shorten the System prompt, or choose a stronger model.            |

## Library and automations

### 19. Knowledge bases

Give chats grounded, citable access to your documents, web pages, and APIs.

**Knowledge** lives in the **Library**: click **Library** in the sidebar, then the **Knowledge** tab. The **Knowledge Bases** panel lists every collection with its status, security posture, and whether it is enabled.

- Click **Add Knowledge Base** and name the collection.
- Choose what feeds it. A knowledge base can index documents you upload, a public web page (**Add a web link** — enter the address and an optional note), or an API (**Connect an API** — enter the endpoint and its authentication).
- **Create** the base, then check its row in the table: it shows the status, the security posture, and whether the base is enabled.

#### GOOD TO KNOW

Only enabled knowledge bases can be searched from chat, and access control is enforced per source — people only ever retrieve what they are allowed to see.

#### Using knowledge in a conversation

- Type **#** in the composer to reference a knowledge source or choose from Files in knowledge sources. A file choice references its name in your prompt and searches that file's knowledge source; the row says which source will be searched.
- Or turn on the **Knowledge** switch in send options to let the reply search your enabled bases.
- Replies grounded in knowledge return citations that link straight back to the source documents.

#### GOOD TO KNOW

The **#** shortcut menu shows Loading files... while retrieving file choices. If Some files could not be loaded appears, use Retry in that menu. Files from sources that loaded successfully remain usable. If access was removed or a source needs configuration, ask your administrator to resolve it.

#### WALKTHROUGH · 2:03 · HELP › LEARN

### Knowledge bases

Create a knowledge base from files, web pages, or an API, watch it index, and ask a question that cites it.

## Before you begin

- To create your own: your administrator grants Can build knowledge bases to one of your groups. Without it, you use the knowledge bases shared with you.
- Files up to 250 MB each: PDF, Word, Excel, PowerPoint, email (.eml, .msg), text, **Markdown**, **CSV**, JSON, images, audio, and video.
- For web pages: public addresses. For pages behind a sign-in, paste the text instead.
- **Connected** sources such as SharePoint, Box, Google Drive, and iManage are set up by administrators and shared with you.

## Step by step

- 1 Choose **Library** in the sidebar; **Knowledge** opens.
- 2 Choose New knowledge base.
- 3 Type a Name, keep Start with on Upload files, and choose **Create**. Who can search it is Only me; administrators manage sharing.
- 4 Drag files onto Drag files here, or choose browse your computer.
- 5 Wait for each file to show Added and appear under Indexed files with its passage count. The base shows Ready.
- 6 In a chat, type # and choose your knowledge base, then type your question and send it.
- 7 Check the reply: it answers from your files and offers View N citations →. Open it to see each file and passage in **Sources gathered**.

## Path: Add web pages

- 1 In the knowledge base, open **Web** pages.
- 2 Type the Page address (and an optional Name), then choose Fetch and add page.
- 3 The page is listed under Indexed web page; Sync now fetches it again later.
- 4 For a page behind a sign-in, tick Paste the text instead of fetching the page, paste it, and choose Add text.

## Path: Connect an API

- 1 Open API. Enter the API address, an optional Path and Name.
- 2 Under Sign-in choose None (public API), API key, Bearer token, or Sign in with provider (OAuth), and fill in its fields.
- 3 Choose Fetch and add. The JSON or text response becomes searchable passages; Sync now refreshes it.

### Path: Ask without #

- 1 **Send options** > **Knowledge**, then choose your knowledge base.
- 2 Send the question; the reply searches only that source.

### Path: Turn it off or tidy it

- 1 The On switch (Assistants can search ...) controls whether chats can search it.
- 2 **Delete** a file with the bin beside it under Indexed files.

### Check that it worked

- ✓ Each file shows Added and is listed under Indexed files with a passage count.
- ✓ The knowledge base shows Ready and On.
- ✓ A question with # returns facts from your files and View N citations; **Sources gathered** names the file (for a PDF, with the page).

### Troubleshooting

| WHAT YOU SEE                                                                                                 | WHAT TO DO                                                                                                    |
|--------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------|
| "You can use the knowledge bases shared with you in chat. Ask an administrator if you need to add your own." | Creating knowledge bases is not granted to your groups. Ask your administrator for Can build knowledge bases. |
| "Skipped ... (save it as .docx first)" or "(unsupported type)" or "(larger than ... MB)"                     | Convert old .doc files to .docx, use a supported type, or split very large files.                             |
| "No readable text — only the name is searchable"                                                             | The file has no text layer (for example a picture-only PDF beyond the OCR limit). Upload a text version.      |
| "Give the knowledge base a name."                                                                            | Type a Name before choosing <b>Create</b> .                                                                   |
| "Enter the page address."                                                                                    | Type the full address, including https://.                                                                    |
| "Some files could not be loaded. Retry" in the # menu                                                        | Choose Retry. Sources that already loaded stay usable.                                                        |

## 20. The Tools library

MCP connections, saved prompts, and skill files inside the Library.

Click **Library** in the sidebar, then the **Tools** tab. The **Tools Library and Connectors** panel has three sections:

| TAB                | WHAT LIVES THERE                                                                          | WHERE IT APPEARS IN CHAT |
|--------------------|-------------------------------------------------------------------------------------------|--------------------------|
| <b>Connections</b> | MCP tools and connections available to models and agents.                                 | Type / in the composer.  |
| <b>Prompts</b>     | Saved prompt templates, including <code>{{variables}}</code> you fill in when using them. | Type / in the composer.  |
| <b>Skills</b>      | Skill files — reusable instructions for recurring workflows.                              | Type \$ in the composer. |

- Click **New Prompt** or **New Skill** to create one: give it a name, a category, a description, and the content itself, then save.
- Click any existing card to review or edit it; the trash icon deletes it permanently.
- Every tool row shows its real status — draft, approval required, or enabled — so you always know whether a connection is actually live before relying on it.

#### WALKTHROUGH · 1:08 · [HELP](#) · [LEARN](#)

### Tools and the Library

Find the connections, saved prompts, and skill files you can use, read them, and add a private connection when your administrator allows it.

#### Before you begin

- **Prompts** and skill files are created by administrators; you can read and use them.
- To add your own connection: your administrator grants Can build tools to one of your groups, and you have the MCP server's address.

#### Step by step

- 1 Choose **Library** in the sidebar, then **Tools**.
- 2 **Connections** lists the MCP connections you can use in **Agent** mode, with On, Off, or Needs setup.
- 3 Choose **Prompts** to see the saved prompts shared with you. Choose a prompt's View to read its content.
- 4 Choose **Skills** to see the skill files shared with you.
- 5 Use them in chat: type / for prompts and connections, and \$ for skill files, or attach them to an agent.

#### Path: Add a private connection

- 1 With Can build tools granted, **Connections** shows Add connection.
- 2 Type a Name, choose how Aperture connects (HTTP or SSE), and enter the Server URL starting with https://.

- 3 Choose Add connection. It starts turned off; set sign-in, access, and approval on the next screen, test it, and turn it on.
- 4 **Connections** you add are private to you; an administrator can share them.

#### Path: Read a connection's status

- 1 On: available under policy. Off: turned off. Needs setup: missing an address, token, or sign-in.
- 2 Last test failed: the latest connection test did not succeed.
- 3 Asks before each use: you approve each run in chat (**Approve** MCP tool run?).

#### Check that it worked

- ✓ **Prompts** and **Skills** list the items shared with you, and View shows each one's content.
- ✓ Typing / in chat lists the same prompts, and \$ lists the same skill files.

#### Troubleshooting

| WHAT YOU SEE                                                                                                 | WHAT TO DO                                                                                             |
|--------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------|
| " <b>Connections are managed by your administrators. These are the ones you can use in Agent mode.</b> "     | You cannot add connections yet. Ask your administrator for Can build tools if you need one.            |
| " <b>Shared prompts are managed by administrators.</b> " or " <b>Skills are managed by administrators.</b> " | Ask your administrator to add or change a prompt or skill file.                                        |
| " <b>Use a full address starting with https://</b> "                                                         | Enter the complete server address.                                                                     |
| " <b>Not signed in to the provider</b> " or " <b>No access token saved</b> "                                 | Open the connection's Sign-in tab and finish signing in, or ask your administrator.                    |
| " <b>MCP tool run denied. No message was sent.</b> "                                                         | You chose Deny when asked to approve a tool run. Send again and choose <b>Approve</b> if you trust it. |

## 21. Scheduled automations

Recurring chat or drafting runs, as one model call or a multi-step chain.

Automations are saved workflows that run a chat or drafting task on a schedule. They are useful for recurring digests, weekly status checks, report refreshes, policy reviews, and any workflow where the same model chain should run the same way each time. Enabled schedules run automatically in the background (times are UTC), and each scheduled run delivers its output as a new chat thread in your sidebar.



## Creating an automation

- Click **Agents** in the sidebar, then the Automations tab, then **New automation**.
- Name it — for example “Monday client digest”.
- Choose what it runs against: chat or draft.
- Pick a trigger: **Daily** (a time each day), **Weekly** (a day and time), **Once** (a specific date and time), or **Cron expression** (such as `0 9 * * 1` for 9:00 every Monday, in UTC).
- Write the initial input — what the first step should work on.
- Build the model chain: each step has a model and an instruction, and each step's output feeds into the next step. Use **Add step** for multi-step chains and the x to remove a step.
- Click **Save automation**.

## What to check before saving

- **Models** — every step must use a model you can actually access. If a model is missing, your admin may need to approve it for your group.
- **Context** — write the initial input as if the automation will run without you watching. Include the matter, audience, date range, source expectations, and desired output format.
- **Cadence** — use **Once** for a single future run, **Daily** or **Weekly** for regular recurring work, or a reviewed **Cron expression** for another schedule. All times are UTC; check the intended local-time equivalent before saving.
- **Ownership** — name the automation so another person can tell what it does later.

## Running and managing

- Each automation card shows its schedule, its steps, and its last run with an honest result — succeeded or failed, with the reason.
- **Run now** executes the whole chain immediately and shows a transcript of every step's output.
- The toggle pauses or resumes the schedule; the pencil edits it; the trash deletes it.
- You can also queue an automation into any chat by typing > in the composer — it runs when you send.
- The list shows the automations you created. Organization administrators can review the automations in their workspace.

### TIP

Always use **Run now** once after creating or editing an automation. If the run fails, the card shows the reason so you can fix the model choice, input, or instructions before depending on the saved schedule.

## Scheduled automations

Build an automation that runs a prompt through one or more models on a schedule, deliver it to a new chat or a new draft, run it now, and read its history.

### Before you begin

- At least one model available to you. Without one, Automations says no approved models are available and **New automation** is off.
- The prompt you want to run, and when it should run (the time zone matters).

### Step by step

- 1 Choose **Agents** in the sidebar, then Automations.
- 2 Choose **New automation**.
- 3 Under What it does, type a Name and the Prompt for the first step.
- 4 Under Steps, choose a model (or one of your agents) for step 1 and an optional instruction. Choose **Add step** to pass the result to another model.
- 5 Under Schedule, choose **Daily**, **Weekly**, **Once**, or Custom, set the time and Time zone, and check Next runs.
- 6 Under Deliver results to, choose **New chat** or New draft.
- 7 Leave On — runs on this schedule turned on, and choose **Save automation**.
- 8 On the automation's card, choose **Run now** to test it. Wait for the result and choose Open chat (or Open draft).

### Path: Daily

- 1 How often › **Daily**, then set Time.
- 2 Next runs lists the coming days at that time.

### Path: Weekly

- 1 How often › **Weekly**, choose the day (Mon to Sun), then set Time.
- 2 Next runs lists the coming weeks.

### Path: Once

- 1 How often › **Once**, then set Run at to a date and time in the future.
- 2 Next runs shows that single run; if the time has passed it says No upcoming runs — this time has passed.

### Path: Custom (cron)

- 1 How often › Custom, then type a **Cron expression** with five fields: minute, hour, day of month, month, day of week.
- 2 For example 30 7 \* \* 1–5 runs at 7:30 AM on weekdays. Check Next runs before saving.

### Path: Deliver to a new draft

- 1 Under Deliver results to, choose New draft.
- 2 The last step is asked for a complete document, saved as a new draft.
- 3 After **Run now**, choose Open draft to edit and export it in **Drafts**.

### Path: Pause, resume, edit

- 1 The switch on the card pauses the automation (Paused) or resumes it (Active).
- 2 Edit changes its steps, schedule, or delivery; choose Save changes.
- 3 Run history (N) lists recent runs: Succeeded, Failed, or Skipped, how each was started, and how long it took.

### Check that it worked

- ✓ After saving, the card shows the schedule, the step chain, and Next run in ....
- ✓ **Run now** reports that the run saved the result to a new chat (or draft), and Output from this run shows each step.
- ✓ Open chat shows the delivered conversation, titled with the automation's name and the run time.
- ✓ Run history lists the run as Succeeded, run now.

## Troubleshooting

| WHAT YOU SEE                                                                                                                  | WHAT TO DO                                                                                                     |
|-------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------|
| "No approved models are available yet. An administrator must enable at least one available model before automations can run." | Ask your administrator for access to a model; <b>New automation</b> stays off until then.                      |
| "Give the automation a name." or "Every step needs a model or agent."                                                         | Type a Name and choose a model or agent for every step.                                                        |
| "Choose a date and time for the one-time run."                                                                                | Set Run at for a <b>Once</b> schedule.                                                                         |
| "Enter a cron expression, for example 0 9 * * 1 -5."                                                                          | Type five fields separated by spaces, and check Next runs.                                                     |
| "Run failed: ..." or Last run failed on the card                                                                              | Read the message in Run history. A model or agent may be unavailable; edit the step and run it again.          |
| Paused after failures                                                                                                         | After 3 scheduled failures in a row the automation pauses itself. Fix the cause, then turn the switch back on. |
| An agent step does not use its tools                                                                                          | <b>Tools</b> and MCP servers need a person to approve them, so automations run agent steps without tools.      |

## Appearance, your account, and help

### 22. Light and dark mode

Switch appearance now or schedule it for this browser.

---


- 1 Find the moon or sun button beside **Help** near the bottom of the sidebar — a moon (**Dark mode**) while in light mode, a sun (**Light mode**) while in dark mode.
- 2 Click it. The entire platform switches immediately — no reload, nothing to save.
- 3 Click it again to switch back.

#### Schedule light and dark mode

Choose the clock button beside it to open **Theme schedule**. Enable **Switch automatically**, set different **Light mode at** and **Dark mode at** times, then choose **Save schedule**. The schedule uses this device's local time every day and is saved in this browser. You can still switch manually until the next scheduled change; disable the switch to stop automatic changes.

## 23. Your account

Your profile, personalization memory, password, role previews, usage, archives, and signing out.



Your account includes profile settings, password controls, and two-step verification. Example workspace with synthetic data.

Click your account card (your initials or photo) at the very bottom of the sidebar. The account drawer opens with everything about you in one place.

- **Profile** — click the card with your name and the pencil icon to edit your display name, firm, website, bio, phone number, and photo (upload an image up to 5 MB, or paste a URL). Click **Save profile** when done, or **Cancel** to discard.
- **Management** — administrators and platform owners see this section, listing the consoles they can open; the sidebar's **Admin console** and Platform console links go to the same places. Regular users do not have it.
- **Personalization memory** — opens your private memory manager when the feature is enabled. Use it to control, review, add, correct, pin, forget, or clear what the assistant remembers about you.
- **Password** — accounts that sign in with a password can change it here (click the pencil, enter the current password, then the new one twice — at least 12 characters). The app continues with a new authenticated session after saving and revokes previous sessions. Accounts that sign in through **SSO** manage their password with the **SSO** provider instead, and the panel says so.
- **Security** — choose **Manage security** to set up an authenticator for a local account, review the remaining recovery-code count, or replace codes after verifying your identity. Turning verification off is available only when organization policy permits it and signs you out. **SSO** accounts follow their identity-provider settings for voluntary enrollment.

- Role and organization — your role, your organization, any personal token caps that apply to you, and how you sign in.
- **View as** — if your role allows it, preview the workspace exactly as a lower role would see it. A note reminds you which role you are previewing; switch back the same way.
- **Usage this month** — your prompts, responses, and estimated tokens.
- **Archived chats** — every chat you have archived, with buttons to restore each one to your sidebar or delete it permanently.
- **Sign out** — ends your session and returns to the sign-in screen.

## WALKTHROUGH · 2:21 · [HELP](#) · [LEARN](#)

### Personalize, use mobile, and get help

Edit your profile, choose light or dark mode and a theme schedule, install the app on a phone, tablet, or computer, and report a problem.

#### Before you begin

- For installing: Safari on iPhone or iPad, Chrome on Android, or Chrome or Edge on a computer.
- For a report: what you did, what you expected, and what happened; optionally a screenshot (**PNG**, **JPEG**, **GIF**, or **WebP**, up to 10 MB) with no passwords or personal data.

#### Step by step

- 1 Select your account card at the bottom of the sidebar, then the pencil beside your name (Edit account profile).
- 2 Change the fields you need, such as Phone number or Firm or organization, and choose **Save profile**. **Profile** saved confirms it.
- 3 Choose the moon button beside **Help** for **Dark mode**, or the sun for **Light mode**.
- 4 Choose the clock button (**Theme schedule**), turn on **Switch automatically**, set **Light mode at** and **Dark mode at** to different times, and choose **Save schedule**.
- 5 Choose **Help** at the bottom of the sidebar for the walkthroughs and the printable User guide (PDF).
- 6 To report a problem, choose Report a problem in **Help**, fill in Subject and Message, optionally add a Screenshot, and choose Send report.
- 7 Wait for Report sent: an administrator can now review it in **Analytics**.

#### Path: iPhone or iPad (Safari)

- 1 Open your workspace address in Safari and sign in. Add Aperture Chat to your home screen opens (or choose **Install app** in the menu).
- 2 Tap the **Share** button in the browser toolbar.

- 3 Scroll the share sheet and choose Add to Home Screen.
- 4 Tap Add. Aperture Chat appears on your home screen and opens full screen.

### Path: Android (Chrome)

- 1 Open your workspace address in Chrome and sign in. The install card opens (or choose **Install app** in the menu).
- 2 If it offers **Install app**, choose it and confirm in the phone's install sheet.
- 3 Otherwise open the browser's : menu and choose Add to Home screen or **Install app**, then confirm.

### Path: Computer (Chrome or Edge)

- 1 Aperture Chat shows no install card on computers.
- 2 Use the browser's own install control: the install icon at the right of the address bar, or the browser menu › Install (Chrome: Cast, save, and share › Install page as app; Edge: Apps › Install this site as an app).
- 3 Confirm Install. The app opens in its own window and appears with your other apps.

### Path: On a narrow screen

- 1 Choose Open menu (top left) to reach navigation, chats, **Help**, **Install app**, and your account.
- 2 Choose where to go; the menu closes. The same permissions apply as on a computer.

### Check that it worked

- ✓ **Profile** saved appears, and your account card shows the new details.
- ✓ The workspace switches between light and dark at the scheduled times on this device.
- ✓ The installed app opens from your home screen or app list.
- ✓ Report sent confirms an administrator can review your report.



## Troubleshooting

| WHAT YOU SEE                                                                                                        | WHAT TO DO                                                             |
|---------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------|
| "Use a complete http(s) website URL."                                                                               | Enter the full website address, including https://.                    |
| "Choose a profile photo that is 5 MB or smaller."                                                                   | Use a smaller image.                                                   |
| "Choose two different times."                                                                                       | <b>Light mode at</b> and <b>Dark mode at</b> must differ.              |
| "Add both a subject and a message."                                                                                 | Fill in both fields before Send report.                                |
| "Choose a screenshot that is 10 MB or smaller." or "Screenshot attachments must be PNG, JPEG, GIF, or WebP images." | Attach a smaller <b>PNG</b> , JPEG, GIF, or WebP image.                |
| "The issue report could not be sent."                                                                               | Your report was not received. Check your connection and send it again. |

## 24. Personalization memory

Save and recall private preferences in plain English across chats and sessions.

Personalization memory belongs to your account, not to one conversation. When service policy and your organization administrator allow it, administrators and regular users can each build their own private memory. A saved item can influence a new chat immediately and still be available after you sign out and return in a later session.

### Save a memory in ordinary English

- In any chat, state what should persist in natural language — for example: “Remember that I prefer a short summary before the detail,” “Please remember that I work in commercial litigation,” or “Remember my project codename is Silver Horizon.”
- Send the message. The assistant confirms the memory after it is saved; you do not need a slash command or special syntax.
- Start a new chat, or return in a later sign-in session, and ask a natural question such as “What do you remember about me?”, “What are my writing preferences?”, or “What is my project codename?” The same account memory is available across those sessions.

#### GOOD TO KNOW

Memory is different from the context window shown in Session details. The context window describes how much of the current chat the selected model can hold at once; personalization memory is durable account context that follows you into other chats and sessions.

## Open and control the memory manager

- Click your account card — your name and role at the bottom-left of the sidebar.
- Click Personalization memory. The dialog titled “What the assistant remembers about you” opens.
- Use memory in my chats controls whether saved items are applied to answers. Learn from my conversations controls whether the assistant may notice durable preferences automatically. Turning a switch off does not expose or silently erase existing items.
- To add something directly, type it under Add something you want remembered, choose Standing instructions, Preferences, About you, Your work, or Other details, and click Add.
- Review the list below. Click a memory's wording to correct it; pin an important item; use Forget to remove one item; or choose Forget everything to clear the list.

### GOOD TO KNOW

Only you can read your memory content. Administrators can see counts and purge memories for compliance, but the platform never shows them what an individual memory says.

### TIP

If Personalization memory is missing from your account drawer, the feature is not currently available at the platform, organization, or group level. Ask your administrator; refreshing the page cannot override that policy.

### CAREFUL

Do not use memory as a password or secret vault. Credentials, keys, and sensitive identifiers are rejected on purpose.

### WALKTHROUGH · 1:28 · [HELP](#) · [LEARN](#)

## Personalization memory

Save preferences in plain English, recall them in later sessions, and stay in control.

### Before you begin

- Memory must be allowed by your platform and your organization. If it is off, the Personalization memory row says why.
- Only you can read your memories. Administrators see how many you have and can delete them, but never read them.

### Step by step

- 1 Select your account card at the bottom of the sidebar and choose Personalization memory.
- 2 Check Use memory in my chats (applies saved memories to answers) and Learn from my conversations (lets the assistant notice lasting preferences itself).

- 3 Type a memory in Add something you want remembered, choose its type (for example Preferences), and choose Add.
- 4 Choose Pin on a memory that should always apply. Choose its text to correct the wording.
- 5 In any chat, say Remember that ... followed by the preference; it is saved for later chats.
- 6 In a new chat, ask a question the memory answers, and check that the reply uses it.

### **Path: Forget**

- 1 Personalization memory › Forget on one memory removes it.
- 2 Forget everything asks Forget all ...? This cannot be undone. before removing them all.

### **Path: Turn memory off**

- 1 Turn off Use memory in my chats to stop applying memories. Saved memories stay private and manageable.
- 2 Turn off Learn from my conversations to save only what you add or ask it to remember.

### **Check that it worked**

- ✓ The memory is listed in Personalization memory under its type, with Pin on if you pinned it.
- ✓ After Remember that ... in chat, the sentence appears in the list.
- ✓ A later chat's reply uses it without being told again.

## Troubleshooting

| WHAT YOU SEE                                                                              | WHAT TO DO                                                                                                                                                                                   |
|-------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "Memories cannot contain credentials, keys, or sensitive identifiers."                    | Passwords, keys, tokens, one-time codes, and card or account numbers are never stored. Keep them out of memory.                                                                              |
| "Write a little more so the memory is useful later."                                      | Write a full phrase of at least a few words.                                                                                                                                                 |
| "Pinned memories already fill your memory limit. Unpin or forget one before saving more." | Unpin or forget a memory, then add the new one.                                                                                                                                              |
| "That memory kind is disabled by your organization's policy."                             | Choose another type, or ask your administrator.                                                                                                                                              |
| The row says memory is turned off                                                         | "Memory is turned off for this platform.", "... for your organization.", or your groups' policy: ask your administrator.<br>"You turned memory off for your account.": turn it back on here. |
| No note under the reply says a memory was saved or used                                   | The reply does not currently show that note. Check Personalization memory to confirm the memory was saved.                                                                                   |

## 25. Install the app on your phone

Put your workspace on the home screen with its own icon and name.

On a phone or tablet, tap **Install app** near the bottom of the sidebar. The "Add ... to your home screen" dialog previews your organization's app name and icon.

- On Android, choose **Install app** if offered, then follow the native install sheet. Otherwise open your browser's menu and choose "Add to Home screen" or "**Install app**".
- On iPhone and iPad, tap the browser's **Share** button, scroll the share sheet, choose "Add to Home Screen", then tap Add.
- Open the new home-screen icon to return to the same workspace in a full-screen app.

## 26. Getting help

Walkthrough videos and this guide, always one click away.


---

- Click **Help** at the bottom of the sidebar for guided videos, captions, transcripts, and this PDF. The guide also covers additional procedures.
- Administrators can also open **Documentation** inside the console for role-specific videos and a printable guide.
- Choose Report a problem in **Help** to describe a problem, the affected screen, and the expected result. **Share** only permitted details with your administrators; omit passwords, keys, and recovery codes.

# The Admin console

## 27. Opening the Admin console

Where the console lives and what its nine tenant-governance tabs control.



The Admin console opens the workspace controls available to your administrator account. Example workspace with synthetic data.

- 1 Click **Admin console** near the bottom of the sidebar. (Only workspace administrators and platform owners see it; the account drawer's **Management** section lists it too.)
- 2 The console opens with nine tabs across the top: **Users**, **Groups**, **Model Access**, **Connections**, **SSO**, **Analytics**, **Policies**, **Audit**, and **Alerts**. **Policies** is always present between **Analytics** and **Audit**; service-wide availability determines which organization controls are active inside it.

Everything you change here writes through the admin API immediately — changes persist across refreshes and restarts, and every action lands in the tenant audit trail. Status messages under the header tell you honestly whether an action synced or failed.

### TIP

The **Documentation** button at the top of the console opens narrated video walkthroughs of every tab, plus this guide as a PDF.

## 28. Users

Create tenant accounts, manage passwords, and remove leavers safely.

The **Users** tab lists the organization accounts you administer, with each person's role, groups, sign-in method, and status.

### Reading and filtering the list

- Use the group filter to work one team at a time.
- Watch the Auth and Status columns: Auth shows sso or local, and Status shows Active, Inactive, or **Pending** for an approved person who has not signed in yet. When the default group is on, new **SSO** accounts start in **Default Users**; otherwise they have no groups until you add them.

### Per-row actions

Every row has an Actions column with account controls:

- **Access** — opens the read-only **Model access trace** for this person, with their groups, each model's status, and the gates that explain it. See [Resolve model requests and trace access](#) for the full workflow.
- **Password** — opens the password dialog for that person: type a password or click **Generate** for a strong random one, optionally mark it a **Temporary password** (they must choose their own at first sign-in), and click **Set password**. The password is shown only here — share it over a safe channel.
- **Deactivate / Activate** — ends or restores sign-in access immediately, keeping the account's audit history intact. You can also select several accounts with their checkboxes and use the **Deactivate** button at the top.
- **Delete** (trash icon) — permanently deletes an eligible account and its chat history. The tooltip spells it out per person: “Permanently delete ... and their chat history”. Administrator-account actions that are unavailable under current service policy remain disabled with an explanation.

#### GOOD TO KNOW

**Password** reset and authenticator recovery are separate. The current **Users** screen has no authenticator-reset action. If someone has lost their authenticator and recovery codes, verify their identity and follow the organization's authorized recovery process; setting a temporary password alone does not clear MFA. Identity-provider recovery remains with the **SSO** provider.

#### GOOD TO KNOW

Administrative continuity rules are enforced by the service. When an account action would violate them, the console blocks the action and explains that it is restricted by administrative continuity policy.

## Approve access and finish sign-in

Take a real access request from the sign-in page to a working first chat, by temporary password or by organization SSO, and decline a request you do not recognize.

### Before you begin

- An Admin account. **Admin console** › **Users** shows the **Access requests** panel only while someone is waiting.
- A model granted to **Default Users** in **Model Access**, so the approved person has something to chat with.
- For the **SSO** path: single sign-on already set up on the **SSO** tab, and the person's account at your identity provider with a verified email on an allowed domain.
- A safe way to share a temporary password, such as in person or your password manager's sharing. Aperture Chat sends no email.

### Step by step

- 1 The person opens your workspace address, chooses **Request access**, enters First name, Last name, and Work email, and chooses Submit access request. They see Request received.
- 2 Choose **Admin console** at the bottom of the sidebar. **Users** opens with an **Access requests** panel that lists each person's name, email, and request time.
- 3 Check that you recognize the person and the email address.
- 4 In **Approve** as, keep User, then choose **Approve**. The status reads "... was approved as User."
- 5 In Finish sign-in setup, choose Set temporary password.
- 6 Choose **Generate** (or type at least 12 characters), keep **Temporary password** on, and choose **Set password**.
- 7 Choose **Copy**, share the password over your safe channel, then choose **Done** in the dialog and **Done** in Finish sign-in setup. The row now shows local under Auth.
- 8 The person chooses Email & password on the sign-in page, signs in with the temporary password, enters their own password twice on Set a new password, and chooses **Set password** and continue.
- 9 They land on the welcome card. Ask them to start a new chat, choose a model, and send a short message.
- 10 Confirm that a reply arrives. That proves the approval, the group, model access, and sign-in all work.



### Path: Organization SSO, no password (shown)

- 1 Make sure single sign-on is set up on the **SSO** tab and the person has an account at your identity provider with a verified email on an allowed domain.
- 2 **Approve** the request as User. In Finish sign-in setup, choose **Done**. Do not set a password.
- 3 Tell the person to open the workspace, enter their work email, and choose **Continue with SSO**.
- 4 They sign in at the identity provider and land on the welcome card.
- 5 **Users** shows the account Active, with sso under Auth.

### Path: Decline a request (shown)

- 1 In **Access requests**, choose **Decline** on the request you do not recognize. There is no confirmation step.
- 2 The status reads "... access request was declined." The pending account is deleted.
- 3 If that person tries to sign in, they see "Unknown local account." They can submit a new request later.

### Path: Approve as Temp User or Admin

- 1 Temp User gives access to the designated Luna model only and stops after 30,000 reported tokens. It needs an enabled Luna model.
- 2 Admin appears in **Approve** as only when service policy allows it: Policies > Administrator accounts reads Available.
- 3 Arrange sign-in the same way: a temporary password or organization **SSO**.

### Check that it worked

- ✓ The status reads "... was approved as User." and the request leaves the queue.
- ✓ **Users** shows the person Active, with local under Auth after a password is set, or sso for the **SSO** path.
- ✓ The person reaches the welcome card and gets a real reply to a first message.
- ✓ A declined person's sign-in attempt shows "Unknown local account."

## Troubleshooting

| WHAT YOU SEE                                                                                         | WHAT TO DO                                                                                                                                                                            |
|------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "Access was not approved. This access request is no longer pending."                                 | Another administrator already approved or declined it. Reload <b>Users</b> to see the current queue.                                                                                  |
| "Access was not approved. Enable a Luna model for this workspace before approving temporary access." | <b>Approve</b> as User instead, or ask your service team to enable a Luna model.                                                                                                      |
| The approved person sees "Unknown local account."                                                    | Approved accounts start as <b>SSO</b> accounts. Set a temporary password (Finish sign-in setup, or <b>Password</b> in the row's Actions), or have them use <b>Continue with SSO</b> . |
| "Invalid local credentials."                                                                         | The temporary password was mistyped. Set a new one; each new password replaces the last.                                                                                              |
| Set password stays unavailable in the password dialog                                                | Passwords need at least 12 characters. Choose <b>Generate</b> .                                                                                                                       |
| The person's model menu reads No models available                                                    | None of their groups carries a model. Grant one to <b>Default Users</b> in <b>Model Access</b> , or choose <b>Access</b> on their <b>Users</b> row to see why.                        |
| "Too many access requests for this email. Try again shortly."                                        | The sign-in page accepts three requests per email in a short window. Wait a minute, then submit once.                                                                                 |

### WALKTHROUGH · 2:36 · ADMIN CONSOLE > DOCUMENTATION

## Users and accounts

Create SSO and password accounts, change roles, reset passwords, deactivate and reactivate, work in bulk, filter by group, explain model access, and delete leavers.

### Before you begin

- An Admin account. You manage User, Temp User, Power User, Auditor, and **Agent Approver** accounts; admin rows are locked.
- A group to start people in. **Default Users** always exists.
- For password accounts: a safe way to share a temporary password.

### Step by step

- 1 Choose **Admin console** at the bottom of the sidebar. **Users** is the first tab.
- 2 Choose **Add User**. Enter Name and Email, then choose a Role and a Starting group.

- 3 Choose **Create Account**. The status reads "... was created through the admin API." The new row shows sso under Auth.
- 4 For someone who signs in with a password, choose **Password** in the row's Actions, then **Generate**, keep **Temporary password** on, and choose **Set password**. Auth changes to local.
- 5 To change a role, choose it in the row's Role menu. It saves immediately.
- 6 To pause access, choose **Deactivate** in the row's Actions. Choose **Activate** to restore it.
- 7 To delete a leaver, choose the trash button in the row's Actions. Deletion is immediate and permanent.

#### Path: An SSO account (shown)

- 1 **Add User** creates an account that signs in through organization **SSO**: Auth reads sso.
- 2 **Share** the workspace address; the person enters their email and chooses **Continue with SSO**.
- 3 Without **SSO** set up, the account cannot sign in until you set a temporary password.

#### Path: A password account (shown)

- 1 **Add User** as above, then choose **Password** in the row's Actions.
- 2 Choose **Generate**, keep **Temporary password** on, choose **Set password**, then **Copy**.
- 3 **Share** the password safely. The person chooses their own at first sign-in.

#### Path: Several accounts at once (shown)

- 1 Tick the checkbox at the start of each row.
- 2 Choose **Deactivate** at the top of the list. The status reads "Deactivated N users through the admin API."
- 3 Bulk deactivation also deletes each person's saved personalization memories. A row's own **Deactivate** button keeps them.
- 4 Reactivate people one at a time with **Activate**.

#### Path: Explain one person's model access (shown)

- 1 Choose **Access** in the person's row.
- 2 Each model shows **Usable**, **Allowed**, **provider offline**, or **Blocked**, with every gate the policy checked.
- 3 The trace explains; it changes nothing. Fix access in **Groups** or **Model Access**.

#### Check that it worked

- ✓ New rows show the role, group, and Auth you chose.

- ✓ A deactivated row reads Inactive; **Activate** returns it to Active.
- ✓ After deletion the row is gone and the status reads "... was permanently deleted."

## Troubleshooting

| WHAT YOU SEE                                                                            | WHAT TO DO                                                                                                                                                                                                      |
|-----------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "User was not created. User email already exists."                                      | Someone already has that address. Set Filter users by group to All groups and find them.                                                                                                                        |
| "Pick a starting group first — create one on the Groups tab if none exist yet."         | <b>Create</b> a group on the <b>Groups</b> tab, then add the user again.                                                                                                                                        |
| The Password button is unavailable                                                      | Passwords can be set only for active accounts other than yours and other administrators'. Reactivate the account first.                                                                                         |
| "Reactivate the account before setting its password."                                   | Choose <b>Activate</b> in the row's Actions, then set the password.                                                                                                                                             |
| "... was not deleted. Tenant admins can only delete regular users in their own tenant." | Administrator accounts are managed by your service team.                                                                                                                                                        |
| "This action is blocked by administrative continuity policy."                           | The change would leave the organization without an administrator. Ask your service team.                                                                                                                        |
| A password reset did not reset two-step verification                                    | Correct: <b>Password</b> does not reset the authenticator, and this screen has no authenticator reset. Verify the person's identity through your organization's recovery process and contact your service team. |

## 29. Groups

Groups carry permissions and model access — everything flows through them.

**Groups** are how access flows in the tenant: model grants, knowledge access controls, and permissions all attach to groups, not to individual people. Someone in no group has no group-based access, and a group that **SSO** group mapping manages follows the identity provider at every sign-in.

Runtime access and authoring permissions

| TOGGLE                    | WHAT IT GRANTS                                                                                              |
|---------------------------|-------------------------------------------------------------------------------------------------------------|
| Can use chat              | Start conversations and use assigned models.                                                                |
| Can use knowledge         | Query approved knowledge bases.                                                                             |
| Can use agents            | Run approved agent workspaces.                                                                              |
| Can use tools             | Invoke enabled tools and MCP actions.                                                                       |
| Can use API               | <b>Create</b> personal keys for approved models when service policy allows API access.                      |
| Can use Hermes companion  | Build and run agent profiles with the <b>Hermes</b> learning companion. Off until approved.                 |
| Can build agents          | <b>Create</b> private agent profiles when service policy permits. Organization publishing stays admin-only. |
| Can build knowledge bases | <b>Create</b> private knowledge bases. Group sharing and organization management stay admin-only.           |
| Can build tools           | <b>Create</b> private tools. Group sharing, stdio commands, and organization management stay admin-only.    |
| Can use memory            | Allow personal preferences to be learned and reused within the workspace memory policy.                     |

**CAREFUL**

Deleting a group removes its members' access that flowed through it. Check what the group grants before deleting.

**WALKTHROUGH · 2:05 · ADMIN CONSOLE > DOCUMENTATION**

### Groups and permissions

Create a group, add people one at a time or by pasting emails, set each permission, and know how the protected Default Users group behaves.

---

#### Before you begin

- An Admin account.
- The people's accounts already exist. **Import** adds existing accounts only.
- **Can use API** stays locked while service policy keeps downstream API access off.

## Step by step

- 1 Choose **Admin console** › **Groups**, then **Add Group**.
- 2 Enter a Group name and choose **Create** group. Leave Optional **SSO** group ID empty unless you use it.
- 3 Choose Manage on the new group's card.
- 4 On the group's **Users** tab, switch on each person you want in the group.
- 5 On the **Import** tab, paste email addresses (one per line, or separated by commas) and choose **Add users to group**.
- 6 On the **Permissions** tab, switch each permission on or off. Each switch saves immediately.
- 7 Open **Users** and choose the group in Filter users by group to confirm its members.

### Path: Add people one at a time (shown)

- 1 Open the group with Manage, then its **Users** tab.
- 2 Switch on Add ... to the group for each person. The pill changes to In group.

### Path: Import by email (shown)

- 1 Open the group's **Import** tab and paste the addresses.
- 2 Choose **Add users to group**. Existing accounts are added; unknown addresses and administrator accounts stay in the box.

### Path: Let SSO manage membership

- 1 On the **SSO** tab, map an identity-provider group value to this group (see Tenant **SSO** and provisioning).
- 2 Membership then follows each person's next **SSO** sign-in.

### Path: Remove a group

- 1 Choose the trash button on its card, or tick several cards and choose Remove selected. There is no confirmation step.
- 2 Members lose the access that came through that group. **Default Users** cannot be removed.

### Check that it worked

- ✓ The group's card shows the member count.
- ✓ **Users** filtered to the group lists exactly its members.
- ✓ Each permission you changed shows its new switch position, and the status reads "... group permissions synced with the admin API."

## Troubleshooting

| WHAT YOU SEE                                                                               | WHAT TO DO                                                                                                                                                      |
|--------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "Group was not created. Group name already exists."                                        | Choose a different name, or manage the existing group.                                                                                                          |
| "Group name is required before creating a group."                                          | Enter a Group name before choosing <b>Create</b> group.                                                                                                         |
| "No users were added to ... . No eligible platform users were found for ...."              | Those addresses have no account, or belong to administrators. <b>Create</b> the accounts on <b>Users</b> first; the addresses stay in the box so you can retry. |
| Can use API reads "Downstream API access is unavailable under the current service policy." | Service policy controls downstream API access. Ask your service team.                                                                                           |
| "Default Users is protected by organization policy and cannot be removed."                 | <b>Default Users</b> is the baseline group. Adjust its permissions and models instead.                                                                          |
| Can build agents is on, but people still cannot build agents                               | It also needs service policy to allow user-built agents. Policies > Default users can build agents shows whether it does.                                       |

## 30. Model Access

Decide which synced models users can see, and which groups can use them.

**Model Access** starts from the catalog available to your organization. You decide what portion of it your users actually see.

### The catalog at a glance

- Click **Sync models** to refresh the catalog whenever service availability changes.
- Read the counters: the All, Enabled, and Disabled status counters carry live totals alongside how many groups are available for scoping.
- Use the search box and the funnel filters in the column headers to cut through the catalog: check off providers, check off model labs, or type text to match runtime routes, and Clear filter resets one. The table has seven columns: **Model**, **Provider**, **User Access**, **Groups**, **Filters**, **Knowledge**, and **Tools**. Everything starts hidden until you decide otherwise.
- Flip the **User Access** toggle on a row to make that model visible, and use **Choose groups** to narrow it to specific teams.

- Use the **Filters** column to attach per-model content filters, and the **Knowledge** and **Tools** columns to scope what each model may search and call.

#### GOOD TO KNOW

Because model access flows through groups, a tenant with no groups yet sees the guard up front — “**Create a group** before enabling models.” — and the access switches stay off until one exists.

#### GOOD TO KNOW

Newly synced models arrive disabled until service availability and organization policy permit them. If a model you expect is missing, it is not currently available to this organization.

#### WALKTHROUGH · 1:51 · ADMIN CONSOLE › DOCUMENTATION

### Tenant model access

Grant a model to the right groups, check the effect for one person in the access trace, and hide a model when you need to.

#### Before you begin

- An Admin account and at least one group.
- **Models** your service team has made available to your organization.
- **Access** is granted to groups only. There is no per-person grant or deny.

#### Step by step

- 1 Choose **Admin console** › **Model Access**, then **Sync models**.
- 2 Find the model with **Search** model names or the funnel filters in the column headings.
- 3 In the model's **Groups** column, choose the groups button (**Choose groups**, or N groups).
- 4 Tick each group that should use the model (Allow ... for ...). Each tick saves immediately.
- 5 On **Users**, choose **Access** for someone in that group. The model reads **Usable**, with group\_grant passed.

#### Path: Grant or remove a group (shown)

- 1 Open the model's groups and tick or untick a group.
- 2 Members of a ticked group can use the model; the trace for someone outside it stops at group\_grant.

#### Path: Hide a model from everyone (shown)

- 1 Turn off the model's **User Access** switch. The row reads Hidden from users.



- 2 Today only your service team can turn it back on: the switch is refused with "This model is not available to this organization." The same happens if you untick its last group.
- 3 To narrow access without hiding it, keep at least one group ticked.

#### Path: Give one person a model

- 1 **Create** a group for them (or use an existing one) and tick it on the model.
- 2 Or approve their request on **Model Access** (see Review model requests and explain access).

#### Check that it worked

- ✓ The **User Access** column reads Visible to N groups.
- ✓ The access trace for a member shows the model **Usable**, or **Allowed, provider offline** if only the provider is down.
- ✓ The person sees the model in their model menu.

#### Troubleshooting

| WHAT YOU SEE                                                                                             | WHAT TO DO                                                                                                                                                                          |
|----------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "... user access could not sync. This model is not available to this organization. Nothing was changed." | The model has no groups from your organization left, so it can no longer be changed here. Ask your service team to restore its groups. Keep one group ticked when narrowing access. |
| "Create a group before enabling models."                                                                 | <b>Create</b> a group on the <b>Groups</b> tab first. <b>Users</b> get models only through groups.                                                                                  |
| The trace reads <b>Allowed, provider offline</b>                                                         | The grant works, but the provider behind the model is not connected. Your service team reconnects it.                                                                               |
| "Model is disabled by platform policy."                                                                  | The model is turned off for the whole service. Ask your service team.                                                                                                               |
| A model is missing from the catalog                                                                      | Choose <b>Sync models</b> . If it is still missing, it has not been made available to your organization. Ask your service team.                                                     |

## 31. Resolve model requests and trace access

Review requests through groups and diagnose each user's actual model access.

#### CAREFUL

Approval changes group membership. When the chosen group does not carry the model yet, it can also widen that model's availability for all group members. Review the group's purpose and members before approving.

## GOOD TO KNOW

The trace is read-only and scoped to users you can administer. It explains existing permissions; opening it does not grant access, change a role, or bypass the organization's model ceiling.

## WALKTHROUGH · 1:37 · ADMIN CONSOLE › DOCUMENTATION

### Review model requests and explain access

Approve a person's model request into the right group, decline another, and confirm both from the person's view and the access trace.

#### Before you begin

- An Admin account and the groups you grant through.
- **Model** requests allowed by service policy. When they are not, people see only the models they can already use and cannot send requests.

#### Step by step

- 1 The person opens the model menu, chooses **Why isn't a model listed?**, and chooses **Request access** on a **Locked** model. They see "Request sent for .... An administrator will review it."
- 2 Choose **Admin console › Model Access. Access requests** lists the person, the model, and the server's reason.
- 3 In **Grant through group**, choose the narrowest group that should carry the model. A group marked (also grant model to group) does not carry it yet; approving grants it to every member.
- 4 Choose **Approve**. The panel reads "... can now use ... through ...."
- 5 Choose **Decline** for requests you will not grant. The panel reads "Declined the request from ...."
- 6 Ask the person to reopen **Why isn't a model listed?**, and check **Users › Access** for them.

#### Path: Approve through a group that already has the model

- 1 Choose a group listed without a suffix. Approving adds the person to it; nobody else gains access.

#### Path: Approve and widen a group (shown)

- 1 Choose a group marked (also grant model to group).
- 2 Read the hint: approving grants the model to everyone in the group. Check its members first.

#### Path: Decline (shown)

- 1 Choose **Decline**. The request leaves the queue, and the person's model reads **Locked** again with **Request access** available.

## Check that it worked

- ✓ The panel reads "... can now use ... through ...." and the request leaves the queue.
- ✓ The person's **Models in your organization** dialog no longer shows the model as **Locked**.
- ✓ **Users > Access** for the person shows group\_grant passed for that model.

## Troubleshooting

| WHAT YOU SEE                                                                          | WHAT TO DO                                                                                                              |
|---------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------|
| "Choose the group that should carry this model before approving."                     | Choose a group in <b>Grant through group</b> . If one is already shown, choose another group and then the one you want. |
| "Approved, but the server still reports: ..."                                         | The approval worked but another gate still blocks the model, such as an offline provider. The trace shows which.        |
| "This request has already been resolved."                                             | Another administrator approved or declined it. Choose <b>Refresh</b> .                                                  |
| <b>Approve is refused because the model is disabled for the whole service</b>         | Ask your service team to enable the model, then approve again.                                                          |
| The person sees " <b>Model access requests are disabled by organization policy.</b> " | Service policy does not allow model requests. Grant access through a group instead.                                     |
| The person sees " <b>A request for this model is already pending.</b> "               | Their earlier request is still in your queue.                                                                           |

## 32. Connections: response actions and connector handoff

Manage response actions and route shared connector setup to the service team.

The **Admin console's Connections** tab contains **Chat output actions**. Shared connector switches, saved credentials, connection tests, and workspace OAuth are managed by the service team outside tenant administration.

### GOOD TO KNOW

**Users** still connect their own Google, Microsoft, Box, or iManage account from the composer's attach menu. Those delegated connections respect each user's source permissions. Shared credentials managed by the service team support connector setup and background knowledge sync; they do not grant users someone else's files.

## Chat output actions

**Chat output actions** adds admin-approved buttons to assistant responses for export, formatting, or handoff. Choose New response action to create one; existing custom actions offer Edit and **Delete**. Each

row shows Enabled or **Draft** and an enable switch. Creating these actions does not configure a shared source connector.

MCP connections and model-callable tools remain in **Library** → **Tools** → **Connections**. **Prompts** and **Skills** also remain in the **Tools** library. Their authoring and use follow the existing permissions and the shared connector availability set by the service team.

#### WALKTHROUGH · 2:30 · ADMIN CONSOLE › DOCUMENTATION

### Response actions and connector responsibilities

Build a script response action, test it, give it to a group, use it on a real reply, edit or draft it, and find where connections are signed in.

#### Before you begin

- An Admin account.
- A short Python script: it reads the reply on standard input and prints the result.
- The group whose members should see the button.

#### Step by step

- 1 Choose **Admin console** › **Connections**. **Chat output actions** lists every tool; response actions carry a Response action pill.
- 2 Choose New response action. Enter Action name, Timeout (seconds, 1–30), Description, and the Python script.
- 3 Under Who can run this action, tick the groups that should see it. With none ticked, everyone can.
- 4 Under Test run, paste sample text and choose Run test. Read "Finished in ... ms." and the output.
- 5 Choose **Create** action. The status reads "... saved and available to models and chat." and the row reads Enabled.
- 6 A member opens any assistant reply and chooses the action's button below it. The result opens with **Copy** and **Download**.
- 7 Choose Edit to change it, or switch it off to set it to **Draft** and hide it from replies.

#### Path: Script response action (shown)

- 1 The script receives the reply text on standard input and prints its result.
- 2 To return files, write them to the folder in APERTURE\_ARTIFACT\_DIR: up to 8 files, 50 MB each and 75 MB per run. Links appear in the output.
- 3 Scripts run isolated, without platform secrets or provider keys, but can reach the network.

#### Path: Sign a connection in (OAuth)

- 1 Open **Library** › **Tools** › **Connections**, choose Edit on the connection, then Sign-in.

- 2 Register the Redirect URL at the provider, then choose Connect with provider and approve there.
- 3 The badge changes from Not signed in to Signed in — token saved.

### Path: Remove an action

- 1 Choose the trash button on its row. It is deleted for the whole workspace at once, with no confirmation.

### Check that it worked

- ✓ Run test shows "Finished in ... ms." with the expected output.
- ✓ Members of the chosen group see the action's button on assistant replies, and it returns the result.
- ✓ After switching to **Draft**, the row reads **Draft** and the button no longer appears on replies.

### Troubleshooting

| WHAT YOU SEE                                                               | WHAT TO DO                                                                                                                                                      |
|----------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "Script has a syntax error on line ..."                                    | Fix the line named in the message and run the test again.                                                                                                       |
| "The script timed out."                                                    | Raise Timeout (up to 30 seconds) or simplify the script.                                                                                                        |
| "The script exited with an error (code 1)."                                | Read the error output under the result and fix the script.                                                                                                      |
| "Script cannot be empty. It reads input from stdin and prints its result." | Enter a script before saving.                                                                                                                                   |
| A member does not see the button                                           | Check that the action is Enabled and that one of their groups is ticked under Who can run this action.                                                          |
| A sign-in page reads "Sign-in link expired"                                | Start sign-in again from the connection's Sign-in tab with Connect with provider. The Authorize shortcut on <b>Connections</b> does not complete sign-in today. |

## 33. Automations: preparing scheduled runs

What admins need to configure before users can rely on recurring model chains.

**Users** create and run automations from the Automations tab of the **Agents** view, and enabled schedules fire automatically in the background (UTC), delivering each run's output as a new chat thread. Admins manage group and model access; the service team manages shared connector availability and credentials outside tenant administration. Check the actual failure details to determine which setup needs attention.

## Admin readiness checklist

- 1 Open **Model Access** and confirm the automation's model is visible to the user's group.
- 2 Open **Groups** and confirm the user belongs to the group that receives the model grant.
- 3 Ask the service team to confirm any required shared connector and inspect its real connection-test result. **Users** complete their own source-account connection when the workflow needs delegated access.
- 4 If the automation uses knowledge, confirm the knowledge base is enabled and its source ACL allows the user's group.
- 5 Ask the user to press **Run now** after saving the automation; the transcript and last-run status will show whether the setup is complete.

### GOOD TO KNOW

Admins do not need to expose every model to make automations work. Grant the smallest approved model set that fits the scheduled workflow, then expand only when a run proves it needs more capability.

## 34. SSO: single sign-on for your tenant

Connect an identity provider, provision on first sign-in, and map IdP groups.

The **SSO** tab lists every identity provider for your tenant as a card. When **Policy Controls** shows **SSO** configuration as available, you can add providers for your domains, test them, map identity-provider groups to tenant groups, and enforce **SSO**. Otherwise the tab is read-only and says so.

### WALKTHROUGH · 4:54 · ADMIN CONSOLE > DOCUMENTATION

#### Tenant SSO and provisioning

Set up your organization's single sign-on from the Admin console: check that policy allows it, create and test the connection, sign a test person in, map groups, enforce it, and recognize the read-only state.

#### Before you begin

- An Admin account, with **SSO** management allowed by service policy: Policies > **SSO** configuration reads Available. Otherwise the **SSO** tab is read-only; ask your service team.
- An OpenID Connect application at your identity provider, with its issuer URL, client ID, and client secret. The paths below list the exact settings for Keycloak, Microsoft Entra ID, Okta, and Google Workspace.
- The redirect URI shown under the **Add SSO configuration** form, registered exactly at the identity provider.
- A groups claim in the ID token if you will map groups. Google sends none.
- A test account at the provider with a verified email on one of your domains.

## Step by step

- 1 Open **Admin console** › Policies › **Policy Controls** and check that **SSO** configuration reads Available. If it reads Read only, ask your service team to allow **SSO** management.
- 2 Choose **Admin console** › **SSO**, then **Add SSO configuration**.
- 3 In **Identity provider**, choose your preset: Microsoft Entra ID, Google Workspace, Okta, or Custom OIDC provider.
- 4 Enter the Issuer URL, **Client ID**, and Allowed email domains. Keep Provision new users on first sign-in (JIT) on.
- 5 Register the redirect URI shown under the form at your identity provider.
- 6 Paste the **Client secret** last, then choose **Create SSO** configuration. Enforcement starts off.
- 7 On the new card, choose **Test connection**. Discovery document, Signing keys (JWKS), and Sign-in readiness must all pass.
- 8 In a private window, enter the test person's email, choose **Continue with SSO**, and sign in at the provider. **Users** then lists them with sso under Auth.
- 9 On the card, choose Add mapping for each identity-provider group value, choose the tenant group, and choose Save mappings.
- 10 Have the test person sign in again, then check **Users** with Filter users by group.
- 11 When everyone can sign in through the provider, turn on Enforce for tenant sign-in.
- 12 Confirm that a local password on the domain is now refused.

### Path: Custom OIDC provider, such as Keycloak (shown)

- 1 Choose Custom OIDC provider. Enter the issuer, such as https://your-keycloak-host/realms/your-realm.
- 2 At Keycloak: a confidential client with Client authentication on, Standard flow, Require PKCE S256, the redirect URI, and a Group Membership mapper named groups with Full group path off.
- 3 Map group values exactly as they appear in the token, such as litigation.

### Path: Microsoft Entra ID

- 1 Choose Microsoft Entra ID. The issuer is prefilled; replace `{tenant-id}` with your Directory (tenant) ID. Never use common or organizations.
- 2 Use the client secret's Value, not its Secret ID.
- 3 At Entra: App registrations › New registration (**Web**, the redirect URI), Certificates & secrets › New client secret, and Token configuration › Add groups claim. Group claims carry object IDs, so map each group's object ID.



### Path: Okta

- 1 Choose Okta. The preset fills `https://{your-domain}.okta.com`; replace it with your Okta domain.
- 2 At Okta: Applications › **Create** App Integration › OIDC, **Web** Application, Authorization **Code**, the redirect URI. For group mapping, use the custom authorization server, `https://{your-org}.okta.com/oauth2/default`, and add a groups claim to the ID token.

### Path: Google Workspace

- 1 Choose Google Workspace. The issuer is `https://accounts.google.com`.
- 2 Google ID tokens carry no groups, so manage group membership on the **Groups** tab instead of mapping.

### Path: When SSO management is read-only (shown)

- 1 The **SSO** tab shows "Organization policy makes **SSO** configuration read-only in this console."
- 2 Existing cards stay visible and keep working, but you cannot add, test, map, enforce, or remove. Ask your service team.

### Check that it worked

- ✓ **Test connection** lists Discovery document, Signing keys (JWKS), and Sign-in readiness as passed.
- ✓ The test person reaches the workspace, and **Users** lists them with the USER role and sso under Auth.
- ✓ After their next sign-in, **Users** filtered to the mapped group lists them.
- ✓ The card reads enforced, and a local password on the domain is refused with "**SSO** is enforced for this email domain; local sign-in is disabled. Use the configured identity provider."



## Troubleshooting

| WHAT YOU SEE                                                                                              | WHAT TO DO                                                                                                                                                                    |
|-----------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "Organization policy makes SSO configuration read-only in this console."                                  | Service policy does not let administrators manage <b>SSO</b> . Policies > <b>SSO</b> configuration reads Read only. Ask your service team.                                    |
| Create SSO configuration stays unavailable                                                                | Issuer URL and <b>Client ID</b> are required.                                                                                                                                 |
| "Issuer verified, but sign-in stays disabled until you add: client secret."                               | The secret was not saved, or allowed domains are missing. Cards cannot be edited: choose Remove and create the configuration again with every field.                          |
| Test connection fails at Discovery document                                                               | The issuer URL is wrong or unreachable from the server. Open the issuer followed by /.well-known/openid-configuration; for Entra use your tenant ID.                          |
| The card shows Custom Oidc instead of the display name you typed                                          | A known limitation: the card is labeled by provider type, and the display name is not kept yet. Identify cards by their issuer.                                               |
| Mapped groups do not appear after sign-in                                                                 | The claim value must match exactly, including case and any leading slash, and the provider must add the groups claim to the ID token. Membership updates at the next sign-in. |
| "... is outside the domains allowed for this SSO provider."                                               | The provider vouched for an address on another domain. Add the domain by recreating the configuration, or keep that person out.                                               |
| "SSO is enforced for this email domain; local sign-in is disabled. Use the configured identity provider." | Expected after enforcement. The person chooses <b>Continue with SSO</b> . Turn enforcement off on the card to allow passwords again.                                          |

### GOOD TO KNOW

The steps inside each identity provider's console (Microsoft Entra ID, Okta, Google Workspace, Keycloak) are spelled out step by step in the **Identity provider** setup topics of the interactive guide at <https://aperturechat.com/guide.html>. The redirect URI to register is the one shown in the **Add SSO configuration** form.

### CAREFUL

Enforcement blocks password sign-in for the allowed domains. A passing **Test connection** checks discovery and keys; it does not prove the client-secret exchange, callback, user session, or group mapping. Validate those with a real sign-in before you enforce.

## 35. Analytics

Runtime metadata, feedback, per-user usage, and token budgets for this tenant.

The **Analytics** tab measures how the organization actually uses the platform. Sections start collapsed behind descriptive headers — click a header to expand one — and each section carries its own filter row pairing a person picker (defaulting to **All admins and users**) with a date range and its presets (All, Today, Week, 30 days). Scoping one section never narrows another.

- **Runtime Clock Metadata** — organization-scoped execution timestamps captured from admin and user **Chat** and **Draft** completion events. The cards total runtime events, chat completions, and draft calls. Its **CSV** includes actor\_id and actor\_name columns, and when one user is selected the filename gains that user's suffix so exports stay unambiguous.
- **Chat Feedback Analytics** — response-level thumbs up and thumbs down signals submitted by tenant admins and users, with totals for overall, positive, and negative feedback.
- **Model Activity** — saved prompt volume by model, date, and person for the selected range, drawn as “**Prompts** by model”, “Prompt trend”, and “**Users** by prompt activity”.
- **User Usage** — durable per-user usage for this organization's admins and users across chat, drafts, agents, automations, and the API gateway. Token counts are provider-reported only and stay blank when a provider reported none. The panel has its own **Filter usage by user** selector for drilling into one person.
- **Workspace Usage Budget** — a read-only view of the organization's service-managed ceiling and its current UTC accounting period. This panel shows where usage stands without exposing service-level configuration.
- **Token Allocations** — per-user and per-group token caps beneath the workspace ceiling, with Per day, Per week, or Per month UTC resets. The most restrictive applicable cap wins. Choose a user or group, enter a cap, pick the period, and click **Set cap**. Until you add one, the table says “No allocations yet. Everyone shares the workspace ceiling.”

### GOOD TO KNOW

Every number here comes from saved audit and usage events, and each **CSV** export button opens its own small date-range popover so a file contains exactly the rows you chose. If a panel has nothing to show, it says so plainly instead of showing sample data.

### WALKTHROUGH · 2:03 · ADMIN CONSOLE › DOCUMENTATION

#### Tenant analytics

Scope each analytics section to a person and dates, export CSV, read usage and the workspace budget, and set a token cap that people actually hit.

#### Before you begin

- An Admin account.

- Some real activity: chats, drafts, or automations from people in your organization.
- The workspace ceiling comes from your service team; your caps sit inside it.

## Step by step

- 1 Choose **Admin console** › **Analytics**. Every section starts collapsed.
- 2 Expand **Runtime Clock Metadata**. Choose a person and dates, or All, Today, Week, or 30 days.
- 3 Choose **CSV**, pick dates or All dates, and choose **Download** N rows.
- 4 Expand **Model Activity**, **User Usage**, and **Workspace Usage Budget** to read volume, tokens, and the ceiling.
- 5 Expand **Token Allocations**. Choose a person or group in the first menu, enter the Token cap, choose Per day, Per week, or Per month, and choose **Set cap**.
- 6 The row shows the cap and a meter of tokens used in the current UTC period.
- 7 Choose Remove allocation for ... to lift the cap.

## Path: Cap one person (shown)

- 1 Choose User · their name, enter the cap, choose the period, and **Set cap**.
- 2 **Once** they pass it, their next message reads "Your daily token budget has been reached. Requests are blocked until the next UTC day."

## Path: Cap a group

- 1 Choose Group · its name instead. Members then see "Your group's daily token budget has been reached. Requests are blocked until the next UTC day."
- 2 When several caps apply, the most restrictive one wins.

## Path: Export other sections

- 1 **Chat** Feedback and **User Usage** have their own **CSV** buttons with the same date choice.
- 2 Choosing one person in a section's filter adds them to the file name.

## Check that it worked

- ✓ Each filter shows "N of M records" for its own section only.
- ✓ The **CSV** downloads with the section and person in its file name, such as aperture-admin-runtime-analytics-....
- ✓ The cap row shows used and capped tokens, and a capped person's next message is refused with the budget message.

## Troubleshooting

| WHAT YOU SEE                                                                | WHAT TO DO                                                                                                    |
|-----------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------|
| "Enter a whole number of tokens (0 means no cap)."                          | Type digits only. 0 records the person or group with no cap.                                                  |
| "Unknown user for this workspace." or "Unknown group for this workspace."   | The person or group was removed. Choose <b>Refresh</b> and pick again.                                        |
| "Allocations total ... tokens, more than the ... -token workspace ceiling." | Allowed, but the ceiling still wins. Lower the caps or ask your service team to raise the ceiling.            |
| The capped person waits several minutes before seeing the refusal           | The chat retries a refused request before showing the message. The refusal itself is immediate on the server. |
| Token columns are blank for some rows                                       | Tokens are counted only when the provider reports them; blank means none was reported.                        |

## 36. Policies: service availability and memory governance

Apply downstream tenant defaults, configure memory, and govern by count without reading content.

The Policies tab is always present. **Policy Controls** starts collapsed. When service policy enables memory, Personalization Memory and Memory by User also appear collapsed. Otherwise, Memory governance explains the restriction; saved organization settings remain intact. Expand the panel you need; administrators can narrow available capabilities, while unavailable settings remain locked.

### Policy Controls: read service availability

- Administrator accounts — shows whether this console can create and manage administrators or whether new administrator accounts require service approval.
- Admin sign-in policy — shows whether admins must use **SSO** or may use an explicitly provisioned local account.
- **SSO** configuration — shows whether tenant **SSO** mappings are delegated or read-only.
- New model defaults — shows whether newly available models begin with **Default Users** or require explicit grants.

Below those status rows are defaults for the protected **Default Users** group: personal API keys, private-agent building, private knowledge-base and tool authoring, and personalization memory. A switch locks when its capability is unavailable under service policy; the saved group grant is preserved rather than silently erased. Use **Groups** for exceptions and **Model Access** for available models. Shared connector availability is managed by the service team outside tenant administration.

## Personalization Memory settings

- When service policy permits memory, expand Personalization Memory and turn on Memory for this organization. This makes the account-level Personalization memory row available to eligible users; existing memories stay saved but are not applied while this switch is off.
- Choose whether to allow Learn from conversations automatically. When off, only explicit requests such as “remember that ...” and direct additions in the memory manager create memories. Every user can still opt out of automatic learning individually.
- Set Retention (days), from 1 through 3650. Older memories retire automatically when they pass this policy.
- Set Maximum memories per user, from 1 through 2000. The default is 200, which is a practical general-purpose limit; raise or lower it only when your retention, compliance, or workload policy calls for a different capacity. When the cap is passed, the least useful unpinned memories retire first.
- Expand Memory by User to review content-free counts. Click **Refresh** to reload the counts, or purge one person's memories when compliance or account cleanup requires it.

### CAREFUL

Memory administration never grants reading access. Administrators see policy, counts, and purge controls only; the API and the interface do not return another person's memory content.

## WALKTHROUGH · 2:29 · ADMIN CONSOLE · DOCUMENTATION

### Policies and memory governance

Read what service policy allows, set the Default Users capabilities, tune personalization memory, purge a person's memories without reading them, and recognize when memory is unavailable.

#### Before you begin

- An Admin account.
- Service policy sets the ceiling: what is locked here is decided by your service team.

#### Step by step

- 1 Choose **Admin console** > Policies. Its sections start collapsed; expand **Policy Controls**.
- 2 Read the four status rows: Administrator accounts, Admin sign-in policy, **SSO** configuration, and New model defaults.
- 3 Switch on the Default users capabilities you want: downstream API, build agents, build knowledge bases, build tools, and use memory. Each saves with "Default user policy saved."
- 4 Confirm the change in **Groups** > **Default Users** > **Permissions**.
- 5 Expand Personalization Memory. Set Memory for this organization and Learn from conversations automatically.

- 6 Type Retention (days), 1 to 3650, or Maximum memories per user, 1 to 2000, and press Enter. The status reads "Memory policy saved."
- 7 Expand Memory by User and choose **Refresh**. To delete one person's memories, choose Purge, then Yes, purge.

#### Path: What each status row means

- 1 Administrator accounts: Available lets you create admins; Service managed means your service team does.
- 2 Admin sign-in policy: **SSO** required or Local allowed for administrator accounts.
- 3 **SSO** configuration: Available lets you manage **SSO** on the **SSO** tab; Read only means your service team manages it.
- 4 New model defaults: **Default Users** means newly available models start granted to **Default Users**; Explicit grants means you grant each one.

#### Path: What each Default users switch changes

- 1 Downstream API: people can create personal API keys for their approved models.
- 2 Build agents, knowledge bases, or tools: people can create their own private ones. Sharing and publishing stay with administrators.
- 3 Use memory: the assistant saves and recalls each person's private preferences.

#### Path: When memory is off for the platform (shown)

- 1 Policies shows a single Memory governance panel: "Personalization memory is unavailable under the current service policy. Saved organization settings remain intact."
- 2 Your saved settings return when service policy allows memory again.

#### Check that it worked

- ✓ The status reads "Default user policy saved." and the matching switch is on in **Groups › Default Users › Permissions**.
- ✓ "Memory policy saved." and the new value remains after reopening Policies.
- ✓ After a purge the status reads "Deleted N memories for ...." and the person leaves Memory by User.

## Troubleshooting

| WHAT YOU SEE                                                                                           | WHAT TO DO                                                                                                   |
|--------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------|
| "Unavailable under the current service policy. The saved group grant is preserved."                    | Service policy has that capability off. Your earlier grant returns when it is turned back on.                |
| "Turn on Memory for this organization below before granting the default group access."                 | Turn on Memory for this organization in Personalization Memory first.                                        |
| A memory number did not change                                                                         | Values outside 1 to 3650 days or 1 to 2000 memories are ignored. Type a value in range and press Enter.      |
| "Memory policy was not saved. Personalization memory is unavailable under the current service policy." | Service policy has personalization memory off. Ask your service team.                                        |
| "Default user policy could not be saved."                                                              | Nothing changed. Reload Policies and try again; if it repeats, check the <b>Audit</b> trail for the refusal. |

## 37. Data retention and conversation tags

Find tagged conversations, inspect their contents, and review batch actions.

### WALKTHROUGH · 2:16 · ADMIN CONSOLE · DOCUMENTATION

#### Data retention and tagging

Keep chats forever by default, label conversations by client or matter, place a legal hold, and preview a finite schedule before anything is deleted.

#### Before you begin

- An Admin account.
- Your organization's retention rules, and the client or matter names and aliases people use.
- Saved chats to label. Labels come only from saved message text.

#### Step by step

- 1 Choose **Admin console** > **Audit** and expand Data Retention.
- 2 Schedule and rules: Keep chats for starts at Forever, with automatic deletion off.
- 3 Under Clients, matters, and regulated records, choose a Source type, enter Name or reference and aliases, and choose Add source to policy.
- 4 Choose Save Forever. The status reads "Forever saved. Automatic deletion is off for all chats."



- 5 Choose Tags and holds, then Scan existing chats. **Search** for the matter and review each suggestion.
- 6 Select the right chats, choose the label, and choose Confirm label.
- 7 Expand Legal holds, enter a Hold name, and choose Hold selected chats.
- 8 To consider a time limit, choose it in Schedule and rules and choose Preview effect. Save only after reviewing the counts.

#### Path: Keep forever (default, shown)

- 1 Leave Keep chats for at Forever. Choose Save Forever after adding sources; nothing is deleted.

#### Path: A finite schedule (previewed, shown)

- 1 Choose 1, 5, 7, or 10 years. Count age from and Review window (at least 7 days) appear.
- 2 Choose Preview effect and read how many chats are old enough, on hold, or kept.
- 3 Choose Save retention policy only when the counts are right. Eligible chats enter the review window first.

#### Path: A rule for one label

- 1 Each source has a Retention rule menu. A longer matching rule wins over the default schedule.
- 2 Tick Apply time limits only to labels with a rule to leave unlabeled chats at Forever.

#### Path: Manual archive or delete

- 1 In Tags and holds, select chats and choose **Archive** selected or **Delete** selected.
- 2 **Delete** asks first and cannot be undone. **Chats** under an active legal hold are skipped.

#### Check that it worked

- ✓ The source appears in the list with its aliases, and Save Forever confirms automatic deletion is off.
- ✓ The confirmed chats show the matter label; held chats show Legal hold.
- ✓ Preview effect reports the totals without changing anything.



## Troubleshooting

| WHAT YOU SEE                                                                           | WHAT TO DO                                                                                                  |
|----------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|
| "Enter a source name of at least two characters."                                      | Enter the client or matter name.                                                                            |
| "Aliases must be 3–160 characters long."                                               | Remove very short aliases, such as initials.                                                                |
| "Preview this exact policy again before saving; the policy or affected chats changed." | Choose Preview effect again, then save.                                                                     |
| Scan suggests nothing for a matter                                                     | Suggestions come from saved message text only. Add the aliases people actually write, save, and scan again. |
| Hold selected chats is unavailable                                                     | Select 1 to 500 chats and enter a Hold name of at least three characters.                                   |

## How retention behaves, in detail

- Open **Audit** and expand Data Retention, then choose Schedule and rules. Forever is the default: automatic deletion is off until an administrator previews and saves a schedule. Choose 1, 5, 7, or 10 years, or keep Forever. One year means 365 days.
- Choose whether age starts at chat creation or the last message change. Renaming or archiving a chat does not restart its clock. Select a review window of at least 7 days. Saving a changed policy restarts that window for eligible chats.
- Add a stable client, matter, or regulated-record source and its aliases. A source can have a longer retention rule or Forever. Choose Apply time limits only to labels with a rule to leave all other chats stored indefinitely. The longest applicable duration wins; a matching Forever rule or active legal hold prevents automatic deletion.
- Save the source definitions, then use Scan existing chats in Data Retention → Tags and holds. New message saves also look for source mentions. Detected mentions are suggestions only: select the correct chats and source, then Confirm label. Remove / dismiss label rejects a false match. The label records who confirmed it and when.
- Optional sensitive-data suggestions check saved message text for email addresses, possible Social **Security** numbers, and payment cards. They do not detect every form of PII or regulated record, or inspect original uploaded files. Raw matched values are never copied into labels. Review before confirming a category.
- Choose Preview effect to see counts across all saved chats, including legal holds and chats without a deadline. Saving the policy recalculates existing and future chats. A shorter duration can make old chats eligible, but they still receive the review window. A longer duration delays deletion; it cannot restore deleted content. Saving Forever stops automatic deletion.

- To inspect chats, open **Audit**, expand Data Retention, and choose Tags and holds. This list includes tagged and untagged conversations available in your administrative scope.
- Use **Search** chats and tags to find titles, people, tags, or client/matter identifiers. Filter by tag type when you need a narrower set.
- Click a conversation title to open Tagged conversation and review its saved prompts and outputs. Close the preview to return to the same list.
- Select the intended chats. **Archive** selected keeps them stored and searchable but removes them from the active list. **Delete** selected opens a permanent-deletion confirmation; it does not delete until you choose Yes, delete.
- Under Legal holds, name a hold and choose Hold selected chats. This hold protects the selected records, including against manual deletion. For ongoing preservation of a client, also set that source's rule to Forever. Load active holds to review them; releasing a hold requires a second confirmation.
- Read the completed action status. **Chats** under an active hold are skipped by deletion. A failure or skipped record needs review; do not assume every selected row was deleted. The background scheduler handles automatic cleanup in bounded batches and writes an audit record per deleted chat.

#### CAREFUL

Permanent deletion removes the selected conversations and their attachments and cannot be undone. Preview contents and check the selected count before confirming. A tagging switch records classification; it is not a scheduled deletion rule or an action that places a hold.

## 38. Review feedback and reported issues

Read response ratings, written feedback, and platform reports in Analytics.

WALKTHROUGH · 1:09 · ADMIN CONSOLE › DOCUMENTATION

### Review feedback and reported issues

Follow a real rating with a note and a reported problem from the person's screen to your review in Analytics.

#### Before you begin

- An Admin account.
- People who rate replies or report problems from **Help**.

#### Step by step

- 1 The person rates a reply with the thumbs buttons below it and may add a note, then chooses Send note.
- 2 The person reports a problem from **Help** › Report a problem with a Subject, Message, and optional screenshot, then Send report.

- 3 Choose **Admin console** > **Analytics** and expand **Chat** Feedback.
- 4 Choose the person and dates, or Today, Week, or 30 days.
- 5 Open a feedback row to read the note and the rated exchange.
- 6 Scroll to Reported platform issues and open a report to read it.
- 7 Follow up through your team's support process.

#### Path: A rating with a note (shown)

- 1 Thumbs up or down, then an optional note. The row shows the sentiment, model, person, and note.

#### Path: A reported problem (shown)

- 1 **Help** > Report a problem. The report shows the subject, message, sender, time, and any screenshot.

#### Check that it worked

- ✓ The summary counts Total feedback, Positive, Negative, and Issue reports for your scope.
- ✓ Feedback and conversation shows the note and the rated exchange.
- ✓ Platform issue report shows the subject and message the person sent.

#### Troubleshooting

| WHAT YOU SEE                                  | WHAT TO DO                                                                               |
|-----------------------------------------------|------------------------------------------------------------------------------------------|
| The list is empty                             | Widen the person and date filter. Each section filters independently.                    |
| The conversation could not be loaded          | The dialog says so and shows the saved reply preview only. Use the evidence that loaded. |
| "The screenshot preview could not be loaded." | The report has a screenshot that failed to load. Read the text and try again later.      |

## 39. Audit

Grouped signals, Audit Insights trends, prompt monitoring, security alerts, and the tenant trail.

The **Audit** tab is the tenant's governance station. It opens on **Admin Audit** and **Audit** Insights, while the sections below start collapsed; click a section header to expand it. **User Prompt Activity**, **Security Alerts**, and the **Audit Trail** each carry their own filter row pairing a person picker with a date range, so one section's scope never narrows another. The **CSV** export buttons produce files of exactly the filtered rows, including actor\_id and actor\_name columns, so every exported row names who acted.

## Admin Audit signals

- The banner at the top reads, for example, “7 of 18 signals need attention”. Expand all and Collapse all open or fold every group. List shows one compact row per signal; Cards shows each signal with its description. The layout you choose is remembered in this browser.
- **Security** signals: audit events, critical events, the prompt watchlist of active DLP or misuse alerts, alert response (the median time to acknowledge an alert), after-hours changes (outside 7 AM–7 PM or on weekends), and failed operations.
- Identity & access: active admins, active users, access requests awaiting review, accounts that have never signed in, role changes, and credential changes such as MFA, password, session, and API key events.
- **Models & workspace**: prompt volume, connector issues, ungrouped models, unassigned users (active users in no group), agent approvals waiting for review, and automation failures.
- **Groups** that need attention open automatically and all-clear groups fold to one summary line. Red rows need attention. Select any signal to open an investigation listing every record behind it, with a filter box to narrow the records.

## Audit Insights

- **Audit** Insights charts the trends behind the signals for your organization. Choose a 7 days, 14 days, or 30 days range; the summary line counts the audit events and security alerts in range.
- **Audit** events by day stacks info, warning, and critical events. **Security** alerts by day draws DLP and behavior alerts as two lines, and Alert breakdown ranks alerts By rule or By person with their open and acknowledged counts.
- Most active people ranks who is acting, Activity by area groups events into areas such as **Chat**, Sign-in & identity, and Administration, and Activity by hour shades the hours outside 7 AM–7 PM in your local time.
- Select any bar, point, or row to open the same investigation view with the records behind it. If the loaded history reaches its page limit, a note says earlier days may be incomplete.

## Sections below the dashboard

- **Recent Governance Activity**: the current tenant snapshot for identity, user, model, and connector posture: how many **SSO** configurations are enforced, how many users are active, and how many connectors are enabled.
- **User Prompt Activity**: drill into saved prompts from this organization's admins and users and their model responses by person, thread, model, and timestamp, scoped by the section's own user and date filter.
- **Security Alerts**: DLP and malicious-behavior flags raised from admin and user prompts, shown with redacted snippets for review. Click **Acknowledge** once an alert is handled, or **Reopen** if it needs another look.

- **Audit Trail:** the tenant's append-only transaction log, newest first. Its toolbar has a severity select, an action-category select, and a search box (**Search actions, people, targets...**) that narrow the rows together; the trail's **CSV** exports exactly the rows you are looking at. Click **Refresh** to reload it straight from the admin audit API.

#### GOOD TO KNOW

**Security** alerts fire on real prompt content, such as payment card numbers, shared credentials, and prompt-injection attempts, but show only redacted snippets, so reviewing an alert never re-exposes the sensitive value itself.

#### WALKTHROUGH · 1:33 · ADMIN CONSOLE › DOCUMENTATION

### Tenant audit

Investigate a real flagged prompt from start to finish: the signal board, Audit Insights, the prompt and reply, the security alert, and the audit trail.

#### Before you begin

- An Admin account.
- A flagged prompt to investigate. This lesson follows a real prompt-injection attempt, the same one the **Alerts** and delivery lesson follows to a delivered email.

#### Step by step

- 1 Choose **Admin console** › **Audit**. The banner counts the signals that need attention.
- 2 Select a red signal, such as Prompt watchlist, to open its records. Close it to return.
- 3 In **Audit** Insights, choose 7, 14, or 30 days, then select a bar or point to list the records behind it.
- 4 Expand **User Prompt Activity**, search for the prompt, and open it to read the prompt and the model's reply.
- 5 Expand **Security Alerts** and choose **Acknowledge** once you have reviewed the alert.
- 6 Expand **Audit Trail**. Filter by severity, category, and text, then choose **CSV** to export what you see.

#### Path: Start from an alert email

- 1 The email names the rule, detection, person, and time, never the flagged text.
- 2 Open **Admin console** › **Audit** › **Security Alerts** and find the alert by detection and time, then follow the steps above.

#### Path: Export evidence

- 1 **User Prompt Activity** and **Audit Trail** each have a **CSV** button. The trail export carries exactly the filtered rows.

## Check that it worked

- ✓ The investigation lists the flagged records with person, model, and time.
- ✓ The prompt dialog shows the exact prompt and the model's saved reply.
- ✓ The alert reads Acknowledged, and the trail shows the security.prompt\_flagged events.

## Troubleshooting

| WHAT YOU SEE                                                      | WHAT TO DO                                                                                                 |
|-------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------|
| A signal opens with no records                                    | The count reflects the current snapshot; check the <b>Audit Insights</b> range or the trail's date filter. |
| The prompt you expect is missing from <b>User Prompt Activity</b> | Each section has its own person and date filter. Choose All and clear the search.                          |
| "Security alert was not updated: ..."                             | Nothing changed. <b>Refresh</b> and acknowledge again; another administrator may have changed it.          |

## 40. Alerts

Watch rules over audit activity, and honest delivery statuses for every alert.

The **Alerts** tab turns audit activity into notifications. It has three panels: **Email Delivery**, **Alert Rules**, and **Alert Deliveries**. Every alert is always logged in-app regardless of email — email is a delivery channel, not the record.

### Email Delivery

For organization administrators this panel is a read-only status: it reports whether email delivery is configured at the service level. If email is not configured, rules still work — their alerts are logged in-app and the delivery log says so honestly.

### Alert Rules

- Click Prompt-injection template to start from a rule that watches prompt-injection, system-prompt extraction, and credential-extraction attempts; **Suspicious-activity template** for security flags and elevated-severity events; or **New rule** to build one from scratch. A template only fills in the form; nothing is saved until you click **Create Rule**.
- A rule has: a name; **Action patterns** — exact audit actions or prefixes such as security.\* or admin.user\_deleted; a **Minimum severity**; an optional **Watched user**; a threshold — how many matches within a time window before it fires; a cooldown between alerts; and email recipients. Leave recipients empty and the alert is in-app only.

- Only these detections narrows a rule to prompts flagged by specific detectors: Prompt injection, System-prompt extraction, Credential extraction, API key or token shared, Private key shared, **Password** shared, US Social **Security** number, and Payment card number. Leave all unchecked to match every event the action patterns allow. A rule that could never fire is refused with an explanation when you save it.
- Alert emails name the rule and the detection but never include the flagged prompt text.
- Rules you create here watch this organization's admin and user audit activity.

## Alert Deliveries

Every alert trigger is listed with its real delivery status: sent, queued, failed with the actual SMTP error, email not configured, or logged in-app. The log exports to **CSV**, archived deliveries included. Click **Archive** on a delivery to clear it from the default view — its history is kept, and **Show archived** reveals archived deliveries so you can review or **Restore** them. Tenant admins can archive their tenant's deliveries only.

**WALKTHROUGH · 1:51 · ADMIN CONSOLE › DOCUMENTATION**

### Alerts and delivery

Check email delivery, create a rule from the Prompt-injection template with a recipient, and follow a real flagged prompt to a delivered email.

#### Before you begin

- An Admin account.
- Email delivery configured by your service team. Without it, alerts are still logged in-app.
- The addresses that should receive alerts.

#### Step by step

- 1 Choose **Admin console › Alerts. Email Delivery** should read Email configured.
- 2 In **Alert Rules**, choose Prompt-injection template.
- 3 Check the action pattern security.prompt\_flagged, the minimum severity, Fire when, and Cooldown.
- 4 Enter Email recipients, separated by commas.
- 5 Under Only these detections, keep or change the ticked detectors.
- 6 Choose **Create Rule**. The rule appears in the list, enabled.
- 7 When a matching prompt is flagged, **Alert Deliveries** shows the delivery and its status, such as sent.



### Path: Prompt-injection template (shown)

- 1 Watches security.prompt\_flagged for Prompt injection, System-prompt extraction, and Credential extraction, at Warning and above, every match, with a 10-minute cooldown.

### Path: Suspicious-activity template

- 1 Watches security.\*: security flags, content-filter hits, and elevated governance events, at Warning and above, with a 15-minute cooldown and no detection filter.

### Path: New rule

- 1 Enter a Rule name. **Action patterns** accepts exact actions or prefixes such as admin.\* or auth.\*; empty matches every action.
- 2 Choose a minimum severity, an optional **Watched user**, Fire when (matches within minutes), and a cooldown.
- 3 Detections apply only to security.prompt\_flagged events.

### Path: In-app only

- 1 Leave Email recipients empty. Deliveries read logged in-app.

### Check that it worked

- ✓ The rule appears enabled, with its patterns, detections, and recipient count.
- ✓ **Alert Deliveries** shows the delivery with status sent and the recipient.
- ✓ The email arrives with the rule and detection in its subject and no flagged text.



## Troubleshooting

| WHAT YOU SEE                                                               | WHAT TO DO                                                                                                               |
|----------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------|
| "A rule name is required."                                                 | Enter a Rule name.                                                                                                       |
| "'...' is not a valid email address."                                      | Separate recipients with commas and check each address.                                                                  |
| "Detections only apply to security.prompt_flagged events. ..."             | Add security.prompt_flagged or security.* to <b>Action patterns</b> , or clear the detections.                           |
| "None of the chosen detections can reach the minimum severity '...' . ..." | Lower <b>Minimum severity</b> or choose other detections.                                                                |
| A delivery reads email not configured                                      | Email delivery is managed at the service level. The alert stays logged in-app; ask your service team to configure email. |
| A delivery reads failed                                                    | Its detail shows the mail server's real error. <b>Share</b> it with your service team.                                   |

## The Platform owner console

### 41. Opening the Platform owner console

The highest access level, its six tabs, and the role ceilings.

- 1 Click Platform console (building icon) near the bottom of the sidebar. (Only platform owners see it; the account drawer's **Management** section lists it as **Platform owner console**.)
- 2 The console has six tabs: **Org Settings**, **Models**, **Providers**, **Analytics**, **Audit**, and **Alerts**. **Org Settings** holds organization configuration; API keys live on each provider's card under **Providers**.

Keep the three role ceilings in mind — they explain who can touch what across the entire platform:

| ROLE                  | WHAT IT CONTROLS                                                                                                                                                                                                    |
|-----------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Platform owner</b> | <b>Providers</b> , API keys and secrets, shared connectors, organization model availability, <b>SSO</b> , branding, policies, usage budgets, platform updates, alert email, and removing admins. The highest level. |
| <b>Admin</b>          | Tenant users, groups, response actions, and tenant model access — always inside the boundaries the owner sets.                                                                                                      |
| <b>User</b>           | <b>Chat</b> , drafts, assigned agents, and whatever models and sources their groups allow.                                                                                                                          |

#### GOOD TO KNOW

First-run setup: when no active platform owner exists yet, the sign-in screen shows “**Create the first platform owner**” — enter a display name, email, and a password of at least 12 characters. This screen never appears again once an owner exists.

#### TIP

The **Documentation** button opens owner walkthroughs and this PDF. Its **Admin documentation** and **Chat help** links open the other role libraries. **Interactive platform guide** opens the public digital walkthrough at <https://aperturechat.com/guide.html> in a new tab for additional platform setup and configuration guidance.

## 42. First-run setup: from owner account to a working team

Create the first owner, connect a model, and verify access before inviting the team.

### GOOD TO KNOW

Before inviting a team, verify the intended provider, enabled models, group memberships, and model grants. Test each intended user role; a working owner account does not prove that a standard user has access.

Use **Providers** and **Models** for service configuration, then Admin Console → **Model Access** and **Groups** for tenant access. Use **Org Settings** for **SSO**, branding, connector availability, policy, and budgets.

### WALKTHROUGH · 3:52 · PLATFORM CONSOLE · DOCUMENTATION

#### Set up the first workspace

Create the first platform owner on a new installation, connect a model, give one person access, and watch that person get a real reply.

#### Before you begin

- A new installation that has never been set up. The first visit shows **Create the first platform owner**.
- The owner's email address, display name, and a password of at least 12 characters.
- A model to connect: a cloud provider API key (see **Providers** and connections), or a local OpenAI-compatible server such as Ollama or LM Studio that the server can reach.
- The email address of the first person who will use the workspace.

#### Step by step

- 1 Open your instance's address. On the **Create the first platform owner** screen, enter Email, Display name, **Create** password, and Confirm password, then choose **Create** platform owner.
- 2 On the Getting started card, choose Set up models. The Platform console opens on **Providers** with Connect your first model provider.
- 3 Choose **Add Provider**. Enter a Name, pick the Kind, enter the Region and Base URL, and give the key a Key name.
- 4 Paste the provider key into API key or secret. A local server that ignores keys still needs a placeholder value here, because **Sync Models** only runs with an active key on file.
- 5 Choose Save **Provider**. Wait for the card to show **Connected** and the message Runtime test passed with the model's name.
- 6 Open **Models**. New models arrive disabled; turn on the switch in Org status for the model your team should use.
- 7 Open **Org Settings** › **Role Boundary**. Enter the person's Display name and Email, keep Role on User, and choose **Create account**.

- 8 Choose the key button on the person's row (Set a password). Choose **Generate**, keep **Temporary password** on, choose **Set password**, and copy the password from the confirmation.
- 9 **Share** the workspace address and the temporary password over a safe channel. Nothing is emailed automatically.
- 10 Optional check: **Admin console** > **Model Access** shows the model as Visible to 1 group, and Edit groups shows **Default Users** ticked.
- 11 In a private window, the person signs in with Email and the temporary password, chooses a new password on Set a new password, and is taken to the workspace.
- 12 The person sends a first message. A reply from the connected model proves the account, the group, the model switch, and the provider all work together.

### Path: Local OpenAI-compatible server (shown in the video)

- 1 Kind: openai-compatible (or ollama for Ollama, local for a generic local runtime).
- 2 Base URL: the server address followed by /v1, for example http://localhost:11434/v1 for Ollama or http://localhost:1234/v1 for LM Studio. Use an address the Aperture Chat server can reach, not just your laptop.
- 3 API key or secret: any placeholder value when the server does not check keys.
- 4 Save **Provider** syncs the model list and runs a short test chat. **Connected** appears only when that test passes.

### Path: Cloud provider (OpenAI, Anthropic, OpenRouter, Gemini, and others)

- 1 **Create** an API key in the vendor's console first. The **Providers** and connections lesson shows where for each vendor.
- 2 Kind: the vendor. The Base URL, Auth type, and Header fill in for most kinds; keep them unless your vendor says otherwise.
- 3 Paste the real key into API key or secret and choose Save **Provider**. The catalog syncs and a test chat runs.
- 4 Azure OpenAI and Azure Foundry cannot list their deployments yet, so they cannot be finished from the console; see the Azure path in **Providers** and connections.

### Path: People sign in with single sign-on instead of passwords

- 1 Skip Set a password. Connect your identity provider first (Single sign-on track).
- 2 With Provision new users on first sign-in on, people are created as users when they first sign in, and join **Default Users**.
- 3 **Share** the workspace address; people choose **Continue with SSO**.

## Check that it worked


- ✓ The provider card shows **Connected** and Runtime test passed with the model's name.
- ✓ **Models** shows the model as Enabled.
- ✓ The person's first sign-in asks them to Set a new password, then opens the workspace.
- ✓ The person's first message gets an answer from the connected model, named in the **Model** menu at the top of the chat.

## Troubleshooting

| WHAT YOU SEE                                                                                   | WHAT TO DO                                                                                                                                                            |
|------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "A platform owner already exists."                                                             | Someone finished setup first. Sign in with that owner account; the setup screen appears only once per installation.                                                   |
| "Password must be at least 12 characters." or "The passwords do not match."                    | Use a password or passphrase of 12 or more characters and type it the same way in both fields.                                                                        |
| <b>Sync Models</b> is greyed out and the card says <b>Needs key</b> .                          | No active key is on file. Open <b>API Keys</b> on the card, choose <b>Add Key</b> , paste the key (a placeholder for a local server), and choose <b>Save Key</b> .    |
| "Key saved; model sync failed:<br><b>Provider model sync failed:</b><br><b>ConnectError.</b> " | The server cannot reach the Base URL. Check the address and port, that the local server is running and listening on the network, and that the URL ends in /v1.        |
| "Key saved; model sync failed:<br><b>Provider model sync failed with HTTP 401.</b> "           | The vendor rejected the key. <b>Create</b> a new key in the vendor's console, then use <b>Replace</b> on the key row.                                                 |
| The person sees <b>No connected models</b> in the <b>Model</b> menu.                           | Turn the model on in <b>Models</b> , and check <b>Admin console &gt; Model Access</b> : the model needs a group the person belongs to, such as <b>Default Users</b> . |
| "Invalid local credentials." when the person signs in.                                         | Check the email and the temporary password. Set a new temporary password from <b>Role Boundary</b> if it was lost; it is shown only once.                             |

## 43. Providers

Register model gateways, read connection health, and sync catalogs.



Provider cards show the actual connection and model-catalog state. Example workspace with synthetic data.

### Kind-driven defaults

- Picking a kind prefills sensible connection defaults. `azure-foundry` uses the inference endpoint `https://{resource}.services.ai.azure.com/models` with api-key authentication; `gcp` routes through Google's OpenAI-compatibility endpoint for Gemini models.
- The **Catalog scope** row appears only on openrouter providers: choose between the zero-data-retention (ZDR) filtered list and the key-scoped model list.

#### GOOD TO KNOW

A provider is not usable until it has an active key in its vault — open **API Keys** on the card and add the key next.

#### WALKTHROUGH · 6:28 · PLATFORM CONSOLE · DOCUMENTATION

### Providers and connections

Connect a model provider of any kind, get its key from the vendor, and read the real result of the connection test.

#### Before you begin

- A platform owner account.
- For a cloud provider: an account with billing or credits at the vendor, and permission to create API keys there.

- For a local server (Ollama, LM Studio, or another OpenAI-compatible server): at least one model downloaded, and the server listening on an address the Aperture Chat server can reach.
- Azure OpenAI: the resource endpoint, KEY 1, and the deployment name from the Azure portal.

## Step by step

- 1 Choose Platform console at the bottom of the sidebar, then open **Providers**.
- 2 Choose **Add Provider**. Enter a Name people will recognise, such as OpenAI production.
- 3 Choose the Kind. For most kinds, Base URL, Auth type, and Header fill in automatically; keep them unless your vendor tells you otherwise.
- 4 Enter the Region if you track it, then a Key name, the Environment, and Expires (the date the vendor key stops working, or Not set).
- 5 Paste the key into API key or secret. It goes straight to the encrypted vault and is never shown on the card.
- 6 Choose Save **Provider**. Aperture Chat saves the connection, stores the key, syncs the model list, and sends one short test chat.
- 7 Read the card: **Connected** with Runtime test passed means the provider works. **Needs key**, **Needs validation**, or an error message means it does not yet; see Troubleshooting.
- 8 Open **Models** and turn on the models your team should use. Newly synced models always arrive disabled.

## Path: Local server: Ollama, LM Studio, or OpenAI-compatible (shown end to end)

- 1 On the model machine, download a model (Ollama: ollama pull, then a model name) and make the server listen on the network (Ollama: OLLAMA\_HOST=0.0.0.0:11434; LM Studio: Developer > Start server and Serve on Local Network).
- 2 Kind: ollama for Ollama (Base URL fills in as http://localhost:11434/v1), or openai-compatible for LM Studio and other servers.
- 3 Base URL: the server's address the Aperture Chat server can reach, ending in /v1. LM Studio's default port is 1234.
- 4 API key or secret: any placeholder if the server does not check keys. Sync needs an active key on file.
- 5 Save **Provider**, and wait for **Connected** and Runtime test passed.

## Path: OpenAI

- 1 At platform.openai.com, open Settings > API keys and choose **Create** new secret key. Pick the Project and **Permissions**, create it, and copy it; it is shown once.
- 2 Kind: openai. Base URL https://api.openai.com/v1, Auth type bearer, Header Authorization fill in.

- 3 Paste the key and choose Save **Provider**.

### Path: Anthropic

- 1 At platform.claude.com, open Settings › API keys, choose **Create** key, set the expiration and Linked account, and copy the key (sk-ant-...).
- 2 Kind: anthropic. Base URL https://api.anthropic.com, Auth type api-key, Header x-api-key fill in.
- 3 Paste the key and choose Save **Provider**.

### Path: Azure OpenAI

- 1 Azure portal › your resource › Resource **Management** › Keys and Endpoint: copy KEY 1 and the Endpoint. In Foundry, copy the deployment name.
- 2 Kind: azure-openai. Base URL: https://your-resource.openai.azure.com/openai. Auth type api-key and Header api-key fill in.
- 3 API version: 2024-10-21. Deployment: your deployment name.
- 4 Save **Provider**. The card stays on Needs validation with ... does not expose a supported model discovery API yet: Azure deployments are not listed automatically, and the console has no control to add one by hand yet. Use another provider kind for live chat until this is supported.

### Path: OpenRouter

- 1 At openrouter.ai, add credits, open Keys, create a key with an optional credit limit, and copy it.
- 2 Kind: openrouter. Base URL https://openrouter.ai/api/v1 fills in.
- 3 **Catalog scope**: ZDR models syncs only zero-data-retention endpoints; Key-scoped models syncs what your key and OpenRouter privacy settings allow.
- 4 Paste the key and choose Save **Provider**.

### Path: Google Gemini

- 1 In Google AI Studio, open **API Keys**, choose **Create** API key, pick the billing project, restrict it to the Gemini API, and copy it.
- 2 Kind: gcp. Base URL https://generativelanguage.googleapis.com/v1beta/openai fills in.
- 3 Paste the key and choose Save **Provider**.

### Path: Other OpenAI-compatible vendors

- 1 Kinds xai, mistral, deepseek, groq, together, fireworks, perplexity, cerebras, sambanova, moonshot, nvidia, deepinfra, cohere, and open-webui fill in the vendor's address and a bearer header.
- 2 **Create** a key in the vendor's console, paste it, and choose Save **Provider**.



## Path: Amazon Bedrock and Azure Foundry

- 1 The connection can be saved, but chat cannot route to it yet. The card stays on Needs validation with does not expose a supported model discovery API yet.
- 2 Use another provider kind for live chat.

### Check that it worked

- ✓ The card shows **Connected**, and its message ends with Runtime test passed and the name of the model that answered.
- ✓ The card's **Models** line counts the synced models (for example 0 of 1 before you enable any).
- ✓ Platform console > **Audit** records platform.provider\_models\_synced for the provider.

## Troubleshooting

| WHAT YOU SEE                                                                                                                                             | WHAT TO DO                                                                                                                                                                         |
|----------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "Provider model sync failed: ConnectError."                                                                                                              | Nothing answered at the Base URL. Check the host and port, that the server is running, and that it listens on the network rather than only on its own machine.                     |
| "Provider model sync did not return a model list." for Ollama                                                                                            | Ollama is running but has no model. Pull one on the model machine (ollama pull, then a model name), then choose <b>Sync Models</b> .                                               |
| "Provider model sync failed with HTTP 401." and the card says Needs key                                                                                  | The vendor rejected the key, and Aperture Chat marked it Inactive. <b>Create</b> a new key in the vendor's console and add it with <b>API Keys › Add Key</b> , or <b>Replace</b> . |
| "... rejected its provider key with HTTP 401. Paste a valid provider-generated key, make sure billing or credits are available, then sync models again." | The model list loaded, but the test chat was refused. Check billing or credits and the key's permissions at the vendor.                                                            |
| "... model sync succeeded, but live chat validation failed ..."                                                                                          | The vendor listed models but could not answer a chat. Check model access for the key, the region, and the Base URL.                                                                |
| "... does not expose a supported model discovery API yet."                                                                                               | Azure OpenAI, Azure Foundry, and Amazon Bedrock cannot sync a model list. Use a different provider kind for live chat.                                                             |
| Sync Models is greyed out.                                                                                                                               | There is no active key. Add one with <b>API Keys › Add Key</b> .                                                                                                                   |
| Delete provider stays greyed out, or "... still has models in use by ..."                                                                                | Type the provider's name exactly to enable <b>Delete</b> provider. If automations or agent profiles use its models, repoint or remove them first.                                  |

## 44. API Key Vault

Each provider card vaults its own secrets; reveal, replace, or delete them.

Each provider's secrets live in a vault attached to its own card — there is no separate keys tab. Open the **Providers** tab and click **API Keys** on a card to expand its **API Key Vault** panel.

- **Reveal** — opens the **Vault reveal** dialog, which shows the secret with a **Copy key** button and a **Done** button to hide it again. Expired keys cannot be revealed.

- **Replace** — swaps in a new provider-generated secret in place.
- **Delete** — removes the key.

Reveals, replacements, and deletes write through the platform API, so the audit trail stays complete.

#### WALKTHROUGH · 1:51 · PLATFORM CONSOLE > DOCUMENTATION

### API Key Vault and replacement

Add, reveal, replace, and delete a provider's keys, and handle a key that has expired.

#### Before you begin

- A platform owner account.
- A provider already added in Platform console > **Providers**.
- A new key created in the vendor's console (see **Providers** and connections for where).

#### Step by step

- 1 Open Platform console > **Providers** and choose **API Keys** on the provider's card. The number on the button is how many keys are stored.
- 2 Read the **API Key Vault** table: Key Name with the masked value (first three and last four characters), Environment, Status, Last Rotated, and Expires.
- 3 To add a key, choose **Add Key**. Enter a Key name, the Environment, and Expires (the date the vendor key stops working, or Not set).
- 4 Paste the key into API key or secret and choose Save Key. Aperture Chat stores it encrypted, then syncs the models and sends a test chat with it.
- 5 To see a stored key, choose the eye button (**Reveal**). The full value opens in a dialog with **Copy key**. Every reveal is recorded in the audit trail as platform.provider\_key\_revealed.
- 6 To replace a key, create the new key at the vendor, then choose the key button (**Replace**) on the old row. The form opens named after the old key; paste the new key and choose Save Key.
- 7 When the new key shows Active and the card still says **Connected**, choose the bin button (**Delete**) on the old row. Then revoke the old key at the vendor.

#### Path: Rotate a key on a schedule

- 1 **Create** the new key at the vendor first.
- 2 **Replace** on the old row > paste the new key > Save Key.
- 3 Confirm the new row is Active and the card is **Connected**.
- 4 **Delete** the old row, then revoke the old key at the vendor.

### Path: A key has expired

- 1 When the Expires date has passed, the row shows Expired in red and **Reveal** is greyed out.
- 2 **Create** a new key at the vendor, add it with **Replace** or **Add Key**, then delete the expired row.

### Path: The vendor rejected a key

- 1 After a failed sync with HTTP 401 or 403, the key is marked Inactive and the card says **Needs key**.
- 2 Add a valid key; a successful sync makes the provider **Connected** again.

### Check that it worked

- ✓ The new row shows Active, and the card shows **Connected** with Runtime test passed.
- ✓ The deleted key no longer appears, and the **API Keys** count on the card went down by one.
- ✓ Platform console > **Audit** lists platform.provider\_key\_created and, for a reveal, platform.provider\_key\_revealed.

### Troubleshooting

| WHAT YOU SEE                                                                   | WHAT TO DO                                                                                                                                                                     |
|--------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Save Key stays greyed out.                                                     | Paste the key into API key or secret; a key value is required.                                                                                                                 |
| Reveal is greyed out and the row says Expired.                                 | Expired keys cannot be revealed or used. Add a replacement and delete the expired row.                                                                                         |
| "Provider key is expired. Add a replacement key before revealing or using it." | The key's Expires date has passed. Add a new key.                                                                                                                              |
| "... key was saved, but model sync failed: ..."                                | The key is stored, but the provider did not accept it or could not be reached. Read the rest of the message and see <b>Providers</b> and <b>connections</b> > Troubleshooting. |
| Copy key says Clipboard access is unavailable.                                 | Select the key text in the dialog and copy it manually.                                                                                                                        |

## 45. Models: organization availability

The organization ceiling — control which synced models tenants can route to.

### GOOD TO KNOW

Newly synced models arrive disabled with no group access. After a sync, switch the filter to Disabled, review the new arrivals, and enable only the routes the organization has actually approved. If the **Default group for enabled models** policy is on (see “Policies”), enabling a model automatically attaches the protected **Default Users** group so users see it without a second step.

### WALKTHROUGH · 2:05 · PLATFORM CONSOLE › DOCUMENTATION

## Organization model availability

Decide which synced models your organization can use, give them clear names, and control what people see when a model is not available to them.

### Before you begin

- A platform owner account.
- At least one provider connected, so models have been synced.
- For the user check: a person who can sign in, or a second browser signed in as one.

### Step by step

- 1 Open Platform console › **Models**. The All, Enabled, and Disabled counters and **Search** models work on the whole catalog.
- 2 Narrow the list with the funnel beside **Provider**, **Model**, or **Runtime route**. Tick values, or type text for **Runtime route**; Clear filter removes it.
- 3 Turn a model on or off with its Org status switch. Off removes it for everyone, whatever their groups; newly synced models start off.
- 4 Choose **Edit details** to set a Display name people will recognise, the **Context window**, **Notes**, and optional System prompt and Meta prompt, then choose Save details.
- 5 Group access is set by administrators in **Admin console** › **Model Access**. With the Default group policy on, enabling a model also gives it to **Default Users**.
- 6 Decide on catalog browsing in **Org Settings** › **Policy Controls** › **Users can browse the model catalog**.
- 7 Check as a user: the **Model** menu lists only models they can use, and **Why isn't a model listed?** explains the rest.

### Path: Open catalog (Users can browse the model catalog on)

- 1 People see every enabled model in **Why isn't a model listed?**, with the reason it is or is not available.
- 2 **Locked** models offer **Request access**, which reaches administrators in **Admin console > Model Access > Access requests**.
- 3 **Provider offline** means the person already has access, but the provider needs reconnecting.

### Path: Closed catalog (Users can browse the model catalog off)

- 1 People see only models they can already use, with the note: Your organization shows only the models you can already use. Ask an administrator about others.
- 2 New access requests are refused: **Model** access requests are disabled by organization policy.

### Check that it worked

- ✓ The Enabled and Disabled counters change, and the confirmation reads ... availability saved through the platform API.
- ✓ A disabled model no longer appears in a user's **Model** menu.
- ✓ The renamed model shows its new Display name in the user's **Model** menu.
- ✓ With browsing off, the user's catalog shows the closed-catalog note.

### Troubleshooting

| WHAT YOU SEE                                                 | WHAT TO DO                                                                                                                                                                  |
|--------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "No models match the current column filters."                | A column filter is still active (the funnel shows a number). Open it and choose Clear filter.                                                                               |
| A user cannot see an enabled model.                          | The model needs a group the person belongs to. Check <b>Admin console &gt; Model Access &gt; Edit groups</b> , and the person's groups in <b>Admin console &gt; Users</b> . |
| The user sees <b>Provider offline</b> .                      | They have access, but the provider is not connected. Fix the provider in <b>Providers</b> , then choose <b>Sync Models</b> .                                                |
| "Model access requests are disabled by organization policy." | <b>Users can browse the model catalog</b> is off. Turn it on, or have an administrator add the person to a group.                                                           |

## 46. Org Settings: users and role boundaries

Create accounts at any level, set passwords, and rely on the account floors.

The **Org Settings** tab gathers organization-level controls: roles and accounts, single sign-on, branding, policies and budgets, platform connectors, and search index readiness. Most sections start collapsed behind a descriptive header; click a header (or its chevron) to expand it. **Search index** starts open. Expand **Role Boundary** for the account tools.

### Setting a password

The key button on a row opens the same dialog admins use — “Set a password for” that person. Type a password or click **Generate** for a strong random one, optionally mark it a **Temporary password** so they must choose their own at first sign-in, and click **Set password**. The password is shown only in that dialog — share it over a safe channel.

#### GOOD TO KNOW

Two floors protect the platform. The last active platform owner can never be removed. And at least one active administrator (owner or admin) must always remain — because owners count as administrators, the sole tenant admin is removable while an active owner exists, but the API refuses any removal that would leave no administrator at all.

Below the account list, the “Clear separation of duties” callout restates the boundary: platform owners manage provider secrets, owner accounts, admin delegation, **SSO** baselines, and platform branding.

#### WALKTHROUGH · 2:10 · PLATFORM CONSOLE > DOCUMENTATION

### Users and role boundaries

Create user, admin, and platform owner accounts, issue temporary passwords, change or remove roles, and know the limits that keep the platform safe.

#### Before you begin

- A platform owner account.
- Each person's display name and work email.
- A safe channel (not the same email thread) for sharing temporary passwords.

#### Step by step

- 1 Open Platform console > **Org Settings** and expand **Role Boundary**.
- 2 Enter the person's Display name and Email, choose the Role (User, Admin, or Platform owner), and choose **Create account**.
- 3 On the new row, choose the key button (Set a password). Choose **Generate**, keep **Temporary password** on, and choose **Set password**.

- 4 **Copy** the password from the confirmation; it is shown only once. **Share** it and the workspace address over a safe channel. Nothing is emailed.
- 5 To change a role, choose a new value in the row's role menu; it saves immediately.
- 6 To remove an account, choose the bin button and confirm. **Users** and admins are deleted permanently; owners are deactivated instead.
- 7 Verify: the person signs in, chooses a new password on Set a new password, and sees the consoles their role allows.

#### Path: User

- 1 Uses chat, drafts, agents, and the library. Joins **Default Users** when the Default group policy is on.
- 2 No console access.

#### Path: Admin

- 1 Everything a user can do, plus the **Admin console**: users, groups, model access, **SSO** mappings, analytics, policies, audit, and alerts for the organization.
- 2 No Platform console: providers, keys, branding, and platform policies stay with owners.

#### Path: Platform owner

- 1 Everything, including the Platform console. Only an owner can create or promote another owner.
- 2 Owners sign in with their own local password; Set a password is unavailable on an owner who already has one.
- 3 Keep at least two owners, so one can recover the other.

#### Path: SSO accounts

- 1 With single sign-on, people are created at their first sign-in as users. Skip Set a password; change their role here afterwards if needed.

#### Check that it worked

- ✓ The confirmation reads ... account created through the admin API, and the row shows the chosen role.
- ✓ At first sign-in the person sees Set a new password, then the workspace.
- ✓ A new admin sees **Admin console** in the sidebar but not Platform console.
- ✓ Platform console > **Audit** records each account creation, role change, password reset, and deletion.



## Troubleshooting

| WHAT YOU SEE                                            | WHAT TO DO                                                                                                           |
|---------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------|
| "... account was not created. User already exists."     | That email already has an account. Find the row in the list and change its role or password instead.                 |
| The owner's role menu and bin button are greyed out.    | At least one active platform owner must remain. <b>Create</b> a second owner first.                                  |
| The bin button is greyed out on your own row.           | Accounts cannot remove themselves. Ask another owner.                                                                |
| The key button is greyed out on an owner's row.         | That owner already has a local password and changes it from their own account panel.                                 |
| The person sees "Invalid local credentials."            | Check the email and the temporary password. Issue a new temporary password if it was lost; it cannot be shown again. |
| Removing an owner deactivates them instead of deleting. | This is by design: owner accounts are never deleted. Another owner can reactivate them.                              |

## 47. Single sign-on: connect any OpenID Connect provider

How sign-in through your identity provider works, and the complete setup loop from app registration to the first signed-in person.

Single sign-on sends people to your organization's identity provider to sign in. The provider returns a signed ID token; Aperture Chat verifies its signature, issuer, audience, nonce, and expiry, then opens a session. OpenID Connect (OIDC) is the working sign-in protocol. SAML appears in the **Protocol** list but is deferred and cannot sign anyone in.

| WHERE                                                         | WHO                                                   | USE IT FOR                                                                                                                               |
|---------------------------------------------------------------|-------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Platform console &gt; Org Settings &gt; Single Sign-On</b> | Platform owners                                       | The organization's main provider: issuer, client, domains, provisioning, the platform authenticator rule, and enforcement.               |
| <b>Admin console &gt; SSO</b>                                 | Owners, and admins when <b>Policy Controls</b> allows | Every provider for the tenant as a card: <b>Test connection</b> , IdP group mapping, enforcement, and extra providers for other domains. |

Plan the rollout in this order: register the application and copy its credentials, save and test the connection, sign in a test account, map groups, then enforce **SSO** only after real sign-ins work. The walkthrough below performs the whole loop with Keycloak; the paths cover Microsoft Entra ID, Okta, and Google Workspace step by step.

## Single sign-on, start to finish

Register the app with an OpenID Connect provider, connect it, test it, and sign the first person in. Shown end to end with Keycloak.

### Before you begin

- A platform owner account that signs in with a local password on a domain you will not enforce. It is your way back in if the identity provider is unavailable.
- Administrator access to your identity provider: Microsoft Entra ID, Okta, Google Workspace, Keycloak, or another OpenID Connect provider.
- The email domains your people sign in with, such as examplecorp.com.
- A test account at the identity provider with a verified email on one of those domains.
- Your instance's public address. The redirect URI is that address followed by `/api/auth/sso/callback`, and the server's `APERTURE_API_BASE_URL` must be the same address.

### Step by step

- 1 Open Platform console › **Org Settings**, expand **Single Sign-On**, and choose **Copy** beside the redirect URI.
- 2 At your identity provider, create a confidential web application that uses OpenID Connect and the authorization code flow. Paste the redirect URI exactly: scheme, host, port, and path must match.
- 3 **Copy** the application's client ID and client secret, and note when the secret expires.
- 4 Make sure the ID token carries the person's email (the email scope). If you will map groups, add a groups claim to the ID token.
- 5 Assign your test account, and later the groups that should use the workspace, to the application.
- 6 Back in **Single Sign-On**, enter a provider name, keep **Protocol** on **OIDC (supported)**, and enter the issuer URL. Aperture Chat reads issuer/.well-known/openid-configuration.
- 7 Paste the client ID and client secret, list the allowed email domains, and keep the group claim name your provider sends (groups by default).
- 8 Leave Provision new users on first sign-in on if accounts should be created automatically. Leave **Enforce SSO** for these domains off for now.
- 9 Choose **Save SSO**, then **Test connection**. Both checks must pass.
- 10 In a private browser window, open the workspace, enter the test account's email, choose **Continue with SSO**, and sign in at the identity provider.
- 11 Confirm the person reaches the workspace, then find the account in **Admin console › Users**.

### Path: Keycloak (shown in the video)

- 1 Manage realms › **Create** realm: name the realm for your organization and choose **Create**. Set a Display name under Realm settings; it appears on the sign-in page.
- 2 Clients › **Create** client: Client type OpenID Connect, **Client ID** aperture-chat, Name Aperture Chat, then Next.
- 3 Capability config: turn on Client authentication, keep Standard flow, clear Direct access grants, turn on Require PKCE, then Next.
- 4 Login settings: paste the redirect URI into Valid redirect URIs (no wildcards), then Save.
- 5 Credentials tab: keep Client Id and Secret and copy the Client Secret. Keep Allowed authentication method on Any.
- 6 Client scopes › aperture-chat-dedicated › **Configure** a new mapper › Group Membership: Name groups, Token Claim Name groups, Full group path off, Add to ID token on, Save.
- 7 **Groups** › **Create** group for each group you will map. **Users** › **Create** new user: Email verified on, username, email, names, Join **Groups**, **Create**; then Credentials › **Set password** with Temporary off.
- 8 Issuer URL for Aperture Chat: <https://your-keycloak-host/realms/your-realm>. Use the same host name the server and browsers both reach.

### Path: Microsoft Entra ID

- 1 Entra ID › App registrations › New registration: name it, choose Single tenant only, then Register. **Copy** the Application (client) ID and Directory (tenant) ID.
- 2 Manage › Authentication › Add Redirect URI › **Web**: paste the redirect URI and **Configure**. Leave the implicit and hybrid ID tokens box unticked.
- 3 Certificates & secrets › New client secret: choose an expiry, Add, and copy the Value (not the Secret ID) right away.
- 4 Token configuration: add the email optional claim to the ID token, and Add groups claim (group object IDs) if you will map groups.
- 5 Enterprise applications › your app: Properties › Assignment required? Yes; **Users** and groups › assign people or groups.
- 6 Issuer URL: <https://login.microsoftonline.com/your-tenant-id/v2.0>. Never use common or organizations.

### Path: Okta

- 1 Applications › Applications › **Create** App Integration (Classic experience if asked): OIDC - OpenID Connect, **Web** Application, Next.

- 2 Grant type Authorization **Code**; Sign-in redirect URIs: the redirect URI; Assignments: Limit access to selected groups; Save.
- 3 General tab: copy the **Client ID** and the **Client secret**.
- 4 For group mapping, use the default custom authorization server: **Security** > API > Authorization Servers > default > Claims > Add Claim named groups, ID Token, Always, value type **Groups**, a narrow filter, Any scope.
- 5 Issuer URL: https://your-org.okta.com/oauth2/default (custom server) or https://your-org.okta.com (org server, no group claim).

### Path: Google Workspace

- 1 In Google Cloud, choose a project in your organization. Google Auth Platform > Branding: app name and support email. Audience: Internal.
- 2 Clients > **Create** client > **Web** application: add the redirect URI under Authorized redirect URIs and **Create**.
- 3 **Copy** the **Client ID** and **Client secret** from the dialog; the full secret is shown only once.
- 4 Issuer URL: https://accounts.google.com. Google ID tokens carry no groups, so manage workspace groups in the **Admin console**.

### Check that it worked

- ✓ **Test connection** lists Discovery document and Signing keys (JWKS) as passed.
- ✓ A private-window sign-in with the test account returns to the workspace already signed in.
- ✓ **Admin console** > **Users** lists the person with the USER role and sso under Auth.
- ✓ **Admin console** > **Audit** records auth.user\_jit\_created for an account created at first sign-in.

## Troubleshooting

| WHAT YOU SEE                                                                                                                                                                       | WHAT TO DO                                                                                                                                                                                    |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| The identity provider rejects the redirect URI (Entra AADSTS50011, Okta "must be a Login redirect URI", Google redirect_uri_mismatch, Keycloak "Invalid parameter: redirect_uri"). | Register the exact URI shown in the panel. Scheme, host, port, path, and any trailing slash must match. If the panel shows an address people do not use, fix APERTURE_API_BASE_URL.           |
| "Could not fetch OIDC discovery document"                                                                                                                                          | The issuer URL is wrong, or the server cannot reach it. Open the issuer followed by /.well-known/openid-configuration from the server's network. For Entra, use your tenant ID, never common. |
| "Token endpoint returned 401: invalid_client"                                                                                                                                      | The client secret is wrong or expired. For Entra, copy the secret's Value rather than its Secret ID. <b>Create</b> a new secret at the provider, paste it, and choose <b>Save SSO</b> .       |
| "ID token validation failed" mentioning the issuer or audience                                                                                                                     | The issuer URL must equal the issuer in the provider's discovery document, character for character, and the client ID must be the one the token was issued to.                                |
| "The identity provider did not return an email claim"                                                                                                                              | Grant the email scope or optional claim, and make sure the account has an email address at the provider.                                                                                      |
| "The identity provider reports ... as unverified"                                                                                                                                  | Mark the address as verified at the provider (in Keycloak, turn on Email verified), then sign in again.                                                                                       |
| "... is outside the domains allowed for this SSO provider"                                                                                                                         | Add the domain to Allowed email domains and save, or sign in with an account on an allowed domain.                                                                                            |
| "JIT provisioning is disabled for this SSO provider"                                                                                                                               | <b>Create</b> the account first in <b>Admin console &gt; Users</b> , or turn on Provision new users on first sign-in.                                                                         |

### GOOD TO KNOW

The redirect URI is your instance's public address followed by /api/auth/sso/callback. The server builds it from APERTURE\_API\_BASE\_URL, so that setting must be the address people use. If the panel shows a different host, fix the setting before registering the application.

## 48. Single sign-on: Microsoft Entra ID, Okta, and Google Workspace

Provider-specific registration, claims, assignment, and the exact issuer to enter for each preset.

Each provider names things differently, so each has its own walkthrough. The steps inside the provider's console are shown as instruction cards; the Aperture Chat steps are shown on real screens. Steps that depend on a provider license are marked.

| PROVIDER                          | ISSUER URL TO ENTER                                                                                                                                                     | GROUPS IN THE ID TOKEN                                                       |
|-----------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------|
| Microsoft Entra ID                | <a href="https://login.microsoftonline.com/your-tenant-id/v2.0">https://login.microsoftonline.com/your-tenant-id/v2.0</a>                                               | Group object IDs, from Add groups claim (200-group limit).                   |
| Okta                              | <a href="https://your-org.okta.com/oauth2/default">https://your-org.okta.com/oauth2/default</a> (or <a href="https://your-org.okta.com">https://your-org.okta.com</a> ) | Group names, from a groups claim on the default custom authorization server. |
| Google Workspace                  | <a href="https://accounts.google.com">https://accounts.google.com</a>                                                                                                   | None. Manage workspace groups in the <b>Admin console</b> .                  |
| Keycloak and other OIDC providers | <a href="https://your-host/realms/your-realm">https://your-host/realms/your-realm</a> (or the provider's issuer)                                                        | Whatever the provider's group mapper sends, such as group names.             |

### WALKTHROUGH · 2:37 · PLATFORM CONSOLE · DOCUMENTATION

#### Single sign-on with Microsoft Entra ID

Register the app in the Entra admin center, add the claims, assign people, and connect it with the Entra preset.

##### Before you begin

- An Entra account with at least the Application Developer role. Token configuration and assignment need Cloud Application Administrator.
- Your Directory (tenant) ID, shown on the Entra overview page.
- Microsoft Entra ID P1 or P2 if you will assign groups to the app or use Conditional **Access**.
- The redirect URI copied from Platform console > **Org Settings** > **Single Sign-On**.
- A test account whose email or user principal name is on an allowed domain.

##### Step by step

- 1 In Aperture Chat, open **Org Settings** > **Single Sign-On** and copy the redirect URI.
- 2 Entra admin center > Entra ID > App registrations > New registration. Name the app, choose Single tenant only (older screens: Accounts in this organizational directory only), and select Register.
- 3 On Overview, copy the Application (client) ID and the Directory (tenant) ID.

- 4 Manage › Authentication › Add Redirect URI › **Web**. Paste the redirect URI and select **Configure**. Leave ID tokens (used for implicit and hybrid flows) unticked.
- 5 Certificates & secrets › Client secrets › New client secret. Add a description, choose an expiry (Microsoft recommends under 12 months), and select Add. **Copy** the Value now; it is never shown again.
- 6 Token configuration › Add optional claim › ID › email › Add.
- 7 To map groups: Token configuration › Add groups claim › **Groups** assigned to the application (or **Security** groups) › Save. Values arrive as group object IDs.
- 8 Enterprise applications › your app › Properties: set Assignment required? to Yes and Save. **Users** and groups › Add user/group: assign your test account and groups.
- 9 In Aperture Chat, choose the Microsoft Entra ID preset and replace `{tenant-id}` in the issuer with your Directory (tenant) ID.
- 10 Paste the Application (client) ID and the secret Value, list your domains, keep the group claim as groups, and leave **Enforce SSO** off.
- 11 Choose **Save SSO**, then **Test connection**. Sign in as the test account from a private window, then check **Admin console** › **Users**.
- 12 To map groups, copy each group's Object ID from Entra ID › **Groups** and map it in **Admin console** › **SSO**.

### Check that it worked

- ✓ **Test connection** passes the discovery document and signing keys for login.microsoftonline.com.
- ✓ The test account signs in through the Microsoft sign-in page and reaches the workspace.
- ✓ **Admin console** › **Users** shows the account with sso under Auth.
- ✓ After group mapping, the account's groups match its Entra groups at the next sign-in.



## Troubleshooting

| WHAT YOU SEE                                                                            | WHAT TO DO                                                                                                                                                            |
|-----------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>AADSTS50011: the redirect URI does not match</b>                                     | Add the exact redirect URI under Authentication on the <b>Web</b> platform, not as a single-page application.                                                         |
| <b>AADSTS7000215: invalid client secret</b>                                             | You copied the Secret ID. <b>Create</b> a new secret and copy its Value.                                                                                              |
| <b>AADSTS7000222: the client secret has expired</b>                                     | <b>Create</b> a new secret, paste it in <b>Single Sign-On</b> , and choose <b>Save SSO</b> . Track the new expiry date.                                               |
| <b>AADSTS700016 (app not found) or AADSTS50194 (single-tenant app used with common)</b> | Check the client ID, and use <a href="https://login.microsoftonline.com/your-tenant-id/v2.0">https://login.microsoftonline.com/your-tenant-id/v2.0</a> as the issuer. |
| <b>AADSTS50105: the user is not assigned to the app</b>                                 | Assign the person or one of their groups under Enterprise applications > <b>Users</b> and groups.                                                                     |
| <b>AADSTS65001 or AADSTS90094: consent is required</b>                                  | An administrator grants consent under API permissions > Grant admin consent for your tenant.                                                                          |
| <b>AADSTS53003: blocked by Conditional Access</b>                                       | Review the Conditional <b>Access</b> policy that targets the app; test policies in Report-only first.                                                                 |
| <b>A person signs in but receives no mapped groups</b>                                  | Entra omits the groups claim above 200 groups. Use <b>Groups</b> assigned to the application, and remember it covers direct members only.                             |

### WALKTHROUGH · 2:05 · PLATFORM CONSOLE > DOCUMENTATION

## Single sign-on with Okta

Create an OIDC web app integration in Okta, choose the authorization server, add a groups claim, and connect it with the Okta preset.

### Before you begin

- An Okta administrator account that can create app integrations.
- Your Okta domain, such as your-org.okta.com.
- API **Access Management**, if you will use the default custom authorization server in production (needed for a groups claim without an extra scope).
- The redirect URI copied from Platform console > **Org Settings** > **Single Sign-On**.
- A test account assigned to the app, on an allowed email domain.



## Step by step

- 1 In Aperture Chat, open **Org Settings** > **Single Sign-On** and copy the redirect URI.
- 2 Okta Admin Console > Applications > Applications > **Create** App Integration. If asked, choose Classic experience.
- 3 Sign-in method: OIDC - OpenID Connect. Application type: **Web** Application. Select Next.
- 4 Name the integration, keep Grant type Authorization **Code**, paste the redirect URI under Sign-in redirect URIs, and choose Limit access to selected groups. Select Save.
- 5 On the General tab, copy the **Client ID** and the **Client secret**. Client authentication stays **Client secret**; Aperture Chat uses client\_secret\_basic first.
- 6 Choose the issuer. Org authorization server: https://your-org.okta.com (no groups claim with the default scopes). Default custom server: https://your-org.okta.com/oauth2/default.
- 7 For group mapping on the custom server: **Security** > API > Authorization Servers > default > Claims > Add Claim. Name groups, include in ID Token Always, Value type **Groups**, a narrow filter such as Starts with aperture-, include in Any scope.
- 8 On the same server, make sure an access policy covers this app with a rule that allows the Authorization **Code** grant.
- 9 Assignments tab > Assign > Assign to **Groups**: assign the groups that should use Aperture Chat.
- 10 In Aperture Chat, choose the Okta preset, enter your issuer, paste the client ID and secret, list domains, and keep the group claim as groups.
- 11 Choose **Save SSO**, then **Test connection**. Sign in as the test account from a private window, then check **Admin console** > **Users**.

## Check that it worked

- ✓ **Test connection** passes for your Okta issuer.
- ✓ The test account signs in through Okta and reaches the workspace.
- ✓ With the custom server's groups claim, mapped workspace groups appear at the next sign-in.

## Troubleshooting

| WHAT YOU SEE                                                                           | WHAT TO DO                                                                                                                                                                                     |
|----------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "The 'redirect_uri' parameter must be a Login redirect URI in the client app settings" | Add the exact redirect URI under Sign-in redirect URIs on the app's General tab.                                                                                                               |
| "User is not assigned to the client application"                                       | Assign the person or one of their groups on the app's Assignments tab.                                                                                                                         |
| "Policy evaluation failed for this request"                                            | The custom authorization server has no active access policy and rule covering this app and the Authorization <b>Code</b> grant.                                                                |
| "Token endpoint returned 401: invalid_client"                                          | The client secret is wrong or inactive. <b>Copy</b> an active secret from the General tab and save it again.                                                                                   |
| "The groups claim matched too many groups"                                             | Narrow the claim's filter, for example to groups that start with aperture-.                                                                                                                    |
| People sign in but no mapped groups arrive                                             | Use the default custom authorization server with a groups claim included in Any scope. The org server sends groups only for the groups scope, which Aperture Chat does not request by default. |

### WALKTHROUGH · 1:38 · PLATFORM CONSOLE > DOCUMENTATION

## Single sign-on with Google Workspace

Create an internal OAuth client in Google Auth Platform, connect it with the Google preset, and test Google's public issuer.

### Before you begin

- A Google Cloud project that belongs to your Workspace or Cloud Identity organization (required for an Internal audience).
- Permission to configure Google Auth Platform in that project.
- The redirect URI copied from Platform console > **Org Settings** > **Single Sign-On**. Google requires HTTPS for production redirect URIs.
- A Workspace test account on an allowed email domain.

### Step by step

- 1 In Aperture Chat, open **Org Settings** > **Single Sign-On** and copy the redirect URI.
- 2 Google Cloud console > select your organization's project > Google Auth Platform > Branding. If asked, choose Get started and enter the app name and user support email.
- 3 Audience: User type Internal, so only your organization can sign in.

- 4 Clients > **Create** client > Application type **Web** application. Name it, add the redirect URI under Authorized redirect URIs, and select **Create**.
- 5 **Copy** the **Client ID** and **Client secret** from the dialog (or download the JSON). Google shows the full secret only once.
- 6 In Aperture Chat, choose the Google Workspace preset. The issuer is <https://accounts.google.com>.
- 7 Paste the client ID and secret, list your Workspace domains, and leave **Enforce SSO** off.
- 8 Choose **Save SSO**, then **Test connection**. Sign in as the test account from a private window and check **Admin console** > **Users**.
- 9 Google ID tokens include no groups. Assign workspace groups to people in **Admin console** > **Users** or **Groups**.

### Check that it worked

- ✓ **Test connection** passes the discovery document and signing keys for <https://accounts.google.com>.
- ✓ The test account signs in with Google and reaches the workspace.
- ✓ Someone outside your organization is refused by Google because the audience is Internal.

### Troubleshooting

| WHAT YOU SEE                                                           | WHAT TO DO                                                                                                                 |
|------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------|
| <b>Error 400: redirect_uri_mismatch</b>                                | Add the exact redirect URI under the client's Authorized redirect URIs.                                                    |
| <b>Access blocked: org_internal</b>                                    | The account is outside your organization. Internal clients accept only your Workspace accounts.                            |
| <b>"Token endpoint returned 401: invalid_client" or deleted_client</b> | The secret is wrong, or Google deleted a client unused for six months. Add a new secret or client and save it.             |
| <b>admin_policy_enforced</b>                                           | A Workspace administrator blocked the app. Trust it under <b>Security</b> > <b>Access</b> and data control > API controls. |
| <b>Mapped groups never change</b>                                      | Expected with Google: its ID tokens carry no groups claim. Manage groups in the <b>Admin console</b> .                     |

## 49. Single sign-on: groups, MFA, enforcement, and recovery

Map identity-provider groups, choose the MFA rule, enforce SSO for your domains, and keep a way back in.

- Provision new users on first sign-in (JIT) creates an account the first time someone on an allowed domain signs in, with the USER role and the default group. Turn it off to admit only accounts an administrator created first.
- Existing accounts on an allowed domain link to the provider the first time that person signs in through **SSO**. If the provider later presents the same email for a different identity, sign-in stops until an administrator resets the account.
- Group mapping makes **SSO** the owner of membership in each mapped workspace group: the person's membership follows the token's group claim at every sign-in. **Groups** that are not mapped stay admin-managed.
- **Require the platform authenticator after SSO**: off trusts the identity provider's own MFA; on adds the Aperture Chat authenticator after every **SSO** sign-in. The **Authenticator app**, MFA methods, and QR enrollment fields only record information for people; they enforce nothing.
- **Enforce SSO** for these domains refuses local passwords for accounts on the allowed domains. **Require SSO for admins** under **Policy Controls** does the same for every tenant admin.

WALKTHROUGH · 2:08 · PLATFORM CONSOLE › DOCUMENTATION

### Go live: groups, MFA, and enforcement

Map identity-provider groups, prove they sync, choose the MFA rule, enforce SSO, and recognize the refusals people may see.

#### Before you begin

- A working **SSO** connection: **Test connection** passes and a test account has signed in.
- Workspace groups that match how you want to grant access.
- A groups claim in the ID token (not available with Google).
- A platform owner account that signs in with a local password on a domain you will not enforce.

#### Step by step

- 1 **Admin console** › **SSO**: on the provider's card, choose Add mapping for each identity-provider group. Enter the claim value exactly as the provider sends it (a name, path, or object ID) and choose the workspace group.
- 2 Choose Save mappings. **SSO** now owns membership in the mapped workspace groups; other groups stay admin-managed.
- 3 Sign in again as the test account, then check **Admin console** › **Users**: the mapped group appears.
- 4 Decide on **Require the platform authenticator after SSO**. Off trusts the provider's MFA; on adds the Aperture Chat authenticator after every **SSO** sign-in.

- 5 Tell people on the domain that password sign-in is ending, and confirm everyone can sign in through the provider.
- 6 Platform console > **Org Settings** > **Single Sign-On**: turn on **Enforce SSO** for these domains and choose **Save SSO**.
- 7 Confirm that a local password on the domain is now refused, and that your owner recovery account still signs in.

### Check that it worked

- ✓ After a fresh sign-in, the person's groups in **Admin console** > **Users** match their identity-provider groups.
- ✓ **Audit** records `auth.sso_groups_synced` when membership changes.
- ✓ A local password on the enforced domain is refused with "**SSO** is enforced for this email domain; local sign-in is disabled."
- ✓ Your owner recovery account still signs in with its password.

### Troubleshooting

| WHAT YOU SEE                                                                 | WHAT TO DO                                                                                                                                                                            |
|------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Mapped groups do not appear</b>                                           | The claim value must match exactly, including case and any leading slash. Check the group claim name in <b>Single Sign-On</b> , and that the provider adds the claim to the ID token. |
| <b>"SSO is enforced for this email domain; local sign-in is disabled."</b>   | Expected after enforcement. Use <b>Continue with SSO</b> . To restore password sign-in, an owner turns enforcement off.                                                               |
| <b>"... is outside the domains allowed for this SSO provider."</b>           | The provider vouched for an address on another domain. Add the domain, or keep that person out.                                                                                       |
| <b>"This email now belongs to a different identity at the provider."</b>     | The provider reissued the address to someone new. An administrator resets the account's password, which re-binds it at the next <b>SSO</b> sign-in.                                   |
| <b>"Organization policy requires admin accounts to sign in through SSO."</b> | <b>Policy Controls</b> requires <b>SSO</b> for admins. Admins use <b>Continue with SSO</b> ; owners keep their local recovery path.                                                   |

### Provisioning with SCIM

**SCIM** 2.0 user provisioning is served at `https://your-instance.example/scim/v2` (**Users**, plus a read-only **Groups** list), and it refuses every request until a bearer token exists. There are two ways to create one, neither in the console yet. For any deployment, a platform owner mints a tenant token with POST

/api/platform/tenants/{tenant}/scim-tokens; the response shows the token once, and DELETE on the same path revokes it. A single-tenant deployment may instead set APERTURE\_SCIM\_BEARER\_TOKEN and restart the API. Point the identity provider's **SCIM** app at the base URL with that bearer token. **SCIM** creates, updates, and deactivates accounts; people still sign in through OIDC.

## 50. Org Settings: platform branding

Rename the product, swap the logos, and recolor the theme everywhere.

### GOOD TO KNOW

Branding reaches further than the shell: the same identity feeds the runtime theme colors and the per-tenant install manifest and icons, so the app people add to their phone's home screen carries your name and logo, not a generic one.

### WALKTHROUGH · 1:58 · PLATFORM CONSOLE · DOCUMENTATION

#### Platform branding

Give the platform your organization's name, logo, and colors, apply them everywhere, and check each place the brand appears.

#### Before you begin

- A platform owner account.
- A square **PNG** logo under 4 MB, or public HTTPS addresses for a logo and a browser icon.
- Your colors as 6-digit hex values, such as #2f5d8a. Pick dark sidebar colors so the light sidebar text stays readable.

#### Step by step

- 1 Open Platform console > **Org Settings** and expand **Platform Branding**. The preview at the top shows the logo, name, and an Interface text sample.
- 2 Enter the Platform name. It replaces Aperture Chat in the sidebar, the sign-in page, the browser tab, and the installed app.
- 3 Choose **Upload PNG** and pick your logo. It fills both **Platform logo URL** and **Browser icon URL**. Or paste HTTPS addresses into those fields instead.
- 4 Enter the **Platform domain** people use. It is recorded for administrators and the API; it does not change DNS, TLS, or routing.
- 5 Under **Theme colors**, set **Accent color**, **Sidebar gradient start**, **Sidebar gradient end**, and **Interface text color**. Leave a field empty to keep the default for that surface.
- 6 Choose **Apply branding**. The message ... branding saved through the platform API and will persist across reloads confirms it.

- 7 Check the sidebar, open the sign-in page in a private window, and on a phone add the workspace to the home screen to see the new icon and name.
- 8 To undo everything, choose **Reset defaults**. It saves the Aperture Chat defaults immediately.

#### Path: Upload a PNG (shown in the video)

- 1 **Upload PNG** accepts **PNG** files up to 4 MB and shrinks them for the sidebar and icons.
- 2 The upload fills **Platform logo URL** and **Browser icon URL** with the stored image; nothing needs hosting elsewhere.
- 3 Choose **Apply branding** to save it.

#### Path: Use hosted image addresses

- 1 Paste an HTTPS address into **Platform logo URL** for the sidebar and sign-in mark.
- 2 Paste an HTTPS address into **Browser icon URL** for the tab icon (it falls back to the logo when empty).
- 3 Choose **Apply branding**.

#### Check that it worked

- ✓ The confirmation reads ... branding saved through the platform API and will persist across reloads, and the fields keep their values after a reload.
- ✓ The sidebar shows the new name, logo, and gradient.
- ✓ A signed-out visitor sees the new name and logo on the sign-in page, including New to ... under the form.
- ✓ The installed app icon (/api/pwa/icon-512.png) shows the uploaded logo, and the manifest's name is the new platform name.

## Troubleshooting

| WHAT YOU SEE                                                                               | WHAT TO DO                                                                                                                      |
|--------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------|
| "Set both gradient colors (or clear both) so the sidebar gradient has a start and an end." | Fill in both <b>Sidebar gradient start</b> and <b>Sidebar gradient end</b> , or clear both.                                     |
| "... is not a 6-digit hex color like #087d8b. Fix it before applying."                     | Use a # followed by exactly six hex digits, or use the color picker beside the field.                                           |
| "Upload a PNG image for the platform logo and browser icon."                               | <b>Export</b> the logo as <b>PNG</b> ; other formats are refused.                                                               |
| "That PNG is larger than 4 MB. Export a smaller logo and try again."                       | <b>Export</b> a smaller <b>PNG</b> ; a few hundred pixels square is plenty.                                                     |
| The sidebar text is hard to read.                                                          | Choose darker gradient colors, or clear both gradient fields to restore the default.                                            |
| A phone still shows the old icon.                                                          | Icons are cached for a few minutes, and an installed app keeps its icon until it is removed and added to the home screen again. |

## 51. Org Settings: policies and budgets

Set the organization ceiling and the workspace usage budget.

### Policy Controls

Expand **Policy Controls** in **Org Settings** when you need it. The first row is not a toggle at all: "Only owners can create platform owners" carries an **Always on** pill because the platform enforces it unconditionally. Below it are eight real switches:



| POLICY                                       | WHAT IT ALLOWS WHEN ON                                                                                                                                                                                                                                                       |
|----------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Downstream API access</b>                 | Owners and admins can use the platform API downstream; standard users still need the group grant <b>Can use API</b> from their admin.                                                                                                                                        |
| <b>Tenant admins can create admins</b>       | Delegates admin creation to tenant admins; off means owner-only.                                                                                                                                                                                                             |
| <b>Require SSO for admins</b>                | Admin accounts must sign in through <b>SSO</b> , not passwords.                                                                                                                                                                                                              |
| <b>Tenant admins can manage SSO mappings</b> | Lets tenant admins edit their tenant's <b>SSO</b> configuration and group mappings.                                                                                                                                                                                          |
| <b>Default group for enabled models</b>      | Newly enabled models automatically include the protected <b>Default Users</b> group, so admins do not need a second step per model.                                                                                                                                          |
| <b>Users can build their own agents</b>      | The first half of a two-part gate: with this on, admins can grant <b>Can build agents</b> to a group, and those users can build private, self-owned agents. Publishing to the organization stays admin-only. Both halves default off.                                        |
| <b>Personalization memory</b>                | Enables the tenant-admin handoff. Admins may then turn memory on for their organization and grant <b>Default Users</b> access. Off locks both downstream controls while preserving saved grants. Admins and owners see counts and purge controls only, never memory content. |
| <b>Users can browse the model catalog</b>    | <b>Users</b> can inspect all models enabled for their organization, read availability reasons, and request access where permitted. Off shows only already-usable models and refuses access requests. This does not reveal private prompts or notes.                          |

#### GOOD TO KNOW

Active policies define the organization ceiling. Tenant admins can only operate inside these boundaries — nothing they configure can exceed them.

## Workspace Usage Budget

- The **Workspace Usage Budget** panel sets a hard spending ceiling for the whole workspace. Choose the **Budget measure** — **Token allowance** (exact provider-reported tokens) or **Dollar amount (USD)** (exact provider-reported cost) — and the **Reset period: Every day, Every week, or Every month**, on the UTC calendar (weeks begin Monday).
- Enter the limit and click **Save budget policy**. A limit of 0 means unlimited. The usage card beside the form shows the live count for the current period, split into provider-reported and unreported completions so the number stays honest.

- **Once** the ceiling is spent, further completions are refused until the period resets. Admins can add per-user and per-group allocations beneath this ceiling from their **Analytics** tab.

## WALKTHROUGH · 6:34 · PLATFORM CONSOLE > DOCUMENTATION

### Policies, budget, and connectors

Set the organization's policy ceiling, cap workspace usage with a budget, and configure and test the shared connectors people use for files and web search.

#### Before you begin

- A platform owner account.
- For connectors: an administrator at the vendor (Google Cloud, Microsoft Entra, Box, or iManage) who can register an application and grant consent.
- Your instance's API address. Connector redirect URLs are that address followed by `/api/connector-oauth/callback`.

#### Step by step

- 1 Open Platform console > **Org Settings**. Each area is a collapsed panel; expand only the one you need.
- 2 Expand **Policy Controls**. Read each switch's description, which states what is true right now, then turn on or off the switches you need. Each change saves immediately and shows a confirmation.
- 3 Expand **Workspace Usage Budget**. Choose the **Budget measure** (**Token allowance** or **Dollar amount (USD)**), the **Reset period** (**Every day**, **Every week**, or **Every month**, in UTC), and the limit. 0 means unlimited.
- 4 Choose **Save budget policy** and read the confirmation, for example Workspace ceiling saved at 1,000 tokens daily.
- 5 Expand **Connectors**. Turn each source on or off with its switch; off removes it for everyone, including the API.
- 6 For a connector with credentials, choose **Configure**, pick the Authentication method, fill in the vendor's values and secret, and choose Save configuration.
- 7 Choose **Test connection** and read every line of the result. For Google Drive, finish with **Connect Google Drive** to authorize the workspace account.

#### Path: Policy Controls: what each switch changes

- 1 Only owners can create platform owners: always on; no switch can loosen it.
- 2 **Downstream API access**. On: owners and tenant admins can create personal pass-through keys; standard users still need an administrator's **Can use API** grant. Off: keys cannot authenticate and admins cannot grant new access.

- 3 **Tenant admins can create admins.** On: tenant admins may create and manage other tenant admins. Off: they can create standard users only.
- 4 **Require SSO for admins.** On: password fallback is rejected for admin accounts. Off: admins may sign in with a local password when one is provisioned.
- 5 **Tenant admins can manage SSO mappings.** On: tenant admins can create and update **SSO** configurations. Off: only platform owners may change them.
- 6 **Default group for enabled models.** On: newly enabled models include **Default Users**. Off: admins grant groups per model.
- 7 **Users can build their own agents** (on: groups granted **Can build agents** may create private agents), **Personalization memory** (off: no memory anywhere), and **Users can browse the model catalog** (see Organization model availability).

#### Path: Budget in tokens

- 1 **Budget measure: Token allowance.** It counts exact provider-reported tokens.
- 2 **Reset period: Every day, Every week** (weeks begin Monday), or **Every month**, all in UTC.
- 3 Token limit: the ceiling; 0 for unlimited. **Save budget policy.**
- 4 When the ceiling is reached, people see: The workspace daily token budget has been reached. Requests are blocked until the next UTC day.

#### Path: Budget in dollars (USD)

- 1 **Budget measure: Dollar amount (USD).** It uses exact provider-reported cost; unreported cost is shown separately.
- 2 Enter the Dollar limit and choose **Save budget policy.**
- 3 Administrators can add per-user and per-group allocations beneath this ceiling in the **Admin console.**

#### Path: Web Search

- 1 **Configure > Search** engine: DuckDuckGo (keyless), SearXNG (your own instance; enter its URL and allow JSON output), or OpenAI, Anthropic, or OpenRouter web search (reuses that provider's key; each search is billed to that account).
- 2 Results per search: how many results are added to the prompt as cited context.
- 3 Save configuration, then **Test connection** runs a real search.

#### Path: Google Drive

- 1 In Google Cloud, enable the Drive API, set up the Google Auth Platform, add the drive.readonly scope, and create a **Web** application client with the redirect URI.

- 2 **Configure** › Google OAuth (recommended): OAuth client ID, optional Drive folder ID and Source label, and the OAuth client secret. Save configuration.
- 3 Choose **Connect Google** Drive and sign in with the workspace account used for knowledge sync. Then **Test connection**.
- 4 **Chat** users connect their own Google account from the attach menu and see only their own files.

### Path: OneDrive / SharePoint / Outlook (Microsoft Graph)

- 1 In Entra, register a single-tenant app, add Files.Read.All (and Sites.Read.All) as application permissions, grant admin consent, and create a client secret.
- 2 **Configure** › App-only (client credentials, recommended): Directory (tenant) ID, Application (client) ID, optional SharePoint site, Drive, and Root folder IDs, and the **Client secret**.
- 3 Save configuration, then **Test connection** requests a token and calls Graph.
- 4 For chat attachments, also add the scopes as delegated permissions and the **Web** redirect URI.

### Path: Box

- 1 In the Box Developer Console, create a Server app with Client Credentials Grant, App + Enterprise **Access**, and Read all files and folders; submit it, and have a Box admin authorize it in Admin Console › Platform › Platform Apps.
- 2 **Configure** › Client Credentials Grant (recommended): **Client ID**, Enterprise ID, optional Folder ID, and the **Client secret**.
- 3 Save configuration, then **Test connection**.

### Path: iManage

- 1 In iManage Control Center, add a **Web**, confidential application with the redirect URI and refresh tokens allowed.
- 2 **Configure** › Each user signs in (recommended): Instance URL, API key (client ID), Customer ID, **Library** ID, optional Workspace ID, and the OAuth client secret.
- 3 Save configuration. **Test connection** confirms the fields; each person completes verification by opening iManage from chat and signing in.
- 4 Service account for background sync is only for knowledge sync; chat always uses each person's own sign-in.

### Check that it worked

- ✓ Each policy change shows its own confirmation, and the switch's description now describes the new state.
- ✓ With **Downstream API access** on, an administrator's **Account** panel shows API access with **Create** key.

- ✓ The budget panel shows the saved ceiling and the current period's usage; a person who exceeds it gets the budget message instead of a reply.
- ✓ A connector's **Test connection** lists Required fields, Authentication, and API access as passed.

## Troubleshooting

| WHAT YOU SEE                                                                                                                             | WHAT TO DO                                                                                                                                                            |
|------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "Web search test failed: DuckDuckGo returned a page with no parseable results (it rate-limits automated queries)."                       | DuckDuckGo limits automated searches from busy servers. Retry later, or switch to SearXNG or a provider's web search.                                                 |
| "OpenAI web search reuses the workspace's OpenAI provider key. Add an active OpenAI key under Providers, or switch to a keyless engine." | Add an OpenAI key in <b>Providers</b> , or choose DuckDuckGo or SearXNG.                                                                                              |
| "Google Drive is not connected: no refresh token is stored ..."                                                                          | Save the client ID and secret, then choose <b>Connect Google Drive</b> and finish Google's consent screen.                                                            |
| "Microsoft Graph token request failed: token endpoint returned HTTP 400: AADSTS90002: Tenant '...' not found."                           | The Directory (tenant) ID is wrong. <b>Copy</b> it from the app's Overview page in Entra.                                                                             |
| "Box token request failed: token endpoint returned HTTP 400: The client credentials are invalid"                                         | Check the <b>Client ID</b> and secret, and that a Box admin authorized the app in Platform Apps.                                                                      |
| People keep seeing Aperture Chat is working for several minutes, then the budget message.                                                | The workspace budget is spent. Raise or clear the ceiling, or wait for the next UTC period. The chat currently retries for a few minutes before it shows the message. |
| "Policy was not saved: ..."                                                                                                              | The change was not stored and the switch returns to its previous state. Read the reason and try again.                                                                |

## Elastic Analytics

The **Elastic Analytics** panel sends platform activity to your Elastic cluster so it can be searched and monitored in Kibana. Its status line says whether export is connected, paused, or off.

[WALKTHROUGH · 3:04 · PLATFORM CONSOLE > DOCUMENTATION](#)

### Elastic Analytics export

Send Aperture Chat's audit trail, usage, chats, documents, and users to Elastic Cloud or your own Elasticsearch, and search them in Kibana.

## Before you begin

- A platform owner account.
- An Elasticsearch cluster: an Elastic Cloud deployment (its Cloud ID or endpoint) or a self-managed cluster the server can reach.
- Permission to create an API key in that cluster (Kibana Dev **Tools** or Stack **Management** › API keys).
- Kibana access with permission to import saved objects.

## Step by step

- 1 Open Platform console › **Org Settings** and expand **Elastic Analytics**. Read the status line at the top.
- 2 Paste the Elasticsearch endpoint (for example `https://your-cluster.example:9200`) or your Elastic Cloud deployment's Cloud ID.
- 3 Open How do I create an API key?, choose **Copy** request, and run it in Kibana Dev **Tools**. **Copy** the response's encoded value.
- 4 Paste the key into API key, and keep or change the Index prefix (aperture by default).
- 5 Under What to send, turn each stream on or off: **Audit** trail, **Model** usage, **Chats**, Documents, and **Users**. Leave Include message and document text off unless you need the text in Elastic.
- 6 Choose Save and check. Every check must pass: Reach the cluster, API key accepted, Cluster details, and Can write Aperture indices.
- 7 Choose Sync now. The Delivery table shows Sent counts and Last sent for each index.
- 8 Choose Kibana data views, then in Kibana open Stack **Management** › Saved objects › **Import**, choose the downloaded file, and **Import**.
- 9 Open Discover, pick an Aperture data view, and widen the time range to see the records.

### Path: Elastic Cloud (Cloud ID)

- 1 In the Elastic Cloud console, open your deployment and copy its Cloud ID (or the Elasticsearch endpoint).
- 2 **Create** the API key in Kibana: Dev **Tools** with the panel's request, or Stack **Management** › API keys with the same privileges.
- 3 Paste the Cloud ID and key into **Elastic Analytics**, then Save and check.

### Path: Self-managed Elasticsearch (shown in the video)

- 1 Use the cluster's HTTP address, usually port 9200, reachable from the Aperture Chat server.
- 2 Run the panel's API key request in Kibana Dev **Tools** (or with curl as a user who can manage API keys).

- 3 Paste the endpoint and key, then Save and check.

### Path: Pause, resume, or resend

- 1 **Export** on pauses or resumes delivery without losing settings.
- 2 Re-send everything sends all history again, for example after switching clusters.
- 3 Check connection tests the form as typed, without saving it.

### Check that it worked

- ✓ Save and check shows Saved and verified, and all four checks are green.
- ✓ The Delivery table shows 0 waiting, a Sent count, and Last sent for each stream.
- ✓ Kibana's import reports the Aperture data views imported, and Discover lists the exported records, including the Elastic changes you just made.

### Troubleshooting


| WHAT YOU SEE                                                                                                | WHAT TO DO                                                                                                                           |
|-------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------|
| <b>Delivery failing: "Could not connect to ...: [Errno 61] Connection refused."</b>                         | Nothing answered at that address and port. Self-managed Elasticsearch usually listens on 9200; Elastic Cloud uses 443.               |
| <b>"Could not connect to ...: [Errno 8] nodename nor servname provided, or not known." after a Cloud ID</b> | The Cloud ID does not point to a reachable deployment. <b>Copy</b> it again from the deployment's page in the Elastic Cloud console. |
| <b>API key accepted fails</b>                                                                               | Paste the encoded value from the API key response (or the whole response), not the key id.                                           |
| <b>Can write Aperture indices fails</b>                                                                     | Recreate the key with the panel's request: create_index, index, and read on aperture-* (or your prefix), plus monitor.               |
| <b>Sync now is greyed out.</b>                                                                              | Save your changes first, and make sure <b>Export</b> on is turned on.                                                                |
| <b>Discover shows no results.</b>                                                                           | Widen the time range; Discover starts at the last 15 minutes. Check the Delivery table for errors.                                   |

**Chat** records carry their retention tags and legal holds, and tag, hold, archive, and matter changes follow within about a minute. Deleted chats, users, and documents stay in Elastic flagged as deleted. Operators can instead set APERTURE\_ELASTIC\_URL or APERTURE\_ELASTIC\_CLOUD\_ID with APERTURE\_ELASTIC\_API\_KEY on the server.



## 52. Org Settings: shared connectors and source credentials

Configure deployment-wide availability, test real connections, and distinguish shared access from personal sign-in.



Shared connector configuration is in the Platform Owner console under Org Settings. Example workspace with synthetic data.

Open **Platform owner console** → **Org Settings** and expand **Connectors**. This owner-only panel controls shared source settings and availability. Turning a connector off removes that capability across chat, source pickers, the command palette, the **Tools** library, and the API. Tenant administrators do not configure these shared connections.

Credential-backed sources show **Credentials saved**, **Saved · disabled**, or **Needs credentials**. A saved credential is configuration evidence; read **Test connection** for the live result. Switch-only capabilities, such as MCP Servers and Prompt **Library**, have an enable switch without a vendor credential form.

### Configure and test a source

- Choose **Configure** on the source row. Select Authentication method and fill in the fields shown for that method. Read the source's setup notes for its required permissions and redirect URI.
- Choose Save configuration. For an existing saved secret, leave its password field blank to retain it, or enter a new value to replace it. Wait for the save result before testing.
- For Google OAuth, save the client ID and secret first, then choose **Connect Google** Drive to authorize the workspace account used for knowledge sync. Complete the consent flow before checking the connection.
- Choose **Test connection** and read its result and individual checks. An incomplete, failed, or missing result is not proof of access. Resolve the reported issue before relying on the source.



- Set the source's enable switch for the deployment and verify the resulting status. **Configure** the intended users' source-account access separately when they will attach files in chat.

| SOURCE                | CONFIGURATION SHOWN BY THE FORM                                                                                                                                                                               |
|-----------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Google Drive          | Google OAuth client ID and secret, with an optional Drive folder ID and source label. Paste access token is a testing option.                                                                                 |
| OneDrive / SharePoint | Microsoft Graph directory ID, application ID, and client secret for app-only access, with optional site, drive, and root-folder IDs.                                                                          |
| Box                   | <b>Client ID</b> , enterprise ID, and client secret for Client Credentials Grant; an optional folder ID limits the starting location. Developer token is a testing option.                                    |
| iManage               | Instance URL, API key (client ID), and OAuth client secret for Each user signs in. Service account for background sync adds the service username and password; chat still requires each user's OAuth sign-in. |

#### GOOD TO KNOW

Google, Microsoft, Box, and iManage chat attachments use each signed-in user's delegated source account. Shared knowledge-sync credentials do not replace that sign-in or bypass source permissions. **Users** keep their Connect action in the attach menu.

## Clear saved configuration

Clear configuration removes the source's saved fields, stored secret, shared OAuth data, and service-account password, and disables its saved configuration. Use it only when that removal is intended. An empty secret field alone retains the saved secret; clearing all visible fields and saving an existing configuration also performs a clear.

## Web Search

Choose **Configure** on **Web Search**, then select **Search** engine and Results per search (1 to 10). DuckDuckGo is keyless. SearXNG requires an instance URL with JSON search output enabled. OpenAI, Anthropic, and OpenRouter choices use the corresponding saved provider key and bill searches to that provider account. Choose Save configuration, then **Test connection** to run a real query.

The engine choice applies to models without hosted web search of their own. OpenRouter-backed models use OpenRouter's built-in search regardless of this selection. The **Web Search** enable switch still governs availability for the deployment.

## 53. Org Settings: search index readiness

Review index coverage and rebuild from live records when search needs maintenance.

### GOOD TO KNOW

If indexing is disabled for this deployment, the panel explains that searches scan records on each request and disables **Rebuild index**. An index entry count is a coverage signal, not a count of items every user may access.

### WALKTHROUGH · 1:02 · PLATFORM CONSOLE > DOCUMENTATION

#### Review workspace search readiness

Check that saved chats and drafts are searchable, rebuild the index when results are missing, and confirm what people can find.

#### Before you begin

- A platform owner account.
- A person's account to check search from, or a second browser signed in as one.

#### Step by step

- 1 Open Platform console > **Org Settings. Search index** is the first panel.
- 2 Read each organization's line: Ready or **Backfilling**, Indexed items, and Last rebuilt.
- 3 If saved content is missing from search results, choose **Rebuild index** and wait for Rebuilt ... index entries from live records.
- 4 Ask a person to press Ctrl + K (⌘ + K on a Mac), or choose **Search** in the sidebar, and search for a word from one of their chats.
- 5 Confirm their results come only from content they can open. Another person's private chat never appears.

#### Path: Index disabled on this server

- 1 The panel says The search index is disabled. **Search** still works by checking saved records directly.
- 2 **Rebuild index** is greyed out; search still works, only more slowly on large workspaces.

#### Check that it worked

- ✓ The status reads Ready, with a current Last rebuilt time.
- ✓ The rebuild message reads Rebuilt ... index entries from live records.
- ✓ A person's search lists their own chats, automations, and drafts that contain the word, and nothing private to anyone else.

## Troubleshooting

| WHAT YOU SEE                                                      | WHAT TO DO                                                                                                                                   |
|-------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------|
| Status shows <b>Backfilling</b> .                                 | The index is still catching up, so results can be incomplete for a while. Wait, or choose <b>Rebuild index</b> .                             |
| "The rebuild did not complete." or another error under the title. | Nothing was confirmed. <b>Try again</b> ; if it repeats, check the server logs.                                                              |
| A deleted chat still appears in results.                          | <b>Search</b> checks every hit against the live record and drops deleted ones; refresh the search. A rebuild never restores deleted content. |
| A person cannot find a chat someone else owns.                    | That is expected: search shows only content the person has permission to open. Rebuilding does not change permissions.                       |

## 54. Review and install a platform release

Use the owner-only sidebar update notice to review release notes, install when available, and verify the reported result.

For platform owners, the sidebar says Release available followed by a version when the updater is unavailable, and Update to followed by a version only when it is ready. It can also show an active update or a recent undismissed result. Tenant administrators and users do not receive this control. The current version comes from the running build; a stale deployment version setting does not establish which build is running.

- 1 When the update row appears, hover over or focus it to read release highlights and the current version. Click it to open the update dialog.
- 2 Review What this update brings, and use Show full release notes or Release page when offered. Check the Release list checked timestamp and any error. Check again requests a fresh release list when no update is running.
- 3 Read How the update runs before proceeding. Installation restarts the API and web app, so arrange an appropriate interruption window and follow the deployment's backup procedure.
- 4 Choose Install followed by the target version only when ready. The button is available when the updater service is configured and connected. If the dialog shows Manual install on this deployment or an offline-updater message, follow those deployment instructions instead. A fresh VPS can use the release bundle installer to prepare private configuration and a stable project, then pull, start, and health-check the stack. Existing installations need the documented one-time API and updater-sidecar configuration in the same project; do not run a fresh-install workflow over existing data. Forks use their own repository release source by default. See the repository Docker release guide for the exact commands and prerequisites.
- 5 Follow the recorded progress through downloading, restarting, and verifying. Show updater log opens available detail. **Closing** the dialog does not cancel an update that has already started.

- 6 After Update installed, choose Reload now to load the new web build, then verify the workspace. If the result is Update failed or Update rolled back, inspect the reported details before using an offered Retry action. Dismiss hides a finished result; it does not install or repair a release.

#### GOOD TO KNOW

A release check or accepted install request is not a completed update. During an active run the API may be temporarily unreachable; the dialog continues polling. Wait for the reported outcome and verify the refreshed application before treating the release as complete.

## 55. Organization retention and tagging

Set the tagging policy and govern conversations within owner access.

In **Platform owner console**, open **Audit**, expand Data Retention, and choose Schedule and rules. The default is Forever, with automatic deletion off. Use the duration slider, review window, client and matter sources, and Preview effect exactly as described in the administrator retention chapter. Deployment never activates an existing metadata-only policy.

#### CAREFUL

A suggested client or sensitive-data label is not a confirmed retention source. Automatic cleanup covers saved chats, linked uploads and image previews, search entries, and chat feedback. **Draft** documents, learned memories, audit and usage records, generated-media storage, backups, external providers, and exported files have separate lifecycles; this control is not a complete client-erasure or regulatory-compliance guarantee.

#### WALKTHROUGH · 2:35 · PLATFORM CONSOLE › DOCUMENTATION

### Data retention and tagging

Keep chats forever or preview a time limit safely, label a client's or matter's chats, and protect them with a legal hold.

#### Before you begin

- A platform owner account (administrators have the same panel in the **Admin console**).
- Your organization's retention requirements, and the client or matter names and aliases people use in chats.
- Approval from whoever owns records policy before you save any time limit.

#### Step by step

- 1 Open Platform console › **Audit** and expand Data Retention, below **Recent Governance Activity**.
- 2 Schedule and rules: Keep chats for starts at Forever, and Automatic deletion off. Forever never deletes anything.

- 3 Under Clients, matters, and regulated records, choose the Source type, enter the Name or reference and its Aliases, separated by commas, and choose Add source to policy.
- 4 To consider a time limit, choose 1, 5, 7, or 10 years and choose Preview effect. Read how many chats are old enough, on legal hold, or without a deadline. Do not save until it is approved.
- 5 Choose Forever again and Save Forever to keep the new source without any deletion.
- 6 Tags and holds: choose Scan existing chats, search for the source, select the right chats, choose the label in Choose a label, and Confirm label.
- 7 Expand Legal holds, enter a Hold name, and choose Hold selected chats.
- 8 Select a chat's title to preview its conversation before deciding anything.

### **Path: Keep everything (Forever, the default)**

- 1 Keep chats for: Forever. Choose Save Forever.
- 2 Sources and labels still work, so chats stay identifiable for later policies or holds.

### **Path: A time limit (after approval)**

- 1 Choose 1, 5, 7, or 10 years, Count age from (Last message change or **Chat** creation), and a Review window of at least 7 days.
- 2 Choose Preview effect and record the counts.
- 3 Only then choose Save retention policy. Eligible chats enter the review window before automatic deletion, and changing the policy restarts that window.
- 4 Longer matching rules and legal holds always win. Increasing the duration never restores deleted chats.

### **Path: Per-source rules**

- 1 Each source has a Retention rule. Use default schedule, or set a longer or shorter period for that client or matter.
- 2 Apply time limits only to labels with a rule leaves unlabeled chats untouched.

### **Check that it worked**

- ✓ Preview effect shows Effect on ... saved chats with the old enough, on legal hold, and without a deletion deadline counts.
- ✓ After Save Forever, the message reads Forever saved. Automatic deletion is off for all chats.
- ✓ The labeled chat shows the source's label, and the held chat shows Legal hold.
- ✓ A second preview counts the held chat under on legal hold.

## Troubleshooting

| WHAT YOU SEE                               | WHAT TO DO                                                                                                                    |
|--------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------|
| "Aliases must be 3–160 characters long."   | Make each alias at least three characters, and separate aliases with commas.                                                  |
| Scan existing chats finds nothing.         | Sources match saved message text and titles. Add the names and aliases people actually type, then scan again.                 |
| Hold selected chats is greyed out.         | Select at least one chat and enter a Hold name.                                                                               |
| A chat you expected to delete was skipped. | Legal holds and longer matching rules always win. Release the hold first; releasing gives eligible chats a new review window. |

## 56. Analytics: runtime, feedback, activity, and usage

Authoritative execution timestamps, feedback, prompt volume, and per-user usage.

The **Analytics** tab carries four panels, and each panel scopes itself: a filter row at the top of every section pairs a user picker with a date range and its presets (All, Today, Week, 30 days), so narrowing one section never hides data in another. The filter's counter reads how many records the scope selected — for example “65 of 65 records”. Three panels carry their own **CSV** export controls, and each **CSV** button opens an independent date-range popover so a file contains exactly the rows you chose.

- **Runtime Clock Metadata** — authoritative execution timestamps captured from chat and draft completion audit events, not client guesses. The scorecards total the runtime events, split between main chat completions and draft generation and revision calls, and each event lists who ran it, which provider served it, how many messages were involved, and the exact client start time.
- **Chat Feedback Analytics** — response-level thumbs up and thumbs down signals submitted from chat actions, totaled by sentiment.
- **Model Activity** — saved prompt volume by model, date, and user for its filter's scope, drawn as “**Prompts by model**”, “Prompt trend”, and “**Users by prompt activity**”.
- **User Usage** — durable per-user usage from real completions across chat, drafts, agents, automations, and the API gateway, with provider-reported token counts only. The owner view is the complete one: its user picker reads **All owners, admins, and users** — unlike the admin console's version, which excludes platform-owner usage. The Usage by user section at the bottom ranks everyone in a contained, scrollable list; pick a person from its selector — or click their row — to focus the whole panel on them.

## Analytics: runtime, activity, and usage

Scope each analytics section to a person and a date range, export it as CSV, open the conversation behind a feedback record, and read the activity and usage charts.

### Before you begin

- A platform owner account.
- Some real activity in the workspace: chats, drafts, feedback, or API calls.

### Step by step

- 1 Open Platform console > **Analytics**. Four sections start collapsed: **Runtime Clock Metadata**, **Chat Feedback**, **Model Activity**, and **User Usage**.
- 2 Expand a section. Each has its own filter: a person picker, From and Through dates, and the presets All, Today, Week, and 30 days. A filter changes only its own section.
- 3 **Runtime Clock Metadata**: read the Runtime events, **Chat**, and **Drafts** totals, and the list of executions with who, which provider, and when.
- 4 Choose **CSV** on a section, set the **CSV** date range (or All dates), and choose **Download**. The file name records the person and range.
- 5 **Chat Feedback**: read the totals, then choose Preview on a rating or issue report to open the conversation or report behind it.
- 6 **Model Activity**: **Prompts** by model, Prompt trend, and **Users** by prompt activity, for the section's filter.
- 7 **User Usage**: Messages, Prompt tokens, Completion tokens, Total tokens, and **Models** used, then Messages by model, Message trend, and Usage by user. Pick a person in Usage by user to focus the section.

### Path: Export a section as CSV

- 1 **Runtime Clock Metadata**, **Chat Feedback**, and **User Usage** each have **CSV**. The export uses its own date range, separate from the section's filter.
- 2 Choose All dates or set From and Through, then **Download**. The button counts the rows it will export.
- 3 Columns include the actor id and name, so the file can be filtered by person.

### Path: Focus on one person

- 1 Pick the person in the section's filter, or in **User Usage** > Usage by user.
- 2 Totals, charts, and lists in that section all follow the choice; other sections do not change.



## Check that it worked

- ✓ The filter line reads, for example, 52 of 119 records, and the totals change to match.
- ✓ The **CSV** downloads with one row per record and a header row (the runtime file starts id, actor\_id, actor\_name, surface...).
- ✓ Preview opens the rated conversation with the person's prompt and the saved model output.

## Troubleshooting

| WHAT YOU SEE                                             | WHAT TO DO                                                                                                   |
|----------------------------------------------------------|--------------------------------------------------------------------------------------------------------------|
| A section shows fewer records than expected.             | Check that section's own filter: person, dates, and preset. Choose All to clear it.                          |
| Token counts are blank.                                  | Token counts are provider-reported only; a provider that reports none leaves them blank rather than guessed. |
| The CSV has a different number of rows from the list.    | The <b>CSV</b> uses its own date range. Set it to match, or choose All dates.                                |
| Chat Feedback says ratings are recorded in this browser. | The platform API is not connected for feedback; server-side feedback loads when it is.                       |

## 57. Automation readiness and governance

The owner-level controls that make scheduled model chains reliable.

Automations depend on the same governance stack as chat: provider health, vaulted keys, enabled models, tenant group grants, connector policy, and the workspace usage budget. Owners do not usually build every tenant's automation, but they set the ceiling that determines whether scheduled runs can work at all.

- 1 In **Providers**, confirm each gateway a schedule depends on shows **Connected** — not **Needs key**, and not **Adapter needed** (a Bedrock provider registers but cannot serve runs until its adapter ships).
- 2 Open **API Keys** on each provider card and replace invalid or expired keys before teams build scheduled workflows around them.
- 3 In **Models**, enable only the routes the organization has reviewed — newly synced models start disabled, and disabled models never flow to tenant admins or user automations.
- 4 In **Org Settings** → **Connectors**, configure and test the shared sources a schedule needs and keep their enable switches aligned with approved use. Check the **Workspace Usage Budget**: once the ceiling is spent, completions — scheduled ones included — are refused (the API answers with HTTP 429) until the period resets, and a refused run is recorded as failed rather than retried automatically.



- 5 In **Analytics** and **Audit**, review automation-related chat/runtime events the same way you review normal completions: who ran it, which provider served it, and whether the run produced the expected output.

**TIP**

For a new deployment, include one automation smoke test: create a weekly one-step run on an approved model, press **Run now**, confirm the transcript succeeds, then restart the API and verify the automation card and last-run status persist.

## 58. Audit: owner governance signals

Grouped signals, Audit Insights trends, prompt monitoring, security alerts, and the platform trail.

The **Audit** tab opens on **Owner Audit** and **Audit** Insights; the sections below start collapsed, and clicking a section header expands it. **User Prompt Activity**, **Security Alerts**, and the **Audit Trail** each carry their own filter row pairing a user picker with a date range, so one section's scope never narrows another. The **CSV** export buttons produce files of exactly the filtered rows, with actor columns, so every exported row names who acted.

### Owner Audit signals

The banner reads, for example, “8 of 24 signals need attention”. Expand all, Collapse all, and the List and Cards switch work as described in the Administrator Guide. The owner board has four groups:

| GROUP                                       | SIGNALS                                                                                                                                                              |
|---------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Security signals</b>                     | <b>Critical events</b> , warning events, after-hours changes, the prompt watchlist, alert response (median time to acknowledge), and failed operations.              |
| <b>Identity &amp; access</b>                | Privileged owners, role changes, password-only admins (tenant admins not using <b>SSO</b> ), access requests, credential changes, and accounts that never signed in. |
| <b>Providers &amp; secrets</b>              | <b>Provider</b> posture, expired keys, keys expiring soon, vault metadata, stale syncs, and providers failing their live validation.                                 |
| <b>Models, connectors &amp; automations</b> | The model ceiling, <b>Connectors</b> , unscoped models, agent approvals awaiting review, connector issues, and automation failures.                                  |

**Groups** with issues open automatically and red rows need attention. Select any signal to open an investigation with every record behind it.

### Audit Insights

**Audit** Insights charts platform-wide trends over 7, 14, or 30 days: audit events by day and severity, security alerts by day, an alert breakdown by rule or by person, the most active people, activity by area, and activity by hour with after-hours time shaded. Selecting any bar, point, or row opens the investigation view with the matching records and a filter box.

## Sections below the dashboard

- **Recent Governance Activity** lists the latest owner-relevant events from the current snapshot — model availability reviews, provider catalog status, vault metadata, **Connector availability**, and **Agent approval activity** — so exceptions become follow-ups instead of surprises.
- **User Prompt Activity** drills into saved user prompts by person, thread, model, and timestamp — the owner-scope view of what is actually being asked across the platform, narrowed by its own user and date filter.
- **Security Alerts** lists DLP and malicious-behavior flags raised from actual prompts, with redacted snippets for review and **Acknowledge / Reopen** actions, scoped by its own filter row.
- The **Audit Trail** is the append-only transaction log of platform and tenant mutations, newest first. On top of its user and date filter, its toolbar has a severity select, an action-category select, and a search box (**Search actions, people, targets...**) that narrow the rows together.

### TIP

Start with the groups that open on their own. Anything unexpected, such as an expired key or an unscoped model, has a matching tab in this console where you can fix it, and this guide's matching section tells you how.

## WALKTHROUGH · 2:16 · PLATFORM CONSOLE › DOCUMENTATION

### Owner audit signals

Triage what needs attention, open the records behind any number, read the Audit Insights trends, act on security alerts, and export the audit trail.

#### Before you begin

- A platform owner account.

#### Step by step

- 1 Open Platform console › **Audit**. **Owner Audit** opens first, with a banner counting the signals that need attention.
- 2 Read the four groups: **Security** signals, Identity & access, **Providers** & secrets, and **Models**, connectors & automations. Red rows need attention.
- 3 Select any row to open an **Audit** investigation listing the exact records behind the number. Use Filter investigation records, then press Escape to close.
- 4 In **Audit** Insights, choose 7 days, 14 days, or 30 days. Select any bar or point to open its records the same way.
- 5 Expand **Security Alerts**. Review each redacted snippet and choose **Acknowledge** once it is handled, or **Reopen** to bring it back.
- 6 Expand **Audit Trail**. Filter by severity and category, or search actions, people, and targets.

- 7 Choose **CSV**, set the **CSV** date range (or All dates), and choose **Download**. The file contains exactly the rows that match your filters.

### Path: Layout choices

- 1 Expand all or Collapse all opens or folds every signal group.
- 2 List or Cards switches between compact rows and full cards. The choice is remembered in this browser.

### Path: Charts in Audit Insights

- 1 **Audit** events by day stacks info, warning, and critical events.
- 2 **Security** alerts by day, and Alert breakdown By rule or By person.
- 3 Most active people, Activity by area, and Activity by hour (hours outside 7 AM to 7 PM in your local time are shaded).

### Check that it worked

- ✓ Each investigation's count matches the number on the signal or bar you selected.
- ✓ The acknowledged alert shows Acknowledged with a **Reopen** button, and the active count goes down by one.
- ✓ The audit **CSV** downloads with a header row (id, created\_at, tenant\_id, actor\_id, actor\_name, actor\_role, action...) and one row per filtered event.

### Troubleshooting

| WHAT YOU SEE                                        | WHAT TO DO                                                                                                                                           |
|-----------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------|
| An investigation lists fewer records than expected. | It reflects the currently loaded audit range. Widen the <b>Audit</b> Insights range, or open the <b>Audit Trail</b> with its own dates.              |
| The CSV is missing events you can see elsewhere.    | The export follows the trail's filters and its own date range. Clear the severity or category filter, or choose All dates.                           |
| A security alert keeps reappearing.                 | <b>Acknowledge</b> marks one alert as handled; new prompts that trip the same detector raise new alerts. Use an alert rule to be emailed about them. |

## 59. Alerts: email delivery and platform-wide rules

Configure SMTP once, watch platform-wide activity, and read honest delivery logs.

The owner **Alerts** tab is the full version of the alerting station: you own the email configuration, your rules watch platform-wide audit activity, and every rule in the organization is listed with its scope. **Alerts** are

always logged in-app regardless of email.

## Email Delivery: configuring SMTP

### GOOD TO KNOW

The relay's TLS certificate is verified, so a relay presenting a self-signed or mismatched certificate is refused rather than trusted. Failed sends are retried with growing delays instead of every scheduler pass.

## Alert Rules

- Rules work exactly as described in the Administrator Guide — Prompt-injection template, **Suspicious-activity template**, or **New rule**, with action patterns, a **Minimum severity**, an optional **Watched user**, a threshold within a window, a cooldown, email recipients (empty recipients means in-app only), and Only these detections to narrow a rule to specific prompt detectors.
- Owner rules are platform-wide: they watch audit activity across the whole organization, including owner actions. Tenant rules created by admins appear in the same list, each labeled with its scope, so you always see the full alerting picture.

## Alert Deliveries

Every alert trigger is listed with its real delivery status — sent, queued, failed with the actual SMTP error, email not configured, or logged in-app — and the log exports to **CSV**, archived deliveries included. Click **Archive** on a delivery to clear it from the default view without deleting its history; **Show archived** reveals archived deliveries for review or **Restore**. Owners can archive any delivery, platform-scope ones included.

### WALKTHROUGH · 3:29 · PLATFORM CONSOLE › DOCUMENTATION

## Alerts and email delivery

Connect a real SMTP relay, send a test email, create an alert rule, and confirm a real alert email arrives when the rule fires.

### Before you begin

- A platform owner account.
- An SMTP relay your server can reach: Microsoft 365, Google Workspace, SendGrid, Amazon SES, or your own, with a username and password or app password.
- A from address the relay allows you to send as, and a mailbox that receives the alerts (for example a security team list).

### Step by step

- 1 Open Platform console › **Alerts. Email Delivery** is at the top.
- 2 Enter the **SMTP host**, **Port**, and **Security** (**STARTTLS** on 587, or **SSL/TLS** on 465), the Username, the **Password**, and the **From address**.

- 3 Choose **Save Email Settings** and wait for SMTP settings saved. The password is stored in the encrypted vault and never shown again.
- 4 Enter a recipient beside **Send test email** and choose it. Read the result: Test email sent to ... or the relay's exact error.
- 5 In **Alert Rules**, choose Prompt-injection template, **Suspicious-activity template**, or **New rule**.
- 6 Review the action patterns, **Minimum severity**, **Watched user**, Fire when threshold, and Cooldown. Name the rule and enter the Email recipients, separated by commas.
- 7 Under Only these detections, keep the detectors this rule should watch. Choose **Create Rule**.
- 8 When a matching event happens, **Alert Deliveries** lists it as queued, then sent (or failed with the SMTP error). The email names the detection but never includes the flagged text.

#### Path: Microsoft 365 / Exchange Online

- 1 Host smtp.office365.com, **Port 587**, **Security STARTTLS**. Username and **From address**: a licensed mailbox (or one the account can Send As).
- 2 Turn on Authenticated SMTP for that mailbox: Microsoft 365 admin center > **Users** > Active users > the user > Mail > Manage email apps.
- 3 Microsoft turns basic authentication for SMTP off by default at the end of December 2026. Plan to use another relay, or confirm with your Microsoft 365 administrator.

#### Path: Google Workspace or Gmail

- 1 Host smtp.gmail.com, **Port 587**, **Security STARTTLS** (or 465 with **SSL/TLS**).
- 2 Username: the full address. **Password**: an app password from myaccount.google.com/apppasswords (2-Step Verification required).
- 3 Alternative: the Workspace SMTP relay, smtp-relay.gmail.com, configured in **Admin console** > Apps > Google Workspace > Gmail > Routing.

#### Path: SendGrid

- 1 Host smtp.sendgrid.net, **Port 587**, **Security STARTTLS**.
- 2 Username: the word apikey. **Password**: a SendGrid API key with Mail Send permission.
- 3 **From address**: a verified sender or authenticated domain (Settings > Sender Authentication).

#### Path: Amazon SES

- 1 Host email-smtp.<region>.amazonaws.com, **Port 587** with **STARTTLS** (or 465 with **SSL/TLS**).
- 2 Username and **Password**: SMTP credentials from SES > SMTP settings > **Create SMTP credentials** (not AWS access keys).
- 3 **From address**: a verified identity. In the SES sandbox, recipients must be verified too.

### Path: Your own relay (shown in the video)

- 1 Host and **Port** of your relay, with **STARTTLS** or **SSL/TLS**. Its certificate must be valid for that host name and trusted by the server; a self-signed certificate is refused.
- 2 Username, **Password**, and a **From address** the relay accepts.

### Path: Rule templates

- 1 Prompt-injection template: security.prompt\_flagged events from the prompt-injection, system-prompt-extraction, and credential-extraction detectors, at warning and above.
- 2 **Suspicious-activity template**: security flags and elevated-severity events.
- 3 **New rule**: start from scratch. Leave Email recipients empty to log alerts in-app only.

### Check that it worked

- ✓ **Send test email** shows Test email sent to ..., and Last test reads sent.
- ✓ The new rule is listed as platform and enabled, with its detectors and recipient count.
- ✓ **Alert Deliveries** lists the rule with sent and the recipient, and the email arrives in the recipient's inbox.

## Troubleshooting

| WHAT YOU SEE                                                                                       | WHAT TO DO                                                                                                                                                   |
|----------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------|
| "(535, b'5.7.8 Authentication credentials invalid')\" or another 535 reply                         | The relay rejected the username or password. Re-enter the password (or app password) and choose <b>Save Email Settings</b> before testing again.             |
| <b>Send test email is greyed out and Unsaved changes. Save them before sending a test. appears</b> | The test uses the saved settings. Choose <b>Save Email Settings</b> first.                                                                                   |
| <b>"Email delivery is not configured. Set the SMTP host and from address first."</b>               | Enter the <b>SMTP host</b> and <b>From address</b> and save.                                                                                                 |
| <b>A certificate verify failed error</b>                                                           | The relay's certificate is not trusted or does not match the host name. Use the provider's official host name, or install a valid certificate on your relay. |
| <b>A delivery shows email not configured or logged in-app</b>                                      | SMTP was not configured when the alert fired, or the rule has no recipients. <b>Configure</b> SMTP and add recipients; later alerts are emailed.             |
| <b>A delivery stays failed</b>                                                                     | Read the SMTP error on the row. Failed sends are retried with growing delays; fix the relay settings and the next retry uses them.                           |