



For workspace administrators

Administrator Guide

Aperture Chat

The complete User Guide, plus the Admin console: accounts, groups, model access requests and diagnostics, response actions, single sign-on, analytics, token budgets, personal data protection, training datasets, the tenant audit trail, and alerts. No prior knowledge assumed.

Contents

PART 1 · GETTING STARTED

1	Welcome	4
2	Signing in	5
3	Request access and finish your first sign-in	8
4	Two-step verification and recovery codes	9
5	Finding your way around	14
6	Organizing chats: folders, pins, and the chat menu	16
7	Searching everything	19

PART 2 · CHAT

8	Your first chat	22
9	Understand model availability and request access	27
10	Reading and acting on replies	30
11	Symbol shortcuts: / @ # \$ >	32
12	Attaching files, web pages, and connected sources	35
13	Send options: Knowledge, Web, Agent, and Reasoning	38
14	Session details and sources	41
15	Personal data in your chats	43

PART 3 · DRAFTS, DECKS, AND AGENTS

16	Drafting documents	46
17	Building a slide deck	52
18	Recover work kept only on this device	57
19	Agent profiles	59

PART 4 · LIBRARY AND AUTOMATIONS

20	Knowledge bases	62
21	The Tools library	64
22	Scheduled automations	66

PART 5 · APPEARANCE, YOUR ACCOUNT, AND HELP

23	Light and dark mode	71
24	Your account	72
25	Personalization memory	75
26	Install the app on your phone	78
27	Getting help	79

PART 6 · THE ADMIN CONSOLE

28	Opening the Admin console	80
29	Users	81
30	Groups	86
31	Model Access	89
32	Resolve model requests and trace access	91
33	Connections: response actions and connector handoff	93
34	Automations: preparing scheduled runs	95
35	SSO: single sign-on for your tenant	96
36	Analytics	100
37	Policies: service availability and memory governance	102
38	Personal Data Protection	105
39	Training datasets	109
40	Data retention and conversation tags	114
41	Review feedback and reported issues	117
42	Audit	119
43	Alerts	121

Getting started

1. Welcome

What Aperture Chat is and how to read this guide.

Aperture Chat is your organization's workspace assistant. You can chat with approved AI models, draft documents and slide decks, search your organization's knowledge, and schedule recurring work. Review the sources and work trace when a response uses connected knowledge or tools.

This guide explains where to look and what to choose. Words in bold — like **Save version** — match the labels on screen. If a step says “hover”, rest your pointer on the item without clicking; controls that offer extra help show it nearby.

How each task is taught

Each task appears in three matching forms: a narrated walkthrough video in the product, a written procedure in this guide, and the same procedure on the website guide. They come from one source, so they describe the same steps. Every walkthrough is a complete loop, from where you start to a result you can see, and it covers each way the product offers to do the task.

IN EACH WALKTHROUGH BOX	WHAT IT GIVES YOU
Before you begin	The role, access, earlier setup, and outside accounts or values you need first.
Step by step	Every action in order, with the exact labels on screen. Where there are several ways to do the task (an identity provider, a delivery target, a device), each path has its own complete list.
Check that it worked	What you should see when the task succeeded.
Troubleshooting	Messages you may meet, quoted as the product shows them, and what to do about each.

Walkthrough screens are real captures of a running workspace with sample accounts. Steps that happen in another product, such as an identity provider's console, are shown as labeled instruction cards rather than imitation screenshots.

GUIDE	WHO IT IS FOR	WHAT IT COVERS
User Guide	Everyone	Chat , model access requests, personalization memory, documents and decks, account sync, agents, knowledge and tools, search and commands, appearance, and account security.
Administrator Guide	Workspace admins	Everything in the User Guide, plus the Admin console : users, groups, model access, response actions, SSO , analytics, policies and memory governance, personal data protection, training datasets, audit, and alerts.
Platform Owner Guide	Platform owners	Everything in the other guides, plus setup readiness, providers and keys, organization policy, shared connectors, branding, releases, search indexing, and platform governance.

TIP

If a feature is not configured — a connector without credentials, a model that is not enabled — the screen reports that state. Ask your administrator about account or model access; shared connector settings and credentials are managed by the service team outside tenant administration.

2. Signing in

Reach the sign-in screen and get into your workspace.

- Open the web address your organization gave you for Aperture Chat in a modern browser (Chrome, Edge, Safari, or Firefox).
- On the sign-in screen, type your work email address in the field marked `you@company.com`.
- If your organization uses single sign-on (**SSO**), the screen recognizes your email domain and offers **Continue with SSO**. Click it — you are sent to your organization's identity provider (for example Microsoft or Google), sign in there as usual, and return to Aperture Chat already signed in.
- When both methods are offered, choose Organization **SSO** or Email & password. The password field appears in Email & password mode; enter your password and click Sign in. Pressing Enter submits the method you selected.

GOOD TO KNOW

If both **SSO** and a password are available, either works. If your organization enforces **SSO** for your email domain, password sign-in is blocked on purpose — use the **SSO** button.

GOOD TO KNOW

See a message you do not expect? The sign-in screen reports real problems plainly (for example, a domain that is not allowed for **SSO**). **Copy** the message and send it to your administrator.

Request access and enter your workspace

Ask for an account, then sign in with a temporary password, your own password, Organization SSO, or an authenticator code.

Before you begin

- Your workspace address, which your organization gives you (for example <https://your-instance.example>).
- Your work email address.
- For Organization **SSO**: an account at your organization's identity provider, such as Microsoft Entra ID, Okta, or Google.
- For a password account: the temporary password your administrator gives you after approving your request.
- If two-step verification is on for your account: your authenticator app, or one unused recovery code.

Step by step

- 1 Open your workspace address in a web browser. The sign-in page opens.
- 2 Under New to Aperture Chat?, choose **Request access**.
- 3 Enter First name, Last name, and Work email, then choose Submit access request.
- 4 Read Request received. The form does not send an email or set a password. Contact your administrator, who approves the request and tells you how to sign in.
- 5 When your administrator gives you a temporary password, return to the sign-in page and choose Email & password (it appears beside Organization **SSO** when your workspace offers both).
- 6 Enter your work email and the temporary password, then choose Sign in.
- 7 On Set a new password, type a password of at least 12 characters into New password and Confirm password, then choose **Set password** and continue.
- 8 Check that the welcome card, A good place to begin., says You're signed in.

Path: Organization SSO

- 1 Open your workspace address. Organization **SSO** is selected first when your workspace offers it.
- 2 Optionally type your work email. It selects your organization's provider, shown as Sign in with followed by the provider's name.
- 3 Choose **Continue with SSO**.
- 4 Sign in on your identity provider's page, including any multifactor check it asks for. Aperture Chat never sees that password.

- 5 You return to the workspace signed in. On your first visit your account is created automatically if your organization allows it; otherwise ask your administrator to add you.

Path: Email & password

- 1 Choose Email & password.
- 2 Enter your work email and your password, then choose Sign in.
- 3 If two-step verification is on, follow the **Authenticator code** path.
- 4 The workspace opens where you left off.

Path: Authenticator code

- 1 After Sign in, Two-step verification asks for a code. It shows the attempts remaining and when the challenge expires (about five minutes).
- 2 Open your authenticator app and type the current 6-digit code into **Authenticator code**.
- 3 Choose Verify and continue.
- 4 No phone? Choose **Use a recovery code instead**, type one unused recovery code, and choose Verify and continue. Each code works once.
- 5 If the challenge expires, choose Back to sign in and start again.

Path: Trouble signing in

- 1 Choose Trouble signing in? under the sign-in button.
- 2 For Organization **SSO**, use your organization's own password recovery.
- 3 For an email-and-password account, ask your workspace administrator for a temporary password.
- 4 If you lost your authenticator and have no recovery codes, ask your administrator to reset two-step verification for your account.

Check that it worked

- ✓ The welcome card A good place to begin. shows You're signed in.
- ✓ Your name appears on the account card at the bottom of the sidebar.
- ✓ With Organization **SSO**, the identity provider's page handled your password and you returned signed in.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Invalid local credentials."	The password is wrong. Retype it (Show password helps), or ask your administrator for a new temporary password.
"Unknown local account."	There is no active password account for that email yet: the request may still be pending, or your account signs in with Organization SSO . Choose Organization SSO , or ask your administrator.
"SSO is enforced for this email domain; local sign-in is disabled. Use the configured identity provider."	Your organization requires single sign-on for your domain. Choose Organization SSO and Continue with SSO .
"Use a password with at least 12 characters." or "The passwords do not match."	On Set a new password, use 12 or more characters and type the same password in both fields.
"Choose a new password."	The new password matches the temporary one. Pick a different password.
"The MFA code is invalid." or "Too many MFA attempts."	Wait for the next code in your authenticator app and check that your phone's clock is set automatically. After too many attempts, wait for the timer, or choose Back to sign in.
"Verification unavailable"	The sign-in challenge expired or was used up. Choose Back to sign in and sign in again.
"Too many failed sign-in attempts for this account. Try again shortly."	Wait a minute before trying again, or use Trouble signing in?.
"Too many access requests for this email. Try again shortly."	Your request was already recorded. Contact your administrator instead of submitting again.

3. Request access and finish your first sign-in

Submit a request, understand the administrator handoff, and enter the workspace.

- 1 On the sign-in screen, choose **Request access**. Enter your first name, last name, and work email, then submit the form.
- 2 The Request received screen confirms submission for your email. If this email is new to the workspace, an administrator reviews the request. This form does not email you or create a password.

- 3 Contact your administrator for an update. After approval, they provide the workspace address and confirm whether to use organization **SSO** or a temporary password.
- 4 Choose **Back** to sign in. Your submitted email remains filled in. Use the sign-in method your administrator arranged.
- 5 If you received a temporary password, sign in with it and complete **Set a new password**. Enter a new password of at least 12 characters twice. After saving, the app continues with a new authenticated session and revokes your older sessions.
- 6 Complete any required authenticator verification. **Once** in the workspace, select an available model and send a short first message. If no model is available, ask your administrator to check your account, group, and model access.

GOOD TO KNOW

Submitting another access request does not approve an account or recover an existing password. Trouble signing in explains the recovery path: contact your administrator for local password access, or use your organization identity provider for **SSO**. If sign-in settings cannot load, choose **Retry connection**.

4. Two-step verification and recovery codes

Set up an authenticator, complete required verification, and manage recovery codes.

To add an authenticator after signing in with a local password, open your account drawer, find **Security**, choose **Manage security**, and select **Set up authenticator**. Confirm your current password, add the QR code or setup key to your authenticator, acknowledge that you added it, and choose **Verify authenticator** with a current six-digit code. **Copy** the one-time recovery codes into secure storage, acknowledge that you saved them, and choose **Done**. Copying alone does not store them for you; finish verification and code storage before leaving the panel.

- If sign-in asks you to set up an authenticator, choose **Begin authenticator setup**. Add the displayed QR code or setup secret to your authenticator app.
- Confirm that you added the account, enter the current six-digit **Authenticator code**, and choose **Verify and enable MFA** before the setup expires.
- Use **Copy recovery codes** and save the codes privately when they are shown. They are not saved to a file automatically. Each recovery code can be used once when your authenticator is unavailable; do not include these codes in screenshots or issue reports.
- On later sign-ins, enter your current authenticator code. **Use a recovery code instead** switches to a single-use recovery code when necessary.
- If you cancel or the challenge expires, return to sign-in and start again. A cancelled setup or failed verification does not finish sign-in.

When verification is enabled, the account card shows how many recovery codes remain. **Replace recovery codes** requires a fresh authenticator code or an unused recovery code; after replacement, all previous

recovery codes stop working. Save the new set before leaving the dialog.

Turn off verification appears only when your organization's policy allows it. Confirm with an authenticator or recovery code; successful removal signs you out. Organization-required verification cannot be turned off here.

If you lose both the authenticator and all unused recovery codes, contact your administrator for the organization's verified recovery process. A new password does not remove the authenticator requirement. For identity-provider MFA, use that provider's recovery process.

TIP

Use the code currently shown for this workspace in your authenticator app. If it expires while you are typing, use the next code. Read the attempts, expiry, and any cooldown shown; wait for a cooldown to finish before trying again. If **Security** cannot load, choose Retry security settings. Keep setup keys and recovery codes out of shared documents, videos, and support attachments.

GOOD TO KNOW

SSO accounts manage voluntary authenticator setup in their identity-provider settings. Your identity provider may also require its own verification during **SSO**. Follow the provider's screen for that check; these instructions describe Aperture Chat's authenticator screens.

WALKTHROUGH · 2:51 · [HELP](#) · [LEARN](#)

Protect your account and recover access

Turn on two-step verification, sign in with a code or a recovery code, replace recovery codes, and change your password.

Before you begin

- An email-and-password account. With Organization **SSO**, your identity provider manages your password and multifactor settings, and the panel says so.
- An authenticator app on your phone, such as **Microsoft Authenticator**, Google Authenticator, or a password manager with one-time codes.
- A safe place for recovery codes, such as a password manager.
- Your current password. If you still have a temporary password, replace it first.

Step by step

- 1 Select your account card at the bottom of the sidebar. In the **Security** card, choose **Manage security**.
- 2 Under Two-step verification (Off), choose **Set up authenticator**.
- 3 Enter your Current password and choose Continue setup.
- 4 In your authenticator app, scan the QR code, or type the Setup key shown under it.

- 5 Tick I added this account to my authenticator., type the 6–digit code from the app into **Authenticator code**, and choose **Verify authenticator**.
- 6 Choose **Copy recovery codes** and store them somewhere safe. Tick I stored these recovery codes somewhere safe., then choose **Done**.
- 7 Check that Two–step verification shows On and Your authenticator is connected. 10 recovery codes remain.

Path: Sign in with a code

- 1 Sign in with your email and password as usual.
- 2 On Two–step verification, type the current 6–digit code from your authenticator app.
- 3 Choose Verify and continue.

Path: Sign in with a recovery code

- 1 On Two–step verification, choose **Use a recovery code instead**.
- 2 Type one unused recovery code into Recovery code and choose Verify and continue.
- 3 Open **Account** › **Manage security**: the count of recovery codes remaining has gone down by one.
- 4 When few codes remain, replace them.

Path: Replace recovery codes

- 1 **Account** › **Manage security** › **Replace recovery codes**.
- 2 Enter a current **Authenticator code** (or choose **Use a recovery code instead**).
- 3 Choose **Create** new recovery codes. The old set stops working at once.
- 4 **Copy** and store the new codes, tick I stored these recovery codes somewhere safe., and choose **Done**.

Path: Change your password

- 1 **Account** › **Password** card › Edit.
- 2 Enter Current password, then a New password of at least 12 characters, and the same in Confirm new password.
- 3 Choose Update password and wait for **Password** updated.
- 4 You stay signed in here; your older sessions on other devices are signed out.

Path: Lost or replaced phone

- 1 Sign in with your password and one recovery code (**Use a recovery code instead**).

- 2 If your organization allows it, choose **Turn off verification**, confirm with another recovery code, sign back in, and set up the new phone.
- 3 If verification is required by your organization, or you have no recovery codes left, ask your administrator to reset two-step verification for your account.
- 4 After the reset, sign in and set up your authenticator again.

Check that it worked


- ✓ Two-step verification shows On, with Your authenticator is connected. and the number of recovery codes that remain.
- ✓ Your next password sign-in asks for an **Authenticator code**.
- ✓ After signing in with a recovery code, the remaining count is one lower.
- ✓ After a password change, the **Password** card shows **Password** updated.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Current password is incorrect. "	Retype your current password. If you have forgotten it, ask your administrator for a temporary password.
"The MFA code is invalid. "	Wait for the next code and type it promptly. Codes depend on the time, so set your phone's clock to update automatically. A code that was already used is refused.
"Change the temporary password before enrolling MFA. "	Sign in and choose your own password first, then set up the authenticator.
"Finish verification or save your recovery codes before leaving this panel. "	Complete Verify authenticator , or copy and store the recovery codes and choose Done , before closing the panel.
"Required by your organization. You cannot turn this off. "	Your organization requires two-step verification, so Turn off verification is not offered. To move to a new phone, ask your administrator to reset it.
"Your organization manages sign-in. Use your identity provider's security settings..."	You sign in with Organization SSO . Change your password and multifactor settings at your identity provider.
"Use a password with at least 12 characters." or "The new passwords do not match. "	Use 12 or more characters and type the same new password in both fields.
"Password changes are only available for local password accounts. "	Your account signs in with Organization SSO . Change the password at your identity provider.

5. Finding your way around

The sidebar, the main area, and what each navigation button opens.



The chat workspace: navigation on the left, model selection above, and the message composer in the main area. Example workspace with synthetic data.

The screen has two areas: a sidebar on the left and the main work area on the right. The sidebar is how you move between the platform's views and your saved chats.

Your first visit

The Getting started card explains what is ready for your account and what comes next. Open quick-start guide opens user help. Administrators see Open admin guide and Manage access. Follow the setup action offered for your role when no model is available. A missing model is an access or setup task, not a successful chat connection.

Choose a guide, take the suggested action, or select I'll explore on my own (or Dismiss welcome) when you are ready. Simply loading the page does not dismiss this card. You can return to guides from **Help** or console **Documentation** later.

SIDEBAR BUTTON	WHAT IT OPENS
New chat	A fresh conversation with your workspace assistant. This is the default view.
Search	Search past work for chats (including archived), documents and decks, agents, and indexed documents, plus commands. Shortcut: Ctrl+K (Windows) or ⌘K (Mac).
Drafts	A document and slide-deck editor with an AI assistant beside it.
Agents	Reusable agent profiles, plus an Automations tab for scheduled runs.
Library	Two tabs: Knowledge (searchable document collections) and Tools (connections, prompts, and skills).

Below the navigation buttons, your chat history is always listed under a small CHATS heading — folders first, then pinned chats, then recent chats. The next section covers organizing them.

Resizing and collapsing the sidebar

- Click the small chevron (◀) near the top of the sidebar to collapse it and give your work more room. Click the chevron again (▶) to bring it back.
- Drag the sidebar's right edge to make it wider or narrower. Double-click that edge to collapse it.
- On a narrow window or a phone, the sidebar hides behind a menu button (three stacked lines, ≡) in the top-left corner. Tap it to open the menu; tap the dimmed area to close it.

Using the keyboard

Use Tab and Shift+Tab to move between controls and Enter or Space to activate the focused control. In a dialog, keyboard focus stays with that dialog. Use its Close or **Cancel** button, or Escape when available, to return to the control that opened it. Confirmations explain when an action is permanent before you commit it.

The bottom of the sidebar

- **Admin console** and Platform console — links shown only to administrators and platform owners (see “Where the consoles live”). Platform owners may also see a Release available notice here.
- **Help** — opens the guided walkthrough videos and this downloadable guide.
- **Dark mode / Light mode** and **Theme schedule** — two small buttons beside **Help**: the moon or sun switches the appearance instantly, and the clock sets a schedule (see “Light and dark mode” later in this guide).
- **Install app** — appears on phones and tablets, and adds the workspace to your home screen (see “Install the app on your phone”).
- Your account card — your initials or photo, name, and role. Click it to open the account drawer, including Personalization memory when your organization has enabled it.

Where the consoles live

Administrators choose **Admin console** near the bottom of the sidebar; platform owners also see Platform console. The account drawer's **Management** section lists the same consoles. Regular users see neither. The account drawer also holds **View as** role previews, **Usage this month**, and your **Archived chats** — all covered in “Your account”.

Returning to a view

The browser address follows the workspace view and console tab. Browser Back and Forward return to those views; opening a saved chat or draft link still checks your account's access. A link does not share private content or grant permission to another person.

6. Organizing chats: folders, pins, and the chat menu

Keep the chats you care about easy to find.

Your chat history is always visible under the CHATS heading: folders first, then a **Pinned** group when something is pinned, then your recent chats (labeled **Recent** only when a group sits above them). Click the CHATS heading to hide or show the whole list; this is remembered for your account on this device. While hidden, it shows a count and a dot for unread replies, and **Search** still finds every chat.

Folders

- Click the folder icon with a plus sign next to the CHATS heading.
- Type a name for the folder — a client, a matter, a project — and click **Create**.
- To file a chat into a folder, hover or focus its row, click the ⋮ (**More** actions) button, and choose Move to folder. Pick a folder, choose New folder to create one right there, or choose Remove from folder to take it out.
- Click a folder to expand or collapse the chats inside it.

GOOD TO KNOW

Deleting a folder does not delete its chats — they move back to the **Recent** list.

Pinned chats

- Hover or focus any chat row, click ⋮ (**More** actions), and choose Pin chat.
- The chat moves to the **Pinned** group, just below your folders, where it stays until you choose Unpin chat from the same menu.

Recent chats

- **Recent** chats — below folders and pinned chats — list your latest conversations. Click one to reopen it exactly where you left off.

- Hover or focus a row to reveal its ⋮ (**More** actions) button. The menu offers Pin chat (or Unpin chat), Move to folder, and **Archive**.
- A dot beside a chat marks a newest assistant reply you have not opened. Read status is saved to your account, so it matches in every browser and device.
- Archiving moves a chat out of the sidebar without deleting it — use it to tidy up. **Archived chats** are listed in your account drawer, where you can restore or permanently delete them.
- Click **View all chats** at the bottom of the list to browse your full history, including everything that no longer fits in the sidebar.

WALKTHROUGH · 1:36 · [HELP](#) > [LEARN](#)

Organize and find your work

Create folders and file chats in them, pin, archive and restore, rename chats, and hide your chat list.

Before you begin

- A few chats in your sidebar.
- **Folders** are kept in this browser: they do not follow you to another browser or device.

Step by step

- 1 Choose **Create** chat folder (the folder button beside CHATS), type a Folder name, and choose **Create**.
- 2 Point at a chat, choose its **More** actions button (...), then Move to folder, and choose the folder.
- 3 Check that the chat now sits inside the folder, with the folder's count updated.
- 4 To keep a chat at the top, choose ... > Pin chat. It moves to **Pinned**.
- 5 To put a chat away, choose ... > **Archive**. Find it later in your account card > **Archived chats** > View, and choose **Restore**.
- 6 To rename the open chat, choose Rename chat (the pencil beside its title), type the name, and choose Save chat name.
- 7 Choose the CHATS heading to hide your chat list; choose it again to show it.

Path: Folders

- 1 **Folders** can hold subfolders, up to four levels: choose **Create** subfolder in ... on a folder.
- 2 ... > Move to folder > Remove from folder takes a chat out again.
- 3 **Delete** ... folder asks first; chats in it stay in **Recent**.

Path: Archive and delete

- 1 **Archive** keeps the chat and removes it from the sidebar (and unpins it).

- 2 **Account** card › **Archived chats** › View lists archived chats, with **Restore** and Permanently delete.
- 3 Permanently delete asks first and cannot be undone.

Path: Rename with AI

- 1 After the first reply, choose Rename chat with AI beside the title for a suggested name.

Path: See everything

- 1 **View all chats** opens All chats, with Filter chats to search titles.
- 2 View all pinned and View all folders appear when you have more than three.

Check that it worked

- ✓ The folder shows the filed chat and its count.
- ✓ The pinned chat appears under **Pinned**.
- ✓ An archived chat is listed under **Archived chats** until you restore it.
- ✓ With the list hidden, CHATS shows how many chats are tucked away.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
My folders are missing in another browser	Folders are stored in the browser where you made them. Create them again there; the chats themselves are in your account.
"Chat title cannot be blank."	Type a name before choosing Save chat name, or choose Cancel renaming chat.
"The AI could not rename this chat. Try again."	Try again , or type a name yourself with Rename chat.
A chat disappeared from the sidebar	It may be archived (account card › Archived chats), inside a folder, or hidden with the CHATS heading. Search finds it either way.

WALKTHROUGH · 0:34 · [HELP](#) › [LEARN](#)

Preview chats at a glance

Hover any listed chat to read its prompts and outputs in a compact, scrollable preview.

Before you begin

- **Chats** in your sidebar, in All chats, or in **Archived chats**.

Step by step

- 1 Point at a chat name in the sidebar and wait a moment; keyboard users can focus the chat row instead.
- 2 Read the **Chat** preview: the title and time, then Prompt 1 and Output 1, in order.
- 3 Scroll inside the preview to read later prompts and outputs. The footer counts prompts and outputs.
- 4 Move away, or press Escape, to close it without opening the chat.
- 5 Choose the chat if you want to open it.

Path: In All chats

- 1 Choose **View all chats** under your recent chats.
- 2 Point at any chat in All chats; the preview opens beside the list.
- 3 Filter chats narrows the list by title first.

Path: In Archived chats

- 1 **Account** card > **Archived chats** > View.
- 2 Point at an archived chat to preview it before you restore or delete it.

Check that it worked

- ✓ The preview shows **Chat** preview, the chat's title, and Prompt 1.
- ✓ The footer counts the prompts and outputs.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"This chat does not have previewable content yet."	The chat has no messages yet.
No preview appears	Rest the pointer on the chat name for a moment, or focus the row with the keyboard.

7. Searching everything

Find saved work, preview a result, or run an available command.

- Click **Search** near the top of the sidebar, just below **New chat**, or press Ctrl+K (Windows) or ⌘K (Mac) from anywhere.

- The **Search past work** box opens. Type a few words from what you remember: a chat title, message text, an agent or draft name, slide text, or text from an indexed document.
- Results are grouped by kind. **Chat** results lead, matching text is highlighted, and archived conversations are labeled explicitly. Saved documents and decks open in their matching editor.
- Hover over or keyboard-focus a supported chat or draft result to preview it. The chat row also offers folder, pin, and archive or restore actions when that conversation is available in the current workspace.
- Click a result to open it, or use Up and Down followed by Enter. Escape closes search. Recently opened results are available when the query is empty.

Commands and saved work

Commands appear in a separate section from search results. Type > at the start of this search box to see all commands available to your role and permissions, then keep typing to narrow them. **Commands** can open workspace views and permitted actions; an unavailable feature does not become accessible through search. The same > symbol inside a chat message has a different purpose: choosing an automation for that message.

GOOD TO KNOW

Search includes only work your account can access. If the panel says Indexing your workspace, results may be incomplete until indexing finishes; try again after it completes. A missing result is not proof the work was deleted. Unsaved device-only changes may need saving to your account before server search can find them.

WALKTHROUGH · 1:20 · HELP › LEARN

Search, commands, and workspace links

Search every chat, draft, agent, and document you can access, reopen recent results, and run commands from the keyboard.

Before you begin

- Some saved work to find. **Archived chats** are included.

Step by step

- 1 Choose **Search** in the sidebar, or press Ctrl+K (⌘K on a Mac).
- 2 Type words from a title or a message.
- 3 Read the results, grouped by kind: Previous chats, Documents & knowledge, **Agents**, Automations, **Drafts**, and more.
- 4 Use the arrow keys and Enter, or choose a result, to open it.
- 5 Press Ctrl+K again. With an empty box, **Recent** lists what you opened from **Search**.
- 6 Type > to list commands, add a word to narrow them (for example >dark), and press Enter to run one.

Path: Act on a chat from Search

- 1 Point at a chat result for its buttons: add to a folder, pin, and archive or restore.
- 2 **Archived chats** are marked **Archived** conversation and can be restored from here.

Path: Commands

- 1 **New chat**; Go to **Chat**, **Drafts**, **Agents**, Automations, **Knowledge**, or **Tools**.
- 2 Switch to dark mode or Switch to light mode; Open **Help**; **Sign out**.
- 3 **Install app** appears on phones and tablets.

Path: Workspace links

- 1 Each screen has its own address (for example /drafts or /chat/ followed by the chat's id).
- 2 Browser Back and Forward move between the screens you visited.
- 3 A link never shares private content: whoever opens it must sign in and have access.

Check that it worked

- ✓ Results list the matching chats and other work.
- ✓ Choosing a result opens it.
- ✓ A command such as >dark changes the workspace at once.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Indexing your workspace... results may be incomplete until it finishes."	Search is still catching up. Try again in a moment.
"No chats or workspace items found for ..."	Try fewer or different words, such as a word from the reply rather than the title.
"No commands match ..."	Type > alone to see every command.
"Search failed."	Check your connection and search again.

8. Your first chat

Send a message, pick a model, and control web search.

- Click **New chat** in the sidebar. A fresh chat greets you by the time of day — “Good morning”, “Good afternoon”, or “Good evening” with your first name (late at night it asks if you are burning the midnight oil) — above the tagline “Your approved models, your sources, your guardrails — ask Aperture Chat anything.” (The name reflects your organization's branding.)
- Click into the message box (the composer) — it reads **Ask anything...** — and type your question.
- Explore an idea, Compare options, and **Draft** a message offer starting points. Selecting one fills an editable prompt; review it before sending.
- Press Enter to send. Press **Shift+Enter** when you want a new line without sending.

Picking the model

The model selector sits in the top bar of every chat. Click it to choose which AI model answers this conversation. Each chat remembers its own choice, and the list only ever shows models your workspace has approved for you.

Use the arrow keys, Home, End, or type a model name to move through available models; Enter selects the highlighted model. Each row has its own star: Set [model name] as default model selects that model and remembers it for new chats. The filled star marks the single current default. Selecting a row without its star changes this conversation's model without changing the default. **Drafts** has a separate default drafting model, set with the star beside a row in its own selector.

Web search and the active-tools chip

- When the selected model supports it, public web search starts turned on. A chip at the bottom of the composer shows what is active: with one tool on it names it — for example **Web search** — and with several on it reads **Tools** with a count.
- Click the small x on that chip to clear the next message's active tools: **Knowledge**, **Web**, and **Agent**, selected MCP connections, and queued automations. Every new chat starts with web search on again when the model supports it.
- Replies that used the web come back with citations you can click, and every source is listed in the session details panel (covered below).

Sending

- The paper-plane button sends your message with the current settings.

- The chevron (v) beside it opens send options — per-reply switches for **Knowledge**, **Web**, and **Agent**, plus the **Reasoning** level (covered in their own section).

Dictate instead of typing

- Click the microphone button in the send controls at the bottom-right of the composer to start dictating. A live waveform shows that the platform is hearing you.
- Click the button again — it becomes a stop square — and your words are transcribed and inserted into the composer as editable text. Nothing sends until you press Enter.
- If the microphone is blocked or no speech was heard, the composer says so plainly so you can fix the problem and try again. The **Drafts** workspace has the same dictation button for drafting instructions.

WALKTHROUGH · 1:38 · [HELP](#) · [LEARN](#)

Start chatting

Send your first message, watch it work, stop a reply, edit and resend your prompt, and copy the answer.

Before you begin

- A signed-in account with at least one available model. If the composer says Connect a model provider to start chatting..., see Choose models and request access.

Step by step

- 1 Choose **New chat** at the top of the sidebar.
- 2 Check the model selector at the top right (**Model:** and a name). Choose another model there if you need one.
- 3 Type your request where it says **Ask anything...**, or choose a starter such as Explore an idea. A starter only fills in the opening words.
- 4 Press Enter, or choose the paper-plane Send message button, to send. **Shift+Enter** adds a new line instead.
- 5 Watch Aperture Chat is working while the reply is prepared; the answer appears under it.
- 6 Read the reply. The work trace above it reads complete, and the chat is listed under **Recent** with a title from your request.
- 7 To keep the answer, choose **Copy** response text under the reply and wait for Copied.

Path: Stop a reply

- 1 While Aperture Chat is working, choose · stop at the top right of the trace.
- 2 The reply stops and keeps any text already written, with You stopped this response before it finished.
- 3 To ask again, use Edit and resend.

Path: Edit and resend

- 1 Point at your message and choose **Edit message** (the pencil).
- 2 Change the text. Escape cancels.
- 3 Choose Send edited message. Everything after that message is removed, and the model answers the new text.

Path: Regenerate

- 1 On a finished reply, choose **Regenerate response** (the circular arrows).
- 2 The model answers the same request again.
- 3 A stopped or failed reply has no actions; use Edit and resend instead.

Path: Reuse a prompt in a new chat

- 1 Point at your message and choose **Load prompt in new chat** (the paper plane beside the pencil).
- 2 A new chat opens with your prompt and its attachments in the message box, ready to send.

Check that it worked

- ✓ Your message and the model's reply both appear, and the trace above the reply reads complete.
- ✓ The chat is listed under **Recent** in the sidebar.
- ✓ Copied appears beside the reply's actions after you copy it.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
The message box says "Connect a model provider to start chatting..."	No model is available to your account. Open the model selector or Why isn't a model listed? , and request access.
"You stopped this response before it finished. Regenerate or resend the message to run it again."	A stopped reply has no Regenerate button. Choose Edit message on your prompt, then Send edited message.
A reply reads "... did not return a completion: ..."	The model service failed. Resend with Edit message . If it repeats, tell your administrator. Some self-hosted models reject Regenerate with "System message must be at the beginning."; resend instead.
The pencil says "Wait for the current response to finish before editing"	Wait for the reply, or choose · stop first.
"The edited message could not be sent. Check the text and try again."	Make sure the edited text is not empty, then choose Send edited message again.
"Copy unavailable"	The browser blocked the clipboard. Select the reply text and copy it with the keyboard instead.

WALKTHROUGH · 1:27 · [HELP](#) · [LEARN](#)

Dictation, images, and diagrams

Dictate a message when your workspace has a transcription model, ask an image model for pictures, and copy, view, and download diagrams.

Before you begin

- For dictation: a microphone, permission for this site to use it, and a transcription model (a Gemini Flash model) set up by your platform owner.
- For images: an image-generation model available to you in the model selector.
- For diagrams: any available model. Ask for a Mermaid diagram.

Step by step

- 1 In the message box, choose Dictate a message (the microphone).
- 2 Allow microphone access if the browser asks.
- 3 Speak. A live waveform shows it is recording; recording stops by itself after 2 minutes.

- 4 Choose Stop dictation and transcribe (the square). The transcript is added to the message box as editable text.
- 5 Review the text, then send it yourself; dictation never sends on its own.

Path: Images

- 1 In the model selector, choose a model that can generate images.
- 2 Describe the picture you want; ask for up to 4.
- 3 The images appear in the reply, labeled Generated image.
- 4 Point at an image and choose **Download** to save it. **Download** links last 7 days.

Path: Diagrams

- 1 Ask for a diagram, for example: show the stages of a process as a Mermaid flowchart.
- 2 The reply shows a live figure with **Copy**, **PNG**, **SVG**, **Code**, and Edit in its header.
- 3 Choose **Code** to read or check the diagram's source; choose Diagram to switch back.
- 4 Choose **PNG** or **SVG** to download the picture, or **Copy** to copy it.

Check that it worked

- ✓ Dictation: the transcript appears in the message box, ready to edit.
- ✓ Images: **Download** saves the image as a file.
- ✓ Diagrams: the figure renders in the reply, and **PNG** downloads it as a picture file.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"No configured Gemini Flash model is available for dictation. ..."	Your workspace has no transcription model. Type your message, and ask your platform owner to set one up.
"Microphone access was denied. Allow microphone use for this site to dictate."	Allow the microphone in your browser's site settings, then try again.
"Microphone capture is not available in this browser."	Use a current browser on a device with a microphone.
"No audio was captured. Check the selected microphone and try again." or "The dictation model heard no speech in the recording."	Check which microphone your device uses and speak closer to it.
"... returned no images for this request. Try rephrasing the prompt or selecting a different image model."	Rephrase the request, or choose another image model.
"This diagram could not be rendered. Open Code to inspect the source, or Edit to fix it."	The model's diagram source has an error. Ask it to correct the Mermaid code, or fix it yourself.

9. Understand model availability and request access

Read the server's reason, ask for a model when permitted, and follow its request status.

- Open the chat model selector and choose **Why isn't a model listed?** If no model is usable, the No models available control opens the same explanation.
- **Models in your organization** separates **Usable now** from **Not available to you**. Read the reason on the model row: **Locked** describes an access restriction; **Provider offline** means you already have access but the provider needs attention.
- Choose **Request access** only where it is enabled. The request goes to an administrator for review; sending it does not grant immediate access.
- While a request is pending, the row shows its date and **Withdraw request**. Use **Refresh model access** after your administrator resolves it, then select the usable model in the chat picker.

GOOD TO KNOW

Your organization may show only models you can already use. Its browsing policy controls whether you can inspect other enabled models and request access. Some restrictions cannot be lifted by a request; ask your administrator about the stated reason. An offline provider needs a platform-owner repair, not another group-access request.

Choose models and request access

Pick the model for a chat, star your default, find out why a model is missing, and request, withdraw, and use a model your administrator approves.

Before you begin

- A signed-in account with at least one model available. If the selector reads **No models available**, start at **Why isn't a model listed?**.
- Your organization must allow access requests. If it shows only models you can already use, ask your administrator about others.
- An administrator to review requests. Approval happens in their **Admin console** > **Model Access**.

Step by step

- 1 In a chat, open the model selector (**Model:** and the model's name) at the top right.
- 2 Choose a model to use it in this chat. Each chat keeps its own model.
- 3 Choose the star beside a model to make it your default for new chats in this browser.
- 4 Choose **Why isn't a model listed?** at the bottom of the selector. **Models in your organization** lists **Usable now** and **Not available to you**, with the server's reason for each.
- 5 Beside a **Locked** model, choose **Request access**. The row changes to Request pending since and the date.
- 6 Wait for your administrator to approve. Approval adds you to a group that carries the model and signs you out of your current sessions.
- 7 Sign in again. The model is listed under **Usable now** and in the selector. Choose it and send a message.

Path: Withdraw a request

- 1 Open the model selector > **Why isn't a model listed?**.
- 2 Beside the pending model, choose **Withdraw request**.
- 3 Withdrew the request for the model appears, and **Request access** is offered again.

Path: Provider offline

- 1 A **Provider offline** badge means the connection behind the model is not working.
- 2 Requesting access does not fix a disconnected provider; the platform owner must reconnect it.
- 3 Choose **Refresh model access** later, or use another model in the meantime.

Path: After approval

- 1 If the dialog or a chat says Session is invalid or expired. Sign in again., sign in again.
- 2 Open the selector: the new model is listed. Choose it for this chat, or star it to make it your default.
- 3 Send a message; the reply comes from the newly granted model.

Check that it worked

- ✓ The selector shows **Model:** and your chosen model's name, and the reply comes from that model.
- ✓ The starred model is filled in, and new chats start with it.
- ✓ After approval and a fresh sign-in, the model is listed under **Usable now** with Available.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Not granted to any of your groups. An administrator can add you to a group that has this model. "	Choose Request access , or ask your administrator to add you to a group that has the model.
"The provider behind this model is not connected right now. " (Provider offline)	A request cannot fix this. The platform owner must reconnect the provider. Use another model meanwhile.
"Disabled by the platform owner for the whole organization. "	The model is switched off for everyone. Ask your administrator whether it will be enabled.
"Session is invalid or expired. Sign in again. " right after a request	Your administrator approved the request, which signs you out. Sign in again; the model is then listed.
"Model access requests are disabled by organization policy. "	Your organization does not take requests here. Contact your administrator directly.
"Too many access requests. Try again shortly. "	Wait a minute. A pending request stays pending; you do not need to resend it.
The selector reads No models available	No model is usable for your account yet. Choose it to open Models in your organization , request access, or contact your administrator.

10. Reading and acting on replies

The work trace, citations, and the actions under every message.

The work trace

- While a reply is being produced, a single status line shows the current step, a live timer, and how many steps remain (for example “4s · step 2 of 5”).
- Click that line at any time — while running or after — to expand the full trace: which model was routed, what context was prepared, whether web search ran, and how the answer was finalized.
- While a reply is still streaming, a stop button appears — its label is **Stop this response**. Clicking it halts the reply and keeps the conversation.

Actions under each response

- **Copy** — copies the response as clean, readable text.
- Thumbs up / thumbs down — sends feedback your administrators can review.
- **Branch** — starts a new chat that continues from that response, leaving the original untouched.
- **Regenerate response** — asks the model to answer the same prompt again. Version arrows appear on the response with a counter (for example “2 / 3”) so you can flip between the versions it has produced.
- **Transfer to Drafts** — turns the response into an editable document in the **Drafts** workspace, where you can format, version, and export it.

Actions under your own messages

- **Edit message** — reopens your prompt for editing. Sending the edit removes the later replies and asks again from that point.
- **Load prompt in new chat** — copies the prompt and its attachments into a fresh chat's composer, ready to adjust and resend.

Diagrams in replies

When a reply contains a diagram (a flowchart, sequence, or similar), it renders as a real figure, not a code block. The figure's toolbar offers **Copy** for the diagram source, **PNG** and **SVG** to download it as an image file, and **Code** to flip between the drawing and its source text.

Generated images

- When the selected model can generate images, just ask for one — the finished images appear directly in the reply.
- Click **Download** under an image to save it as a file.

Citations and token usage

- When a reply used the web or your knowledge bases, numbered citations appear with it. Click a citation to open the original source.
- A small token count appears under a reply only when the model provider actually reported one — the platform never shows an estimate dressed up as a real number.

WALKTHROUGH · 1:43 · [HELP](#) · [LEARN](#)

Follow the work trace and act on replies

Read the steps behind a reply, rate it with a note, branch it into a new chat, and turn it into an editable document.

Before you begin

- A chat with at least one finished reply (see Start chatting).

Step by step

- 1 Under a finished reply's heading, find Work trace and its summary, such as complete · 4 steps.
- 2 Choose Work trace to expand it and read every step: Routing request names the model, Preparing context lists the sources, then Generating answer and Finalizing response.
- 3 Choose Work trace again to collapse it.
- 4 Choose Send positive feedback (thumbs up) or Send negative feedback (thumbs down) under the reply.
- 5 Optionally type a note in the box that opens, then choose Send note. Note sent confirms it.
- 6 Choose **Branch** response into new chat to continue from this reply in a separate chat.
- 7 Choose **Transfer to Drafts** to open the reply as an editable document.

Path: Rate a reply

- 1 Choose the thumbs up or thumbs down under the reply. Positive feedback sent (or Negative feedback sent) appears.
- 2 In What worked well? or What went wrong?, add an optional note of up to 2,000 characters.
- 3 Choose Send note, or close the box to skip the note.

Path: Branch into a new chat

- 1 Choose **Branch** response into new chat under the reply.
- 2 A new chat opens, titled **Branch:** and the original title, holding the prompt and this reply.
- 3 Continue there; the original chat is unchanged.

Path: Transfer to Drafts

- 1 Choose **Transfer to Drafts** at the right of the reply's actions.
- 2 **Drafts** opens a new document titled after the chat, saved as Version 1, with the reply and any citations.
- 3 Edit it, save versions, and export it (see **Draft** documents).

Check that it worked

- ✓ The expanded trace lists the model used in Routing request.
- ✓ Note sent appears beside the reply's actions after sending a note.
- ✓ A chat titled **Branch: ...** appears under **Recent**.
- ✓ **Drafts** opens with the reply as a saved document, and the assistant rail says it was transferred.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
The running trace moves to the next step before that step seems done	While a reply runs, the trace advances through the planned steps on a timer. Treat it as a guide to what the request involves; the finished trace is the record.
"Could not branch"	The branch could not be created, often because the connection dropped. Try again when the chat has saved.
No actions under a reply	Actions appear only on finished replies. A stopped or failed reply has none; resend with Edit message .
Transfer to Drafts is missing	It appears on finished replies with text. Wait for the reply to finish.

11. Symbol shortcuts: / @ # \$ >

Five characters that pull prompts, agents, knowledge, skill files, and automations into a message.

Start a word in the message box with one of five symbols and a menu opens with matching items. You can also choose **Send options** → **Resources** to browse without memorizing a symbol. Forgot one? The bottom of Session details (①) lists all five, and **Help** → Symbol shortcuts: / @ # \$ > shows each menu in action.

TYPE	WHAT IT OPENS
/	Saved prompts (inserted as text) and enabled MCP connections (turned on for this message).
@	Agent profiles — the reply routes through the one you pick.
#	Knowledge bases and the files inside them. Picking one turns on Knowledge for that source.
\$	Saved skill files. The one you pick is attached to the message.
>	Automations, including paused ones. The one you pick runs on your message when you send.

- Type the symbol, then keep typing to filter the list.
- Use the ↑ and ↓ arrow keys to move through the menu.
- Press Enter to insert the highlighted item, Tab to complete it, or Escape to dismiss the menu.

Browse resources

On the **Resources** tab, use **Find a resource** to search by name, or select All, **Knowledge** bases, Files in knowledge sources, MCP connections, **Prompts**, **Agents**, Skill files, or Automations. A selected connection or automation shows a check; choose it again to remove it. Read an empty or failed-loading state before retrying. The **Type shortcuts in chat** reference explains each symbol below the results.

TIP

The expanded composer focuses on editing a longer message. Return to the normal composer to browse resources, adjust send options, and send with the visible controls.

WALKTHROUGH · 1:49 · HELP > LEARN

Symbol shortcuts: / @ # \$ >

Type a symbol in the message box to add a saved prompt, an agent, knowledge, a skill file, or an automation, then send.

Before you begin

- A chat with an available model.
- Items to choose from: saved prompts and skill files are shared by your administrators; agents, knowledge bases, and automations are ones you can use or have built.

Step by step

- 1 Click in the message box. At the start of a word, type one of the five symbols: / @ # \$ or >.
- 2 A list opens above the message box. Keep typing to narrow it.
- 3 Use the up and down arrow keys to move, and Enter (or a click) to choose. Escape closes the list.

- 4 Check what was added: prompt text in the box, or a chip under it (**Agent**, **Knowledge**, a skill file, or an automation).
- 5 Finish your message and send it.

Path: / Prompts and MCP connections

- 1 Type / to list saved prompts and enabled MCP connections.
- 2 Choose a prompt to put its text in the message box; edit it before sending.
- 3 Choose a connection to turn it on for this message.

Path: @ Agents

- 1 Type @ to list the agent profiles you can use.
- 2 Choose one. @ and its name are added, and the **Agent** chip appears.
- 3 The reply runs through that agent's model, instructions, knowledge, and tools.

Path: # Knowledge

- 1 Type # to list knowledge bases and the files in them.
- 2 Choose a knowledge base, or a single file to search its source.
- 3 The **Knowledge** chip appears; the reply searches it and cites what it uses.

Path: \$ Skill files

- 1 Type \$ to list saved skill files, such as citation rules or an approval checklist.
- 2 Choose one. It is attached to the message as a file chip (Name.md).
- 3 The model follows it for this message.

Path: > Automations

- 1 Type > to list your automations, including paused ones.
- 2 Choose one. It is queued in the chip under the message box.
- 3 When you send, the automation runs on your message. x removes it before sending.

Check that it worked

- ✓ After choosing, the message box holds the prompt text or shows the matching chip.
- ✓ A message sent with # and \$ returns a reply that lists its sources and View N citations →.
- ✓ Session details > Symbol shortcuts lists all five symbols at any time.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"No matching commands. Try another name."	Nothing matches what you typed after the symbol. Delete some letters, or check the item's name.
"No saved prompts or enabled MCP connections yet — add them under Tools"	Nothing is shared with you yet. Your administrator adds prompts and connections.
"No agent profiles available yet — create one under Agents"	Build an agent (if your account can) or ask your administrator to share one.
"No knowledge bases are enabled yet — add one under Knowledge"	Ask your administrator to share a knowledge base, or create one if your account can.
"Some files could not be loaded. Retry"	Choose Retry in the list. Sources that already loaded stay usable.
The list does not open	The symbol must start a word: type a space first if it follows other text.

12. Attaching files, web pages, and connected sources

Upload from your computer, fetch a page by link, or pull from your organization's drives.

- Click the paperclip button at the bottom-left of the composer.
- Click **Upload from computer** to attach files from this device. The model reads them as context for your next message.
- Or click **Web page by link** and paste one or more web addresses — the platform fetches up to 3 public pages and attaches them as cited sources for this message.
- Attached items appear as chips above the composer. Click the x on a chip to remove one before sending.

Attach from source

Below those options, under the **Attach from source** heading, the menu lists your organization's workspace sources: Google Drive, OneDrive, SharePoint, Box, and iManage. Availability depends on the service team's connector setup and your own source-account access.

Connecting your own account

- Cloud sources read from your own account. The first time you pick one, the picker may ask you to connect it — click the Connect button for that source and a sign-in window opens.

- **Approve** read access in that window. The file list refreshes on its own once access is granted, and only you can see files from your account.
- The picker shows files at the top level of the drive or the folder configured for the source, within your account's permissions.

GOOD TO KNOW

If a source needs shared configuration, ask your administrator to coordinate with the service team. Shared connector administration is managed outside tenant administration. Your own Connect action in the attach menu remains separate and grants access only through your source account.

WALKTHROUGH · 2:08 · HELP › LEARN

Attach files and sources

Attach a file from your computer, a web page by link, or a file from a connected cloud source, and know what the model actually receives.

Before you begin

- A chat with an available model.
- For cloud sources (Google Drive, OneDrive, SharePoint, Box, iManage): your administrator must set up the connector, and you sign in to your own account the first time.
- Files up to 25 MB each. **Web** pages must be public addresses starting with <http://> or <https://>.

Step by step

- 1 In the message box, choose Add attachment (the paperclip).
- 2 Choose **Upload from computer** and pick one or more files.
- 3 Wait for the file chip to finish uploading. Remove a file with the x on its chip.
- 4 Type your question about the file and send it.
- 5 Check the reply: it uses the file's content, and the trace includes a step for reading attachments.

Path: Web page by link

- 1 Paperclip › **Web page by link**.
- 2 Type or paste the full address into **Web** page address, such as <https://example.com>, and choose Add link (or press Enter).
- 3 The link appears as a chip. Add up to 3 links per message.
- 4 Send your question. The page is fetched when you send and cited in the reply (View 1 citation →).

Path: Cloud sources

- 1 Paperclip › choose a source under **Attach from source**: Google Drive, OneDrive, SharePoint, Box, or iManage.
- 2 The first time, choose Connect and approve read access in the sign-in window. Only you see files from your account.
- 3 Tick up to 10 files and choose Attach selected.
- 4 If the dialog says the source is not ready, your administrator must finish setting up that connector.

Path: What the model receives

- 1 Word (.docx) and text files (.txt, .md, .csv, .json, .xml, .html, .rtf, .eml, .log): the opening text, up to about 2,400 characters.
- 2 Images: the picture itself, if the selected model can read images; otherwise the model says it cannot see them.
- 3 Audio and video: a transcript, when your workspace has a transcription model. Media longer than 45 minutes is refused.
- 4 PDF, Excel, and PowerPoint files from your computer: currently only the file name. To ask about their content, add them to a knowledge base, attach them from a cloud source, or open Word files in **Drafts**.

Check that it worked

- ✓ The file chip finishes uploading without Upload failed.
- ✓ The reply states facts that appear only in your file or page.
- ✓ A web-page reply offers View 1 citation → for the attached page.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Attachment exceeds the 25 MB chat upload limit."	Split or compress the file, or add it to a knowledge base, which accepts larger files.
"Could not upload the attachment. Check your connection and try again."	Check your connection, remove the failed chip, and attach the file again.
The reply says it cannot see the content of a PDF, spreadsheet, or slide file	Files of those types uploaded from your computer currently reach the model by name only. Use a knowledge base or a cloud source instead.
"Enter a full web address starting with http:// or https://."	Paste the complete address, including https://.
"You can attach up to 3 web links per message."	Send the first three, then attach the rest in a follow-up message.
"Web source fetch failed for ..."	The page could not be fetched (it may need a sign-in or be offline). Copy the text into your message instead.
"Google Drive is not ready." (or another source)	Your administrator has not finished that connector. Ask them, or upload the file from your computer.
"Media longer than 45 minutes cannot be transcribed in this version."	Split the recording and upload the part you need.

13. Send options: Knowledge, Web, Agent, and Reasoning

Control what the model may use, one reply at a time.

- Click the chevron (v) beside the paper-plane send button.
- The menu opens with **Send now** at the top — that simply sends with the current settings.
- Below it are three switches you can turn on or off for the next reply, and a **Reasoning** slider at the bottom.

SWITCH	WHAT IT DOES WHEN ON
Knowledge	The reply searches your enabled knowledge bases and cites what it finds, with links back to the sources. A picker lets you choose which bases to use.
Web	The reply uses public web search. Results come back as clickable citations, and the sources are listed in session details. Available when the selected model supports it.
Agent	The reply uses the selected agent profile and its permitted tools. Turning this on can choose the first available profile if none is selected; confirm the Agent profile field before sending, or choose a profile with @.

Reasoning level

The **Reasoning** group at the bottom of the menu is a three-position slider from **Fast** to **Smart**. **Fast** favors quicker answers; **Smart** makes the model think longer for more detailed output; the middle is a balanced default. The slider is only active when the selected model supports reasoning levels — otherwise it is disabled and says so.

Whenever tools are on, the active-tools chip appears next to the paperclip: the single tool's name, or **Tools** with a count. Click its x to turn off **Knowledge**, **Web**, and **Agent** and clear selected MCP connections and queued automations for the next message.

On the **Reply settings** tab, **Stream replies** controls whether text appears as the model produces it or only after the reply finishes. The **Resources** tab opens the searchable resource picker; **MCP connections and resources** also opens it from **Reply settings**. These controls change how the next reply is prepared or displayed; they do not supply missing provider or connector access.

WALKTHROUGH · 2:38 · [HELP](#) · [LEARN](#)

Knowledge, Web, Agent, and reply settings

Ground a reply in a knowledge base and check its citations, and choose Web, an agent, reasoning, streaming, and resources for a message.

Before you begin

- A chat with an available model.
- For **Knowledge**: at least one knowledge base shared with you (see **Knowledge** bases).
- For **Web**: your workspace must connect web search to the model you use. Otherwise **Web** is greyed out.
- For **Agent**: an agent profile you can use (see **Agent** profiles).

Step by step

- 1 In the message box, choose **Send options** (the arrow beside the paper plane). **Send options** opens on **Reply settings**.

- 2 Choose **Knowledge**, then choose the source to search, such as Litigation Playbook. Each source shows how many files it holds.
- 3 Choose Close send options. The **Knowledge** chip under the message box shows it is on.
- 4 Type your question and send it.
- 5 Check the reply: the work trace lists one more step, and View 2 citations → appears under the answer.
- 6 Choose View citations to open Session details › **Sources gathered**, with each source's K number, file name, and the passage used.

Path: Web

- 1 **Send options** › **Web**. When it is available, the chip under the message box reads **Web search**.
- 2 Send your message. Aperture Chat searches public web sources and gives the results, with their addresses, to the model so the reply can cite them.
- 3 If **Web** is greyed out, your workspace has not connected web search to this model. Ask your administrator.

Path: Agent

- 1 **Send options** › **Agent**. An **Agent profile** list appears under it.
- 2 Choose the profile you want; the line under it counts its knowledge, files, and tools.
- 3 Send your message. The trace reads **Agent** work trace. See **Agent** profiles to build your own.

Path: Reasoning and streaming

- 1 Under **Reasoning**, drag the slider toward **Fast** for quicker answers or **Smart** for deeper thinking, on models that support it.
- 2 Clear **Stream replies** to show each reply only when it is complete; leave it ticked to watch the answer arrive.

Path: Resources

- 1 **Send options** › **Resources** (or **MCP connections and resources**).
- 2 **Search** by name, or filter by **Knowledge** bases, Files in knowledge sources, MCP connections, **Prompts**, **Agents**, Skill files, or Automations.
- 3 Choose an item. A prompt's text fills the message box; other items turn on for this message.

Path: Turn everything off

- 1 Choose x on the chip under the message box.

- 2 **Knowledge, Web, and Agent** turn off, selected MCP connections clear, and queued automations are removed.

Check that it worked

- ✓ The chip under the message box names what is on (**Knowledge**, **Web search**, **Agent**, or **Tools** with a count).
- ✓ A knowledge reply shows View N citations →, and **Sources gathered** lists the files and passages.
- ✓ The reply's trace includes a step for each option you turned on.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
Web is greyed out: "Web search is turned off for this model by your workspace configuration"	Your workspace has not connected web search to this model. Choose another model, or ask your administrator.
With Web on, the reply stays on Aperture Chat is working	Some self-hosted models currently reject web-search requests. Choose · stop, turn Web off, and resend; tell your administrator.
"Web search (...) is unavailable: ... Turn off Web for this reply to continue without it."	The search service could not be reached. Turn Web off for this message, or try again later.
Knowledge cannot be chosen	No knowledge base is shared with you. Ask your administrator, or create one if your account can (see Knowledge bases).
"Create an agent profile in Agents to enable this mode."	No agent profile is available to you. Build one under Agents , or ask your administrator to share one.
"... does not support reasoning levels"	The selected model has no reasoning control. The slider has no effect for it.

14. Session details and sources

Real token usage, active tools, every source this chat gathered, and the symbol shortcuts.

- Click the round information button (i) at the right end of the chat's top bar.
- The session details panel opens for the current chat.
- **Current chat** and **Model** — the user-message and response counts, and which model is active.
- **Tokens used** — shown only when the model provider reported real numbers. Otherwise the panel says "Not reported by the provider".

- **Tools** — exactly what is switched on for your next message: web search, knowledge, agent mode, plus any automations or MCP connections you added in the composer. “Off” means nothing is active.
- **Sources gathered** — every web source (as clickable links) and workspace citation collected during this conversation.
- Symbol shortcuts — at the bottom of the panel, the five symbols you can type in the message box and what each one pulls in. See “Symbol shortcuts: / @ # \$ >”.

The **Context window** meter describes how much context this conversation uses. For older conversations without provider counts it can show $\approx$ and estimated from message length. Treat that as an estimate; **Tokens used** remains based on actual provider reports. When the context window fills, start a new chat and include the details it needs to continue reliably.

WALKTHROUGH · 0:59 · HELP > LEARN

Session details and context

Open Session details to check the model, tokens used, how full the context window is, the sources a reply used, and the symbol shortcuts.

Before you begin

- A chat with at least one reply. **Sources gathered** appears after a reply that used knowledge or the web.

Step by step

- 1 In a chat, choose Session info (the i button at the top right). Session details opens on the right.
- 2 Read Session summary: **Current chat**, **Tokens used**, **Model**, **Knowledge** bases, **Tools**, and **Agent profile**.
- 3 Read **Context window**: the meter and how many of the model's tokens this chat already uses.
- 4 Scroll to **Sources gathered** to see every source a reply used, with its K (knowledge) or W (web) number and the passage.
- 5 Scroll to Symbol shortcuts for the five symbols you can type in the message box.
- 6 Choose Close session details when you are done.

Path: From a reply's citations

- 1 Under a reply, choose View N citations →.
- 2 Session details opens at **Sources gathered** for that reply. Hide citations closes it again.

Check that it worked

- ✓ **Tokens used** shows the counts the provider reported (in, out, and total), or Not reported by the provider.

- ✓ The context window meter shows a percentage and tokens used of the model's total.
- ✓ **Sources gathered** lists the same number of sources as View N citations.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
Tokens used reads "Not reported by the provider"	The model service did not return usage. Nothing is wrong with your chat; the context meter may use an estimate instead.
The context meter is marked estimated or ≈	It is based on message length, because the provider did not report usage for some messages.
The context window is nearly full	Start a new chat for a fresh window. Older details may otherwise matter less to the answer.
No Sources gathered section	The replies in this chat did not use knowledge or the web. Turn on Knowledge in Send options and ask again.

15. Personal data in your chats

Locked chips for concealed values, what the model receives, and how rated answers are kept.

Your organization can conceal personal data in chats. When it does, the line under the message box ends with Personal data is concealed. As soon as you send a message, values such as Social Security, card, and account numbers, email addresses, phone numbers, health identifiers, and passwords or keys are replaced with a locked chip such as SSN or Email. The original value is not saved anywhere in the workspace, so it cannot be recovered from the chat later.

YOU SEE	WHAT IT MEANS
A locked chip in your message, such as SSN	The value was concealed before the chat was saved. Reopening the chat shows the chip, not the value.
[SSN] or [EMAIL] in a reply	Your organization keeps values from the model, so the model received a placeholder. Add the real value yourself in the system where it belongs, not in the chat.
A locked chip in a reply	Your organization lets the model read values for work that needs them. A value the model writes is concealed before you see it.
A note in the rating box about a de-identified copy	Your organization captures training examples from rated answers and notes. The copy is de-identified and never sent to a model provider.

GOOD TO KNOW

Concealment recognizes values by their format and check digits. It does not recognize names or free-text descriptions of someone's health, and it does not alter drafts or uploaded files. Follow your organization's policy for those.

WALKTHROUGH · 1:15 · [HELP](#) · [LEARN](#)

Personal data in your chats

Recognize the personal data your organization conceals, know what the model received, and see how rated answers are kept when your organization captures training examples.

Before you begin

- Your administrator has turned on **Personal Data Protection**. If not, none of this appears and chats are stored as typed.
- An available model.

Step by step

- 1 Look under the message box. When your organization protects personal data, the line ends with Personal data is concealed. Point at it for a short explanation.
- 2 Send your message as usual. Social Security, card, and account numbers, contact details, health identifiers, and secrets become locked chips, such as SSN or Email, as soon as the message is sent.
- 3 Read the reply. When your organization keeps values from the model, the model received placeholders only, so its draft shows them, for example [SSN], where a value belongs.
- 4 Add the real values yourself in the system where they belong, not in the chat.
- 5 When you rate a reply with the thumbs buttons, read the note box. If your organization captures training examples, it says so: a de-identified copy of rated answers and notes is kept and never sent to a model provider.

Path: Values kept from the model (shown)

- 1 Your sent message shows locked chips.
- 2 The model received placeholders instead of the values, so its reply refers to them, for example [SSN] or [EMAIL].

Path: Model allowed to read values (shown)

- 1 Your sent message still shows locked chips.
- 2 The model read the values for that turn, so it can use them, for example in a signature block.

- 3 Any value it writes is replaced with a locked chip before you see the reply and before the chat is saved.

Path: Rating answers when training capture is on (shown)

- 1 Choose thumbs up or thumbs down under a reply. The note box says your organization keeps a de-identified copy of rated answers and notes, and that it is never sent to a model provider.
- 2 Type a note and choose Send note, or close the box with x. The rating counts either way.
- 3 Corrections you type in the chat, such as "That's wrong, the deadline is 60 days. Please revise.", can be kept the same way.

Check that it worked

- ✓ The line under the message box ends with Personal data is concealed.
- ✓ Your sent message shows a locked chip where you typed the value.
- ✓ Reopening the chat later still shows the chip, not the value.


Troubleshooting

WHAT YOU SEE	WHAT TO DO
Your message shows the value for a moment after you send it	The concealed copy replaces it within a moment. The original value is not saved.
The reply says [SSN] or similar where you expected a value	The model never received the value. Fill it in yourself where it belongs; do not paste it back into the chat.
A name or a description of someone's health was not concealed	Concealment recognizes numbers, codes, and contact details by their format. It does not recognize names or free text about health, so leave them out when your policy requires it.
A draft still contains a value	Drafts are documents of record and are not altered. Keep personal data out of drafts unless your policy allows it.
You need the original value from an earlier chat	It was never saved, so the chat cannot give it back. Use the system of record it came from.

Drafts, decks, and agents

16. Drafting documents

Generate, edit, version, and export documents with the drafting assistant.



Drafts combines the document editor with a separate assistant and document controls. Example workspace with synthetic data.

Click **Drafts** in the sidebar. The view pairs a full document editor with an AI assistant rail on the left that writes into the editor for you. A **Draft format** switch at the top of the editor has two buttons — **Document** and **Deck**. This section covers document mode; “Building a slide deck” covers deck mode.

GOOD TO KNOW

If the model selector says No models connected, AI drafting, AI rename, inline AI edits, and AI deck prompts are unavailable. You can still edit manually, import, save, use document history, and export. Follow the setup guidance for your role or ask your administrator for access to a working model.

Generating a draft

- Look at the Sources, **Web**, and Templates chips at the top of the assistant rail first: they show which context is on before you generate anything.

- In the assistant's message box, describe the document you need: a memo, an engagement letter, a summary. Be specific about audience and tone if it matters. The suggested requests under the heading fill the box for you.
- The assistant writes the document directly into the editor. You can also send any chat response here with **Transfer to Drafts**.
- Use the model selector in the toolbar area to choose which approved model writes and edits this draft.

Opening a file from this device

Click the paperclip in the assistant's message box. Under From this device are two choices. Attach to chat gives files to the assistant to read with your next request; they become sources, not the document. Open in editor opens one file as the document itself, so you can edit it directly or select text and use **Ask AI**.

- Open in editor accepts Word (.docx, .dotx, .doc), **Markdown** (.md), plain text (.txt), and web page (.html) files. Word and web pages keep their formatting, and **Markdown** keeps its headings, lists, and emphasis. When something could not be carried over, an import note in the assistant rail says what.
- The file opens as a new draft and is saved to your account as one. The draft you had open stays in **Document history**; if it has unsaved changes, the unsaved-changes dialog asks first.
- Open in editor is unavailable while the assistant is still working on a request. In deck mode the same choice opens a PowerPoint file; see “Building a slide deck”.

Editing like a document

The main toolbar keeps undo and redo, block style, bold, italic, underline, and inline AI editing close at hand. **Text** opens font, size, advanced styles, color, highlighting, and **Clear formatting**. **Paragraph** opens alignment, lists, and quotes. **More** opens **Copy document** and the **AI edit trail**. Select the intended text before opening a panel; Escape returns to the editor.

- The plus (+) **Insert** button opens links and citations alongside **Web image**, **Chart**, **Table**, **Divider**, and **Page break**. A manual page break starts a new page at your cursor; the editor also paginates longer documents automatically.
- Selecting text shows a floating toolbar beside it with **Ask AI**, block style, bold, italic, underline, strikethrough, link, and highlight.
- Type / at the start of a line for the command menu: AI actions such as **Continue writing** and **Summarize the document**, headings, lists, quotes, a table, a divider, a page break, a web image, or today's date. Keep typing to filter and press Enter to insert.
- Typing shortcuts format as you go: start a line with “#”, “##”, or “###” and a space for a title, heading, or subheading; “-” or “1.” for a list; or “>” for a quote. Type “---” and press Enter for a divider, and wrap words in “**”, “*”, or “`” for bold, italic, or code.
- Click inside a table to insert or delete rows and columns, turn the header row on or off, or delete the table. Click a picture to set its size (Small, Half, Large, or Full width), alignment, and alt text, or to delete it.

- Pasting from Word, Google Docs, or a web page keeps headings, lists, tables, links, pictures, and emphasis but drops the source's fonts, sizes, and colors so the text matches your document. Ctrl+Shift+V (⌘⇧V on a Mac) pastes plain text.

Editing with AI

- Select a passage and choose **Ask AI** in the floating toolbar, or press Ctrl+J (⌘J on a Mac). With nothing selected, the same command writes new text at the cursor.
- Pick an action — **Improve writing**, **Fix spelling & grammar**, **Make shorter**, **Make longer**, **Simplify language**, **Turn into a bulleted list**, **Turn into a table**, **Summarize**, or **Use active voice** — or a tone or language, or type your own instruction.
- Review the suggestion. **Changes** marks what was added and removed; **Result** shows the finished text. Nothing in the document changes yet.
- Choose **Replace**, **Insert below**, **Try again**, or **Discard**. To adjust the suggestion instead, type a follow-up such as “shorter” or “add a number” in the box underneath.

An accepted AI edit stays highlighted briefly and is recorded with the document; **More → AI edit trail** highlights the recorded edits again so you can review what changed.

Revise the document you already have

Describe the change in the assistant: shorten a section, change tone, apply a template, or transform the existing paper. Revision requests work from the current document. Ask explicitly to start over, replace, or clear it when that is your intention. The editor keeps instructions separate from the deliverable; review the resulting document and use Undo or a saved version when a change needs correction.

For an existing paper, **Assistant settings → Apply MLA layout** applies double-spaced 12-point Times New Roman, a student heading, centered title, body and reference indents, and a separate Works Cited page without asking a model. It preserves the paper's text and supports Undo. Fill missing student details and verify quotations and sources yourself; formatting is not citation verification.

The page navigator shows the current sheet and total pages outside the document text. Older drafts are paginated when opened, and manual edits reflow after you leave the editor. Page labels do not become part of the exported body.

Status bar, outline, find, and zoom

- The status bar under the page shows the current page, word and character counts, and reading time. When text is selected, it counts the selected words too.
- **Outline** lists the document's headings; select one to jump to that section.
- Find (Ctrl+F, or ⌘F on a Mac) searches the document, with **Match case** and whole-word options. **Find and replace** (Ctrl+H, or ⌘⇧H on a Mac) replaces one match at a time or all of them.
- The keyboard button lists every editor shortcut (Ctrl+/, or ⌘/ on a Mac). The zoom controls resize the page on screen from 50% to 200% without changing the document or its exports.

Versions and comparing

- Click **Save version** whenever the draft reaches a good state. You can restore an earlier version at any time from the draft's history.
- **Compare versions** opens a read-only visual redline of two saved versions, so you can see exactly what changed between them. It needs two genuinely different saved versions before it activates.
- Open **Document history** to browse saved documents and decks. Hover or keyboard-focus an entry for a preview. **Archive** moves finished work into **Archived**, where **Unarchive** restores it; **Delete** requires confirmation and permanently removes the saved item and revisions.
- Read the save status: account-backed documents and decks can be reopened after sign-in on another device. A pending or failed save stays a working copy in this browser. If the account version changed elsewhere, resolve the reported conflict before treating the save as complete.

Leaving with unsaved changes

When you leave an edited draft through workspace navigation, open another chat or draft, or choose **Sign out**, the unsaved-changes dialog lets you Keep editing, Save copy and continue, or **Discard** and continue. The saved copy goes into local document history in this browser. If browser storage fails, the app keeps you in the draft so you can recover it. A forced security sign-out can still interrupt work; save important changes regularly.

Exporting

- Click **Export** to open the export panel. In document mode it offers a **Word document** (an editable Word file with preview page breaks and embedded images), **Markdown** (best for plain text or web publishing), and **Print / Save as PDF**, which opens your browser's print dialog with the saved version — choose “Save as PDF” there to keep a PDF copy.
- Exports always run from a saved version. If you have unsaved edits, the panel says **Save your edits first** and offers a one-click **Save version and export** (or **Save version and print**) button.
- When your browser supports it, choose **Choose a location** to select a destination, or **Browser downloads** to use its download folder. Word and print retain document formatting; **Markdown** preserves the text without its typography. Exporting a file does not by itself upload pending work to your account.

TIP

On a narrow window, the assistant rail becomes a slide-out drawer so the editor gets the full screen. Open it with the pen tab on the left edge or Ctrl+. (⌘. on a Mac); resting the mouse on that edge previews it. Press Escape or click the dimmed area to close it.

WALKTHROUGH · 3:34 · [HELP](#) · [LEARN](#)

Draft documents

Start a document from a blank page, a chat reply, or a file; have the assistant write it; edit with AI; save a version; and export it.

Before you begin

- An available model for AI writing and editing. Without one, you can still type, open files, save, and export.
- To open a file in the editor: a Word (.docx or .doc), **Markdown** (.md), text (.txt), or web page (.html) file.

Step by step

- 1 Choose **Drafts** in the sidebar. A blank document opens, with the **Document** Assistant on the left.
- 2 Type a title in the title field at the top of the page.
- 3 In Ask the document assistant, describe the document you need, or choose a suggestion. Choose Apply instruction (the arrow).
- 4 Wait while the assistant writes on the page. Ready to write returns when it is done.
- 5 Check the title. If the assistant renamed the document, type your own title again.
- 6 Select a passage, choose **Ask AI** in the floating toolbar (or press Ctrl+J, ⌘J on a Mac), and choose an action such as **Make shorter**.
- 7 Review the **Changes** or **Result** view, then choose **Replace**, **Insert below**, **Try again**, and **Discard** are the alternatives.
- 8 Choose **Save version** and wait for Saved.
- 9 Choose **Export**, set Save to to **Browser downloads**, and choose **Word document**. The receipt names the downloaded file.

Path: From a chat reply

- 1 Under a finished chat reply, choose **Transfer to Drafts**.
- 2 **Drafts** opens a new document with the reply, saved as Version 1.
- 3 Continue from step 6 of the main steps.

Path: Open a file from this device

- 1 Choose Attach file (the paperclip) in the assistant's message box.
- 2 Under From this device, choose Open in editor, and pick a .docx, .doc, .md, .txt, or .html file.
- 3 The file opens as a new draft with its headings and lists. The draft you had open stays in **Document history**.
- 4 Attach to chat instead gives the file to the assistant to read with your next request.

Path: Write it yourself

- 1 Click on the page and type. **Markdown** shortcuts work: **#** and a space starts a title, **-** and a space starts a list.
- 2 Type **/** at the start of a line for the command menu: headings, lists, tables, page breaks, and AI actions.
- 3 Use **Text**, **Paragraph**, **More**, and **Insert** above the page for formatting.

Path: Export formats

- 1 **Word document**: an editable .docx with page breaks and images.
- 2 **Markdown**: a .md file for plain text or web publishing.
- 3 **Print / Save as PDF**: opens the browser's print dialog; choose Save as PDF there.
- 4 Save to: **Browser downloads** saves straight to your downloads; **Choose a location** asks the browser where to save.

Check that it worked

- ✓ The toolbar shows Saved and the version number after **Save version**.
- ✓ The export receipt names the file, for example synthetic-vendor-onboarding-checklist.docx, and offers **Download** again.
- ✓ An opened file shows a note in the assistant rail: the file opened in the editor.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"AI drafting is unavailable." or Model: No models connected	No model is available to you. You can still type, open files, save, and export. Ask your administrator for model access.
The document title changed to a template name, such as Board Update Draft	The assistant can rename a new document after a template it matched. Type your own title back into the title field, then save.
"Could not open ... Choose a .docx, .doc, .md, .txt, or .html file."	Save the file in one of those formats and open it again. For PDFs, use a knowledge base or attach the file to chat.
"Wait for the assistant to finish before opening another document."	Let the current request finish, then open the file.
Export says "Save your edits first"	Choose Save version and export (or Save version and print). Exports use the saved version.
"Local only — server save failed"	Your changes are kept in this browser. Choose Retry, or export a copy before you leave.
"Could not export ... Try again from the export menu."	Try again , or choose Browser downloads in Save to.

17. Building a slide deck

Turn a draft into slides, edit them, brand them, and export real PowerPoint.

Deck mode turns the **Drafts** workspace into a slide editor with a real PowerPoint export. Click the **Deck** button on the **Draft format** switch. If the draft already has document content, a dialog asks "Turn this draft into slides?" — choose **Start a blank deck** to keep the document untouched, or **Convert into slides** to split its headings and text into slides. Your document versions are kept either way.

The filmstrip

- Slide thumbnails run along the side. Click one to select it, or drag it to reorder the deck.
- Hover a thumbnail for its actions: move up, move down, duplicate, and delete (undo restores a deleted slide).
- **Add slide** at the end of the filmstrip opens the layout menu for a new slide.
- **Slide sorter**, at the right of the deck toolbar, lays every slide out in a grid. Drag to reorder, double-click a slide to edit it, and use each slide's duplicate and delete buttons.

Slide layouts

Above the slide, **Layouts** previews seven arrangements: **Title**, **Title + bullets**, **Two columns**, **Image + caption**, **Section**, **Quote**, and **Closing**. Select one to switch the selected slide; the layout button on the slide and in the toolbar opens the same choices. **Themes** offers five color palettes — Aperture, Midnight, Editorial, Violet, and Forest; a palette change recolors every slide, keeps text, media, and layout, and can be undone.

Editing slide text

- Click any text region on a slide to edit it in place. The deck toolbar carries undo and redo, bold, italic, underline, and **Text** for font, size, and color; **More** copies the deck outline or lists the keyboard shortcuts.
- Drag a text box by its frame to move it. Guides snap it to the slide's edges and center and to other boxes; arrow keys nudge the selected box (Shift moves ten times as far).
- Inside bullet lists, press Tab to indent a bullet one level and Shift+Tab to outdent it.

Editing slides with AI

- **Edit slide with AI** (the AI pen in the deck toolbar, or Ctrl+J / ⌘J) reworks the whole slide: **Improve this slide**, **Make it punchier**, **Cut the text in half**, **Write speaker notes**, **Pick a better layout**, **Split into two slides**, **Fix spelling & grammar**, a tone or language, or your own instruction.
- The review shows the slide before and after, plus any new speaker notes. Choose **Apply to slide**, **Try again**, or **Discard**; undo restores the previous slide.
- To rewrite only part of a slide, highlight the text and choose **Ask AI**, then review the suggestion before replacing it.

Opening a PowerPoint file

- In deck mode, click the paperclip in the assistant's message box and choose Open in editor (Edit slides directly) to open a .pptx file as the deck itself.
- Each slide keeps its title, subtitle, bullets with their indent levels, two-column text, speaker notes, one picture, and its layout artwork as the background. **Text** with no place in the chosen layout moves into that slide's speaker notes instead of disappearing.
- If the presentation is larger than the deck limit, pictures are left out first and then layout artwork, and the assistant rail says so. The opened deck is a new deck; the deck you had open stays in **Document history**.

Starter templates and your brand template

- **Deck starters & brand themes**, above the slide (or the Templates chip in the assistant rail), opens the templates panel with five starter structures: **Pitch deck**, **Quarterly review**, **Project kickoff**, **Client proposal**, and **Training session**. Pick one and start it to get scaffold slides you replace.
- A **Use templates in chat** toggle (off by default) lets the selected template guide the deck assistant when it drafts for you.
- **Upload brand template** accepts your firm's own file — “.pptx or .potx — brand colors, fonts, logo, and every slide's text are extracted.” The extracted theme is stored only on this device.

- **Once** a brand theme is loaded, **Apply to deck** restyles the current deck with its colors, fonts, and logo, and a “Load all N slides” button can replace the deck with every slide from the uploaded template.

Backgrounds and AI images

- The background button in the toolbar opens a per-slide menu: **Upload background...** (a PNG, JPEG, or WebP from this device), **Use on every slide**, **Remove from this slide**, and **Clear from all slides**.
- **Generate AI slide image** writes an image prompt for you, prefilled from the slide's content; press **Generate image** and the result becomes that slide's background. This needs an image-generation model enabled for your workspace — the button says so honestly when none is.
- In the assistant rail, **Toggle AI slide images** makes the assistant also generate an image for every slide whenever it drafts a deck.

Speaker notes

Every slide has a **Speaker notes** field below the stage. **Notes** export with the deck: they land in each slide's notes pane in the PowerPoint file, they appear in the **Markdown outline**, and they follow you into presentation mode.

Presenting from Aperture

- Click **Present deck** in the deck toolbar (the monitor icon), or press F5 (⌘Enter on a Mac); add Shift to start from the current slide. The deck opens full screen.
- Advance by clicking the slide or with the arrow keys or Space; move back with the left arrow. Type a slide number and press Enter to jump to it.
- The **Notes** button (or N) shows your speaker notes under the slide. **Presenter view** (or P) shows the current slide, the next slide, a timer you can pause or reset, and your notes side by side.
- Press B or W to black or white out the screen; press it again to return. Leave with Escape or the **Exit** button in the top corner.

Saving and exporting

- **Save version** stores decks to your account with revision history, like documents. Read the status to confirm the save reached the server. A local working copy or pending save remains only in this browser until synchronization succeeds; use the **Only on this device** reminder to find unfinished saves.
- **Document history** includes decks. Preview, reopen, archive, unarchive, and confirmed deletion use the same account controls as documents. A concurrent account edit can prevent a stale save or deletion; read the conflict message before retrying.
- Click **Export** to open the **Export deck** panel: **PowerPoint deck** produces an editable .pptx that mirrors the slides on screen — speaker notes included in each slide's notes pane — and **Markdown outline** exports slide titles, bullets, and speaker notes as text.
- Limits: a deck holds up to 100 slides and 8 MB of content. The editor tells you plainly if a deck exceeds them.

Build a slide deck

Turn a document into slides, build a deck from an outline or a starter template, refine it, present it, and export a PowerPoint file.

Before you begin

- An available model for the deck assistant and **Edit slide with AI**. Templates, layouts, notes, presenting, and export work without one.
- Optional: your organization's PowerPoint brand template (.pptx or .potx, up to 30 MB, no macros).
- AI slide images need an image-generation model in your workspace.

Step by step

- 1 Choose **Drafts** in the sidebar, then choose **Deck** in the **Draft format** switch. A blank deck opens with the **Deck** Assistant.
- 2 Type a title for the deck.
- 3 In Ask the deck assistant, describe the deck: how many slides and what each covers. Choose Apply instruction.
- 4 Wait while Building slides... runs; on smaller models this can take a few minutes. The filmstrip fills and the rail says how many slides were drafted.
- 5 Review each slide. Use **Layouts**, **Themes**, speaker notes, and **Edit slide with AI** to refine it.
- 6 Choose **Save version** and wait for Saved.
- 7 Choose **Present deck**, then press P for presenter view. Escape exits.
- 8 Choose **Export**, set Save to to **Browser downloads**, and choose **PowerPoint deck**. The receipt names the .pptx file.

Path: Convert a document

- 1 In a document with content, choose **Deck** in the **Draft format** switch.
- 2 Turn this draft into slides? appears. Choose **Convert into slides** (or **Start a blank deck**).
- 3 The headings become slides. Your document versions are kept either way.

Path: Start from a template

- 1 In a deck, choose **Deck starters & brand themes**.
- 2 Choose a starter: **Pitch deck**, **Quarterly review**, **Project kickoff**, **Client proposal**, or **Training session**.
- 3 Choose Start and the template's name. **Replace** the scaffold text with your content.

Path: Your brand template

- 1 In **Deck starters & brand themes**, choose Upload under Brand theme and pick your .pptx or .potx file.
- 2 Choose **Apply to deck**. The colors, fonts, and logo carry through the slides.
- 3 Load all N slides brings in every slide design. The brand theme is stored only on this device.

Path: Present

- 1 Choose **Present deck**. Arrow keys, Space, or a click move forward.
- 2 Press N to show notes, P for presenter view (timer, next slide, notes), and B to black out the screen.
- 3 Press Escape, or choose **Exit**, to leave.

Path: Export formats

- 1 **PowerPoint deck**: an editable .pptx with your slides and speaker notes.
- 2 **Markdown outline**: slide titles, bullets, and notes as text.

Check that it worked

- ✓ The filmstrip lists the slides, and the rail reports how many the assistant drafted.
- ✓ **Presenter view** shows the current slide, a timer, the next slide, and your notes.
- ✓ The export receipt names the .pptx file.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
Building slides... runs for several minutes	Larger decks take longer, especially on smaller models. Ask for fewer slides, or wait; your draft stays intact meanwhile.
Generate AI slide image is greyed out: "No image-generation model is enabled for your workspace"	Your workspace has no image model. Upload your own background from Slide background instead.
"Could not open ... Choose a PowerPoint .pptx file."	Open in editor accepts .pptx and .potx files. Save other formats as .pptx first.
"... exceeds the 30 MB brand-template upload limit." or "Macro-enabled presentations are not accepted. Save as .pptx or .potx first."	Use a smaller template, and save it without macros.
"Brand template upload failed: ..."	Check that the file is a valid PowerPoint template and try again.

18. Recover work kept only on this device

Distinguish a browser working copy from a completed account save.

The sidebar's **Only on this device** reminder appears when work has changes kept in this browser that have not reached your account. It can include chats whose last save failed, documents or decks with unsent changes, and older history from before account sync.

- Open the reminder to review the named items. For chats whose last save failed, choose **Retry all** and check the resulting save state.
- For documents and decks, choose **Open Drafts**, open the item, and save it to your account. Review any load, quota, connection, or conflict error before leaving the browser.
- Older browser history remains separate until you explicitly import each item from **Drafts** → **Document history** with **Import to my account**. Confirm the item belongs in this account before importing.
- **Archive** finished drafts when you want them out of the active list. **Account** archives remain recoverable; archiving a browser-only item does not prove it has been uploaded.

GOOD TO KNOW

Clear list dismisses these reminders without deleting or uploading work. New edits may appear again. **Hide this reminder** turns off this account's sidebar notice in this browser; save status remains in **Drafts** and chats. Downloading an export also leaves pending account saves to be resolved.

Save, organize, and recover your drafts

Save documents and decks to your account, reopen them from Document history, archive and restore them, return to an earlier version, and recover work kept only on this device.

Before you begin

- At least one document or deck in **Drafts**.

Step by step

- 1 In a document or deck, choose **Save version** and wait for Saved in the toolbar.
- 2 Choose **Document history** (the clock at the top of the assistant rail).
- 3 Point at a card to preview it, or choose the card to reopen that draft.
- 4 To put finished work away, point at its card and choose **Archive. Draft** archived. Open **Archived** to restore it. confirms it.
- 5 Choose **Archived** to find it, then choose **Unarchive** to bring it back to Active.
- 6 Under Versions of this draft, choose a version to return to it.
- 7 **Delete** is separate: choose **Delete**, read the confirmation, then choose **Delete** draft or **Cancel**.

Path: When a save fails

- 1 Local only means the latest changes are kept in this browser but have not reached your account.
- 2 Choose Retry, or export a copy before you leave.
- 3 If **Draft** changed elsewhere appears, keep your local changes safe (export), then choose Reload server copy.

Path: Only on this device

- 1 The sidebar reminder counts items only on this device. Choose it to open the list.
- 2 **Retry all** retries chat saves; **Open Drafts** takes you to drafts and decks with unsent changes.
- 3 **Clear list** dismisses the current reminders; **Hide this reminder** hides the notice. Neither uploads nor deletes your work.

Path: Leaving with unsaved edits

- 1 Keep your unsaved edits? appears when you leave or open something else.
- 2 Choose Keep editing, Save copy and continue, or **Discard** and continue.

Check that it worked

- ✓ Saved appears in the toolbar with the version number.

- ✓ The draft is listed under **Archived** after archiving, and under Active after **Unarchive**.
- ✓ Versions of this draft lists each saved version with its time.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Local only — ... Your changes are kept on this device."	The server save failed. Choose Retry, or export a copy before you close the browser.
"Not saved — ..., and browser storage could not keep a copy."	Keep the workspace open, then retry or export your changes.
"this draft exceeds the 2 MB server draft limit" (or the 8 MB deck limit)	Split the document, or remove large images, then save again.
"Draft changed elsewhere"	Another tab or device saved this draft. Export your local version if you need it, then choose Reload server copy.
"Save two different versions of this draft to compare them"	Compare versions needs two saved versions that differ. Save again after an edit.

19. Agent profiles

Reusable bundles of model, prompts, knowledge, and tools.

An agent profile bundles a model route, meta prompts, knowledge bases, and MCP tools into one reusable configuration. A usable profile needs an approved, enabled model backed by a working provider. A saved profile marked not connected is a configuration example, not proof that it can run.

Creating an agent

- If your account has agent authoring permission, click **Agents** in the sidebar, then **New Agent**.
- Give it a name and a short description of what it is for.
- Work through the editor tabs: **Profile** (name, model, and description), **Knowledge** (assign knowledge bases), **Tools** (select MCP tools), **Prompts & Skills** (attach system prompts and skill files), **Access** (who can use it), and **Hermes** (the optional learning companion).
- Click the save button. Standard users create private profiles; administrators control group sharing and organization publishing. Check model readiness and access before trying the profile in chat.

Using an agent in chat

- Type @ in the composer and pick an available profile to use its configured model route, instructions, knowledge, and tools.

- The **Agent** switch in **Send options** can choose the first available profile when none is selected. Check **Agent profile** and select the intended configuration before sending.
- If no profiles appear, ask your administrator to check that a ready profile is available to your account.

From the **Agents** view you can also open a profile to edit it, jump straight into a chat with it, or delete it. The Automations tab at the top of this view holds your scheduled runs — covered in “Scheduled automations”.

WALKTHROUGH · 1:52 · [HELP](#) · [LEARN](#)

Agent profiles

Build a private agent with its own instructions, model, knowledge, and skill files, then use it in chat and check that the reply follows it.

Before you begin

- Permission to build agents: your platform owner allows users to build their own agents, and your administrator grants **Can build agents** to one of your groups. Until then, **Agents** says Building agents is not available to you.
- An available model for the agent to use.
- Optional: knowledge bases, tools, prompts, and skill files shared with you, to attach to the agent.

Step by step

- 1 Choose **Agents** in the sidebar, then **New Agent**.
- 2 On **Profile**, type an **Agent** name and choose the AI model. Optionally start from a template such as Research Analyst.
- 3 In System prompt, write the agent's instructions: its role, what it should produce, and what it must never do. Meta prompt adds standing rules on every turn.
- 4 On **Knowledge**, tick the knowledge bases it should search on every message.
- 5 On **Tools**, tick any tools it may use in agent mode; tools that need approval ask before they run.
- 6 On **Prompts & Skills**, tick saved prompts and skill files to add to its instructions.
- 7 On **Access**, read Private — only you can use this agent.
- 8 Choose **Create** agent and wait for the agent's name followed by created. It is ready to use in chat.
- 9 In a chat, type @ and choose the agent, then type your question and send it.
- 10 Check the reply: the trace reads **Agent** work trace, and the answer follows the agent's instructions and cites its knowledge.

Path: Use it from Send options

- 1 **Send options** > **Agent**, then choose the profile in **Agent profile**.
- 2 Check the selected profile before sending; turning **Agent** on can select the first one.

Path: Use it from the Agents page

- 1 On the agent's card or in its editor, choose **Chat** (or Use in **Chat**).
- 2 A chat opens with **Agent** mode on for that agent.

Path: Edit, duplicate, or delete

- 1 On the agent's card, choose Edit, change the settings, and choose Save **Profile**.
- 2 **More** actions on the card offers Duplicate and **Delete**.

Path: Share it

- 1 **Agents** you build stay private to your account.
- 2 Ask an administrator to share the agent with a group.

Check that it worked

- ✓ The agent appears under **Agents** with your chosen model.
- ✓ In chat, the **Agent** chip shows the agent, and the reply's trace reads **Agent** work trace with Preparing agent tools and the knowledge search.
- ✓ The answer follows the instructions you wrote (here, three bullet points from the playbook) and offers its citations.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Building agents is not available to you."	Your groups lack Can build agents , or the service policy does not allow user-built agents. Ask your administrator.
"Select a base AI model and enter an agent name before saving."	Type a name and choose an available model in AI model.
"... is not connected, so the agent cannot answer."	The agent's model has no working provider. Choose another model in Profile and save.
"Knowledge base ... is turned off, so the agent answers without it."	Ask the knowledge base's owner to turn it back on, or untick it.
"No agent profiles available yet — create one under Agents" when typing @	You have no agents yet, or none is shared with you. Build one, or ask your administrator.
The reply ignores part of the instructions	Smaller models follow short, specific instructions best. Shorten the System prompt, or choose a stronger model.

Library and automations

20. Knowledge bases

Give chats grounded, citable access to your documents, web pages, and APIs.

Knowledge lives in the **Library**: click **Library** in the sidebar, then the **Knowledge** tab. The **Knowledge Bases** panel lists every collection with its status, security posture, and whether it is enabled.

- Click **Add Knowledge Base** and name the collection.
- Choose what feeds it. A knowledge base can index documents you upload, a public web page (**Add a web link** — enter the address and an optional note), or an API (**Connect an API** — enter the endpoint and its authentication).
- **Create** the base, then check its row in the table: it shows the status, the security posture, and whether the base is enabled.

GOOD TO KNOW

Only enabled knowledge bases can be searched from chat, and access control is enforced per source — people only ever retrieve what they are allowed to see.

Using knowledge in a conversation

- Type **#** in the composer to reference a knowledge source or choose from Files in knowledge sources. A file choice references its name in your prompt and searches that file's knowledge source; the row says which source will be searched.
- Or turn on the **Knowledge** switch in send options to let the reply search your enabled bases.
- Replies grounded in knowledge return citations that link straight back to the source documents.

GOOD TO KNOW

The **#** shortcut menu shows Loading files... while retrieving file choices. If Some files could not be loaded appears, use Retry in that menu. Files from sources that loaded successfully remain usable. If access was removed or a source needs configuration, ask your administrator to resolve it.

WALKTHROUGH · 2:03 · HELP › LEARN

Knowledge bases

Create a knowledge base from files, web pages, or an API, watch it index, and ask a question that cites it.

Before you begin

- To create your own: your administrator grants Can build knowledge bases to one of your groups. Without it, you use the knowledge bases shared with you.
- Files up to 250 MB each: PDF, Word, Excel, PowerPoint, email (.eml, .msg), text, **Markdown**, **CSV**, JSON, images, audio, and video.
- For web pages: public addresses. For pages behind a sign-in, paste the text instead.
- **Connected** sources such as SharePoint, Box, Google Drive, and iManage are set up by administrators and shared with you.

Step by step

- 1 Choose **Library** in the sidebar; **Knowledge** opens.
- 2 Choose New knowledge base.
- 3 Type a Name, keep Start with on Upload files, and choose **Create**. Who can search it is Only me; administrators manage sharing.
- 4 Drag files onto Drag files here, or choose browse your computer.
- 5 Wait for each file to show Added and appear under Indexed files with its passage count. The base shows Ready.
- 6 In a chat, type # and choose your knowledge base, then type your question and send it.
- 7 Check the reply: it answers from your files and offers View N citations →. Open it to see each file and passage in **Sources gathered**.

Path: Add web pages

- 1 In the knowledge base, open **Web** pages.
- 2 Type the Page address (and an optional Name), then choose Fetch and add page.
- 3 The page is listed under Indexed web page; Sync now fetches it again later.
- 4 For a page behind a sign-in, tick Paste the text instead of fetching the page, paste it, and choose Add text.

Path: Connect an API

- 1 Open API. Enter the API address, an optional Path and Name.
- 2 Under Sign-in choose None (public API), API key, Bearer token, or Sign in with provider (OAuth), and fill in its fields.
- 3 Choose Fetch and add. The JSON or text response becomes searchable passages; Sync now refreshes it.

Path: Ask without

- 1 **Send options** > **Knowledge**, then choose your knowledge base.
- 2 Send the question; the reply searches only that source.

Path: Turn it off or tidy it

- 1 The On switch (Assistants can search ...) controls whether chats can search it.
- 2 **Delete** a file with the bin beside it under Indexed files.

Check that it worked

- ✓ Each file shows Added and is listed under Indexed files with a passage count.
- ✓ The knowledge base shows Ready and On.
- ✓ A question with # returns facts from your files and View N citations; **Sources gathered** names the file (for a PDF, with the page).

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"You can use the knowledge bases shared with you in chat. Ask an administrator if you need to add your own."	Creating knowledge bases is not granted to your groups. Ask your administrator for Can build knowledge bases.
"Skipped ... (save it as .docx first)" or "(unsupported type)" or "(larger than ... MB)"	Convert old .doc files to .docx, use a supported type, or split very large files.
"No readable text — only the name is searchable"	The file has no text layer (for example a picture-only PDF beyond the OCR limit). Upload a text version.
"Give the knowledge base a name."	Type a Name before choosing Create .
"Enter the page address."	Type the full address, including https://.
"Some files could not be loaded. Retry" in the # menu	Choose Retry. Sources that already loaded stay usable.

21. The Tools library

MCP connections, saved prompts, and skill files inside the Library.

Click **Library** in the sidebar, then the **Tools** tab. The **Tools Library and Connectors** panel has three sections:

TAB	WHAT LIVES THERE	WHERE IT APPEARS IN CHAT
Connections	MCP tools and connections available to models and agents.	Type / in the composer.
Prompts	Saved prompt templates, including <code>{{variables}}</code> you fill in when using them.	Type / in the composer.
Skills	Skill files — reusable instructions for recurring workflows.	Type \$ in the composer.

- Click **New Prompt** or **New Skill** to create one: give it a name, a category, a description, and the content itself, then save.
- Click any existing card to review or edit it; the trash icon deletes it permanently.
- Every tool row shows its real status — draft, approval required, or enabled — so you always know whether a connection is actually live before relying on it.

WALKTHROUGH · 1:08 · [HELP](#) · [LEARN](#)

Tools and the Library

Find the connections, saved prompts, and skill files you can use, read them, and add a private connection when your administrator allows it.

Before you begin

- **Prompts** and skill files are created by administrators; you can read and use them.
- To add your own connection: your administrator grants Can build tools to one of your groups, and you have the MCP server's address.

Step by step

- 1 Choose **Library** in the sidebar, then **Tools**.
- 2 **Connections** lists the MCP connections you can use in **Agent** mode, with On, Off, or Needs setup.
- 3 Choose **Prompts** to see the saved prompts shared with you. Choose a prompt's View to read its content.
- 4 Choose **Skills** to see the skill files shared with you.
- 5 Use them in chat: type / for prompts and connections, and \$ for skill files, or attach them to an agent.

Path: Add a private connection

- 1 With Can build tools granted, **Connections** shows Add connection.
- 2 Type a Name, choose how Aperture connects (HTTP or SSE), and enter the Server URL starting with https://.

- 3 Choose Add connection. It starts turned off; set sign-in, access, and approval on the next screen, test it, and turn it on.
- 4 **Connections** you add are private to you; an administrator can share them.

Path: Read a connection's status

- 1 On: available under policy. Off: turned off. Needs setup: missing an address, token, or sign-in.
- 2 Last test failed: the latest connection test did not succeed.
- 3 Asks before each use: you approve each run in chat (**Approve** MCP tool run?).

Check that it worked

- ✓ **Prompts** and **Skills** list the items shared with you, and View shows each one's content.
- ✓ Typing / in chat lists the same prompts, and \$ lists the same skill files.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Connections are managed by your administrators. These are the ones you can use in Agent mode."	You cannot add connections yet. Ask your administrator for Can build tools if you need one.
"Shared prompts are managed by administrators." or "Skills are managed by administrators."	Ask your administrator to add or change a prompt or skill file.
"Use a full address starting with https://"	Enter the complete server address.
"Not signed in to the provider" or "No access token saved"	Open the connection's Sign-in tab and finish signing in, or ask your administrator.
"MCP tool run denied. No message was sent."	You chose Deny when asked to approve a tool run. Send again and choose Approve if you trust it.

22. Scheduled automations

Recurring chat or drafting runs, as one model call or a multi-step chain.

Automations are saved workflows that run a chat or drafting task on a schedule. They are useful for recurring digests, weekly status checks, report refreshes, policy reviews, and any workflow where the same model chain should run the same way each time. Enabled schedules run automatically in the background (times are UTC), and each scheduled run delivers its output as a new chat thread in your sidebar.

Creating an automation

- Click **Agents** in the sidebar, then the Automations tab, then **New automation**.
- Name it — for example “Monday client digest”.
- Choose what it runs against: chat or draft.
- Pick a trigger: **Daily** (a time each day), **Weekly** (a day and time), **Once** (a specific date and time), or **Cron expression** (such as `0 9 * * 1` for 9:00 every Monday, in UTC).
- Write the initial input — what the first step should work on.
- Build the model chain: each step has a model and an instruction, and each step's output feeds into the next step. Use **Add step** for multi-step chains and the x to remove a step.
- Click **Save automation**.

What to check before saving

- **Models** — every step must use a model you can actually access. If a model is missing, your admin may need to approve it for your group.
- **Context** — write the initial input as if the automation will run without you watching. Include the matter, audience, date range, source expectations, and desired output format.
- **Cadence** — use **Once** for a single future run, **Daily** or **Weekly** for regular recurring work, or a reviewed **Cron expression** for another schedule. All times are UTC; check the intended local-time equivalent before saving.
- **Ownership** — name the automation so another person can tell what it does later.

Running and managing

- Each automation card shows its schedule, its steps, and its last run with an honest result — succeeded or failed, with the reason.
- **Run now** executes the whole chain immediately and shows a transcript of every step's output.
- The toggle pauses or resumes the schedule; the pencil edits it; the trash deletes it.
- You can also queue an automation into any chat by typing > in the composer — it runs when you send.
- The list shows the automations you created. Organization administrators can review the automations in their workspace.

TIP

Always use **Run now** once after creating or editing an automation. If the run fails, the card shows the reason so you can fix the model choice, input, or instructions before depending on the saved schedule.

Scheduled automations

Build an automation that runs a prompt through one or more models on a schedule, deliver it to a new chat or a new draft, run it now, and read its history.

Before you begin

- At least one model available to you. Without one, Automations says no approved models are available and **New automation** is off.
- The prompt you want to run, and when it should run (the time zone matters).

Step by step

- 1 Choose **Agents** in the sidebar, then Automations.
- 2 Choose **New automation**.
- 3 Under What it does, type a Name and the Prompt for the first step.
- 4 Under Steps, choose a model (or one of your agents) for step 1 and an optional instruction. Choose **Add step** to pass the result to another model.
- 5 Under Schedule, choose **Daily**, **Weekly**, **Once**, or Custom, set the time and Time zone, and check Next runs.
- 6 Under Deliver results to, choose **New chat** or New draft.
- 7 Leave On — runs on this schedule turned on, and choose **Save automation**.
- 8 On the automation's card, choose **Run now** to test it. Wait for the result and choose Open chat (or Open draft).

Path: Daily

- 1 How often › **Daily**, then set Time.
- 2 Next runs lists the coming days at that time.

Path: Weekly

- 1 How often › **Weekly**, choose the day (Mon to Sun), then set Time.
- 2 Next runs lists the coming weeks.

Path: Once

- 1 How often › **Once**, then set Run at to a date and time in the future.
- 2 Next runs shows that single run; if the time has passed it says No upcoming runs — this time has passed.

Path: Custom (cron)

- 1 How often › Custom, then type a **Cron expression** with five fields: minute, hour, day of month, month, day of week.
- 2 For example 30 7 * * 1–5 runs at 7:30 AM on weekdays. Check Next runs before saving.

Path: Deliver to a new draft

- 1 Under Deliver results to, choose New draft.
- 2 The last step is asked for a complete document, saved as a new draft.
- 3 After **Run now**, choose Open draft to edit and export it in **Drafts**.

Path: Pause, resume, edit

- 1 The switch on the card pauses the automation (Paused) or resumes it (Active).
- 2 Edit changes its steps, schedule, or delivery; choose Save changes.
- 3 Run history (N) lists recent runs: Succeeded, Failed, or Skipped, how each was started, and how long it took.

Check that it worked

- ✓ After saving, the card shows the schedule, the step chain, and Next run in
- ✓ **Run now** reports that the run saved the result to a new chat (or draft), and Output from this run shows each step.
- ✓ Open chat shows the delivered conversation, titled with the automation's name and the run time.
- ✓ Run history lists the run as Succeeded, run now.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"No approved models are available yet. An administrator must enable at least one available model before automations can run."	Ask your administrator for access to a model; New automation stays off until then.
"Give the automation a name." or "Every step needs a model or agent."	Type a Name and choose a model or agent for every step.
"Choose a date and time for the one-time run."	Set Run at for a Once schedule.
"Enter a cron expression, for example 0 9 * * 1 -5."	Type five fields separated by spaces, and check Next runs.
"Run failed: ..." or Last run failed on the card	Read the message in Run history. A model or agent may be unavailable; edit the step and run it again.
Paused after failures	After 3 scheduled failures in a row the automation pauses itself. Fix the cause, then turn the switch back on.
An agent step does not use its tools	Tools and MCP servers need a person to approve them, so automations run agent steps without tools.

Appearance, your account, and help

23. Light and dark mode

Switch appearance now or schedule it for this browser.


- 1 Find the moon or sun button beside **Help** near the bottom of the sidebar — a moon (**Dark mode**) while in light mode, a sun (**Light mode**) while in dark mode.
- 2 Click it. The entire platform switches immediately — no reload, nothing to save.
- 3 Click it again to switch back.

Schedule light and dark mode

Choose the clock button beside it to open **Theme schedule**. Enable **Switch automatically**, set different **Light mode at** and **Dark mode at** times, then choose **Save schedule**. The schedule uses this device's local time every day and is saved in this browser. You can still switch manually until the next scheduled change; disable the switch to stop automatic changes.

24. Your account

Your profile, personalization memory, password, role previews, usage, archives, and signing out.



Your account includes profile settings, password controls, and two-step verification. Example workspace with synthetic data.

Click your account card (your initials or photo) at the very bottom of the sidebar. The account drawer opens with everything about you in one place.

- **Profile** — click the card with your name and the pencil icon to edit your display name, firm, website, bio, phone number, and photo (upload an image up to 5 MB, or paste a URL). Click **Save profile** when done, or **Cancel** to discard.
- **Management** — administrators and platform owners see this section, listing the consoles they can open; the sidebar's **Admin console** and Platform console links go to the same places. Regular users do not have it.
- **Personalization memory** — opens your private memory manager when the feature is enabled. Use it to control, review, add, correct, pin, forget, or clear what the assistant remembers about you.
- **Password** — accounts that sign in with a password can change it here (click the pencil, enter the current password, then the new one twice — at least 12 characters). The app continues with a new authenticated session after saving and revokes previous sessions. Accounts that sign in through **SSO** manage their password with the **SSO** provider instead, and the panel says so.
- **Security** — choose **Manage security** to set up an authenticator for a local account, review the remaining recovery-code count, or replace codes after verifying your identity. Turning verification off is available only when organization policy permits it and signs you out. **SSO** accounts follow their identity-provider settings for voluntary enrollment.

- Role and organization — your role, your organization, any personal token caps that apply to you, and how you sign in.
- **View as** — if your role allows it, preview the workspace exactly as a lower role would see it. A note reminds you which role you are previewing; switch back the same way.
- **Usage this month** — your prompts, responses, and estimated tokens.
- **Archived chats** — every chat you have archived, with buttons to restore each one to your sidebar or delete it permanently.
- **Sign out** — ends your session and returns to the sign-in screen.

WALKTHROUGH · 2:21 · [HELP](#) · [LEARN](#)

Personalize, use mobile, and get help

Edit your profile, choose light or dark mode and a theme schedule, install the app on a phone, tablet, or computer, and report a problem.

Before you begin

- For installing: Safari on iPhone or iPad, Chrome on Android, or Chrome or Edge on a computer.
- For a report: what you did, what you expected, and what happened; optionally a screenshot (**PNG**, JPEG, GIF, or WebP, up to 10 MB) with no passwords or personal data.

Step by step

- 1 Select your account card at the bottom of the sidebar, then the pencil beside your name (Edit account profile).
- 2 Change the fields you need, such as Phone number or Firm or organization, and choose **Save profile**. **Profile** saved confirms it.
- 3 Choose the moon button beside **Help** for **Dark mode**, or the sun for **Light mode**.
- 4 Choose the clock button (**Theme schedule**), turn on **Switch automatically**, set **Light mode at** and **Dark mode at** to different times, and choose **Save schedule**.
- 5 Choose **Help** at the bottom of the sidebar for the walkthroughs and the printable User guide (PDF).
- 6 To report a problem, choose Report a problem in **Help**, fill in Subject and Message, optionally add a Screenshot, and choose Send report.
- 7 Wait for Report sent: an administrator can now review it in **Analytics**.

Path: iPhone or iPad (Safari)

- 1 Open your workspace address in Safari and sign in. Add Aperture Chat to your home screen opens (or choose **Install app** in the menu).
- 2 Tap the **Share** button in the browser toolbar.

- 3 Scroll the share sheet and choose Add to Home Screen.
- 4 Tap Add. Aperture Chat appears on your home screen and opens full screen.

Path: Android (Chrome)

- 1 Open your workspace address in Chrome and sign in. The install card opens (or choose **Install app** in the menu).
- 2 If it offers **Install app**, choose it and confirm in the phone's install sheet.
- 3 Otherwise open the browser's : menu and choose Add to Home screen or **Install app**, then confirm.

Path: Computer (Chrome or Edge)

- 1 Aperture Chat shows no install card on computers.
- 2 Use the browser's own install control: the install icon at the right of the address bar, or the browser menu > Install (Chrome: Cast, save, and share > Install page as app; Edge: Apps > Install this site as an app).
- 3 Confirm Install. The app opens in its own window and appears with your other apps.

Path: On a narrow screen

- 1 Choose Open menu (top left) to reach navigation, chats, **Help**, **Install app**, and your account.
- 2 Choose where to go; the menu closes. The same permissions apply as on a computer.

Check that it worked

- ✓ **Profile** saved appears, and your account card shows the new details.
- ✓ The workspace switches between light and dark at the scheduled times on this device.
- ✓ The installed app opens from your home screen or app list.
- ✓ Report sent confirms an administrator can review your report.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Use a complete http(s) website URL."	Enter the full website address, including https://.
"Choose a profile photo that is 5 MB or smaller."	Use a smaller image.
"Choose two different times."	Light mode at and Dark mode at must differ.
"Add both a subject and a message."	Fill in both fields before Send report.
"Choose a screenshot that is 10 MB or smaller." or "Screenshot attachments must be PNG, JPEG, GIF, or WebP images."	Attach a smaller PNG , JPEG, GIF, or WebP image.
"The issue report could not be sent."	Your report was not received. Check your connection and send it again.

25. Personalization memory

Save and recall private preferences in plain English across chats and sessions.

Personalization memory belongs to your account, not to one conversation. When service policy and your organization administrator allow it, administrators and regular users can each build their own private memory. A saved item can influence a new chat immediately and still be available after you sign out and return in a later session.

Save a memory in ordinary English

- In any chat, state what should persist in natural language — for example: “Remember that I prefer a short summary before the detail,” “Please remember that I work in commercial litigation,” or “Remember my project codename is Silver Horizon.”
- Send the message. The assistant confirms the memory after it is saved; you do not need a slash command or special syntax.
- Start a new chat, or return in a later sign-in session, and ask a natural question such as “What do you remember about me?”, “What are my writing preferences?”, or “What is my project codename?” The same account memory is available across those sessions.

GOOD TO KNOW

Memory is different from the context window shown in Session details. The context window describes how much of the current chat the selected model can hold at once; personalization memory is durable account context that follows you into other chats and sessions.

Open and control the memory manager

- Click your account card — your name and role at the bottom-left of the sidebar.
- Click Personalization memory. The dialog titled “What the assistant remembers about you” opens.
- Use memory in my chats controls whether saved items are applied to answers. Learn from my conversations controls whether the assistant may notice durable preferences automatically. Turning a switch off does not expose or silently erase existing items.
- To add something directly, type it under Add something you want remembered, choose Standing instructions, Preferences, About you, Your work, or Other details, and click Add.
- Review the list below. Click a memory's wording to correct it; pin an important item; use Forget to remove one item; or choose Forget everything to clear the list.

GOOD TO KNOW

Only you can read your memory content. Administrators can see counts and purge memories for compliance, but the platform never shows them what an individual memory says.

TIP

If Personalization memory is missing from your account drawer, the feature is not currently available at the platform, organization, or group level. Ask your administrator; refreshing the page cannot override that policy.

CAREFUL

Do not use memory as a password or secret vault. Credentials, keys, and sensitive identifiers are rejected on purpose.

WALKTHROUGH · 1:28 · [HELP](#) · [LEARN](#)

Personalization memory

Save preferences in plain English, recall them in later sessions, and stay in control.

Before you begin

- Memory must be allowed by your platform and your organization. If it is off, the Personalization memory row says why.
- Only you can read your memories. Administrators see how many you have and can delete them, but never read them.

Step by step

- 1 Select your account card at the bottom of the sidebar and choose Personalization memory.
- 2 Check Use memory in my chats (applies saved memories to answers) and Learn from my conversations (lets the assistant notice lasting preferences itself).

- 3 Type a memory in Add something you want remembered, choose its type (for example Preferences), and choose Add.
- 4 Choose Pin on a memory that should always apply. Choose its text to correct the wording.
- 5 In any chat, say Remember that ... followed by the preference; it is saved for later chats.
- 6 In a new chat, ask a question the memory answers, and check that the reply uses it.

Path: Forget

- 1 Personalization memory › Forget on one memory removes it.
- 2 Forget everything asks Forget all ...? This cannot be undone. before removing them all.

Path: Turn memory off

- 1 Turn off Use memory in my chats to stop applying memories. Saved memories stay private and manageable.
- 2 Turn off Learn from my conversations to save only what you add or ask it to remember.

Check that it worked

- ✓ The memory is listed in Personalization memory under its type, with Pin on if you pinned it.
- ✓ After Remember that ... in chat, the sentence appears in the list.
- ✓ A later chat's reply uses it without being told again.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Memories cannot contain credentials, keys, or sensitive identifiers."	Passwords, keys, tokens, one-time codes, and card or account numbers are never stored. Keep them out of memory.
"Write a little more so the memory is useful later."	Write a full phrase of at least a few words.
"Pinned memories already fill your memory limit. Unpin or forget one before saving more."	Unpin or forget a memory, then add the new one.
"That memory kind is disabled by your organization's policy."	Choose another type, or ask your administrator.
The row says memory is turned off	"Memory is turned off for this platform.", "... for your organization.", or your groups' policy: ask your administrator. "You turned memory off for your account.": turn it back on here.
No note under the reply says a memory was saved or used	The reply does not currently show that note. Check Personalization memory to confirm the memory was saved.

26. Install the app on your phone

Put your workspace on the home screen with its own icon and name.

On a phone or tablet, tap **Install app** near the bottom of the sidebar. The "Add ... to your home screen" dialog previews your organization's app name and icon.

- On Android, choose **Install app** if offered, then follow the native install sheet. Otherwise open your browser's menu and choose "Add to Home screen" or "**Install app**".
- On iPhone and iPad, tap the browser's **Share** button, scroll the share sheet, choose "Add to Home Screen", then tap Add.
- Open the new home-screen icon to return to the same workspace in a full-screen app.

27. Getting help


Walkthrough videos and this guide, always one click away.

- Click **Help** at the bottom of the sidebar for guided videos, captions, transcripts, and this PDF. The guide also covers additional procedures.
- Administrators can also open **Documentation** inside the console for role-specific videos and a printable guide.
- Choose Report a problem in **Help** to describe a problem, the affected screen, and the expected result. **Share** only permitted details with your administrators; omit passwords, keys, and recovery codes.

The Admin console

28. Opening the Admin console

Where the console lives and what its ten tenant-governance tabs control.



The Admin console opens the workspace controls available to your administrator account. Example workspace with synthetic data.

- 1 Click **Admin console** near the bottom of the sidebar. (Only workspace administrators and platform owners see it; the account drawer's **Management** section lists it too.)
- 2 The console opens with ten tabs across the top: **Users**, **Groups**, **Model Access**, **Connections**, **SSO**, **Analytics**, **Policies**, **Datasets**, **Audit**, and **Alerts**. **Policies** is always present between **Analytics** and **Datasets**; service-wide availability determines which organization controls are active inside it.

Everything you change here writes through the admin API immediately — changes persist across refreshes and restarts, and every action lands in the tenant audit trail. Status messages under the header tell you honestly whether an action synced or failed.

TIP

The **Documentation** button at the top of the console opens narrated video walkthroughs of every tab, plus this guide as a PDF.

29. Users

Create tenant accounts, manage passwords, and remove leavers safely.

The **Users** tab lists the organization accounts you administer, with each person's role, groups, sign-in method, and status.

Reading and filtering the list

- Use the group filter to work one team at a time.
- Watch the Auth and Status columns: Auth shows sso or local, and Status shows Active, Inactive, or **Pending** for an approved person who has not signed in yet. When the default group is on, new **SSO** accounts start in **Default Users**; otherwise they have no groups until you add them.

Per-row actions

Every row has an Actions column with account controls:

- **Access** — opens the read-only **Model access trace** for this person, with their groups, each model's status, and the gates that explain it. See [Resolve model requests and trace access](#) for the full workflow.
- **Password** — opens the password dialog for that person: type a password or click **Generate** for a strong random one, optionally mark it a **Temporary password** (they must choose their own at first sign-in), and click **Set password**. The password is shown only here — share it over a safe channel.
- **Deactivate / Activate** — ends or restores sign-in access immediately, keeping the account's audit history intact. You can also select several accounts with their checkboxes and use the **Deactivate** button at the top.
- **Delete** (trash icon) — permanently deletes an eligible account and its chat history. The tooltip spells it out per person: “Permanently delete ... and their chat history”. Administrator-account actions that are unavailable under current service policy remain disabled with an explanation.

GOOD TO KNOW

Password reset and authenticator recovery are separate. The current **Users** screen has no authenticator-reset action. If someone has lost their authenticator and recovery codes, verify their identity and follow the organization's authorized recovery process; setting a temporary password alone does not clear MFA. Identity-provider recovery remains with the **SSO** provider.

GOOD TO KNOW

Administrative continuity rules are enforced by the service. When an account action would violate them, the console blocks the action and explains that it is restricted by administrative continuity policy.

Approve access and finish sign-in

Take a real access request from the sign-in page to a working first chat, by temporary password or by organization SSO, and decline a request you do not recognize.

Before you begin

- An Admin account. **Admin console** › **Users** shows the **Access requests** panel only while someone is waiting.
- A model granted to **Default Users** in **Model Access**, so the approved person has something to chat with.
- For the **SSO** path: single sign-on already set up on the **SSO** tab, and the person's account at your identity provider with a verified email on an allowed domain.
- A safe way to share a temporary password, such as in person or your password manager's sharing. Aperture Chat sends no email.

Step by step

- 1 The person opens your workspace address, chooses **Request access**, enters First name, Last name, and Work email, and chooses Submit access request. They see Request received.
- 2 Choose **Admin console** at the bottom of the sidebar. **Users** opens with an **Access requests** panel that lists each person's name, email, and request time.
- 3 Check that you recognize the person and the email address.
- 4 In **Approve** as, keep User, then choose **Approve**. The status reads "... was approved as User."
- 5 In Finish sign-in setup, choose Set temporary password.
- 6 Choose **Generate** (or type at least 12 characters), keep **Temporary password** on, and choose **Set password**.
- 7 Choose **Copy**, share the password over your safe channel, then choose **Done** in the dialog and **Done** in Finish sign-in setup. The row now shows local under Auth.
- 8 The person chooses Email & password on the sign-in page, signs in with the temporary password, enters their own password twice on Set a new password, and chooses **Set password** and continue.
- 9 They land on the welcome card. Ask them to start a new chat, choose a model, and send a short message.
- 10 Confirm that a reply arrives. That proves the approval, the group, model access, and sign-in all work.

Path: Organization SSO, no password (shown)

- 1 Make sure single sign-on is set up on the **SSO** tab and the person has an account at your identity provider with a verified email on an allowed domain.
- 2 **Approve** the request as User. In Finish sign-in setup, choose **Done**. Do not set a password.
- 3 Tell the person to open the workspace, enter their work email, and choose **Continue with SSO**.
- 4 They sign in at the identity provider and land on the welcome card.
- 5 **Users** shows the account Active, with sso under Auth.

Path: Decline a request (shown)

- 1 In **Access requests**, choose **Decline** on the request you do not recognize. There is no confirmation step.
- 2 The status reads "... access request was declined." The pending account is deleted.
- 3 If that person tries to sign in, they see "Unknown local account." They can submit a new request later.

Path: Approve as Temp User or Admin

- 1 Temp User gives access to the designated Luna model only and stops after 30,000 reported tokens. It needs an enabled Luna model.
- 2 Admin appears in **Approve** as only when service policy allows it. Otherwise Policies > **Policy Controls** notes "Administrator accounts are created by your service team."
- 3 Arrange sign-in the same way: a temporary password or organization **SSO**.

Check that it worked

- ✓ The status reads "... was approved as User." and the request leaves the queue.
- ✓ **Users** shows the person Active, with local under Auth after a password is set, or sso for the **SSO** path.
- ✓ The person reaches the welcome card and gets a real reply to a first message.
- ✓ A declined person's sign-in attempt shows "Unknown local account."

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Access was not approved. This access request is no longer pending."	Another administrator already approved or declined it. Reload Users to see the current queue.
"Access was not approved. Enable a Luna model for this workspace before approving temporary access."	Approve as User instead, or ask your service team to enable a Luna model.
The approved person sees "Unknown local account."	Approved accounts start as SSO accounts. Set a temporary password (Finish sign-in setup, or Password in the row's Actions), or have them use Continue with SSO .
"Invalid local credentials."	The temporary password was mistyped. Set a new one; each new password replaces the last.
Set password stays unavailable in the password dialog	Passwords need at least 12 characters. Choose Generate .
The person's model menu reads No models available	None of their groups carries a model. Grant one to Default Users in Model Access , or choose Access on their Users row to see why.
"Too many access requests for this email. Try again shortly."	The sign-in page accepts three requests per email in a short window. Wait a minute, then submit once.

WALKTHROUGH · 2:36 · ADMIN CONSOLE > DOCUMENTATION

Users and accounts

Create SSO and password accounts, change roles, reset passwords, deactivate and reactivate, work in bulk, filter by group, explain model access, and delete leavers.

Before you begin

- An Admin account. You manage User, Temp User, Power User, Auditor, and **Agent Approver** accounts; admin rows are locked.
- A group to start people in. **Default Users** always exists.
- For password accounts: a safe way to share a temporary password.

Step by step

- 1 Choose **Admin console** at the bottom of the sidebar. **Users** is the first tab.
- 2 Choose **Add User**. Enter Name and Email, then choose a Role and a Starting group.

- 3 Choose **Create Account**. The status reads "... was created through the admin API." The new row shows sso under Auth.
- 4 For someone who signs in with a password, choose **Password** in the row's Actions, then **Generate**, keep **Temporary password** on, and choose **Set password**. Auth changes to local.
- 5 To change a role, choose it in the row's Role menu. It saves immediately.
- 6 To pause access, choose **Deactivate** in the row's Actions. Choose **Activate** to restore it.
- 7 To delete a leaver, choose the trash button in the row's Actions. Deletion is immediate and permanent.

Path: An SSO account (shown)

- 1 **Add User** creates an account that signs in through organization **SSO**: Auth reads sso.
- 2 **Share** the workspace address; the person enters their email and chooses **Continue with SSO**.
- 3 Without **SSO** set up, the account cannot sign in until you set a temporary password.

Path: A password account (shown)

- 1 **Add User** as above, then choose **Password** in the row's Actions.
- 2 Choose **Generate**, keep **Temporary password** on, choose **Set password**, then **Copy**.
- 3 **Share** the password safely. The person chooses their own at first sign-in.

Path: Several accounts at once (shown)

- 1 Tick the checkbox at the start of each row.
- 2 Choose **Deactivate** at the top of the list. The status reads "Deactivated N users through the admin API."
- 3 Bulk deactivation also deletes each person's saved personalization memories. A row's own **Deactivate** button keeps them.
- 4 Reactivate people one at a time with **Activate**.

Path: Explain one person's model access (shown)

- 1 Choose **Access** in the person's row.
- 2 Each model shows **Usable**, **Allowed**, **provider offline**, or **Blocked**, with every gate the policy checked.
- 3 The trace explains; it changes nothing. Fix access in **Groups** or **Model Access**.

Check that it worked

- ✓ New rows show the role, group, and Auth you chose.

- ✓ A deactivated row reads Inactive; **Activate** returns it to Active.
- ✓ After deletion the row is gone and the status reads "... was permanently deleted."

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"User was not created. User email already exists."	Someone already has that address. Set Filter users by group to All groups and find them.
"Pick a starting group first — create one on the Groups tab if none exist yet."	Create a group on the Groups tab, then add the user again.
The Password button is unavailable	Passwords can be set only for active accounts other than yours and other administrators'. Reactivate the account first.
"Reactivate the account before setting its password."	Choose Activate in the row's Actions, then set the password.
"... was not deleted. Tenant admins can only delete regular users in their own tenant."	Administrator accounts are managed by your service team.
"This action is blocked by administrative continuity policy."	The change would leave the organization without an administrator. Ask your service team.
A password reset did not reset two-step verification	Correct: Password does not reset the authenticator, and this screen has no authenticator reset. Verify the person's identity through your organization's recovery process and contact your service team.

30. Groups

Groups carry permissions and model access — everything flows through them.

Groups are how access flows in the tenant: model grants, knowledge access controls, and permissions all attach to groups, not to individual people. Someone in no group has no group-based access, and a group that **SSO** group mapping manages follows the identity provider at every sign-in.

Runtime access and authoring permissions

TOGGLE	WHAT IT GRANTS
Can use chat	Start conversations and use assigned models.
Can use knowledge	Query approved knowledge bases.
Can use agents	Run approved agent workspaces.
Can use tools	Invoke enabled tools and MCP actions.
Can use API	Create personal keys for approved models when service policy allows API access.
Can use Hermes companion	Build and run agent profiles with the Hermes learning companion. Off until approved.
Can build agents	Create private agent profiles when service policy permits. Organization publishing stays admin-only.
Can build knowledge bases	Create private knowledge bases. Group sharing and organization management stay admin-only.
Can build tools	Create private tools. Group sharing, stdio commands, and organization management stay admin-only.
Can use memory	Allow personal preferences to be learned and reused within the workspace memory policy.

CAREFUL

Deleting a group removes its members' access that flowed through it. Check what the group grants before deleting.

WALKTHROUGH · 2:05 · ADMIN CONSOLE > DOCUMENTATION

Groups and permissions

Create a group, add people one at a time or by pasting emails, set each permission, and know how the protected Default Users group behaves.

Before you begin

- An Admin account.
- The people's accounts already exist. **Import** adds existing accounts only.
- **Can use API** stays locked while service policy keeps downstream API access off.

Step by step

- 1 Choose **Admin console** › **Groups**, then **Add Group**.
- 2 Enter a Group name and choose **Create** group. Leave Optional **SSO** group ID empty unless you use it.
- 3 Choose Manage on the new group's card.
- 4 On the group's **Users** tab, switch on each person you want in the group.
- 5 On the **Import** tab, paste email addresses (one per line, or separated by commas) and choose **Add users to group**.
- 6 On the **Permissions** tab, switch each permission on or off. Each switch saves immediately.
- 7 Open **Users** and choose the group in Filter users by group to confirm its members.

Path: Add people one at a time (shown)

- 1 Open the group with Manage, then its **Users** tab.
- 2 Switch on Add ... to the group for each person. The pill changes to In group.

Path: Import by email (shown)

- 1 Open the group's **Import** tab and paste the addresses.
- 2 Choose **Add users to group**. Existing accounts are added; unknown addresses and administrator accounts stay in the box.

Path: Let SSO manage membership

- 1 On the **SSO** tab, map an identity-provider group value to this group (see Tenant **SSO** and provisioning).
- 2 Membership then follows each person's next **SSO** sign-in.

Path: Remove a group

- 1 Choose the trash button on its card, or tick several cards and choose Remove selected. There is no confirmation step.
- 2 Members lose the access that came through that group. **Default Users** cannot be removed.

Check that it worked

- ✓ The group's card shows the member count.
- ✓ **Users** filtered to the group lists exactly its members.
- ✓ Each permission you changed shows its new switch position, and the status reads "... group permissions synced with the admin API."

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Group was not created. Group name already exists."	Choose a different name, or manage the existing group.
"Group name is required before creating a group."	Enter a Group name before choosing Create group.
"No users were added to No eligible platform users were found for"	Those addresses have no account, or belong to administrators. Create the accounts on Users first; the addresses stay in the box so you can retry.
Can use API reads "Downstream API access is unavailable under the current service policy."	Service policy controls downstream API access. Ask your service team.
"Default Users is protected by organization policy and cannot be removed."	Default Users is the baseline group. Adjust its permissions and models instead.
Can build agents is on, but people still cannot build agents	It also needs service policy to allow user-built agents. Policies > Default users can build agents shows whether it does.

31. Model Access

Decide which synced models users can see, and which groups can use them.

Model Access starts from the catalog available to your organization. You decide what portion of it your users actually see.

The catalog at a glance

- Click **Sync models** to refresh the catalog whenever service availability changes.
- Read the counters: the All, Enabled, and Disabled status counters carry live totals alongside how many groups are available for scoping.
- Use the search box and the funnel filters in the column headers to cut through the catalog: check off providers, check off model labs, or type text to match runtime routes, and Clear filter resets one. The table has seven columns: **Model**, **Provider**, **User Access**, **Groups**, **Filters**, **Knowledge**, and **Tools**. Everything starts hidden until you decide otherwise.
- Flip the **User Access** toggle on a row to make that model visible, and use **Choose groups** to narrow it to specific teams.

- Use the **Filters** column to attach per-model content filters, and the **Knowledge** and **Tools** columns to scope what each model may search and call.

GOOD TO KNOW

Because model access flows through groups, a tenant with no groups yet sees the guard up front — “**Create a group** before enabling models.” — and the access switches stay off until one exists.

GOOD TO KNOW

Newly synced models arrive disabled until service availability and organization policy permit them. If a model you expect is missing, it is not currently available to this organization.

WALKTHROUGH · 1:51 · ADMIN CONSOLE › DOCUMENTATION

Tenant model access

Grant a model to the right groups, check the effect for one person in the access trace, and hide a model when you need to.

Before you begin

- An Admin account and at least one group.
- **Models** your service team has made available to your organization.
- **Access** is granted to groups only. There is no per-person grant or deny.

Step by step

- 1 Choose **Admin console** › **Model Access**, then **Sync models**.
- 2 Find the model with **Search** model names or the funnel filters in the column headings.
- 3 In the model's **Groups** column, choose the groups button (**Choose groups**, or N groups).
- 4 Tick each group that should use the model (Allow ... for ...). Each tick saves immediately.
- 5 On **Users**, choose **Access** for someone in that group. The model reads **Usable**, with group_grant passed.

Path: Grant or remove a group (shown)

- 1 Open the model's groups and tick or untick a group.
- 2 Members of a ticked group can use the model; the trace for someone outside it stops at group_grant.

Path: Hide a model from everyone (shown)

- 1 Turn off the model's **User Access** switch. The row reads Hidden from users.

- 2 Today only your service team can turn it back on: the switch is refused with "This model is not available to this organization." The same happens if you untick its last group.
- 3 To narrow access without hiding it, keep at least one group ticked.

Path: Give one person a model

- 1 **Create** a group for them (or use an existing one) and tick it on the model.
- 2 Or approve their request on **Model Access** (see Review model requests and explain access).

Check that it worked

- ✓ The **User Access** column reads Visible to N groups.
- ✓ The access trace for a member shows the model **Usable**, or **Allowed, provider offline** if only the provider is down.
- ✓ The person sees the model in their model menu.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"... user access could not sync. This model is not available to this organization. Nothing was changed."	The model has no groups from your organization left, so it can no longer be changed here. Ask your service team to restore its groups. Keep one group ticked when narrowing access.
"Create a group before enabling models."	Create a group on the Groups tab first. Users get models only through groups.
The trace reads Allowed, provider offline	The grant works, but the provider behind the model is not connected. Your service team reconnects it.
"Model is disabled by platform policy."	The model is turned off for the whole service. Ask your service team.
A model is missing from the catalog	Choose Sync models . If it is still missing, it has not been made available to your organization. Ask your service team.

32. Resolve model requests and trace access

Review requests through groups and diagnose each user's actual model access.

CAREFUL

Approval changes group membership. When the chosen group does not carry the model yet, it can also widen that model's availability for all group members. Review the group's purpose and members before approving.

GOOD TO KNOW

The trace is read-only and scoped to users you can administer. It explains existing permissions; opening it does not grant access, change a role, or bypass the organization's model ceiling.

WALKTHROUGH · 1:37 · ADMIN CONSOLE › DOCUMENTATION

Review model requests and explain access

Approve a person's model request into the right group, decline another, and confirm both from the person's view and the access trace.

Before you begin

- An Admin account and the groups you grant through.
- **Model** requests allowed by service policy. When they are not, people see only the models they can already use and cannot send requests.

Step by step

- 1 The person opens the model menu, chooses **Why isn't a model listed?**, and chooses **Request access** on a **Locked** model. They see "Request sent for An administrator will review it."
- 2 Choose **Admin console › Model Access. Access requests** lists the person, the model, and the server's reason.
- 3 In **Grant through group**, choose the narrowest group that should carry the model. A group marked (also grant model to group) does not carry it yet; approving grants it to every member.
- 4 Choose **Approve**. The panel reads "... can now use ... through"
- 5 Choose **Decline** for requests you will not grant. The panel reads "Declined the request from"
- 6 Ask the person to reopen **Why isn't a model listed?**, and check **Users › Access** for them.

Path: Approve through a group that already has the model

- 1 Choose a group listed without a suffix. Approving adds the person to it; nobody else gains access.

Path: Approve and widen a group (shown)

- 1 Choose a group marked (also grant model to group).
- 2 Read the hint: approving grants the model to everyone in the group. Check its members first.

Path: Decline (shown)

- 1 Choose **Decline**. The request leaves the queue, and the person's model reads **Locked** again with **Request access** available.

Check that it worked

- ✓ The panel reads "... can now use ... through" and the request leaves the queue.
- ✓ The person's **Models in your organization** dialog no longer shows the model as **Locked**.
- ✓ **Users > Access** for the person shows group_grant passed for that model.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Choose the group that should carry this model before approving."	Choose a group in Grant through group . If one is already shown, choose another group and then the one you want.
"Approved, but the server still reports: ..."	The approval worked but another gate still blocks the model, such as an offline provider. The trace shows which.
"This request has already been resolved."	Another administrator approved or declined it. Choose Refresh .
Approve is refused because the model is disabled for the whole service	Ask your service team to enable the model, then approve again.
The person sees " Model access requests are disabled by organization policy. "	Service policy does not allow model requests. Grant access through a group instead.
The person sees " A request for this model is already pending. "	Their earlier request is still in your queue.

33. Connections: response actions and connector handoff

Manage response actions and route shared connector setup to the service team.

The **Admin console's Connections** tab contains **Chat output actions**. Shared connector switches, saved credentials, connection tests, and workspace OAuth are managed by the service team outside tenant administration.

GOOD TO KNOW

Users still connect their own Google, Microsoft, Box, or iManage account from the composer's attach menu. Those delegated connections respect each user's source permissions. Shared credentials managed by the service team support connector setup and background knowledge sync; they do not grant users someone else's files.

Chat output actions

Chat output actions adds admin-approved buttons to assistant responses for export, formatting, or handoff. Choose New response action to create one; existing custom actions offer Edit and **Delete**. Each

row shows Enabled or **Draft** and an enable switch. Creating these actions does not configure a shared source connector.

MCP connections and model-callable tools remain in **Library** → **Tools** → **Connections**. **Prompts** and **Skills** also remain in the **Tools** library. Their authoring and use follow the existing permissions and the shared connector availability set by the service team.

WALKTHROUGH · 2:30 · ADMIN CONSOLE › DOCUMENTATION

Response actions and connector responsibilities

Build a script response action, test it, give it to a group, use it on a real reply, edit or draft it, and find where connections are signed in.

Before you begin

- An Admin account.
- A short Python script: it reads the reply on standard input and prints the result.
- The group whose members should see the button.

Step by step

- 1 Choose **Admin console** › **Connections**. **Chat output actions** lists every tool; response actions carry a Response action pill.
- 2 Choose New response action. Enter Action name, Timeout (seconds, 1–30), Description, and the Python script.
- 3 Under Who can run this action, tick the groups that should see it. With none ticked, everyone can.
- 4 Under Test run, paste sample text and choose Run test. Read "Finished in ... ms." and the output.
- 5 Choose **Create** action. The status reads "... saved and available to models and chat." and the row reads Enabled.
- 6 A member opens any assistant reply and chooses the action's button below it. The result opens with **Copy** and **Download**.
- 7 Choose Edit to change it, or switch it off to set it to **Draft** and hide it from replies.

Path: Script response action (shown)

- 1 The script receives the reply text on standard input and prints its result.
- 2 To return files, write them to the folder in APERTURE_ARTIFACT_DIR: up to 8 files, 50 MB each and 75 MB per run. Links appear in the output.
- 3 Scripts run isolated, without platform secrets or provider keys, but can reach the network.

Path: Sign a connection in (OAuth)

- 1 Open **Library** › **Tools** › **Connections**, choose Edit on the connection, then Sign-in.

- 2 Register the Redirect URL at the provider, then choose Connect with provider and approve there.
- 3 The badge changes from Not signed in to Signed in — token saved.

Path: Remove an action

- 1 Choose the trash button on its row. It is deleted for the whole workspace at once, with no confirmation.

Check that it worked

- ✓ Run test shows "Finished in ... ms." with the expected output.
- ✓ Members of the chosen group see the action's button on assistant replies, and it returns the result.
- ✓ After switching to **Draft**, the row reads **Draft** and the button no longer appears on replies.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Script has a syntax error on line ..."	Fix the line named in the message and run the test again.
"The script timed out."	Raise Timeout (up to 30 seconds) or simplify the script.
"The script exited with an error (code 1)."	Read the error output under the result and fix the script.
"Script cannot be empty. It reads input from stdin and prints its result."	Enter a script before saving.
A member does not see the button	Check that the action is Enabled and that one of their groups is ticked under Who can run this action.
A sign-in page reads "Sign-in link expired"	Start sign-in again from the connection's Sign-in tab with Connect with provider. The Authorize shortcut on Connections does not complete sign-in today.

34. Automations: preparing scheduled runs

What admins need to configure before users can rely on recurring model chains.

Users create and run automations from the Automations tab of the **Agents** view, and enabled schedules fire automatically in the background (UTC), delivering each run's output as a new chat thread. Admins manage group and model access; the service team manages shared connector availability and credentials outside tenant administration. Check the actual failure details to determine which setup needs attention.

Admin readiness checklist

- 1 Open **Model Access** and confirm the automation's model is visible to the user's group.
- 2 Open **Groups** and confirm the user belongs to the group that receives the model grant.
- 3 Ask the service team to confirm any required shared connector and inspect its real connection-test result. **Users** complete their own source-account connection when the workflow needs delegated access.
- 4 If the automation uses knowledge, confirm the knowledge base is enabled and its source ACL allows the user's group.
- 5 Ask the user to press **Run now** after saving the automation; the transcript and last-run status will show whether the setup is complete.

GOOD TO KNOW

Admins do not need to expose every model to make automations work. Grant the smallest approved model set that fits the scheduled workflow, then expand only when a run proves it needs more capability.

35. SSO: single sign-on for your tenant

Connect an identity provider, provision on first sign-in, and map IdP groups.

The **SSO** tab lists every identity provider for your tenant as a card. When **Policy Controls** shows **SSO** configuration as available, you can add providers for your domains, test them, map identity-provider groups to tenant groups, and enforce **SSO**. Otherwise the tab is read-only and says so.

WALKTHROUGH · 4:54 · ADMIN CONSOLE > DOCUMENTATION

Tenant SSO and provisioning

Set up your organization's single sign-on from the Admin console: check that policy allows it, create and test the connection, sign a test person in, map groups, enforce it, and recognize the read-only state.

Before you begin

- An Admin account, with **SSO** management allowed by service policy: Policies > **SSO** configuration reads Available. Otherwise the **SSO** tab is read-only; ask your service team.
- An OpenID Connect application at your identity provider, with its issuer URL, client ID, and client secret. The paths below list the exact settings for Keycloak, Microsoft Entra ID, Okta, and Google Workspace.
- The redirect URI shown under the **Add SSO configuration** form, registered exactly at the identity provider.
- A groups claim in the ID token if you will map groups. Google sends none.
- A test account at the provider with a verified email on one of your domains.

Step by step

- 1 Open **Admin console** › Policies › **Policy Controls** and check that **SSO** configuration reads Available. If it reads Read only, ask your service team to allow **SSO** management.
- 2 Choose **Admin console** › **SSO**, then **Add SSO configuration**.
- 3 In **Identity provider**, choose your preset: Microsoft Entra ID, Google Workspace, Okta, or Custom OIDC provider.
- 4 Enter the Issuer URL, **Client ID**, and Allowed email domains. Keep Provision new users on first sign-in (JIT) on.
- 5 Register the redirect URI shown under the form at your identity provider.
- 6 Paste the **Client secret** last, then choose **Create SSO** configuration. Enforcement starts off.
- 7 On the new card, choose **Test connection**. Discovery document, Signing keys (JWKS), and Sign-in readiness must all pass.
- 8 In a private window, enter the test person's email, choose **Continue with SSO**, and sign in at the provider. **Users** then lists them with sso under Auth.
- 9 On the card, choose Add mapping for each identity-provider group value, choose the tenant group, and choose Save mappings.
- 10 Have the test person sign in again, then check **Users** with Filter users by group.
- 11 When everyone can sign in through the provider, turn on Enforce for tenant sign-in.
- 12 Confirm that a local password on the domain is now refused.

Path: Custom OIDC provider, such as Keycloak (shown)

- 1 Choose Custom OIDC provider. Enter the issuer, such as https://your-keycloak-host/realms/your-realm.
- 2 At Keycloak: a confidential client with Client authentication on, Standard flow, Require PKCE S256, the redirect URI, and a Group Membership mapper named groups with Full group path off.
- 3 Map group values exactly as they appear in the token, such as litigation.

Path: Microsoft Entra ID

- 1 Choose Microsoft Entra ID. The issuer is prefilled; replace `{tenant-id}` with your Directory (tenant) ID. Never use common or organizations.
- 2 Use the client secret's Value, not its Secret ID.
- 3 At Entra: App registrations › New registration (**Web**, the redirect URI), Certificates & secrets › New client secret, and Token configuration › Add groups claim. Group claims carry object IDs, so map each group's object ID.

Path: Okta

- 1 Choose Okta. The preset fills `https://{your-domain}.okta.com`; replace it with your Okta domain.
- 2 At Okta: Applications › **Create** App Integration › OIDC, **Web** Application, Authorization **Code**, the redirect URI. For group mapping, use the custom authorization server, `https://{your-org}.okta.com/oauth2/default`, and add a groups claim to the ID token.

Path: Google Workspace

- 1 Choose Google Workspace. The issuer is `https://accounts.google.com`.
- 2 Google ID tokens carry no groups, so manage group membership on the **Groups** tab instead of mapping.

Path: When SSO management is read-only (shown)

- 1 The **SSO** tab shows "Organization policy makes **SSO** configuration read-only in this console."
- 2 Existing cards stay visible and keep working, but you cannot add, test, map, enforce, or remove. Ask your service team.

Check that it worked

- ✓ **Test connection** lists Discovery document, Signing keys (JWKS), and Sign-in readiness as passed.
- ✓ The test person reaches the workspace, and **Users** lists them with the USER role and sso under Auth.
- ✓ After their next sign-in, **Users** filtered to the mapped group lists them.
- ✓ The card reads enforced, and a local password on the domain is refused with "**SSO** is enforced for this email domain; local sign-in is disabled. Use the configured identity provider."

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Organization policy makes SSO configuration read-only in this console."	Service policy does not let administrators manage SSO . Policies > SSO configuration reads Read only. Ask your service team.
Create SSO configuration stays unavailable	Issuer URL and Client ID are required.
"Issuer verified, but sign-in stays disabled until you add: client secret."	The secret was not saved, or allowed domains are missing. Cards cannot be edited: choose Remove and create the configuration again with every field.
Test connection fails at Discovery document	The issuer URL is wrong or unreachable from the server. Open the issuer followed by /.well-known/openid-configuration; for Entra use your tenant ID.
The card shows Custom Oidc instead of the display name you typed	A known limitation: the card is labeled by provider type, and the display name is not kept yet. Identify cards by their issuer.
Mapped groups do not appear after sign-in	The claim value must match exactly, including case and any leading slash, and the provider must add the groups claim to the ID token. Membership updates at the next sign-in.
"... is outside the domains allowed for this SSO provider."	The provider vouched for an address on another domain. Add the domain by recreating the configuration, or keep that person out.
"SSO is enforced for this email domain; local sign-in is disabled. Use the configured identity provider."	Expected after enforcement. The person chooses Continue with SSO . Turn enforcement off on the card to allow passwords again.

GOOD TO KNOW

The steps inside each identity provider's console (Microsoft Entra ID, Okta, Google Workspace, Keycloak) are spelled out step by step in the **Identity provider** setup topics of the interactive guide at <https://aperturechat.com/guide.html>. The redirect URI to register is the one shown in the **Add SSO configuration** form.

CAREFUL

Enforcement blocks password sign-in for the allowed domains. A passing **Test connection** checks discovery and keys; it does not prove the client-secret exchange, callback, user session, or group mapping. Validate those with a real sign-in before you enforce.

36. Analytics

Runtime metadata, feedback, per-user usage, and token budgets for this tenant.

The **Analytics** tab measures how the organization actually uses the platform. Sections start collapsed behind descriptive headers — click a header to expand one — and each section carries its own filter row pairing a person picker (defaulting to **All admins and users**) with a date range and its presets (All, Today, Week, 30 days). Scoping one section never narrows another.

- **Runtime Clock Metadata** — organization-scoped execution timestamps captured from admin and user **Chat** and **Draft** completion events. The cards total runtime events, chat completions, and draft calls. Its **CSV** includes actor_id and actor_name columns, and when one user is selected the filename gains that user's suffix so exports stay unambiguous.
- **Chat Feedback Analytics** — response-level thumbs up and thumbs down signals submitted by tenant admins and users, with totals for overall, positive, and negative feedback.
- **Model Activity** — saved prompt volume by model, date, and person for the selected range, drawn as “**Prompts** by model”, “Prompt trend”, and “**Users** by prompt activity”.
- **User Usage** — durable per-user usage for this organization's admins and users across chat, drafts, agents, automations, and the API gateway. Token counts are provider-reported only and stay blank when a provider reported none. The panel has its own **Filter usage by user** selector for drilling into one person.
- **Workspace Usage Budget** — a read-only view of the organization's service-managed ceiling and its current UTC accounting period. This panel shows where usage stands without exposing service-level configuration.
- **Token Allocations** — per-user and per-group token caps beneath the workspace ceiling, with Per day, Per week, or Per month UTC resets. The most restrictive applicable cap wins. Choose a user or group, enter a cap, pick the period, and click **Set cap**. Until you add one, the table says “No allocations yet. Everyone shares the workspace ceiling.”

GOOD TO KNOW

Every number here comes from saved audit and usage events, and each **CSV** export button opens its own small date-range popover so a file contains exactly the rows you chose. If a panel has nothing to show, it says so plainly instead of showing sample data.

WALKTHROUGH · 2:03 · ADMIN CONSOLE › DOCUMENTATION

Tenant analytics

Scope each analytics section to a person and dates, export CSV, read usage and the workspace budget, and set a token cap that people actually hit.

Before you begin

- An Admin account.

- Some real activity: chats, drafts, or automations from people in your organization.
- The workspace ceiling comes from your service team; your caps sit inside it.

Step by step

- 1 Choose **Admin console** › **Analytics**. Every section starts collapsed.
- 2 Expand **Runtime Clock Metadata**. Choose a person and dates, or All, Today, Week, or 30 days.
- 3 Choose **CSV**, pick dates or All dates, and choose **Download** N rows.
- 4 Expand **Model Activity**, **User Usage**, and **Workspace Usage Budget** to read volume, tokens, and the ceiling.
- 5 Expand **Token Allocations**. Choose a person or group in the first menu, enter the Token cap, choose Per day, Per week, or Per month, and choose **Set cap**.
- 6 The row shows the cap and a meter of tokens used in the current UTC period.
- 7 Choose Remove allocation for ... to lift the cap.

Path: Cap one person (shown)

- 1 Choose User · their name, enter the cap, choose the period, and **Set cap**.
- 2 **Once** they pass it, their next message reads "Your daily token budget has been reached. Requests are blocked until the next UTC day."

Path: Cap a group

- 1 Choose Group · its name instead. Members then see "Your group's daily token budget has been reached. Requests are blocked until the next UTC day."
- 2 When several caps apply, the most restrictive one wins.

Path: Export other sections

- 1 **Chat** Feedback and **User Usage** have their own **CSV** buttons with the same date choice.
- 2 Choosing one person in a section's filter adds them to the file name.

Check that it worked

- ✓ Each filter shows "N of M records" for its own section only.
- ✓ The **CSV** downloads with the section and person in its file name, such as aperture-admin-runtime-analytics-....
- ✓ The cap row shows used and capped tokens, and a capped person's next message is refused with the budget message.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Enter a whole number of tokens (0 means no cap)."	Type digits only. 0 records the person or group with no cap.
"Unknown user for this workspace." or "Unknown group for this workspace."	The person or group was removed. Choose Refresh and pick again.
"Allocations total ... tokens, more than the ... -token workspace ceiling."	Allowed, but the ceiling still wins. Lower the caps or ask your service team to raise the ceiling.
The capped person waits several minutes before seeing the refusal	The chat retries a refused request before showing the message. The refusal itself is immediate on the server.
Token columns are blank for some rows	Tokens are counted only when the provider reports them; blank means none was reported.

37. Policies: service availability and memory governance

Apply downstream tenant defaults, configure memory, and govern by count without reading content.

The Policies tab is always present. **Policy Controls** and **Personal Data Protection** start collapsed. When service policy enables memory, Personalization Memory and Memory by User also appear collapsed. Otherwise, Memory governance explains the restriction; saved organization settings remain intact. Expand the panel you need; administrators can narrow available capabilities, while unavailable settings remain locked. **Personal Data Protection** has its own section in this guide.

Policy Controls

When your service team limits something you cannot change here, a Service policy note at the bottom of **Policy Controls** says so: administrator accounts created by the service team, administrators required to sign in with **SSO**, or newly available models starting without access until you grant a group in **Model Access**. No note means nothing is limited. Whether you can edit **SSO** is shown on the **SSO** tab itself.

Policy Controls sets defaults for the protected **Default Users** group: personal API keys, private-agent building, private knowledge-base and tool authoring, and personalization memory. A switch locks when its capability is unavailable under service policy; the saved group grant is preserved rather than silently erased. Use **Groups** for exceptions and **Model Access** for available models. Shared connector availability is managed by the service team outside tenant administration.

Personalization Memory settings

- When service policy permits memory, expand Personalization Memory and turn on Memory for this organization. This makes the account-level Personalization memory row available to eligible users; existing memories stay saved but are not applied while this switch is off.

- Choose whether to allow Learn from conversations automatically. When off, only explicit requests such as “remember that ...” and direct additions in the memory manager create memories. Every user can still opt out of automatic learning individually.
- Set Retention (days), from 1 through 3650. Older memories retire automatically when they pass this policy.
- Set Maximum memories per user, from 1 through 2000. The default is 200, which is a practical general-purpose limit; raise or lower it only when your retention, compliance, or workload policy calls for a different capacity. When the cap is passed, the least useful unpinned memories retire first.
- Expand Memory by User to review content-free counts. Click **Refresh** to reload the counts, or purge one person's memories when compliance or account cleanup requires it.

CAREFUL

Memory administration never grants reading access. Administrators see policy, counts, and purge controls only; the API and the interface do not return another person's memory content.

WALKTHROUGH · 2:29 · ADMIN CONSOLE › DOCUMENTATION

Policies and memory governance

Read what service policy allows, set the Default Users capabilities, tune personalization memory, purge a person's memories without reading them, and recognize when memory is unavailable.

Before you begin

- An Admin account.
- Service policy sets the ceiling: what is locked here is decided by your service team.

Step by step

- 1 Choose **Admin console** › Policies. Its sections start collapsed; expand **Policy Controls**.
- 2 If your service team limits anything you cannot change here, a Service policy note at the bottom of **Policy Controls** says so. No note means nothing is limited.
- 3 Switch on the Default users capabilities you want: downstream API, build agents, build knowledge bases, build tools, and use memory. Each saves with "Default user policy saved."
- 4 Confirm the change in **Groups** › **Default Users** › **Permissions**.
- 5 Expand Personalization Memory. Set Memory for this organization and Learn from conversations automatically.
- 6 Type Retention (days), 1 to 3650, or Maximum memories per user, 1 to 2000, and press Enter. The status reads "Memory policy saved."
- 7 Expand Memory by User and choose **Refresh**. To delete one person's memories, choose Purge, then Yes, purge.

Path: What the Service policy note can say

- 1 "Administrator accounts are created by your service team." Admin is not offered when you add or approve people.
- 2 "Administrators must sign in with **SSO**." Administrator accounts cannot use email and password.
- 3 "Newly available models start without access until you grant a group in **Model Access**." Grant each new model yourself.
- 4 Whether you can edit **SSO** is shown on the **SSO** tab itself.

Path: What each Default users switch changes

- 1 Downstream API: people can create personal API keys for their approved models.
- 2 Build agents, knowledge bases, or tools: people can create their own private ones. Sharing and publishing stay with administrators.
- 3 Use memory: the assistant saves and recalls each person's private preferences.

Path: When memory is off for the platform (shown)

- 1 Policies shows a single Memory governance panel: "Personalization memory is unavailable under the current service policy. Saved organization settings remain intact."
- 2 Your saved settings return when service policy allows memory again.

Check that it worked

- ✓ The status reads "Default user policy saved." and the matching switch is on in **Groups > Default Users > Permissions**.
- ✓ "Memory policy saved." and the new value remains after reopening Policies.
- ✓ After a purge the status reads "Deleted N memories for" and the person leaves Memory by User.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Unavailable under the current service policy. The saved group grant is preserved."	Service policy has that capability off. Your earlier grant returns when it is turned back on.
"Turn on Memory for this organization below before granting the default group access."	Turn on Memory for this organization in Personalization Memory first.
A memory number did not change	Values outside 1 to 3650 days or 1 to 2000 memories are ignored. Type a value in range and press Enter.
"Memory policy was not saved. Personalization memory is unavailable under the current service policy."	Service policy has personalization memory off. Ask your service team.
"Default user policy could not be saved."	Nothing changed. Reload Policies and try again; if it repeats, check the Audit trail for the refusal.

38. Personal Data Protection

Conceal personal data in chats, records, and exports, and decide whether the model may read it.

Personal Data Protection is a panel on the Policies tab. It is off until an administrator turns it on, and it applies to everyone in the organization at once. When it is on, each detected value is replaced with a labeled placeholder, such as SSN or Card number, before a chat or record is saved or shown. People see the placeholder as a locked chip, and the line under their message box ends with Personal data is concealed.

Controls

CONTROL	WHAT IT DOES
Conceal personal data	Turns concealment on or off for the organization. Off: chats are stored and shown exactly as typed; content filters attached to individual models still apply.
Hide values from the model too	On by default. Typed prompts and attached file text reach the model provider as placeholders. Off: the model reads the original value for that turn, and the value is still concealed everywhere it is stored or shown, including the model's reply.
What to conceal	Government and personal IDs (Social Security number, ITIN, passport, driver's license, date of birth, vehicle identification number); Contact details (email, phone, street address); Financial accounts (payment card, card security code, bank account, ABA routing number, IBAN); Health identifiers (medical record number, health plan member or subscriber ID, Medicare beneficiary identifier); Secrets and credentials (private keys, API keys and access tokens, disclosed passwords or PINs); Network identifiers (IP addresses). At least one stays selected.
Try it with sample text	Preview concealment runs the selected kinds on sample text and reports how many values it concealed. The preview is not saved or logged.

Where it applies

- **Chat** messages, titles, and regenerated answers, before they are saved and every time they are shown.
- **Model** replies, including streamed text, before they reach the browser.
- **User Prompt Activity**, feedback notes, security alert snippets, and retention tags.
- Memories, issue reports, search results, the Elastic export, and training datasets.
- **Chats** saved before protection was turned on are concealed whenever they are shown or exported, and stored concealed the next time they are saved.

CAREFUL

Concealment is permanent for what it stores. Turning protection off does not bring values back, because the originals were never saved.

GOOD TO KNOW

Detection runs inside your deployment and checks formats and check digits: card numbers must pass the Luhn check, Social Security numbers must fall in issued ranges, and routing numbers, IBANs, and vehicle identification numbers must pass their checks. Look-alike characters, such as full-width digits or hidden spaces, cannot hide a value. It does not recognize names or free-text health details. **Drafts** are documents of record and are not altered; a draft request that contained personal data is recorded as DRAFT_NOT_CONCEALED. Uploaded files are stored as uploaded; only the text extracted for the model is concealed.

Evidence in Audit

- **User Prompt Activity** shows concealed prompts with the same chips. Choose **Refresh monitor** to load new activity.
- The **Audit Trail** records PROMPT_CONCEALED for each concealed message, with the sender, the model, and whether the model saw placeholders, and POLICY_UPDATED for each change to the panel. **Search** privacy to list them.

WALKTHROUGH · 3:20 · ADMIN CONSOLE › DOCUMENTATION

Protect personal data

Turn on Personal Data Protection, choose what to conceal and whether the model may read it, test it on sample text, and confirm the result in a real chat and in Audit.

Before you begin

- An Admin account.
- Your organization's decision on which kinds of personal data must not be stored in chats, and whether the model may read them.
- Synthetic test values, such as the sample already in the panel. Never test with a real person's data.

Step by step

- 1 Choose **Admin console** › Policies. **Personal Data Protection** sits below **Policy Controls**; its header reads Off or On.
- 2 Expand it and switch on **Conceal personal data**. It saves at once for everyone in the organization, and the header reads On.
- 3 Leave **Hide values from the model too** on unless the work needs the model to read the values. When on, typed prompts and attached file text reach the model provider as placeholders.
- 4 Under **What to conceal**, switch off any kind of personal data you do not want concealed. At least one stays on.
- 5 Under **Try it with sample text**, keep or edit the synthetic sample and choose **Preview concealment**. Read the count, the concealed text, and the pill for each detector. The preview is not saved or logged.
- 6 Read **Where it applies**, and the limits below it.
- 7 Have a person send a message with a synthetic value. It shows a locked chip such as SSN, and the line under their message box ends with Personal data is concealed.
- 8 Choose **Audit** › **User Prompt Activity** › **Refresh monitor**: the prompt shows the same chip.
- 9 In the **Audit Trail** choose **Refresh** and search privacy. PROMPT_CONCEALED records each concealed message; POLICY_UPDATED records your change.

Path: Keep values from the model (default, shown)

- 1 **Hide values from the model too:** on.
- 2 The model receives placeholders such as `[[SSN]]` instead of the value, so its reply refers to the placeholder, for example `[SSN]`.
- 3 Use this when the model does not need the value to do the work, such as drafting a note or summarizing a record.

Path: Let the model read the values (shown)

- 1 Switch off **Hide values from the model too**. The row reads "The model reads the original value for that turn. It is still concealed everywhere it is stored or shown."
- 2 Use this when the work needs the value, such as filling in a form or a signature block.
- 3 Values the model writes are concealed before the reply reaches the browser and before it is saved.

Path: Narrow what is concealed

- 1 Under **What to conceal**, switch a kind of personal data off or on. Each switch saves at once.
- 2 Each card lists the values it finds. For example, Contact details covers email addresses, phone numbers, and street addresses.
- 3 The last switch that is on cannot be turned off. To stop concealing, switch off **Conceal personal data** instead.

Path: Turn protection off (shown)

- 1 Switch off **Conceal personal data**. The header reads Off, and new chats are stored and shown exactly as typed.
- 2 Content filters attached to individual models still apply.
- 3 **Chats** saved while protection was on stay concealed: their values were never stored, so they cannot come back.

Check that it worked

- ✓ The panel header reads On.
- ✓ **Preview concealment** reports the values concealed, for example "6 values concealed", with a pill for each detector.
- ✓ A test message shows locked chips in the chat and in **Audit > User Prompt Activity**.
- ✓ The **Audit Trail** lists `PROMPT_CONCEALED` for the message and `POLICY_UPDATED` for your change.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
A kind of personal data cannot be switched off	At least one stays selected. Switch another one on first, or switch off Conceal personal data .
Hide values from the model too is dimmed	It applies only while Conceal personal data is on.
"Nothing detected"	The sample has no value the selected kinds recognize. Detection checks formats and checksums, so a look-alike such as a card number that fails its check digit is left alone.
A name or a description of someone's health was not concealed	Detection does not recognize names or free-text health details. Ask people to leave them out, or add a content filter to the model.
A draft still contains a value	Drafts are documents of record and are not altered. A draft request that contained personal data is recorded in the Audit Trail as DRAFT_NOT_CONCEALED.
"Personal data protection settings could not be loaded."	Reload Policies. If it repeats, confirm your account is still an administrator.
"The setting was not saved."	Nothing changed. Try again ; if it repeats, check the Audit Trail and your connection.

39. Training datasets

Capture de-identified ratings and corrections, route them into datasets, review them, and download them for fine-tuning.

The **Datasets** tab keeps a private, de-identified record of how your people rate and correct answers, so you can fine-tune an open-weight model on your own work later. It is off until you turn it on. Your model providers keep running under zero data retention: captured examples are never sent to a provider. They stay in your deployment until an administrator downloads a dataset, and every download is audited.

What is captured

SIGNAL	CAPTURED WHEN	BECOMES
Rated helpful	A person gives a reply thumbs up, with or without a note.	An approved answer.
Rated unhelpful	A person gives a reply thumbs down, with any note.	A rejected answer, with the note.
Corrected	The person's next message pushes back, such as "That's wrong, the deadline is 60 days. Please revise." A regenerated reply counts too.	The first answer as rejected, and the reply that followed as preferred unless it was rated down or corrected again.

Corrections are recognized by a transparent rule: factual pushback such as "that's wrong", "you missed", or "should be" counts more than a style request such as "shorter" or "rewrite", and an opener such as "No," or "Actually" adds weight. Ordinary follow-up questions are not captured as corrections.

Safeguards

- Every captured text is de-identified with all kinds of personal data, whatever the **Personal Data Protection** settings are.
- **Conceal people and client names** (on by default) also replaces the full names of people in your workspace and your configured client and matter names. Other names are not detected.
- **Skip sensitive or regulated chats** (on by default): chats with a sensitive or regulated retention tag, confirmed or suggested, are never captured.
- **Never capture from these groups**: members of a selected group are never captured.
- **Review before export** (on by default): new examples wait for an administrator to approve them, and downloads include approved examples only unless you choose otherwise.
- Examples follow their chat: deleting a chat, a retention purge, or deleting the person or organization removes its examples.

How examples are labeled and routed

- Practice area: from the chat's subject retention tag when subject tagging is on, otherwise from keywords in the person's own words, for example Legal · Litigation or Financial · Tax.
- Kind of work: Drafting, Review & redlining, Research & Q&A, Summarization, Analysis, Extraction, Translation, Coding, or General.
- Department: the person's groups, other than **Default Users**.
- A dataset is a set of routing rules (signals, practice areas, kinds of work, departments, and models; an empty rule accepts everything) plus a training format. Membership is computed, so editing a rule re-routes every example, and one example can feed several datasets.

- **Suggested datasets** appear for practice areas and departments with three or more examples and no dataset yet.

Training formats

FORMAT	TRAIN.JSONL CONTAINS	USES
Supervised fine-tuning	Chat messages ending in the answer people approved, or the revised answer after a correction.	Helpful ratings and accepted corrections.
Preference pairs (DPO)	prompt, chosen, and rejected.	Corrections with an accepted revision.
Binary feedback (KTO)	prompt, completion, and a true or false label.	Every judged answer; works with thumbs ratings alone.

A download is a ZIP with a folder named after the dataset, holding train.jsonl in the dataset's format, metadata.jsonl with line-aligned labels and no user identity, and a README.md dataset card. Identical content captured twice, for example from a forked chat, is exported once. The **Audit Trail** records each capture setting change, scan, dataset change, review, and download; set its category filter to training to list them.

WALKTHROUGH · 4:36 · ADMIN CONSOLE › DOCUMENTATION

Build training datasets from ratings and corrections

Capture de-identified ratings and corrections, route them into datasets by practice area, department, and kind of work, review them, and download a file for fine-tuning an open-weight model.

Before you begin

- An Admin account.
- Your organization's approval to keep de-identified copies of rated answers for model training, and a decision on which groups are never captured.
- People who rate replies with the thumbs buttons or correct the assistant in chat.
- For the download: a place inside your organization to keep the file, such as the environment where you will fine-tune an open-weight model.

Step by step

- 1 Choose **Admin console › Datasets. Training Data Capture** reads Off, and nothing is collected.
- 2 Switch on **Capture training signals**. The header reads Capturing.
- 3 Under **Capture these signals**, keep or clear Helpful ratings, Unhelpful ratings and notes, and Corrections and regenerations.

- 4 Check the safeguards. **Review before export**, **Skip sensitive or regulated chats**, and **Conceal people and client names** are on by default.
- 5 Under **Never capture from these groups**, choose any group whose work must never be captured.
- 6 Choose **Scan existing chats** to capture chats saved before capture was on. The notice reads "Scanned N chats and captured N new examples."
- 7 From now on, ratings and corrections are captured as chats are saved. In **Captured Signals**, choose **Refresh** to see the counts and the work mix.
- 8 Under **Suggested datasets**, choose **Create** on a suggestion, or choose **New dataset**.
- 9 Check the Name, **Training format**, and routing rules, then choose **Create dataset**.
- 10 Choose Review on the dataset card. Read each example as it would be exported, then choose **Approve** or **Exclude**, or **Approve all shown**.
- 11 Choose **Download**, then **Download** N examples in the dialog. Keep the ZIP inside your organization.
- 12 Choose **Audit** > **Audit Trail** > **Refresh** and set the category filter to training to see each step recorded.

Path: Preference pairs (DPO) (shown)

- 1 Uses corrections whose revised answer was accepted: not rated down and not corrected again.
- 2 train.jsonl holds prompt, chosen (the revised answer), and rejected (the original answer).

Path: Supervised fine-tuning

- 1 Uses answers rated helpful and accepted revisions.
- 2 train.jsonl holds chat messages that end in the approved or revised answer.

Path: Binary feedback (KTO)

- 1 Uses every judged answer, so it works with thumbs ratings alone.
- 2 train.jsonl holds prompt, completion, and a true or false label.

Path: Create a dataset from scratch

- 1 Choose **New dataset** in the **Datasets** panel.
- 2 Type a Name and an optional Description, and choose a **Training format**.
- 3 Choose rule chips under Signals, Practice areas (a broad area includes its specialties), Kinds of work, Departments, and **Models**. Leave a rule empty to accept everything.
- 4 Optionally type a System prompt for every example, then choose **Create dataset**.

Path: Skip review

- 1 Switch off **Review before export**. New examples are approved as they are captured, and you can still exclude any of them.
- 2 Or keep review on and switch on **Include examples waiting for review** in the **Download** dialog, for one download.

Path: Edit, archive, or delete a dataset

- 1 The pencil edits the name, format, and rules. Every example is re-routed at once.
- 2 **Archive** stops routing examples to the dataset; **Restore** brings it back.
- 3 **Delete** removes the dataset definition. Its examples stay captured and keep routing to other datasets.

Check that it worked

- ✓ **Training Data Capture** reads Capturing, and **Captured Signals** counts examples.
- ✓ The dataset card shows approved examples and none waiting.
- ✓ The download is a ZIP with train.jsonl, metadata.jsonl, and README.md.
- ✓ The **Audit Trail** lists DATASET_CREATED, EXAMPLES_REVIEWED, and DATASET_EXPORTED.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
Scan existing chats is dimmed	Switch on Capture training signals first.
"No examples captured yet"	Capture is off, or nobody has rated or corrected an answer since it was turned on. Turn capture on and choose Scan existing chats to include earlier work.
A rated chat was not captured	It may carry a sensitive or regulated retention tag; a chat in which a Social Security or card number was found is usually tagged. Or its owner is in a group under Never capture from these groups . Very short replies are also skipped.
A follow-up was not treated as a correction	A correction needs clear pushback, such as "that's wrong", "you missed", or "should be". A new question, or a style request such as "make it shorter" on its own, is not enough.
An example has the wrong practice area	Without a subject retention tag, the practice area comes from keywords, and the example says Practice area from keywords. Turn on subject tagging in Data Retention, or route by department instead.
Download is dimmed	No examples match the dataset yet. Check its rules and format: preference pairs need corrections whose revised answer was accepted.
The dialog says Download 0 examples	Nothing is approved yet. Approve examples, or switch on Include examples waiting for review .
"Training data settings could not be loaded."	Reload Datasets . If it repeats, confirm your account is still an administrator.

40. Data retention and conversation tags

Find tagged conversations, inspect their contents, and review batch actions.

[WALKTHROUGH](#) · [2:16](#) · [ADMIN CONSOLE](#) · [DOCUMENTATION](#)

Data retention and tagging

Keep chats forever by default, label conversations by client or matter, place a legal hold, and preview a finite schedule before anything is deleted.

Before you begin

- An Admin account.

- Your organization's retention rules, and the client or matter names and aliases people use.
- Saved chats to label. Labels come only from saved message text.

Step by step

- 1 Choose **Admin console** > **Audit** and expand Data Retention.
- 2 Schedule and rules: Keep chats for starts at Forever, with automatic deletion off.
- 3 Under Clients, matters, and regulated records, choose a Source type, enter Name or reference and aliases, and choose Add source to policy.
- 4 Choose Save Forever. The status reads "Forever saved. Automatic deletion is off for all chats."
- 5 Choose Tags and holds, then **Scan existing chats**. **Search** for the matter and review each suggestion.
- 6 Select the right chats, choose the label, and choose Confirm label.
- 7 Expand Legal holds, enter a Hold name, and choose Hold selected chats.
- 8 To consider a time limit, choose it in Schedule and rules and choose Preview effect. Save only after reviewing the counts.

Path: Keep forever (default, shown)

- 1 Leave Keep chats for at Forever. Choose Save Forever after adding sources; nothing is deleted.

Path: A finite schedule (previewed, shown)

- 1 Choose 1, 5, 7, or 10 years. Count age from and Review window (at least 7 days) appear.
- 2 Choose Preview effect and read how many chats are old enough, on hold, or kept.
- 3 Choose Save retention policy only when the counts are right. Eligible chats enter the review window first.

Path: A rule for one label

- 1 Each source has a Retention rule menu. A longer matching rule wins over the default schedule.
- 2 Tick Apply time limits only to labels with a rule to leave unlabeled chats at Forever.

Path: Manual archive or delete

- 1 In Tags and holds, select chats and choose **Archive** selected or **Delete** selected.
- 2 **Delete** asks first and cannot be undone. **Chats** under an active legal hold are skipped.

Check that it worked

- ✓ The source appears in the list with its aliases, and Save Forever confirms automatic deletion is off.
- ✓ The confirmed chats show the matter label; held chats show Legal hold.

- ✓ Preview effect reports the totals without changing anything.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"Enter a source name of at least two characters."	Enter the client or matter name.
"Aliases must be 3–160 characters long."	Remove very short aliases, such as initials.
"Preview this exact policy again before saving; the policy or affected chats changed."	Choose Preview effect again, then save.
Scan suggests nothing for a matter	Suggestions come from saved message text only. Add the aliases people actually write, save, and scan again.
Hold selected chats is unavailable	Select 1 to 500 chats and enter a Hold name of at least three characters.

How retention behaves, in detail

- Open **Audit** and expand Data Retention, then choose Schedule and rules. Forever is the default: automatic deletion is off until an administrator previews and saves a schedule. Choose 1, 5, 7, or 10 years, or keep Forever. One year means 365 days.
- Choose whether age starts at chat creation or the last message change. Renaming or archiving a chat does not restart its clock. Select a review window of at least 7 days. Saving a changed policy restarts that window for eligible chats.
- Add a stable client, matter, or regulated-record source and its aliases. A source can have a longer retention rule or Forever. Choose Apply time limits only to labels with a rule to leave all other chats stored indefinitely. The longest applicable duration wins; a matching Forever rule or active legal hold prevents automatic deletion.
- Save the source definitions, then use **Scan existing chats** in Data Retention → Tags and holds. New message saves also look for source mentions. Detected mentions are suggestions only: select the correct chats and source, then Confirm label. Remove / dismiss label rejects a false match. The label records who confirmed it and when.
- Optional sensitive-data suggestions check saved message text for email addresses, possible Social Security numbers, and payment cards. They do not detect every form of PII or regulated record, or inspect original uploaded files. Raw matched values are never copied into labels. Review before confirming a category.

- Choose Preview effect to see counts across all saved chats, including legal holds and chats without a deadline. Saving the policy recalculates existing and future chats. A shorter duration can make old chats eligible, but they still receive the review window. A longer duration delays deletion; it cannot restore deleted content. Saving Forever stops automatic deletion.
- To inspect chats, open **Audit**, expand Data Retention, and choose Tags and holds. This list includes tagged and untagged conversations available in your administrative scope.
- Use **Search** chats and tags to find titles, people, tags, or client/matter identifiers. Filter by tag type when you need a narrower set.
- Click a conversation title to open Tagged conversation and review its saved prompts and outputs. Close the preview to return to the same list.
- Select the intended chats. **Archive** selected keeps them stored and searchable but removes them from the active list. **Delete** selected opens a permanent-deletion confirmation; it does not delete until you choose Yes, delete.
- Under Legal holds, name a hold and choose Hold selected chats. This hold protects the selected records, including against manual deletion. For ongoing preservation of a client, also set that source's rule to Forever. Load active holds to review them; releasing a hold requires a second confirmation.
- Read the completed action status. **Chats** under an active hold are skipped by deletion. A failure or skipped record needs review; do not assume every selected row was deleted. The background scheduler handles automatic cleanup in bounded batches and writes an audit record per deleted chat.

CAREFUL

Permanent deletion removes the selected conversations and their attachments and cannot be undone. Preview contents and check the selected count before confirming. A tagging switch records classification; it is not a scheduled deletion rule or an action that places a hold.

41. Review feedback and reported issues

Read response ratings, written feedback, and platform reports in Analytics.

WALKTHROUGH · 1:09 · ADMIN CONSOLE > DOCUMENTATION

Review feedback and reported issues

Follow a real rating with a note and a reported problem from the person's screen to your review in Analytics.

Before you begin

- An Admin account.
- People who rate replies or report problems from **Help**.

Step by step

- 1 The person rates a reply with the thumbs buttons below it and may add a note, then chooses Send note.
- 2 The person reports a problem from **Help** › Report a problem with a Subject, Message, and optional screenshot, then Send report.
- 3 Choose **Admin console** › **Analytics** and expand **Chat** Feedback.
- 4 Choose the person and dates, or Today, Week, or 30 days.
- 5 Open a feedback row to read the note and the rated exchange.
- 6 Scroll to Reported platform issues and open a report to read it.
- 7 Follow up through your team's support process.

Path: A rating with a note (shown)

- 1 Thumbs up or down, then an optional note. The row shows the sentiment, model, person, and note.

Path: A reported problem (shown)

- 1 **Help** › Report a problem. The report shows the subject, message, sender, time, and any screenshot.

Check that it worked

- ✓ The summary counts Total feedback, Positive, Negative, and Issue reports for your scope.
- ✓ Feedback and conversation shows the note and the rated exchange.
- ✓ Platform issue report shows the subject and message the person sent.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
The list is empty	Widen the person and date filter. Each section filters independently.
The conversation could not be loaded	The dialog says so and shows the saved reply preview only. Use the evidence that loaded.
"The screenshot preview could not be loaded."	The report has a screenshot that failed to load. Read the text and try again later.

42. Audit

Grouped signals, Audit Insights trends, prompt monitoring, security alerts, and the tenant trail.

The **Audit** tab is the tenant's governance station. It opens on **Admin Audit** and **Audit** Insights, while the sections below start collapsed; click a section header to expand it. **User Prompt Activity**, **Security Alerts**, and the **Audit Trail** each carry their own filter row pairing a person picker with a date range, so one section's scope never narrows another. The **CSV** export buttons produce files of exactly the filtered rows, including actor_id and actor_name columns, so every exported row names who acted.

Admin Audit signals

- The banner at the top reads, for example, “7 of 18 signals need attention”. Expand all and Collapse all open or fold every group. List shows one compact row per signal; Cards shows each signal with its description. The layout you choose is remembered in this browser.
- **Security** signals: audit events, critical events, the prompt watchlist of active DLP or misuse alerts, alert response (the median time to acknowledge an alert), after-hours changes (outside 7 AM–7 PM or on weekends), and failed operations.
- Identity & access: active admins, active users, access requests awaiting review, accounts that have never signed in, role changes, and credential changes such as MFA, password, session, and API key events.
- **Models** & workspace: prompt volume, connector issues, ungrouped models, unassigned users (active users in no group), agent approvals waiting for review, and automation failures.
- **Groups** that need attention open automatically and all-clear groups fold to one summary line. Red rows need attention. Select any signal to open an investigation listing every record behind it, with a filter box to narrow the records.

Audit Insights

- **Audit** Insights charts the trends behind the signals for your organization. Choose a 7 days, 14 days, or 30 days range; the summary line counts the audit events and security alerts in range.
- **Audit** events by day stacks info, warning, and critical events. **Security** alerts by day draws DLP and behavior alerts as two lines, and Alert breakdown ranks alerts By rule or By person with their open and acknowledged counts.
- Most active people ranks who is acting, Activity by area groups events into areas such as **Chat**, Sign-in & identity, and Administration, and Activity by hour shades the hours outside 7 AM–7 PM in your local time.
- Select any bar, point, or row to open the same investigation view with the records behind it. If the loaded history reaches its page limit, a note says earlier days may be incomplete.

Sections below the dashboard

- **Recent Governance Activity:** the current tenant snapshot for identity, user, model, and connector posture: how many **SSO** configurations are enforced, how many users are active, and how many connectors are enabled.
- **User Prompt Activity:** drill into saved prompts from this organization's admins and users and their model responses by person, thread, model, and timestamp, scoped by the section's own user and date filter.
- **Security Alerts:** DLP and malicious-behavior flags raised from admin and user prompts, shown with redacted snippets for review. Click **Acknowledge** once an alert is handled, or **Reopen** if it needs another look.
- **Audit Trail:** the tenant's append-only transaction log, newest first. Its toolbar has a severity select, an action-category select, and a search box (**Search actions, people, targets...**) that narrow the rows together; the trail's **CSV** exports exactly the rows you are looking at. Click **Refresh** to reload it straight from the admin audit API.

GOOD TO KNOW

Security alerts fire on real prompt content, such as payment card numbers, shared credentials, and prompt-injection attempts, but show only redacted snippets, so reviewing an alert never re-exposes the sensitive value itself.

WALKTHROUGH · 1:33 · ADMIN CONSOLE › DOCUMENTATION

Tenant audit

Investigate a real flagged prompt from start to finish: the signal board, Audit Insights, the prompt and reply, the security alert, and the audit trail.

Before you begin

- An Admin account.
- A flagged prompt to investigate. This lesson follows a real prompt-injection attempt, the same one the **Alerts** and delivery lesson follows to a delivered email.

Step by step

- 1 Choose **Admin console › Audit**. The banner counts the signals that need attention.
- 2 Select a red signal, such as Prompt watchlist, to open its records. Close it to return.
- 3 In **Audit Insights**, choose 7, 14, or 30 days, then select a bar or point to list the records behind it.
- 4 Expand **User Prompt Activity**, search for the prompt, and open it to read the prompt and the model's reply.
- 5 Expand **Security Alerts** and choose **Acknowledge** once you have reviewed the alert.
- 6 Expand **Audit Trail**. Filter by severity, category, and text, then choose **CSV** to export what you see.

Path: Start from an alert email

- 1 The email names the rule, detection, person, and time, never the flagged text.
- 2 Open **Admin console** > **Audit** > **Security Alerts** and find the alert by detection and time, then follow the steps above.

Path: Export evidence

- 1 **User Prompt Activity** and **Audit Trail** each have a **CSV** button. The trail export carries exactly the filtered rows.

Check that it worked

- ✓ The investigation lists the flagged records with person, model, and time.
- ✓ The prompt dialog shows the exact prompt and the model's saved reply.
- ✓ The alert reads Acknowledged, and the trail shows the security.prompt_flagged events.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
A signal opens with no records	The count reflects the current snapshot; check the Audit Insights range or the trail's date filter.
The prompt you expect is missing from User Prompt Activity	Each section has its own person and date filter. Choose All and clear the search.
"Security alert was not updated: ..."	Nothing changed. Refresh and acknowledge again; another administrator may have changed it.

43. Alerts

Watch rules over audit activity, and honest delivery statuses for every alert.

The **Alerts** tab turns audit activity into notifications. It has three panels: **Email Delivery**, **Alert Rules**, and **Alert Deliveries**. Every alert is always logged in-app regardless of email — email is a delivery channel, not the record.

Email Delivery

For organization administrators this panel is a read-only status: it reports whether email delivery is configured at the service level. If email is not configured, rules still work — their alerts are logged in-app and the delivery log says so honestly.

Alert Rules

- Click Prompt-injection template to start from a rule that watches prompt-injection, system-prompt extraction, and credential-extraction attempts; **Suspicious-activity template** for security flags and elevated-severity events; or **New rule** to build one from scratch. A template only fills in the form; nothing is saved until you click **Create Rule**.
- A rule has: a name; **Action patterns** — exact audit actions or prefixes such as security.* or admin.user_deleted; a **Minimum severity**; an optional **Watched user**; a threshold — how many matches within a time window before it fires; a cooldown between alerts; and email recipients. Leave recipients empty and the alert is in-app only.
- Only these detections narrows a rule to prompts flagged by specific detectors: Prompt injection, System-prompt extraction, Credential extraction, API key or token shared, Private key shared, **Password** shared, US Social Security number, and Payment card number. Leave all unchecked to match every event the action patterns allow. A rule that could never fire is refused with an explanation when you save it.
- Alert emails name the rule and the detection but never include the flagged prompt text.
- Rules you create here watch this organization's admin and user audit activity.

Alert Deliveries

Every alert trigger is listed with its real delivery status: sent, queued, failed with the actual SMTP error, email not configured, or logged in-app. The log exports to **CSV**, archived deliveries included. Click **Archive** on a delivery to clear it from the default view — its history is kept, and **Show archived** reveals archived deliveries so you can review or **Restore** them. Tenant admins can archive their tenant's deliveries only.

WALKTHROUGH · 1:51 · ADMIN CONSOLE › DOCUMENTATION

Alerts and delivery

Check email delivery, create a rule from the Prompt-injection template with a recipient, and follow a real flagged prompt to a delivered email.

Before you begin

- An Admin account.
- Email delivery configured by your service team. Without it, alerts are still logged in-app.
- The addresses that should receive alerts.

Step by step

- 1 Choose **Admin console › Alerts. Email Delivery** should read Email configured.
- 2 In **Alert Rules**, choose Prompt-injection template.
- 3 Check the action pattern security.prompt_flagged, the minimum severity, Fire when, and Cooldown.
- 4 Enter Email recipients, separated by commas.

- 5 Under Only these detections, keep or change the ticked detectors.
- 6 Choose **Create** Rule. The rule appears in the list, enabled.
- 7 When a matching prompt is flagged, **Alert Deliveries** shows the delivery and its status, such as sent.

Path: Prompt-injection template (shown)

- 1 Watches security.prompt_flagged for Prompt injection, System-prompt extraction, and Credential extraction, at Warning and above, every match, with a 10-minute cooldown.

Path: Suspicious-activity template

- 1 Watches security.*: security flags, content-filter hits, and elevated governance events, at Warning and above, with a 15-minute cooldown and no detection filter.

Path: New rule

- 1 Enter a Rule name. **Action patterns** accepts exact actions or prefixes such as admin.* or auth.*; empty matches every action.
- 2 Choose a minimum severity, an optional **Watched user**, Fire when (matches within minutes), and a cooldown.
- 3 Detections apply only to security.prompt_flagged events.

Path: In-app only

- 1 Leave Email recipients empty. Deliveries read logged in-app.

Check that it worked

- ✓ The rule appears enabled, with its patterns, detections, and recipient count.
- ✓ **Alert Deliveries** shows the delivery with status sent and the recipient.
- ✓ The email arrives with the rule and detection in its subject and no flagged text.

Troubleshooting

WHAT YOU SEE	WHAT TO DO
"A rule name is required."	Enter a Rule name.
"'...' is not a valid email address."	Separate recipients with commas and check each address.
"Detections only apply to security.prompt_flagged events. ..."	Add security.prompt_flagged or security.* to Action patterns , or clear the detections.
"None of the chosen detections can reach the minimum severity '...'"	Lower Minimum severity or choose other detections.
A delivery reads email not configured	Email delivery is managed at the service level. The alert stays logged in-app; ask your service team to configure email.
A delivery reads failed	Its detail shows the mail server's real error. Share it with your service team.