



For platform owners

Platform Owner Guide

Aperture Chat

The complete User and Administrator Guides, plus first-run guidance, providers and shared connectors, the API key vault, organization model availability, SSO and MFA policy, branding, search indexing, budgets, releases, analytics, audit, and alerts. No prior knowledge assumed.

Contents

PART 1 · GETTING STARTED

1	Welcome	4
2	Signing in	4
3	Request access and finish your first sign-in	5
4	Two-step verification and recovery codes	6
5	Finding your way around	7
6	Organizing chats: folders, pins, and recent	9
7	Searching everything	10

PART 2 · CHAT

8	Your first chat	11
9	Understand model availability and request access	12
10	Reading and acting on replies	12
11	Composer symbol shortcuts	14
12	Attaching files, web pages, and connected sources	14
13	Send options: Knowledge, Web, Agent, and Reasoning	15
14	Session details and sources	16

PART 3 · DRAFTS, DECKS, AND AGENTS

15	Drafting documents	17
16	Building a slide deck	19
17	Recover work kept only on this device	21
18	Agent profiles	22

PART 4 · KNOWLEDGE/TOOLS AND AUTOMATIONS

19	Knowledge bases	23
20	The Tools library	23
21	Scheduled automations	24

PART 5 · APPEARANCE, YOUR ACCOUNT, AND HELP

22	Light and dark mode	26
23	Your account	27
24	Personalization memory	28
25	Install the app on your phone	29
26	Getting help	29

PART 6 · THE ADMIN CONSOLE

27	Opening the Admin console	30
28	Users	31
29	Groups	32
30	Model Access	33
31	Resolve model requests and trace access	34
32	Connections: response actions and connector handoff	35
33	Automations: preparing scheduled runs	35
34	SSO: single sign-on for your tenant	36
35	Analytics	37
36	Policies: service availability and memory governance	38
37	Data retention and conversation tags	39
38	Review feedback and reported issues	40
39	Audit	40
40	Alerts	41

PART 7 · THE PLATFORM OWNER CONSOLE

41	Opening the Platform owner console	43
42	First-run setup: from owner account to a working team	44
43	Providers	45
44	API Key Vault	46
45	Models: organization availability	46
46	Org Settings: users and role boundaries	47
47	Org Settings: single sign-on setup	48
48	Org Settings: SSO claims, MFA, and go-live	48
49	Org Settings: platform branding	49
50	Org Settings: policies and budgets	50
51	Org Settings: shared connectors and source credentials	52
52	Org Settings: search index readiness	54
53	Review and install a platform release	54
54	Organization retention and tagging	55
55	Analytics: runtime, feedback, activity, and usage	56
56	Automation readiness and governance	56
57	Audit: owner governance signals	57
58	Alerts: email delivery and platform-wide rules	58

Getting started

1. Welcome

What Aperture Chat is and how to read this guide.

Aperture Chat is your organization's workspace assistant. You can chat with approved AI models, draft documents and slide decks, search your organization's knowledge, and schedule recurring work. Review the sources and work trace when a response uses connected knowledge or tools.

This guide explains where to look and what to choose. Words in bold — like **Save version** — match the labels on screen. If a step says “hover”, rest your pointer on the item without clicking; controls that offer extra help show it nearby.

GUIDE	WHO IT IS FOR	WHAT IT COVERS
User Guide	Everyone	Chat , model access requests, personalization memory, documents and decks, account sync, agents, knowledge and tools, search and commands, appearance, and account security.
Administrator Guide	Workspace admins	Everything in the User Guide, plus the Admin console : users, groups, model access, response actions, SSO , analytics, policies and memory governance, audit, and alerts.
Platform Owner Guide	Platform owners	Everything in the other guides, plus setup readiness, providers and keys, organization policy, shared connectors, branding, releases, search indexing, and platform governance.

TIP

If a feature is not configured — a connector without credentials, a model that is not enabled — the screen reports that state. Ask your administrator about account or model access; shared connector settings and credentials are managed by the service team outside tenant administration.

2. Signing in

Reach the sign-in screen and get into your workspace.

- 1 Open the web address your organization gave you for Aperture Chat in a modern browser (Chrome, Edge, Safari, or Firefox).
- 2 On the sign-in screen, type your work email address in the field marked `you@company.com`.

- 3 If your organization uses single sign-on (**SSO**), the screen recognizes your email domain and offers **Continue with SSO**. Click it — you are sent to your organization's identity provider (for example Microsoft or Google), sign in there as usual, and return to Aperture Chat already signed in.
- 4 When both methods are offered, choose Organization **SSO** or Email & password. The password field appears in Email & password mode; enter your password and click Sign in. Pressing Enter submits the method you selected.

GOOD TO KNOW

If both **SSO** and a password are available, either works. If your organization enforces **SSO** for your email domain, password sign-in is blocked on purpose — use the **SSO** button.

GOOD TO KNOW

See a message you do not expect? The sign-in screen reports real problems plainly (for example, a domain that is not allowed for **SSO**). **Copy** the message and send it to your administrator.

3. Request access and finish your first sign-in

Submit a request, understand the administrator handoff, and enter the workspace.

- 1 On the sign-in screen, choose **Request access**. Enter your first name, last name, and work email, then submit the form.
- 2 The Request received screen confirms submission for your email. If this email is new to the workspace, an administrator reviews the request. This form does not email you or create a password.
- 3 Contact your administrator for an update. After approval, they provide the workspace address and confirm whether to use organization **SSO** or a temporary password.
- 4 Choose Back to sign in. Your submitted email remains filled in. Use the sign-in method your administrator arranged.
- 5 If you received a temporary password, sign in with it and complete Set a new password. Enter a new password of at least 12 characters twice. After saving, the app continues with a new authenticated session and revokes your older sessions.
- 6 Complete any required authenticator verification. **Once** in the workspace, select an available model and send a short first message. If no model is available, ask your administrator to check your account, group, and model access.

GOOD TO KNOW

Submitting another access request does not approve an account or recover an existing password. Trouble signing in explains the recovery path: contact your administrator for local password access, or use your organization identity provider for **SSO**. If sign-in settings cannot load, choose Retry connection.

4. Two-step verification and recovery codes

Set up an authenticator, complete required verification, and manage recovery codes.

To add an authenticator after signing in with a local password, open your account drawer, find **Security**, choose **Manage security**, and select **Set up authenticator**. Confirm your current password, add the QR code or setup key to your authenticator, acknowledge that you added it, and choose **Verify authenticator** with a current six-digit code. **Copy** the one-time recovery codes into secure storage, acknowledge that you saved them, and choose **Done**. Copying alone does not store them for you; finish verification and code storage before leaving the panel.

- 1 If sign-in asks you to set up an authenticator, choose **Begin authenticator setup**. Add the displayed QR code or setup secret to your authenticator app.
- 2 Confirm that you added the account, enter the current six-digit **Authenticator code**, and choose **Verify and enable MFA** before the setup expires.
- 3 Use **Copy recovery codes** and save the codes privately when they are shown. They are not saved to a file automatically. Each recovery code can be used once when your authenticator is unavailable; do not include these codes in screenshots or issue reports.
- 4 On later sign-ins, enter your current authenticator code. **Use a recovery code instead** switches to a single-use recovery code when necessary.
- 5 If you cancel or the challenge expires, return to sign-in and start again. A cancelled setup or failed verification does not finish sign-in.

When verification is enabled, the account card shows how many recovery codes remain. **Replace recovery codes** requires a fresh authenticator code or an unused recovery code; after replacement, all previous recovery codes stop working. Save the new set before leaving the dialog.

Turn off verification appears only when your organization's policy allows it. Confirm with an authenticator or recovery code; successful removal signs you out. Organization-required verification cannot be turned off here.

If you lose both the authenticator and all unused recovery codes, contact your administrator for the organization's verified recovery process. A new password does not remove the authenticator requirement. For identity-provider MFA, use that provider's recovery process.

TIP


Use the code currently shown for this workspace in your authenticator app. If it expires while you are typing, use the next code. Read the attempts, expiry, and any cooldown shown; wait for a cooldown to finish before trying again. If **Security** cannot load, choose Retry security settings. Keep setup keys and recovery codes out of shared documents, videos, and support attachments.

GOOD TO KNOW

SSO accounts manage voluntary authenticator setup in their identity-provider settings. Your identity provider may also require its own verification during **SSO**. Follow the provider's screen for that check; these instructions describe Aperture Chat's authenticator screens.

5. Finding your way around

The sidebar, the main area, and what each navigation button opens.



The chat workspace: navigation on the left, model selection above, and the message composer in the main area. Example workspace with synthetic data.

The screen has two areas: a sidebar on the left and the main work area on the right. The sidebar is how you move between the platform's views and your saved chats.

Your first visit

The Getting started card explains what is ready for your account and what comes next. Open quick-start guide opens user help. Administrators see Open admin guide and Manage access. Follow the setup action offered for your role when no model is available. A missing model is an access or setup task, not a successful chat connection.

Choose a guide, take the suggested action, or select I'll explore on my own (or Dismiss welcome) when you are ready. Simply loading the page does not dismiss this card. You can return to guides from **Help** or console **Documentation** later.

SIDEBAR BUTTON	WHAT IT OPENS
New chat	A fresh conversation with your workspace assistant. This is the default view.
Drafts	A document and slide-deck editor with an AI assistant beside it.
Agents/Automations	Reusable agent profiles, plus an Automations tab for scheduled runs.
Knowledge/Tools	Two tabs: Knowledge (searchable document collections) and Tools (connections, prompts, and skills).

Below the navigation buttons, a **Chats** row expands into your organized chat history — **Folders**, **Pinned**, and **Recent** sections, covered in the next section.

Resizing and collapsing the sidebar

- Click the small chevron (◀) near the top of the sidebar to collapse it and give your work more room. Click the chevron again (▶) to bring it back.
- Drag the sidebar's right edge to make it wider or narrower. Double-click that edge to collapse it.
- On a narrow window or a phone, the sidebar hides behind a menu button (three stacked lines, ≡) in the top-left corner. Tap it to open the menu; tap the dimmed area to close it.

Using the keyboard

Use Tab and Shift+Tab to move between controls and Enter or Space to activate the focused control. In a dialog, keyboard focus stays with that dialog. Use its Close or **Cancel** button, or Escape when available, to return to the control that opened it. Confirmations explain when an action is permanent before you commit it.

The bottom of the sidebar

- **Search** — opens **Search past work** for chats, documents and decks, agents, and indexed documents. It also offers commands. The keyboard shortcut is Ctrl+K (Windows) or ⌘K (Mac).
- **Help** — opens the guided walkthrough videos and this downloadable guide.
- **Dark mode / Light mode** — switches the appearance instantly (see “Light and dark mode” later in this guide).
- **Install app** — appears on phones and tablets, and adds the workspace to your home screen (see “Install the app on your phone”).
- Your account card — your initials or photo, name, and role. Click it to open the account drawer, including Personalization memory when your organization has enabled it.

Where the consoles live

Administrators open the **Admin console** from the account drawer: click your account card, expand the **Management** section, and choose **Admin console**. Regular users simply do not have a **Management**

section. The account drawer also holds **View as** role previews, **Usage this month**, and your **Archived chats** — all covered in “Your account”.

Returning to a view

The browser address follows the workspace view and console tab. Browser Back and Forward return to those views; opening a saved chat or draft link still checks your account's access. A link does not share private content or grant permission to another person.

6. Organizing chats: folders, pins, and recent

Keep the chats you care about easy to find.

Click the **Chats** row in the sidebar to expand your organized history. Inside it are three sections — **Folders**, **Pinned**, and **Recent** — and each one expands or collapses with a click.

Folders

- 1 Expand the **Folders** section and click the folder icon with a plus sign next to its heading.
- 2 Type a name for the folder — a client, a matter, a project — and click **Create**.
- 3 To file a chat into the folder, hover the chat's row in the sidebar and click the folder-plus icon that appears, then pick the folder (or create a new one right there).
- 4 Click a folder to expand or collapse the chats inside it.

GOOD TO KNOW

Deleting a folder does not delete its chats — they move back to the **Recent** list.

Pinned chats

- 1 Hover any chat row and click the pin icon that appears.
- 2 The chat moves to the **Pinned** section, where it stays until you unpin it (click the pin icon again).

Recent chats

- The **Recent** section lists your latest conversations. Click one to reopen it exactly where you left off.
- Hover a row to reveal three quick actions: add to folder (folder-plus icon), pin (pin icon), and archive (box icon).
- Archiving moves a chat out of the sidebar without deleting it — use it to tidy up. **Archived chats** are listed in your account drawer, where you can restore or permanently delete them.
- Click **View all chats** at the bottom of the list to browse your full history, including everything that no longer fits in the sidebar.

7. Searching everything

Find saved work, preview a result, or run an available command.

- 1 Click **Search** near the bottom of the sidebar, or press Ctrl+K (Windows) or ⌘K (Mac) from anywhere.
- 2 The **Search past work** box opens. Type a few words from what you remember: a chat title, message text, an agent or draft name, slide text, or text from an indexed document.
- 3 Results are grouped by kind. **Chat** results lead, matching text is highlighted, and archived conversations are labeled explicitly. Saved documents and decks open in their matching editor.
- 4 Hover over or keyboard-focus a supported chat or draft result to preview it. The chat row also offers folder, pin, and archive or restore actions when that conversation is available in the current workspace.
- 5 Click a result to open it, or use Up and Down followed by Enter. Escape closes search. Recently opened results are available when the query is empty.

Commands and saved work

Commands appear in a separate section from search results. Type > at the start of this search box to see all commands available to your role and permissions, then keep typing to narrow them. **Commands** can open workspace views and permitted actions; an unavailable feature does not become accessible through search. The same > symbol inside a chat message has a different purpose: choosing an automation for that message.

GOOD TO KNOW

Search includes only work your account can access. If the panel says Indexing your workspace, results may be incomplete until indexing finishes; try again after it completes. A missing result is not proof the work was deleted. Unsaved device-only changes may need saving to your account before server search can find them.

8. Your first chat

Send a message, pick a model, and control web search.

- 1 Click **New chat** in the sidebar. A fresh chat greets you by the time of day — “Good morning”, “Good afternoon”, or “Good evening” with your first name (late at night it asks if you are burning the midnight oil) — above the tagline “Your approved models, your sources, your guardrails — ask Aperture Chat anything.” (The name reflects your organization's branding.)
- 2 Click into the message box (the composer) — it reads **Ask anything...** — and type your question.
- 3 Explore an idea, Compare options, and **Draft** a message offer starting points. Selecting one fills an editable prompt; review it before sending.
- 4 Press Enter to send. Press **Shift+Enter** when you want a new line without sending.

Picking the model

The model selector sits in the top bar of every chat. Click it to choose which AI model answers this conversation. Each chat remembers its own choice, and the list only ever shows models your workspace has approved for you.

Use the arrow keys, Home, End, or type a model name to move through available models; Enter selects the highlighted model. Each row has its own star: Set [model name] as default model selects that model and remembers it for new chats. The filled star marks the single current default. Selecting a row without its star changes this conversation's model without changing the default. **Drafts** has a separate default drafting model, set with the star beside a row in its own selector.

Web search and the active-tools chip

- When the selected model supports it, public web search starts turned on. A chip at the bottom of the composer shows what is active: with one tool on it names it — for example **Web search** — and with several on it reads **Tools** with a count.
- Click the small x on that chip to clear the next message's active tools: **Knowledge**, **Web**, and **Agent**, selected MCP connections, and queued automations. Every new chat starts with web search on again when the model supports it.
- Replies that used the web come back with citations you can click, and every source is listed in the session details panel (covered below).

Sending

- The paper-plane button sends your message with the current settings.

- The chevron (v) beside it opens send options — per-reply switches for **Knowledge**, **Web**, and **Agent**, plus the **Reasoning** level (covered in their own section).

Dictate instead of typing

- Click the microphone button in the send controls at the bottom-right of the composer to start dictating. A live waveform shows that the platform is hearing you.
- Click the button again — it becomes a stop square — and your words are transcribed and inserted into the composer as editable text. Nothing sends until you press Enter.
- If the microphone is blocked or no speech was heard, the composer says so plainly so you can fix the problem and try again. The **Drafts** workspace has the same dictation button for drafting instructions.

9. Understand model availability and request access

Read the server's reason, ask for a model when permitted, and follow its request status.

- 1 Open the chat model selector and choose **Why isn't a model listed?** If no model is usable, the No models available control opens the same explanation.
- 2 **Models in your organization** separates **Usable now** from **Not available to you**. Read the reason on the model row: **Locked** describes an access restriction; **Provider offline** means you already have access but the provider needs attention.
- 3 Choose **Request access** only where it is enabled. The request goes to an administrator for review; sending it does not grant immediate access.
- 4 While a request is pending, the row shows its date and **Withdraw request**. Use **Refresh model access** after your administrator resolves it, then select the usable model in the chat picker.

GOOD TO KNOW

Your organization may show only models you can already use. Its browsing policy controls whether you can inspect other enabled models and request access. Some restrictions cannot be lifted by a request; ask your administrator about the stated reason. An offline provider needs a platform-owner repair, not another group-access request.

10. Reading and acting on replies

The work trace, citations, and the actions under every message.

The work trace

- While a reply is being produced, a single status line shows the current step, a live timer, and how many steps remain (for example “4s · step 2 of 5”).

- Click that line at any time — while running or after — to expand the full trace: which model was routed, what context was prepared, whether web search ran, and how the answer was finalized.
- While a reply is still streaming, a stop button appears — its label is **Stop this response**. Clicking it halts the reply and keeps the conversation.

Actions under each response

- **Copy** — copies the response as clean, readable text.
- Thumbs up / thumbs down — sends feedback your administrators can review.
- **Branch** — starts a new chat that continues from that response, leaving the original untouched.
- **Regenerate response** — asks the model to answer the same prompt again. Version arrows appear on the response with a counter (for example “2 / 3”) so you can flip between the versions it has produced.
- **Transfer to Drafts** — turns the response into an editable document in the **Drafts** workspace, where you can format, version, and export it.

Actions under your own messages

- **Edit message** — reopens your prompt for editing. Sending the edit removes the later replies and asks again from that point.
- **Load prompt in new chat** — copies the prompt and its attachments into a fresh chat's composer, ready to adjust and resend.

Diagrams in replies

When a reply contains a diagram (a flowchart, sequence, or similar), it renders as a real figure, not a code block. The figure's toolbar offers **Copy** for the diagram source, **PNG** and **SVG** to download it as an image file, and **Code** to flip between the drawing and its source text.

Generated images

- When the selected model can generate images, just ask for one — the finished images appear directly in the reply.
- Click **Download** under an image to save it as a file.

Citations and token usage

- When a reply used the web or your knowledge bases, numbered citations appear with it. Click a citation to open the original source.
- A small token count appears under a reply only when the model provider actually reported one — the platform never shows an estimate dressed up as a real number.

11. Composer symbol shortcuts

Five characters that insert prompts, agents, knowledge, skills, and automations.

Start a word in the composer with one of five symbols and a menu opens with matching items. You can also choose **Send options** → **Resources** to browse without memorizing a symbol. The **Reply settings** tab's **MCP connections and resources** button opens the same resource browser.

TYPE	WHAT IT OPENS
/	Your saved prompts and enabled MCP connections.
@	Agent profiles — the reply routes through the one you pick.
#	Knowledge bases and the files inside them, to ground your question.
\$	Saved skill files, inserted into the message.
>	Automations — the one you pick queues and runs when you send.

- 1 Type the symbol, then keep typing to filter the list.
- 2 Use the ↑ and ↓ arrow keys to move through the menu.
- 3 Press Enter to insert the highlighted item, Tab to complete it, or Escape to dismiss the menu.

Browse resources

On the **Resources** tab, use **Find a resource** to search by name, or select All, **Knowledge** bases, Files in knowledge sources, MCP connections, **Prompts**, **Agents**, Skill files, or Automations. A selected connection or automation shows a check; choose it again to remove it. Read an empty or failed-loading state before retrying. The **Type shortcuts in chat** reference explains each symbol below the results.

TIP

The expanded composer focuses on editing a longer message. Return to the normal composer to browse resources, adjust send options, and send with the visible controls.

12. Attaching files, web pages, and connected sources

Upload from your computer, fetch a page by link, or pull from your organization's drives.

- 1 Click the paperclip button at the bottom-left of the composer.
- 2 Click **Upload from computer** to attach files from this device. The model reads them as context for your next message.
- 3 Or click **Web page by link** and paste one or more web addresses — the platform fetches up to 3 public pages and attaches them as cited sources for this message.

- 4 Attached items appear as chips above the composer. Click the x on a chip to remove one before sending.

Attach from source

Below those options, under the **Attach from source** heading, the menu lists your organization's workspace sources: Google Drive, OneDrive, SharePoint, Box, and iManage. Availability depends on the service team's connector setup and your own source-account access.

Connecting your own account

- Cloud sources read from your own account. The first time you pick one, the picker may ask you to connect it — click the Connect button for that source and a sign-in window opens.
- **Approve** read access in that window. The file list refreshes on its own once access is granted, and only you can see files from your account.
- The picker shows files at the top level of the drive or the folder configured for the source, within your account's permissions.

GOOD TO KNOW

If a source needs shared configuration, ask your administrator to coordinate with the service team. Shared connector administration is managed outside tenant administration. Your own Connect action in the attach menu remains separate and grants access only through your source account.

13. Send options: Knowledge, Web, Agent, and Reasoning

Control what the model may use, one reply at a time.

- 1 Click the chevron (v) beside the paper-plane send button.
- 2 The menu opens with **Send now** at the top — that simply sends with the current settings.
- 3 Below it are three switches you can turn on or off for the next reply, and a **Reasoning** slider at the bottom.

SWITCH	WHAT IT DOES WHEN ON
Knowledge	The reply searches your enabled knowledge bases and cites what it finds, with links back to the sources. A picker lets you choose which bases to use.
Web	The reply uses public web search. Results come back as clickable citations, and the sources are listed in session details. Available when the selected model supports it.
Agent	The reply uses the selected agent profile and its permitted tools. Turning this on can choose the first available profile if none is selected; confirm the Agent profile field before sending, or choose a profile with @.

Reasoning level

The **Reasoning** group at the bottom of the menu is a three-position slider from **Fast** to **Smart**. **Fast** favors quicker answers; **Smart** makes the model think longer for more detailed output; the middle is a balanced default. The slider is only active when the selected model supports reasoning levels — otherwise it is disabled and says so.

Whenever tools are on, the active-tools chip appears next to the paperclip: the single tool's name, or **Tools** with a count. Click its x to turn off **Knowledge**, **Web**, and **Agent** and clear selected MCP connections and queued automations for the next message.

On the **Reply settings** tab, **Stream replies** controls whether text appears as the model produces it or only after the reply finishes. The **Resources** tab opens the searchable resource picker; **MCP connections and resources** also opens it from **Reply settings**. These controls change how the next reply is prepared or displayed; they do not supply missing provider or connector access.

14. Session details and sources

Real token usage, active tools, and every source this chat gathered.


- 1 Click the round information button (i) at the right end of the chat's top bar.
 - 2 The session details panel opens for the current chat.
- **Current chat** and **Model** — the user-message and response counts, and which model is active.
 - **Tokens used** — shown only when the model provider reported real numbers. Otherwise the panel says “Not reported by the provider”.
 - **Tools** — exactly what is switched on for your next message: web search, knowledge, agent mode, plus any automations or MCP connections you added in the composer. “Off” means nothing is active.
 - **Sources gathered** — every web source (as clickable links) and workspace citation collected during this conversation.

The **Context window** meter describes how much context this conversation uses. For older conversations without provider counts it can show $\approx$ and estimated from message length. Treat that as an estimate; **Tokens used** remains based on actual provider reports. When the context window fills, start a new chat and include the details it needs to continue reliably.

Drafts, decks, and agents

15. Drafting documents

Generate, edit, version, and export documents with the drafting assistant.



Drafts combines the document editor with a separate assistant and document controls. Example workspace with synthetic data.

Click **Drafts** in the sidebar. The view pairs a full document editor with an AI assistant rail on the left that writes into the editor for you. A **Draft format** switch at the top of the editor has two buttons — **Document** and **Deck**. This section covers document mode; “Building a slide deck” covers deck mode.

GOOD TO KNOW

If the model selector says No models connected, AI drafting, AI rename, inline AI edits, and AI deck prompts are unavailable. You can still edit manually, import, save, use document history, and export. Follow the setup guidance for your role or ask your administrator for access to a working model.

Generating a draft

- 1 Look at the assistant rail first: it shows which context sources are on — workspace files, web, and templates — before you generate anything.
- 2 In the assistant's message box, describe the document you need: a memo, an engagement letter, a summary. Be specific about audience and tone if it matters.

- 3 The assistant writes the document directly into the editor. You can also send any chat response here with **Transfer to Drafts**.
- 4 Use the model selector in the toolbar area to choose which approved model writes and edits this draft.

Editing like a document

The main toolbar keeps undo and redo, block style, bold, italic, underline, and inline AI editing close at hand. **Text** opens font, size, advanced styles, color, highlighting, and **Clear formatting**. **Paragraph** opens alignment, lists, and quotes. **More** opens **Copy**, **AI edit trail**, and word count. Select the intended text before opening a panel; Escape returns to the editor.

- The plus (+) **Insert** button opens links and citations alongside **Web image**, **Chart**, **Table**, **Divider**, and **Page break**. A manual page break starts a new page at your cursor; the editor also paginates longer documents automatically.
- **Inline AI edit** rewrites just the text you have selected: select a passage, click the AI pen, type an instruction, and review the change.

Revise the document you already have

Describe the change in the assistant: shorten a section, change tone, apply a template, or transform the existing paper. Revision requests work from the current document. Ask explicitly to start over, replace, or clear it when that is your intention. The editor keeps instructions separate from the deliverable; review the resulting document and use Undo or a saved version when a change needs correction.

For an existing paper, **Assistant settings** → **Apply MLA layout** applies double-spaced 12-point Times New Roman, a student heading, centered title, body and reference indents, and a separate Works Cited page without asking a model. It preserves the paper's text and supports Undo. Fill missing student details and verify quotations and sources yourself; formatting is not citation verification.

The page navigator shows the current sheet and total pages outside the document text. Older drafts are paginated when opened, and manual edits reflow after you leave the editor. Page labels do not become part of the exported body.

Versions and comparing

- Click **Save version** whenever the draft reaches a good state. You can restore an earlier version at any time from the draft's history.
- **Compare versions** opens a read-only visual redline of two saved versions, so you can see exactly what changed between them. It needs two genuinely different saved versions before it activates.
- Open **Document history** to browse saved documents and decks. Hover or keyboard-focus an entry for a preview. **Archive** moves finished work into **Archived**, where **Unarchive** restores it; **Delete** requires confirmation and permanently removes the saved item and revisions.
- Read the save status: account-backed documents and decks can be reopened after sign-in on another device. A pending or failed save stays a working copy in this browser. If the account version changed elsewhere, resolve the reported conflict before treating the save as complete.

Leaving with unsaved changes

When you leave an edited draft through workspace navigation, open another chat or draft, or choose **Sign out**, the unsaved-changes dialog lets you Keep editing, Save copy and continue, or Discard and continue. The saved copy goes into local document history in this browser. If browser storage fails, the app keeps you in the draft so you can recover it. A forced security sign-out can still interrupt work; save important changes regularly.

Exporting

- Click **Export** to open the export panel. In document mode it offers a **Word document** (an editable Word file with preview page breaks and embedded images), **Markdown** (best for plain text or web publishing), and **Print / Save as PDF**, which opens your browser's print dialog with the saved version — choose “Save as PDF” there to keep a PDF copy.
- Exports always run from a saved version. If you have unsaved edits, the panel says **Save your edits first** and offers a one-click **Save version and export** (or **Save version and print**) button.
- When your browser supports it, choose **Choose a location** to select a destination, or **Browser downloads** to use its download folder. Word and print retain document formatting; **Markdown** preserves the text without its typography. Exporting a file does not by itself upload pending work to your account.

TIP

On a narrow window, the assistant rail becomes a slide-out drawer so the editor gets the full screen. Open it with the floating pen button on the left edge; press Escape or click the dimmed area to close it.

16. Building a slide deck

Turn a draft into slides, edit them, brand them, and export real PowerPoint.

Deck mode turns the **Drafts** workspace into a slide editor with a real PowerPoint export. Click the **Deck** button on the **Draft format** switch. If the draft already has document content, a dialog asks “Turn this draft into slides?” — choose **Start a blank deck** to keep the document untouched, or **Convert into slides** to split its headings and text into slides. Your document versions are kept either way.

The filmstrip

- Slide thumbnails run along the side. Click one to select it, or drag it to reorder the deck.
- Each thumbnail has its own actions: move-up and move-down arrows, a duplicate button, and a delete button (undo restores a deleted slide).
- **Add slide** at the end of the filmstrip opens the layout menu for a new slide.

Slide layouts

Seven layout previews are available: **Title**, **Title + bullets**, **Two columns**, **Image + caption**, **Quote**, **Section**, and **Closing**. Open the layout gallery to scroll through them and switch the selected slide's arrangement.

The deck color controls offer five palettes; palette changes keep slide text, media, and layout and can be undone.

Editing slide text

- Click any text region on a slide to edit it in place. The deck toolbar carries the same rich formatting controls as documents — undo and redo, bold, italics, and more.
- Inside bullet lists, press Tab to indent a bullet one level and Shift+Tab to outdent it.
- **Edit selection with AI** rewrites highlighted slide text: select the text, click the AI pen, type what should change, and choose **Replace highlight**.

Starter templates and your brand template

- The templates drawer (the slides icon in the assistant rail) offers five starter structures: **Pitch deck**, **Quarterly review**, **Project kickoff**, **Client proposal**, and **Training session**. Pick one and start it to get scaffold slides you replace.
- A **Use templates in chat** toggle (off by default) lets the selected template guide the deck assistant when it drafts for you.
- **Upload brand template** accepts your firm's own file — “.pptx or .potx — brand colors, fonts, logo, and every slide's text are extracted.” The extracted theme is stored only on this device.
- **Once** a brand theme is loaded, **Apply to deck** restyles the current deck with its colors, fonts, and logo, and a “Load all N slides” button can replace the deck with every slide from the uploaded template.

Backgrounds and AI images

- The background button in the toolbar opens a per-slide menu: **Upload background...** (a PNG, JPEG, or WebP from this device), **Use on every slide**, **Remove from this slide**, and **Clear from all slides**.
- **Generate AI slide image** writes an image prompt for you, prefilled from the slide's content; press **Generate image** and the result becomes that slide's background. This needs an image-generation model enabled for your workspace — the button says so honestly when none is.
- In the assistant rail, **Toggle AI slide images** makes the assistant also generate an image for every slide whenever it drafts a deck.

Speaker notes

Every slide has a **Speaker notes** field below the stage. **Notes** export with the deck: they land in each slide's notes pane in the PowerPoint file, they appear in the **Markdown outline**, and they follow you into presentation mode.

Presenting from Aperture

- 1 Click the **Present deck** button at the end of the deck toolbar (the monitor icon). The deck opens full screen.
- 2 Advance by clicking anywhere on the slide, or with the arrow keys; move back with the left arrow.

- 3 Keep your speaker notes open in the panel beneath the slide and read from them as you go — the **Notes** button (or the N key) shows and hides the panel.
- 4 Leave the presentation with Escape or the **Exit** button in the top corner.

Saving and exporting

- **Save version** stores decks to your account with revision history, like documents. Read the status to confirm the save reached the server. A local working copy or pending save remains only in this browser until synchronization succeeds; use the **Only on this device** reminder to find unfinished saves.
- **Document history** includes decks. Preview, reopen, archive, unarchive, and confirmed deletion use the same account controls as documents. A concurrent account edit can prevent a stale save or deletion; read the conflict message before retrying.
- Click **Export** to open the **Export deck** panel: **PowerPoint deck** produces an editable .pptx that mirrors the slides on screen — speaker notes included in each slide's notes pane — and **Markdown outline** exports slide titles, bullets, and speaker notes as text.
- Limits: a deck holds up to 100 slides and 8 MB of content. The editor tells you plainly if a deck exceeds them.

17. Recover work kept only on this device

Distinguish a browser working copy from a completed account save.

The sidebar's **Only on this device** reminder appears when work has changes kept in this browser that have not reached your account. It can include chats whose last save failed, documents or decks with unsent changes, and older history from before account sync.

- 1 Open the reminder to review the named items. For chats whose last save failed, choose **Retry all** and check the resulting save state.
- 2 For documents and decks, choose **Open Drafts**, open the item, and save it to your account. Review any load, quota, connection, or conflict error before leaving the browser.
- 3 Older browser history remains separate until you explicitly import each item from **Drafts** → **Draft history** with **Import to my account**. Confirm the item belongs in this account before importing.
- 4 **Archive** finished drafts when you want them out of the active list. **Account** archives remain recoverable; archiving a browser-only item does not prove it has been uploaded.

GOOD TO KNOW

Clear list dismisses these reminders without deleting or uploading work. New edits may appear again. **Hide this reminder** turns off this account's sidebar notice in this browser; save status remains in **Drafts** and chats. Downloading an export also leaves pending account saves to be resolved.

18. Agent profiles

Reusable bundles of model, prompts, knowledge, and tools.

An agent profile bundles a model route, meta prompts, knowledge bases, and MCP tools into one reusable configuration. A usable profile needs an approved, enabled model backed by a working provider. A saved profile marked not connected is a configuration example, not proof that it can run.

Creating an agent

- 1 If your account has agent authoring permission, click **Agents**/Automations in the sidebar, then **New Agent**.
- 2 Give it a name and a short description of what it is for.
- 3 Work through the editor tabs: **Profile** (name, model, and description), **Knowledge** (assign knowledge bases), **Tools** (select MCP tools), **Prompts & Skills** (attach system prompts and skill files), **Access** (who can use it), and **Hermes** (the optional learning companion).
- 4 Click the save button. Standard users create private profiles; administrators control group sharing and organization publishing. Check model readiness and access before trying the profile in chat.

Using an agent in chat

- Type @ in the composer and pick an available profile to use its configured model route, instructions, knowledge, and tools.
- The **Agent** switch in **Send options** can choose the first available profile when none is selected. Check **Agent profile** and select the intended configuration before sending.
- If no profiles appear, ask your administrator to check that a ready profile is available to your account.

From the **Agents** view you can also open a profile to edit it, jump straight into a chat with it, or delete it. The Automations tab at the top of this view holds your scheduled runs — covered in “Scheduled automations”.

Knowledge/Tools and automations

19. Knowledge bases

Give chats grounded, citable access to your documents, web pages, and APIs.

Knowledge lives in the **Library**: click **Knowledge/Tools** in the sidebar, then the **Knowledge** tab. The **Knowledge Bases** panel lists every collection with its status, security posture, and whether it is enabled.

- 1 Click **Add Knowledge Base** and name the collection.
- 2 Choose what feeds it. A knowledge base can index documents you upload, a public web page (**Add a web link** — enter the address and an optional note), or an API (**Connect an API** — enter the endpoint and its authentication).
- 3 **Create** the base, then check its row in the table: it shows the status, the security posture, and whether the base is enabled.

GOOD TO KNOW

Only enabled knowledge bases can be searched from chat, and access control is enforced per source — people only ever retrieve what they are allowed to see.

Using knowledge in a conversation

- Type # in the composer to reference a knowledge source or choose from Files in knowledge sources. A file choice references its name in your prompt and searches that file's knowledge source; the row says which source will be searched.
- Or turn on the **Knowledge** switch in send options to let the reply search your enabled bases.
- Replies grounded in knowledge return citations that link straight back to the source documents.

GOOD TO KNOW

The # shortcut menu shows Loading files... while retrieving file choices. If Some files could not be loaded appears, use Retry in that menu. Files from sources that loaded successfully remain usable. If access was removed or a source needs configuration, ask your administrator to resolve it.

20. The Tools library

MCP connections, saved prompts, and skill files inside the Library.

Click **Knowledge/Tools** in the sidebar, then the **Tools** tab. The **Tools Library and Connectors** panel has three sections:

TAB	WHAT LIVES THERE	WHERE IT APPEARS IN CHAT
Connections	MCP tools and connections available to models and agents.	Type / in the composer.
Prompts	Saved prompt templates, including <code>{{variables}}</code> you fill in when using them.	Type / in the composer.
Skills	Skill files — reusable instructions for recurring workflows.	Type \$ in the composer.

- Click **New Prompt** or **New Skill** to create one: give it a name, a category, a description, and the content itself, then save.
- Click any existing card to review or edit it; the trash icon deletes it permanently.
- Every tool row shows its real status — draft, approval required, or enabled — so you always know whether a connection is actually live before relying on it.

21. Scheduled automations

Recurring chat or drafting runs, as one model call or a multi-step chain.

Automations are saved workflows that run a chat or drafting task on a schedule. They are useful for recurring digests, weekly status checks, report refreshes, policy reviews, and any workflow where the same model chain should run the same way each time. Enabled schedules run automatically in the background (times are UTC), and each scheduled run delivers its output as a new chat thread in your sidebar.

Creating an automation

- 1 Click **Agents/Automations** in the sidebar, then the Automations tab, then **New automation**.
- 2 Name it — for example “Monday client digest”.
- 3 Choose what it runs against: chat or draft.
- 4 Pick a trigger: **Daily** (a time each day), **Weekly** (a day and time), **Once** (a specific date and time), or **Cron expression** (such as `0 9 * * 1` for 9:00 every Monday, in UTC).
- 5 Write the initial input — what the first step should work on.
- 6 Build the model chain: each step has a model and an instruction, and each step's output feeds into the next step. Use **Add step** for multi-step chains and the x to remove a step.
- 7 Click **Save automation**.

What to check before saving

- **Models** — every step must use a model you can actually access. If a model is missing, your admin may need to approve it for your group.

- Context — write the initial input as if the automation will run without you watching. Include the matter, audience, date range, source expectations, and desired output format.
- Cadence — use **Once** for a single future run, **Daily** or **Weekly** for regular recurring work, or a reviewed **Cron expression** for another schedule. All times are UTC; check the intended local-time equivalent before saving.
- Ownership — name the automation so another person can tell what it does later.

Running and managing

- Each automation card shows its schedule, its steps, and its last run with an honest result — succeeded or failed, with the reason.
- **Run now** executes the whole chain immediately and shows a transcript of every step's output.
- The toggle pauses or resumes the schedule; the pencil edits it; the trash deletes it.
- You can also queue an automation into any chat by typing > in the composer — it runs when you send.
- The list shows the automations you created. Organization administrators can review the automations in their workspace.

TIP

Always use **Run now** once after creating or editing an automation. If the run fails, the card shows the reason so you can fix the model choice, input, or instructions before depending on the saved schedule.

Appearance, your account, and help

22. Light and dark mode

Switch appearance now or schedule it for this browser.


- 1 Find the appearance row near the bottom of the sidebar — it reads **Dark mode** with a moon icon in light mode, and **Light mode** with a sun icon in dark mode.
- 2 Click it. The entire platform switches immediately — no reload, nothing to save.
- 3 Click it again to switch back.

Schedule light and dark mode

Choose the clock button beside the appearance row to open **Theme schedule**. Enable **Switch automatically**, set different **Light mode at** and **Dark mode at** times, then choose **Save schedule**. The schedule uses this device's local time every day and is saved in this browser. You can still switch manually until the next scheduled change; disable the switch to stop automatic changes.

23. Your account

Your profile, personalization memory, password, role previews, usage, archives, and signing out.



Your account includes profile settings, password controls, and two-step verification. Example workspace with synthetic data.

Click your account card (your initials or photo) at the very bottom of the sidebar. The account drawer opens with everything about you in one place.

- **Profile** — click the card with your name and the pencil icon to edit your display name, firm, website, bio, phone number, and photo (upload an image up to 5 MB, or paste a URL). Click **Save profile** when done, or **Cancel** to discard.
- **Management** — organization administrators see this section; expanding it opens the **Admin console**. Regular users do not have it.
- **Personalization memory** — opens your private memory manager when the feature is enabled. Use it to control, review, add, correct, pin, forget, or clear what the assistant remembers about you.
- **Password** — accounts that sign in with a password can change it here (click the pencil, enter the current password, then the new one twice — at least 12 characters). The app continues with a new authenticated session after saving and revokes previous sessions. Accounts that sign in through **SSO** manage their password with the **SSO** provider instead, and the panel says so.
- **Security** — choose **Manage security** to set up an authenticator for a local account, review the remaining recovery-code count, or replace codes after verifying your identity. Turning verification off is available only when organization policy permits it and signs you out. **SSO** accounts follow their identity-provider settings for voluntary enrollment.

- Role and organization — your role, your organization, any personal token caps that apply to you, and how you sign in.
- **View as** — if your role allows it, preview the workspace exactly as a lower role would see it. A note reminds you which role you are previewing; switch back the same way.
- **Usage this month** — your prompts, responses, and estimated tokens.
- **Archived chats** — every chat you have archived, with buttons to restore each one to your sidebar or delete it permanently.
- **Sign out** — ends your session and returns to the sign-in screen.

24. Personalization memory

Save and recall private preferences in plain English across chats and sessions.

Personalization memory belongs to your account, not to one conversation. When service policy and your organization administrator allow it, administrators and regular users can each build their own private memory. A saved item can influence a new chat immediately and still be available after you sign out and return in a later session.

Save a memory in ordinary English

- 1 In any chat, state what should persist in natural language — for example: “Remember that I prefer a short summary before the detail,” “Please remember that I work in commercial litigation,” or “Remember my project codename is Silver Horizon.”
- 2 Send the message. The assistant confirms the memory after it is saved; you do not need a slash command or special syntax.
- 3 Start a new chat, or return in a later sign-in session, and ask a natural question such as “What do you remember about me?”, “What are my writing preferences?”, or “What is my project codename?” The same account memory is available across those sessions.

GOOD TO KNOW

Memory is different from the context window shown in Session details. The context window describes how much of the current chat the selected model can hold at once; personalization memory is durable account context that follows you into other chats and sessions.

Open and control the memory manager

- 1 Click your account card — your name and role at the bottom-left of the sidebar.
- 2 Click Personalization memory. The dialog titled “What the assistant remembers about you” opens.
- 3 Use memory in my chats controls whether saved items are applied to answers. Learn from my conversations controls whether the assistant may notice durable preferences automatically. Turning a switch off does not expose or silently erase existing items.

- 4 To add something directly, type it under Add something you want remembered, choose Standing instructions, Preferences, About you, Your work, or Other details, and click Add.
- 5 Review the list below. Click a memory's wording to correct it; pin an important item; use Forget to remove one item; or choose Forget everything to clear the list.

GOOD TO KNOW

Only you can read your memory content. Administrators can see counts and purge memories for compliance, but the platform never shows them what an individual memory says.

TIP

If Personalization memory is missing from your account drawer, the feature is not currently available at the platform, organization, or group level. Ask your administrator; refreshing the page cannot override that policy.

CAREFUL

Do not use memory as a password or secret vault. Credentials, keys, and sensitive identifiers are rejected on purpose.

25. Install the app on your phone

Put your workspace on the home screen with its own icon and name.

On a phone or tablet, tap **Install app** near the bottom of the sidebar. The “Add ... to your home screen” dialog previews your organization's app name and icon.

- On Android, choose **Install app** if offered, then follow the native install sheet. Otherwise open your browser's menu and choose “Add to Home screen” or “**Install app**”.
- On iPhone and iPad, tap the browser's **Share** button, scroll the share sheet, choose “Add to Home Screen”, then tap Add.
- Open the new home-screen icon to return to the same workspace in a full-screen app.

26. Getting help


Walkthrough videos and this guide, always one click away.

- Click **Help** at the bottom of the sidebar for guided videos, captions, transcripts, and this PDF. The guide also covers additional procedures.
- Administrators can also open **Documentation** inside the console for role-specific videos and a printable guide.
- Choose Report a problem in **Help** to describe a problem, the affected screen, and the expected result. **Share** only permitted details with your administrators; omit passwords, keys, and recovery codes.

The Admin console

27. Opening the Admin console

Where the console lives and what its nine tenant-governance tabs control.



The Admin console opens the workspace controls available to your administrator account. Example workspace with synthetic data.

- 1 Click your account card at the bottom of the sidebar to open the account drawer.
- 2 Expand the **Management** section and click **Admin console**. (This section appears only for workspace administrators.)
- 3 The console opens with nine tabs across the top: **Users**, **Groups**, **Model Access**, **Connections**, **SSO**, **Analytics**, **Policies**, **Audit**, and **Alerts**. **Policies** is always present between **Analytics** and **Audit**; service-wide availability determines which organization controls are active inside it.

Everything you change here writes through the admin API immediately — changes persist across refreshes and restarts, and every action lands in the tenant audit trail. Status messages under the header tell you honestly whether an action synced or failed.

TIP

The **Documentation** button at the top of the console opens narrated video walkthroughs of every tab, plus this guide as a PDF.

28. Users

Create tenant accounts, manage passwords, and remove leavers safely.

The **Users** tab lists the organization accounts you administer, with each person's role, groups, sign-in method, and status.

Adding a user

- 1 Click **Add User**.
- 2 Enter the person's name and work email, then pick from the roles allowed by current organization policy: User, Power User, Auditor, and **Agent** Approver, plus Admin when administrator creation is available.
- 3 Optionally pick a starting group, then create the account.

Reading and filtering the list

- Use the group filter to work one team at a time.
- Watch the auth and status columns: users who arrived through **SSO** stay **Pending** until you assign them to a group — that is the signal they are waiting on you.

Reviewing access requests

- 1 Open **Users** and review the **Access requests** queue. Check the requester's name, email, and request time before choosing their access level under **Approve** as.
- 2 Choose User for standard group-based access, Temp User for the restricted Luna-only 30,000–reported-token allowance, or Admin only when that option is permitted and appropriate. Click **Approve**, or **Decline** to reject the request.
- 3 After approval, use Finish sign-in setup for the approved person. **Share** the workspace's sign-in address and confirm the sign-in method. Approval does not send an email.
- 4 For organization **SSO**, confirm the account exists with the identity provider. For email and password, choose Set temporary password, save a generated or entered password, and share it securely. The person must choose their own password at first sign-in.
- 5 Check the person's group membership and model access, then ask them to sign in and complete a first message. Approval alone does not establish every resource permission.

Per-row actions

Every row has an Actions column with account controls:

- **Access** — opens the read-only **Model access trace** for this person, with their groups, each model's status, and the gates that explain it. See Resolve model requests and trace access for the full workflow.

- **Password** — opens the password dialog for that person: type a password or click **Generate** for a strong random one, optionally mark it a **Temporary password** (they must choose their own at first sign-in), and click **Set password**. The password is shown only here — share it over a safe channel.
- **Deactivate / Activate** — ends or restores sign-in access immediately, keeping the account's audit history intact. You can also select several accounts with their checkboxes and use the **Deactivate** button at the top.
- **Delete** (trash icon) — permanently deletes an eligible account and its chat history. The tooltip spells it out per person: “Permanently delete ... and their chat history”. Administrator-account actions that are unavailable under current service policy remain disabled with an explanation.

GOOD TO KNOW

Password reset and authenticator recovery are separate. The current **Users** screen has no authenticator-reset action. If someone has lost their authenticator and recovery codes, verify their identity and follow the organization's authorized recovery process; setting a temporary password alone does not clear MFA. Identity-provider recovery remains with the **SSO** provider.

GOOD TO KNOW

Administrative continuity rules are enforced by the service. When an account action would violate them, the console blocks the action and explains that it is restricted by administrative continuity policy.

29. Groups

Groups carry permissions and model access — everything flows through them.

Groups are how access flows in the tenant: model grants, knowledge access controls, and permissions all attach to groups, not to individual people. New **SSO** users stay **Pending** until you put them in one.

- 1 Click **Add Group** and name it for the team or matter it represents.
- 2 Select the group to open its editor. It has three tabs: **Users**, **Permissions**, and **Import**.
- 3 On the **Users** tab, tick the people who belong in it.
- 4 On the **Import** tab — headed “**Import Users to**” the group's name — paste addresses into the **User emails** box and click **Add users to group** to bulk-add them.
- 5 On the **Permissions** tab, switch each capability on or off.

Runtime access and authoring permissions

TOGGLE	WHAT IT GRANTS
Can use chat	Start conversations and use assigned models.
Can use knowledge	Query approved knowledge bases.
Can use agents	Run approved agent workspaces.
Can use tools	Invoke enabled tools and MCP actions.
Can use API	Create personal keys for approved models when service policy allows API access.
Can use Hermes companion	Build and run agent profiles with the Hermes learning companion. Off until approved.
Can build agents	Create private agent profiles when service policy permits. Organization publishing stays admin-only.
Can build knowledge bases	Create private knowledge bases. Group sharing and organization management stay admin-only.
Can build tools	Create private tools. Group sharing, stdio commands, and organization management stay admin-only.
Can use memory	Allow personal preferences to be learned and reused within the workspace memory policy.

CAREFUL

Deleting a group removes its members' access that flowed through it. Check what the group grants before deleting.

30. Model Access

Decide which synced models users can see, and which groups can use them.

Model Access starts from the catalog available to your organization. You decide what portion of it your users actually see.

- 1 Click **Sync models** to refresh the catalog whenever service availability changes.
- 2 Read the counters: the All, Enabled, and Disabled status counters carry live totals alongside how many groups are available for scoping.

- 3 Use the search box and the funnel filters in the column headers to cut through the catalog: check off providers, check off model labs, or type text to match runtime routes, and Clear filter resets one. The table has seven columns: **Model**, **Provider**, **User Access**, **Groups**, **Filters**, **Knowledge**, and **Tools**. Everything starts hidden until you decide otherwise.
- 4 Flip the **User Access** toggle on a row to make that model visible, and use **Choose groups** to narrow it to specific teams.
- 5 Use the **Filters** column to attach per-model content filters, and the **Knowledge** and **Tools** columns to scope what each model may search and call.

GOOD TO KNOW

Because model access flows through groups, a tenant with no groups yet sees the guard up front — “**Create** a group before enabling models.” — and the access switches stay off until one exists.

GOOD TO KNOW

Newly synced models arrive disabled until service availability and organization policy permit them. If a model you expect is missing, it is not currently available to this organization.

31. Resolve model requests and trace access

Review requests through groups and diagnose each user's actual model access.

Review a pending request

- 1 **Open Model Access** and review **Access requests**. **Refresh** reloads pending requests for your tenant. Each row names the requester and model and shows the server's access reason.
- 2 Choose **Grant through group** carefully. An eligible existing group already carries the model. If the option says also grant model to group, approval grants the model to every member of that group as well as adding the requester.
- 3 Choose **Approve** or **Decline**, then read the recorded outcome. A permission or provider restriction can still require attention after the request is resolved; a submitted decision is not a successful model response.

CAREFUL

Approval changes group membership. When the chosen group does not carry the model yet, it can also widen that model's availability for all group members. Review the group's purpose and members before approving.

Trace one person's access

- 1 Open **Users** and choose **Access** in that person's Actions column. On a narrow table, scroll to expose the actions.

- 2 **Model access trace** lists the person's groups and each enabled model. Expand a model to inspect the policy gates in order and the reason for its result.
- 3 **Usable** means both access and provider readiness pass. **Allowed**, **provider offline** needs a platform-owner connection repair. **Blocked** requires the indicated account, group, model, or policy issue to be resolved.
- 4 Close the trace, make the authorized correction in **Users**, **Groups**, or **Model Access**, and reopen **Access** to verify the current server result. Ask the user to refresh model access and try the intended model.

GOOD TO KNOW

The trace is read-only and scoped to users you can administer. It explains existing permissions; opening it does not grant access, change a role, or bypass the organization's model ceiling.

32. Connections: response actions and connector handoff

Manage response actions and route shared connector setup to the service team.

The **Admin console**'s **Connections** tab contains **Chat output actions**. Shared connector switches, saved credentials, connection tests, and workspace OAuth are managed by the service team outside tenant administration.

GOOD TO KNOW

Users still connect their own Google, Microsoft, Box, or iManage account from the composer's attach menu. Those delegated connections respect each user's source permissions. Shared credentials managed by the service team support connector setup and background knowledge sync; they do not grant users someone else's files.

Chat output actions

Chat output actions adds admin-approved buttons to assistant responses for export, formatting, or handoff. Choose New response action to create one; existing custom actions offer Edit and **Delete**. Each row shows Enabled or **Draft** and an enable switch. Creating these actions does not configure a shared source connector.

MCP connections and model-callable tools remain in **Knowledge/Tools** → **Tools** → **Connections**. **Prompts** and **Skills** also remain in the **Tools** library. Their authoring and use follow the existing permissions and the shared connector availability set by the service team.

33. Automations: preparing scheduled runs

What admins need to configure before users can rely on recurring model chains.

Users create and run automations from the Automations tab of the **Agents** view, and enabled schedules fire automatically in the background (UTC), delivering each run's output as a new chat thread. Admins manage

group and model access; the service team manages shared connector availability and credentials outside tenant administration. Check the actual failure details to determine which setup needs attention.

Admin readiness checklist

- 1 **Open Model Access** and confirm the automation's model is visible to the user's group.
- 2 **Open Groups** and confirm the user belongs to the group that receives the model grant.
- 3 Ask the service team to confirm any required shared connector and inspect its real connection-test result. **Users** complete their own source-account connection when the workflow needs delegated access.
- 4 If the automation uses knowledge, confirm the knowledge base is enabled and its source ACL allows the user's group.
- 5 Ask the user to press **Run now** after saving the automation; the transcript and last-run status will show whether the setup is complete.

GOOD TO KNOW

Admins do not need to expose every model to make automations work. Grant the smallest approved model set that fits the scheduled workflow, then expand only when a run proves it needs more capability.

34. SSO: single sign-on for your tenant

Connect an identity provider, provision on first sign-in, and map IdP groups.

GOOD TO KNOW

If organization policy does not permit **SSO** management in this console, the tab is read-only and says so plainly.

- 1 Open the **SSO** tab and click **Add SSO configuration**.
- 2 Pick a preset — Microsoft Entra ID, Google Workspace, Okta — or choose custom OIDC and enter any issuer. Presets prefill the issuer URL; for Entra, replace the `{tenant-id}` placeholder with your directory ID.
- 3 Paste the **Client ID** and **Client secret** from your identity provider's app registration. The secret is vaulted server-side and never shown again.
- 4 List the email domains allowed to sign in through this provider.
- 5 Decide on just-in-time provisioning: when on, the first sign-in from an allowed domain creates the account automatically with the User role. Remember — they stay **Pending** until you assign a group.
- 6 **Copy** the redirect URI shown at the bottom of the form and register it with your identity provider. Sign-ins cannot complete until the provider trusts that exact callback address.
- 7 Click create. Enforcement always starts off, so nothing can lock the tenant out.

- 8 On the new configuration's card, click **Test connection** — it performs a real discovery and key check against the provider.
- 9 Map identity-provider group values to tenant groups on the card, so JIT users land with the right access.
- 10 After the discovery and key test passes, complete a real sign-in with a permitted test account. Verify the redirect back to this workspace, the resulting account and group access, and any required MFA. Only then consider enforcement.

CAREFUL

Enforcement blocks password sign-in for the allowed domains. A passing **Test connection** checks discovery and keys; it does not prove the client-secret exchange, callback, user session, or group mapping. Validate those in a real login before making enforcement the last step.

35. Analytics

Runtime metadata, feedback, per-user usage, and token budgets for this tenant.

The **Analytics** tab measures how the organization actually uses the platform. Sections start collapsed behind descriptive headers — click a header to expand one — and each section carries its own filter row pairing a person picker (defaulting to **All admins and users**) with a date range and its presets (All, Today, Week, 30 days). Scoping one section never narrows another.

- **Runtime Clock Metadata** — organization-scoped execution timestamps captured from admin and user **Chat** and **Draft** completion events. The cards total runtime events, chat completions, and draft calls. Its **CSV** includes `actor_id` and `actor_name` columns, and when one user is selected the filename gains that user's suffix so exports stay unambiguous.
- **Chat Feedback Analytics** — response-level thumbs up and thumbs down signals submitted by tenant admins and users, with totals for overall, positive, and negative feedback.
- **Model Activity** — saved prompt volume by model, date, and person for the selected range, drawn as “**Prompts** by model”, “Prompt trend”, and “**Users** by prompt activity”.
- **User Usage** — durable per-user usage for this organization's admins and users across chat, drafts, agents, automations, and the API gateway. Token counts are provider-reported only and stay blank when a provider reported none. The panel has its own **Filter usage by user** selector for drilling into one person.
- **Workspace Usage Budget** — a read-only view of the organization's service-managed ceiling and its current UTC accounting period. This panel shows where usage stands without exposing service-level configuration.
- **Token Allocations** — per-user and per-group token caps beneath the workspace ceiling, with Per day, Per week, or Per month UTC resets. The most restrictive applicable cap wins. Choose a user or group, enter a cap, pick the period, and click **Set cap**. Until you add one, the table says “No allocations yet. Everyone shares the workspace ceiling.”

GOOD TO KNOW

Every number here comes from saved audit and usage events, and each **CSV** export button opens its own small date-range popover so a file contains exactly the rows you chose. If a panel has nothing to show, it says so plainly instead of showing sample data.

36. Policies: service availability and memory governance

Apply downstream tenant defaults, configure memory, and govern by count without reading content.

The Policies tab is always present. **Policy Controls** starts collapsed. When service policy enables memory, Personalization Memory and Memory by User also appear collapsed. Otherwise, Memory governance explains the restriction; saved organization settings remain intact. Expand the panel you need; administrators can narrow available capabilities, while unavailable settings remain locked.

Policy Controls: read service availability

- Administrator accounts — shows whether this console can create and manage administrators or whether new administrator accounts require service approval.
- Admin sign-in policy — shows whether admins must use **SSO** or may use an explicitly provisioned local account.
- **SSO** configuration — shows whether tenant **SSO** mappings are delegated or read-only.
- New model defaults — shows whether newly available models begin with **Default Users** or require explicit grants.

Below those status rows are defaults for the protected **Default Users** group: personal API keys, private-agent building, private knowledge-base and tool authoring, and personalization memory. A switch locks when its capability is unavailable under service policy; the saved group grant is preserved rather than silently erased. Use **Groups** for exceptions and **Model Access** for available models. Shared connector availability is managed by the service team outside tenant administration.

Personalization Memory

- 1 When service policy permits memory, expand Personalization Memory and turn on Memory for this organization. This makes the account-level Personalization memory row available to eligible users; existing memories stay saved but are not applied while this switch is off.
- 2 Choose whether to allow Learn from conversations automatically. When off, only explicit requests such as “remember that ...” and direct additions in the memory manager create memories. Every user can still opt out of automatic learning individually.
- 3 Set Retention (days), from 1 through 3650. Older memories retire automatically when they pass this policy.

- 4 Set Maximum memories per user, from 1 through 2000. The default is 200, which is a practical general-purpose limit; raise or lower it only when your retention, compliance, or workload policy calls for a different capacity. When the cap is passed, the least useful unpinned memories retire first.
- 5 Expand Memory by User to review content-free counts. Click **Refresh** to reload the counts, or purge one person's memories when compliance or account cleanup requires it.

CAREFUL

Memory administration never grants reading access. Administrators see policy, counts, and purge controls only; the API and the interface do not return another person's memory content.

37. Data retention and conversation tags

Find tagged conversations, inspect their contents, and review batch actions.

- 1 Open **Audit** and expand Data Retention, then choose Schedule and rules. Forever is the default: automatic deletion is off until an administrator previews and saves a schedule. Choose 1, 5, 7, or 10 years, or keep Forever. One year means 365 days.
- 2 Choose whether age starts at chat creation or the last message change. Renaming or archiving a chat does not restart its clock. Select a review window of at least 7 days. Saving a changed policy restarts that window for eligible chats.
- 3 Add a stable client, matter, or regulated-record source and its aliases. A source can have a longer retention rule or Forever. Choose Apply time limits only to labels with a rule to leave all other chats stored indefinitely. The longest applicable duration wins; a matching Forever rule or active legal hold prevents automatic deletion.
- 4 Save the source definitions, then use Scan existing chats in Data Retention → Tags and holds. New message saves also look for source mentions. Detected mentions are suggestions only: select the correct chats and source, then Confirm label. Remove / dismiss label rejects a false match. The label records who confirmed it and when.
- 5 Optional sensitive-data suggestions check saved message text for email addresses, possible Social **Security** numbers, and payment cards. They do not detect every form of PII or regulated record, or inspect original uploaded files. Raw matched values are never copied into labels. Review before confirming a category.
- 6 Choose Preview effect to see counts across all saved chats, including legal holds and chats without a deadline. Saving the policy recalculates existing and future chats. A shorter duration can make old chats eligible, but they still receive the review window. A longer duration delays deletion; it cannot restore deleted content. Saving Forever stops automatic deletion.
- 7 To inspect chats, open **Audit**, expand Data Retention, and choose Tags and holds. This list includes tagged and untagged conversations available in your administrative scope.

- 8 Use **Search** chats and tags to find titles, people, tags, or client/matter identifiers. Filter by tag type when you need a narrower set.
- 9 Click a conversation title to open Tagged conversation and review its saved prompts and outputs. Close the preview to return to the same list.
- 10 Select the intended chats. **Archive** selected keeps them stored and searchable but removes them from the active list. **Delete** selected opens a permanent-deletion confirmation; it does not delete until you choose Yes, delete.
- 11 Under Legal holds, name a hold and choose Hold selected chats. This hold protects the selected records, including against manual deletion. For ongoing preservation of a client, also set that source's rule to Forever. Load active holds to review them; releasing a hold requires a second confirmation.
- 12 Read the completed action status. **Chats** under an active hold are skipped by deletion. A failure or skipped record needs review; do not assume every selected row was deleted. The background scheduler handles automatic cleanup in bounded batches and writes an audit record per deleted chat.

CAREFUL

Permanent deletion removes the selected conversations and their attachments and cannot be undone. Preview contents and check the selected count before confirming. A tagging switch records classification; it is not a scheduled deletion rule or an action that places a hold.

38. Review feedback and reported issues

Read response ratings, written feedback, and platform reports in Analytics.

- 1 Open **Analytics** and expand **Chat** Feedback. Set its person and date filters; these apply to that feedback panel.
- 2 Select a response rating to preview its feedback and saved conversation. Read any written comment alongside the model output.
- 3 Under Reported platform issues, select a report to review its subject, description, reporter, time, and optional screenshot. If a screenshot cannot load, read the error rather than treating it as an empty attachment.
- 4 Use the report to reproduce the problem with permitted test data, then follow your organization's support process. Viewing a report does not send a reply or mark the problem resolved.

39. Audit

Governance signals, prompt monitoring, security alerts, and the tenant trail.

The **Audit** tab is the tenant's governance station. It opens on the **Admin Audit** posture dashboard while the sections below start collapsed — click a section header to expand it. **User Prompt Activity**, **Security Alerts**, and the **Audit Trail** each carry their own filter row pairing a person picker with a date range, so one section's

scope never narrows another. The **CSV** export buttons produce files of exactly the filtered rows — including actor_id and actor_name columns, so every exported row names who acted.

- **Admin Audit** — summary cards for tenant security and governance signals: audit events, the prompt watchlist of active DLP or misuse alerts, prompt volume, active admins, active users, connector issues, and ungrouped models. Cards that need attention are highlighted; hover one to see what it counts.
- **Recent Governance Activity** — the current tenant snapshot for identity, user, model, and connector posture: how many **SSO** configurations are enforced, how many users are active, and how many connectors are enabled.
- **User Prompt Activity** — drill into saved prompts from this organization's admins and users and their model responses by person, thread, model, and timestamp, scoped by the section's own user and date filter.
- **Security Alerts** — DLP and malicious-behavior flags raised from admin and user prompts, shown with redacted snippets for review. Click **Acknowledge** once an alert is handled, or **Reopen** if it needs another look.
- **Audit Trail** — the tenant's append-only transaction log, newest first. Its toolbar has a severity select, an action-category select, and a search box (**Search actions, people, targets...**) that narrow the rows together; the trail's **CSV** exports exactly the rows you are looking at. Click **Refresh** to reload it straight from the admin audit API.

GOOD TO KNOW

Security alerts fire on real prompt content — payment card numbers, shared credentials, prompt-injection attempts — but show only redacted snippets, so reviewing an alert never re-exposes the sensitive value itself.

40. Alerts

Watch rules over audit activity, and honest delivery statuses for every alert.

The **Alerts** tab turns audit activity into notifications. It has three panels: **Email Delivery**, **Alert Rules**, and **Alert Deliveries**. Every alert is always logged in-app regardless of email — email is a delivery channel, not the record.

Email Delivery

For organization administrators this panel is a read-only status: it reports whether email delivery is configured at the service level. If email is not configured, rules still work — their alerts are logged in-app and the delivery log says so honestly.

Alert Rules

- Click **Suspicious-activity template** to start from a sensible security rule, or **New rule** to build one from scratch.

- A rule has: a name; **Action patterns** — exact audit actions or prefixes such as security.* or admin.user_deleted; a **Minimum severity**; an optional **Watched user**; a threshold — how many matches within a time window before it fires; a cooldown between alerts; and email recipients. Leave recipients empty and the alert is in-app only.
- Rules you create here watch this organization's admin and user audit activity.

Alert Deliveries

Every alert trigger is listed with its real delivery status: sent, queued, failed with the actual SMTP error, email not configured, or logged in-app. The log exports to **CSV**, archived deliveries included. Click **Archive** on a delivery to clear it from the default view — its history is kept, and **Show archived** reveals archived deliveries so you can review or **Restore** them. Tenant admins can archive their tenant's deliveries only.

The Platform owner console

41. Opening the Platform owner console

The highest access level, its six tabs, and the role ceilings.

- 1 Click your account card at the bottom of the sidebar to open the account drawer.
- 2 Expand the **Management** section and click **Platform owner console**. (Only platform owners see it.)
- 3 The console has six tabs: **Org Settings**, **Models**, **Providers**, **Analytics**, **Audit**, and **Alerts**. **Org Settings** holds organization configuration; API keys live on each provider's card under **Providers**.

Keep the three role ceilings in mind — they explain who can touch what across the entire platform:

ROLE	WHAT IT CONTROLS
Platform owner	Providers , API keys and secrets, shared connectors, organization model availability, SSO , branding, policies, usage budgets, platform updates, alert email, and removing admins. The highest level.
Admin	Tenant users, groups, response actions, and tenant model access — always inside the boundaries the owner sets.
User	Chat , drafts, assigned agents, and whatever models and sources their groups allow.

GOOD TO KNOW

First-run setup: when no active platform owner exists yet, the sign-in screen shows “**Create the first platform owner**” — enter a display name, email, and a password of at least 12 characters. This screen never appears again once an owner exists.

TIP

The **Documentation** button opens owner walkthroughs and this PDF. Its **Admin documentation** and **Chat help** links open the other role libraries. **Interactive platform guide** opens the public digital walkthrough at <https://aperturechat.com/guide.html> in a new tab for additional platform setup and configuration guidance.

42. First-run setup: from owner account to a working team

Create the first owner, connect a model, and verify access before inviting the team.

- 1 On a new installation with no active owner, complete **Create the first platform owner** with your display name, work email, and a password of at least 12 characters. Confirm the password and choose **Create platform owner**.
- 2 The Getting started card offers the setup action for your role. Open owner guide opens this role's documentation. Choose an action or explicitly dismiss the card; merely loading the workspace does not mark it reviewed.
- 3 Open the account drawer, expand **Management**, and choose **Platform owner console**. **Documentation** opens the owner guide, narrated lessons, and the **Interactive platform guide** for step-by-step configuration advice.
- 4 Open **Providers**, register the intended gateway, and save its real credential in **API Keys**. A successful save confirms configuration storage; Needs validation still requires a successful runtime check.
- 5 For a provider that supports catalog discovery, choose **Sync Models** on its card. This refreshes the catalog and tests a small live model request. Read the returned status and correct credential or runtime failures before relying on the connection.
- 6 In **Models**, review the resulting catalog and enable only the models you intend to offer. **Providers** with manually managed catalogs need their model configuration checked separately.
- 7 Open **Admin console** to configure the team's groups, model access, and any permitted knowledge or tools. **Approve** access requests or create accounts, then complete the sign-in handoff described in **Users**.
- 8 Verify a real first message with a synthetic standard-user account. Confirm the intended model works, restricted resources stay unavailable, and any temporary password or authenticator requirement completes correctly.
- 9 If you will use **SSO**, complete discovery testing and a real end-to-end sign-in before enforcing it. Keep an authorized administrative sign-in path available while validating the setup.


GOOD TO KNOW

Before inviting a team, verify the intended provider, enabled models, group memberships, and model grants. Test each intended user role; a working owner account does not prove that a standard user has access.

Use **Providers** and **Models** for service configuration, then Admin Console → **Model Access** and **Groups** for tenant access. Use **Org Settings** for **SSO**, branding, connector availability, policy, and budgets.

43. Providers

Register model gateways, read connection health, and sync catalogs.



Provider cards show the actual connection and model-catalog state. Example workspace with synthetic data.

- 1 Open the **Providers** tab and click **Add Provider**.
- 2 Pick the kind. Eleven are available: openai, anthropic, azure-openai, azure-foundry, gcp, amazon-bedrock, open-webui, openrouter, ollama, openai-compatible, and local. Save the base URL, auth type, region, and the label for its first key.
- 3 Read the provider's card after saving. It shows the provider's brand logo, the kind, region, and a **Models** counter that reads enabled "of" total — for example "12 of 40" — so you can see the ceiling at a glance.
- 4 Read the status badge separately from the save confirmation: **Needs key** means no active credential is available; **Needs validation** means the saved connection and credential still need a successful runtime check; **Adapter needed** means the required runtime adapter is unavailable; **Connected** reflects a validated usable route. Saving a provider or key is not a successful model test.
- 5 Click **API Keys** on the card — the button shows how many keys are stored — to manage that provider's vaulted credentials. The next section covers the vault in detail.
- 6 Click **Sync Models** to pull the provider's catalog. The button stays disabled until an active key exists, and says so.
- 7 Use **Edit Connection** whenever a base URL, auth header, or catalog scope needs correction.

Kind-driven defaults

- Picking a kind prefills sensible connection defaults. azure-foundry uses the inference endpoint `https://{resource}.services.ai.azure.com/models` with api-key authentication; gcp routes through Google's OpenAI-compatibility endpoint for Gemini models.
- The **Catalog scope** row appears only on openrouter providers: choose between the zero-data-retention (ZDR) filtered list and the key-scoped model list.

GOOD TO KNOW

A provider is not usable until it has an active key in its vault — open **API Keys** on the card and add the key next.

44. API Key Vault

Each provider card vaults its own secrets; reveal, replace, or delete them.

Each provider's secrets live in a vault attached to its own card — there is no separate keys tab. Open the **Providers** tab and click **API Keys** on a card to expand its **API Key Vault** panel.

- 1 Click **Add Key** and enter the key's name, environment, expiry, and the secret itself. The secret is saved to the backend vault — only a masked value ever reaches the browser.
 - 2 Read each row's metadata: the key's name, environment, active status, when it was last rotated, and its expiry — stale credentials are visible before they fail.
- **Reveal** — opens the **Vault reveal** dialog, which shows the secret with a **Copy key** button and a **Done** button to hide it again. Expired keys cannot be revealed.
 - **Replace** — swaps in a new provider-generated secret in place.
 - **Delete** — removes the key.

Reveals, replacements, and deletes write through the platform API, so the audit trail stays complete.

45. Models: organization availability

The organization ceiling — control which synced models tenants can route to.

- 1 Open the **Models** tab. It lists every model synced from your providers.
- 2 Narrow the list three ways: the search box, the All / Enabled / Disabled status filter (each option shows a live count), and per-column filters on the **Provider**, **Model lab**, and **Runtime route** columns — click a column's filter icon to tick specific providers, labs (OpenAI, Anthropic, Google, and so on), or routes.
- 3 Read each row: the provider, the model, the exact runtime route requests will use, and whether the organization has it enabled.
- 4 Flip the org status toggle to enable or disable a model. This is the ceiling: models disabled here never reach tenant admins or users.

- 5 Click **Edit details** on any row you allow, to record its display name, runtime route, context window, notes, system prompt, and meta prompt.

GOOD TO KNOW

Newly synced models arrive disabled with no group access. After a sync, switch the filter to Disabled, review the new arrivals, and enable only the routes the organization has actually approved. If the **Default group for enabled models** policy is on (see “Policies”), enabling a model automatically attaches the protected **Default Users** group so users see it without a second step.

46. Org Settings: users and role boundaries

Create accounts at any level, set passwords, and rely on the account floors.

The **Org Settings** tab gathers organization-level controls: roles and accounts, single sign-on, branding, policies and budgets, platform connectors, and search index readiness. Most sections start collapsed behind a descriptive header; click a header (or its chevron) to expand it. **Search index** starts open. Expand **Role Boundary** for the account tools.

- 1 In the create form, enter a display name and email, and pick a role: User, Admin, or Platform owner.
- 2 Click **Create account**. It writes through the platform admin API, so the account survives refresh and restart.
- 3 Manage existing accounts in the list below: change a role from its dropdown, set a password with the key button, or remove an account with the trash button.

Setting a password

The key button on a row opens the same dialog admins use — “Set a password for” that person. Type a password or click **Generate** for a strong random one, optionally mark it a **Temporary password** so they must choose their own at first sign-in, and click **Set password**. The password is shown only in that dialog — share it over a safe channel.

GOOD TO KNOW

Two floors protect the platform. The last active platform owner can never be removed. And at least one active administrator (owner or admin) must always remain — because owners count as administrators, the sole tenant admin is removable while an active owner exists, but the API refuses any removal that would leave no administrator at all.

Below the account list, the “Clear separation of duties” callout restates the boundary: platform owners manage provider secrets, owner accounts, admin delegation, **SSO** baselines, and platform branding.

47. Org Settings: single sign-on setup

Choose the protocol, connect a real OIDC provider, and register the redirect URI.

The **Single Sign-On** panel configures live OIDC. **Users** are redirected to your identity provider, and the returned ID tokens are cryptographically verified before a session is issued.

- 1 In **Org Settings**, expand the **Single Sign-On** section by clicking its header.
- 2 Pick the **Protocol** first. **OIDC (supported)** is the live sign-in path. SAML appears but is disabled — labeled “SAML — Deferred, not a working sign-in path” — so a stored SAML configuration stays honest without pretending to sign anyone in. **SCIM** stores a **SCIM** provisioning base URL and token instead of a sign-in flow.
- 3 Pick a preset — Microsoft Entra ID, Google Workspace, or Okta — or enter any OIDC issuer URL. For Entra, replace the `{tenant-id}` placeholder in the issuer with your directory ID.
- 4 The issuer drives everything else: the platform fetches the provider's discovery document from `{issuer}/.well-known/openid-configuration` to learn its endpoints and keys.
- 5 Paste the **Client ID** and **Client secret** from your app registration (for **SCIM**, the secret field doubles as the **SCIM** token). Secrets are vaulted server-side only — never returned to the browser.
- 6 For **SCIM**, fill the **SCIM base URL**; if you document Duo MFA, the **Duo API hostname** field records it alongside.
- 7 **Copy** the redirect URI shown in the panel and register it with your identity provider. Sign-ins cannot complete until the provider trusts this exact callback address.
- 8 List the allowed email domains — only accounts on those domains can sign in through this provider.

48. Org Settings: SSO claims, MFA, and go-live

Separate identity-provider MFA from the platform authenticator and verify the complete sign-in path.

- Role and group claims — tell the platform which token attributes carry your identity provider's role and group assignments.
- MFA documentation — the **Authenticator app** select records which app your organization uses (**Microsoft Authenticator**, **Duo Mobile**, or **Identity provider**), alongside the MFA methods the provider enforces, so the sign-in experience is documented where admins look for it.
- Enrollment QR — add an enrollment URI (a standards-based TOTP link or a Duo enrollment URL) and the panel renders it as a scannable QR code for authenticator setup.
- **Require the platform authenticator after SSO** — off trusts the identity provider's MFA; on also requires users to enroll and verify with Aperture Chat's authenticator. Choose deliberately to avoid an unintended second challenge. Recording an authenticator name or enrollment link alone does not enable this requirement.

- Just-in-time provisioning — creates new accounts automatically the first time someone on an allowed domain signs in. They arrive with the User role until you promote them.

- 1 Click **Save SSO**.
- 2 Click **Test connection** — it performs a real discovery and key check against the provider. The test requires the OIDC protocol; with SAML or **SCIM** selected the button is disabled and its tooltip says to switch to OIDC.
- 3 Complete a real sign-in with a permitted test account, including client-secret exchange, redirect back to the workspace, session creation, role/group mapping, and any required authenticator challenge.
- 4 Only after that full sign-in succeeds, consider enforcement and confirm users know which sign-in and recovery path to use.

CAREFUL

Enforce SSO blocks local password sign-in for every allowed domain. **Test connection** proves discovery and keys, not a completed login. Treat enforcement as the last step after the real login succeeds; a misconfigured provider can lock out administrators on those domains too.

49. Org Settings: platform branding

Rename the product, swap the logos, and recolor the theme everywhere.

- 1 In **Org Settings**, expand **Platform Branding**. It starts with a live preview of the current name and logo, exactly as users see them.
- 2 Fill in the identity fields: the platform name, the **Platform logo URL**, and the **Browser icon URL** (the tab icon).
- 3 Record the **Platform domain** if you use one. As the note under the field says, it is “Recorded for admins and the API” — you still point DNS and TLS at this deployment separately; saving here does not move traffic.
- 4 Set the **Theme colors**: **Accent color** for buttons and highlights, **Sidebar gradient start** and **Sidebar gradient end** for the rail (pick darker stops so the light sidebar text stays readable), and **Interface text color**, which applies to the light theme.
- 5 No image URL handy? Click **Upload PNG** and the platform hosts the image for you — **PNG** files up to 4 MB.
- 6 Click **Apply branding** to roll the change out everywhere the default brand appears, or **Reset defaults** to restore the original identity in one click.

GOOD TO KNOW

Branding reaches further than the shell: the same identity feeds the runtime theme colors and the per-tenant install manifest and icons, so the app people add to their phone's home screen carries your name and logo, not a generic one.

50. Org Settings: policies and budgets

Set the organization ceiling and the workspace usage budget.

Policy Controls

Expand **Policy Controls** in **Org Settings** when you need it. The first row is not a toggle at all: “Only owners can create platform owners” carries an **Always on** pill because the platform enforces it unconditionally. Below it are eight real switches:

POLICY	WHAT IT ALLOWS WHEN ON
Downstream API access	Owners and admins can use the platform API downstream; standard users still need the group grant Can use API from their admin.
Tenant admins can create admins	Delegates admin creation to tenant admins; off means owner-only.
Require SSO for admins	Admin accounts must sign in through SSO , not passwords.
Tenant admins can manage SSO mappings	Lets tenant admins edit their tenant's SSO configuration and group mappings.
Default group for enabled models	Newly enabled models automatically include the protected Default Users group, so admins do not need a second step per model.
Users can build their own agents	The first half of a two-part gate: with this on, admins can grant Can build agents to a group, and those users can build private, self-owned agents. Publishing to the organization stays admin-only. Both halves default off.
Personalization memory	Enables the tenant-admin handoff. Admins may then turn memory on for their organization and grant Default Users access. Off locks both downstream controls while preserving saved grants. Admins and owners see counts and purge controls only, never memory content.
Users can browse the model catalog	Users can inspect all models enabled for their organization, read availability reasons, and request access where permitted. Off shows only already-usable models and refuses access requests. This does not reveal private prompts or notes.

GOOD TO KNOW

Active policies define the organization ceiling. Tenant admins can only operate inside these boundaries — nothing they configure can exceed them.

Workspace Usage Budget


- The **Workspace Usage Budget** panel sets a hard spending ceiling for the whole workspace. Choose the **Budget measure** — **Token allowance** (exact provider-reported tokens) or **Dollar amount (USD)** (exact provider-reported cost) — and the **Reset period: Every day, Every week, or Every month**, on the UTC calendar (weeks begin Monday).
- Enter the limit and click **Save budget policy**. A limit of 0 means unlimited. The usage card beside the form shows the live count for the current period, split into provider-reported and unreported completions so the number stays honest.
- **Once** the ceiling is spent, further completions are refused until the period resets. Admins can add per-user and per-group allocations beneath this ceiling from their **Analytics** tab.

Elastic Analytics

The **Elastic Analytics** panel reports whether backend audit export is configured, its status, and buffered events. **Configure** `APERTURE_ELASTIC_URL` or `APERTURE_ELASTIC_CLOUD_ID`, together with `APERTURE_ELASTIC_API_KEY`, in the backend environment and restart the API. The browser never receives this API key. Buffered events are delivered once the configured cluster is reachable; the panel has no connection form to save.

51. Org Settings: shared connectors and source credentials

Configure deployment-wide availability, test real connections, and distinguish shared access from personal sign-in.



Shared connector configuration is in the Platform Owner console under Org Settings. Example workspace with synthetic data.

Open **Platform owner console** → **Org Settings** and expand **Connectors**. This owner-only panel controls shared source settings and availability. Turning a connector off removes that capability across chat, source pickers, the command palette, the **Tools** library, and the API. Tenant administrators do not configure these shared connections.

Credential-backed sources show **Credentials saved**, **Saved · disabled**, or **Needs credentials**. A saved credential is configuration evidence; read **Test connection** for the live result. Switch-only capabilities, such as MCP Servers and Prompt **Library**, have an enable switch without a vendor credential form.

Configure and test a source

- 1 Choose **Configure** on the source row. Select Authentication method and fill in the fields shown for that method. Read the source's setup notes for its required permissions and redirect URI.
- 2 Choose Save configuration. For an existing saved secret, leave its password field blank to retain it, or enter a new value to replace it. Wait for the save result before testing.
- 3 For Google OAuth, save the client ID and secret first, then choose **Connect Google** Drive to authorize the workspace account used for knowledge sync. Complete the consent flow before checking the connection.
- 4 Choose **Test connection** and read its result and individual checks. An incomplete, failed, or missing result is not proof of access. Resolve the reported issue before relying on the source.

- 5 Set the source's enable switch for the deployment and verify the resulting status. **Configure** the intended users' source-account access separately when they will attach files in chat.

SOURCE	CONFIGURATION SHOWN BY THE FORM
Google Drive	Google OAuth client ID and secret, with an optional Drive folder ID and source label. Paste access token is a testing option.
OneDrive / SharePoint	Microsoft Graph directory ID, application ID, and client secret for app-only access, with optional site, drive, and root-folder IDs.
Box	Client ID , enterprise ID, and client secret for Client Credentials Grant; an optional folder ID limits the starting location. Developer token is a testing option.
iManage	Instance URL, API key (client ID), and OAuth client secret for Each user signs in. Service account for background sync adds the service username and password; chat still requires each user's OAuth sign-in.

GOOD TO KNOW

Google, Microsoft, Box, and iManage chat attachments use each signed-in user's delegated source account. Shared knowledge-sync credentials do not replace that sign-in or bypass source permissions. **Users** keep their Connect action in the attach menu.

Clear saved configuration

Clear configuration removes the source's saved fields, stored secret, shared OAuth data, and service-account password, and disables its saved configuration. Use it only when that removal is intended. An empty secret field alone retains the saved secret; clearing all visible fields and saving an existing configuration also performs a clear.

Web Search

Choose **Configure** on **Web Search**, then select **Search** engine and Results per search (1 to 10). DuckDuckGo is keyless. SearXNG requires an instance URL with JSON search output enabled. OpenAI, Anthropic, and OpenRouter choices use the corresponding saved provider key and bill searches to that provider account. Choose Save configuration, then **Test connection** to run a real query.

The engine choice applies to models without hosted web search of their own. OpenRouter-backed models use OpenRouter's built-in search regardless of this selection. The **Web Search** enable switch still governs availability for the deployment.

52. Org Settings: search index readiness

Review index coverage and rebuild from live records when search needs maintenance.

- 1 Open **Org Settings** and find **Search index**, which starts open. This supports **Search past work** for chats and saved documents and decks; it is separate from knowledge-source ingestion.
- 2 Read the per-organization status: Ready or **Backfilling**, entry count, index mode, and completion time when available. While backfilling, the user search panel may warn that results are incomplete.
- 3 Use **Rebuild index** when the stored index needs rebuilding from current records. Wait for the completed result and refreshed entry count. A failure message requires review; a button click is not a completed rebuild.
- 4 Verify representative search terms with an account that is allowed to see the underlying work. Rebuilding does not broaden permissions: every search hit is checked against its live record and the requesting account.

GOOD TO KNOW

If indexing is disabled for this deployment, the panel explains that searches scan records on each request and disables **Rebuild index**. An index entry count is a coverage signal, not a count of items every user may access.

53. Review and install a platform release

Use the owner-only sidebar update notice to review release notes, install when available, and verify the reported result.

For platform owners, the sidebar says Release available followed by a version when the updater is unavailable, and Update to followed by a version only when it is ready. It can also show an active update or a recent undismissed result. Tenant administrators and users do not receive this control. The current version comes from the running build; a stale deployment version setting does not establish which build is running.

- 1 When the update row appears, hover over or focus it to read release highlights and the current version. Click it to open the update dialog.
- 2 Review What this update brings, and use Show full release notes or Release page when offered. Check the Release list checked timestamp and any error. Check again requests a fresh release list when no update is running.
- 3 Read How the update runs before proceeding. Installation restarts the API and web app, so arrange an appropriate interruption window and follow the deployment's backup procedure.

- 4 Choose Install followed by the target version only when ready. The button is available when the updater service is configured and connected. If the dialog shows Manual install on this deployment or an offline-updater message, follow those deployment instructions instead. A fresh VPS can use the release bundle installer to prepare private configuration and a stable project, then pull, start, and health-check the stack. Existing installations need the documented one-time API and updater-sidecar configuration in the same project; do not run a fresh-install workflow over existing data. Forks use their own repository release source by default. See the repository Docker release guide for the exact commands and prerequisites.
- 5 Follow the recorded progress through downloading, restarting, and verifying. Show updater log opens available detail. **Closing** the dialog does not cancel an update that has already started.
- 6 After Update installed, choose Reload now to load the new web build, then verify the workspace. If the result is Update failed or Update rolled back, inspect the reported details before using an offered Retry action. Dismiss hides a finished result; it does not install or repair a release.

GOOD TO KNOW

A release check or accepted install request is not a completed update. During an active run the API may be temporarily unreachable; the dialog continues polling. Wait for the reported outcome and verify the refreshed application before treating the release as complete.

54. Organization retention and tagging

Set the tagging policy and govern conversations within owner access.

In **Platform owner console**, open **Audit**, expand Data Retention, and choose Schedule and rules. The default is Forever, with automatic deletion off. Use the duration slider, review window, client and matter sources, and Preview effect exactly as described in the administrator retention chapter. Deployment never activates an existing metadata-only policy.

- 1 Open **Audit**, expand Data Retention, and choose Tags and holds. **Search** and filter the conversations before selecting rows.
- 2 Preview the saved prompts and model outputs for a conversation. If the data cannot load, resolve the error before deciding on an action.
- 3 Review the selected count and choose **Archive** selected or **Delete** selected. Read the confirmation carefully; archival preserves stored conversations, while confirmed deletion permanently removes eligible chats and attachments.
- 4 Inspect the resulting action status and skipped records. Under Legal holds, create a named hold for the selected chats or review active holds. Releasing a hold requires confirmation and starts a new review window before eligible records can be deleted.

CAREFUL

A suggested client or sensitive-data label is not a confirmed retention source. Automatic cleanup covers saved chats, linked uploads and image previews, search entries, and chat feedback. **Draft** documents, learned memories, audit and usage records, generated-media storage, backups, external providers, and exported files have separate lifecycles; this control is not a complete client-erasure or regulatory-compliance guarantee.

55. Analytics: runtime, feedback, activity, and usage

Authoritative execution timestamps, feedback, prompt volume, and per-user usage.

The **Analytics** tab carries four panels, and each panel scopes itself: a filter row at the top of every section pairs a user picker with a date range and its presets (All, Today, Week, 30 days), so narrowing one section never hides data in another. The filter's counter reads how many records the scope selected — for example “65 of 65 records”. Three panels carry their own **CSV** export controls, and each **CSV** button opens an independent date-range popover so a file contains exactly the rows you chose.

- **Runtime Clock Metadata** — authoritative execution timestamps captured from chat and draft completion audit events, not client guesses. The scorecards total the runtime events, split between main chat completions and draft generation and revision calls, and each event lists who ran it, which provider served it, how many messages were involved, and the exact client start time.
- **Chat Feedback Analytics** — response-level thumbs up and thumbs down signals submitted from chat actions, totaled by sentiment.
- **Model Activity** — saved prompt volume by model, date, and user for its filter's scope, drawn as “**Prompts** by model”, “Prompt trend”, and “**Users** by prompt activity”.
- **User Usage** — durable per-user usage from real completions across chat, drafts, agents, automations, and the API gateway, with provider-reported token counts only. The owner view is the complete one: its user picker reads **All owners, admins, and users** — unlike the admin console's version, which excludes platform-owner usage. The Usage by user section at the bottom ranks everyone in a contained, scrollable list; pick a person from its selector — or click their row — to focus the whole panel on them.

56. Automation readiness and governance

The owner-level controls that make scheduled model chains reliable.

Automations depend on the same governance stack as chat: provider health, vaulted keys, enabled models, tenant group grants, connector policy, and the workspace usage budget. Owners do not usually build every tenant's automation, but they set the ceiling that determines whether scheduled runs can work at all.

- 1 In **Providers**, confirm each gateway a schedule depends on shows **Connected** — not **Needs key**, and not **Adapter needed** (a Bedrock provider registers but cannot serve runs until its adapter ships).

- 2 Open **API Keys** on each provider card and replace invalid or expired keys before teams build scheduled workflows around them.
- 3 In **Models**, enable only the routes the organization has reviewed — newly synced models start disabled, and disabled models never flow to tenant admins or user automations.
- 4 In **Org Settings** → **Connectors**, configure and test the shared sources a schedule needs and keep their enable switches aligned with approved use. Check the **Workspace Usage Budget**: once the ceiling is spent, completions — scheduled ones included — are refused (the API answers with HTTP 429) until the period resets, and a refused run is recorded as failed rather than retried automatically.
- 5 In **Analytics** and **Audit**, review automation-related chat/runtime events the same way you review normal completions: who ran it, which provider served it, and whether the run produced the expected output.

TIP

For a new deployment, include one automation smoke test: create a weekly one-step run on an approved model, press **Run now**, confirm the transcript succeeds, then restart the API and verify the automation card and last-run status persist.

57. Audit: owner governance signals

Posture tiles, prompt monitoring, security alerts, and the platform trail.

The **Audit** tab opens on the **Owner Audit** posture dashboard; the sections below it start collapsed — click a section header to expand it. **User Prompt Activity**, **Security Alerts**, and the **Audit Trail** each carry their own filter row pairing a user picker with a date range, so one section's scope never narrows another. The **CSV** export buttons produce files of exactly the filtered rows — with actor columns, so every exported row names who acted.

- **Owner Audit** condenses governance posture into tiles: **Critical events**, provider posture, the model ceiling, vault metadata, approvals awaiting review, **Connectors**, expired keys, connector issues, unscoped models, privileged owners, stale syncs, and the prompt watchlist of active DLP or misuse alerts. Hover a tile to see exactly what it counts.
- **Recent Governance Activity** lists the latest owner-relevant events from the current snapshot — model availability reviews, provider catalog status, vault metadata, **Connector availability**, and **Agent approval activity** — so exceptions become follow-ups instead of surprises.
- **User Prompt Activity** drills into saved user prompts by person, thread, model, and timestamp — the owner-scope view of what is actually being asked across the platform, narrowed by its own user and date filter.
- **Security Alerts** lists DLP and malicious-behavior flags raised from actual prompts, with redacted snippets for review and **Acknowledge / Reopen** actions, scoped by its own filter row.

- The **Audit Trail** is the append-only transaction log of platform and tenant mutations, newest first. On top of its user and date filter, its toolbar has a severity select, an action-category select, and a search box (**Search actions, people, targets...**) that narrow the rows together.

TIP

Read the tiles first. Anything unexpected — an expired key, an unscoped model — has a matching tab in this console where you can fix it, and this guide's matching section tells you how.

58. Alerts: email delivery and platform-wide rules

Configure SMTP once, watch platform-wide activity, and read honest delivery logs.

The owner **Alerts** tab is the full version of the alerting station: you own the email configuration, your rules watch platform-wide audit activity, and every rule in the organization is listed with its scope. **Alerts** are always logged in-app regardless of email.

Email Delivery: configuring SMTP

- 1 Fill in the **SMTP host** and **Port**, and pick the security mode — **STARTTLS**, **SSL/TLS**, or unencrypted.
- 2 Enter the username and password. The password is stored in the encrypted vault and never shown again; the field's placeholder confirms when one is stored.
- 3 Set the **From address** alert emails will come from, then click **Save Email Settings**.
- 4 Prove it works: enter a test recipient and click **Send test email**. It sends a real message through the saved settings and reports the genuine result — including the SMTP error if it fails.

Alert Rules

- Rules work exactly as described in the Administrator Guide — **Suspicious-activity template** or **New rule**, with action patterns, a **Minimum severity**, an optional **Watched user**, a threshold within a window, a cooldown, and email recipients (empty recipients means in-app only).
- Owner rules are platform-wide: they watch audit activity across the whole organization, including owner actions. Tenant rules created by admins appear in the same list, each labeled with its scope, so you always see the full alerting picture.

Alert Deliveries

Every alert trigger is listed with its real delivery status — sent, queued, failed with the actual SMTP error, email not configured, or logged in-app — and the log exports to **CSV**, archived deliveries included. Click **Archive** on a delivery to clear it from the default view without deleting its history; **Show archived** reveals archived deliveries for review or **Restore**. Owners can archive any delivery, platform-scope ones included.